

Computer Graphics (COMP0027) 2023/24

Rasterization

Tobias Ritschel

Ray-tracing: photo-realistic, but not usually real-time



Goal: Photo-realism at interactive rate for VR/games



Summary

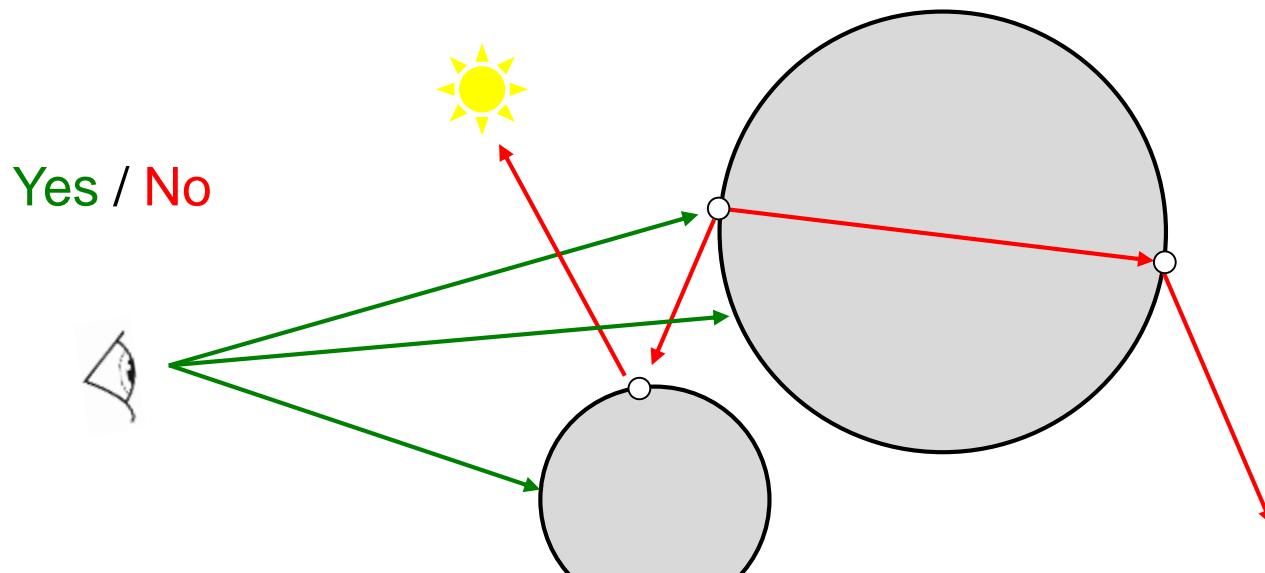
- Ray-tracing is elegant and general but slow
- Now we want to speed this up
- Main idea:
 - Don't **trace** rays
 - **Project** primitives
- Creates many new challenges

Ray-tracing cost

- The process of casting rays is very slow
- Example
 - 10,000 triangles,
 - 1000x1000 pixels
 - Result in 1,000,000 primary rays to cast
 - Each one testing for intersection with the 10,000 triangles (and then reflections, shadow rays, etc...)

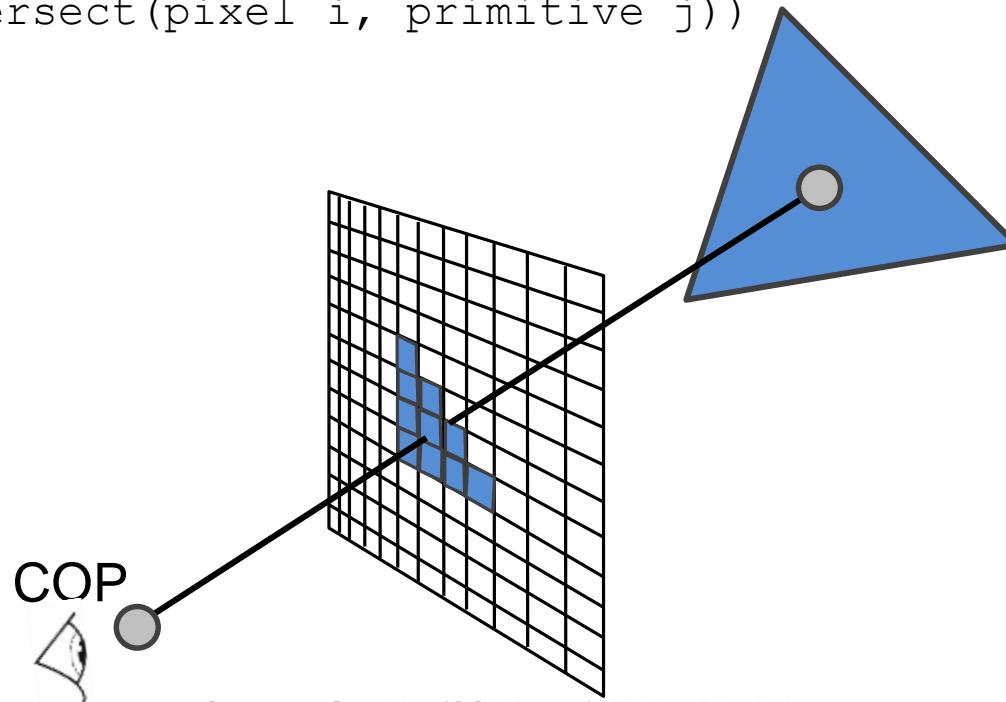
Simplifications from ray-tracing

- Rays only from one or a few specific points
- Only convex polygons
- No global illumination part (i.e. no recursion)



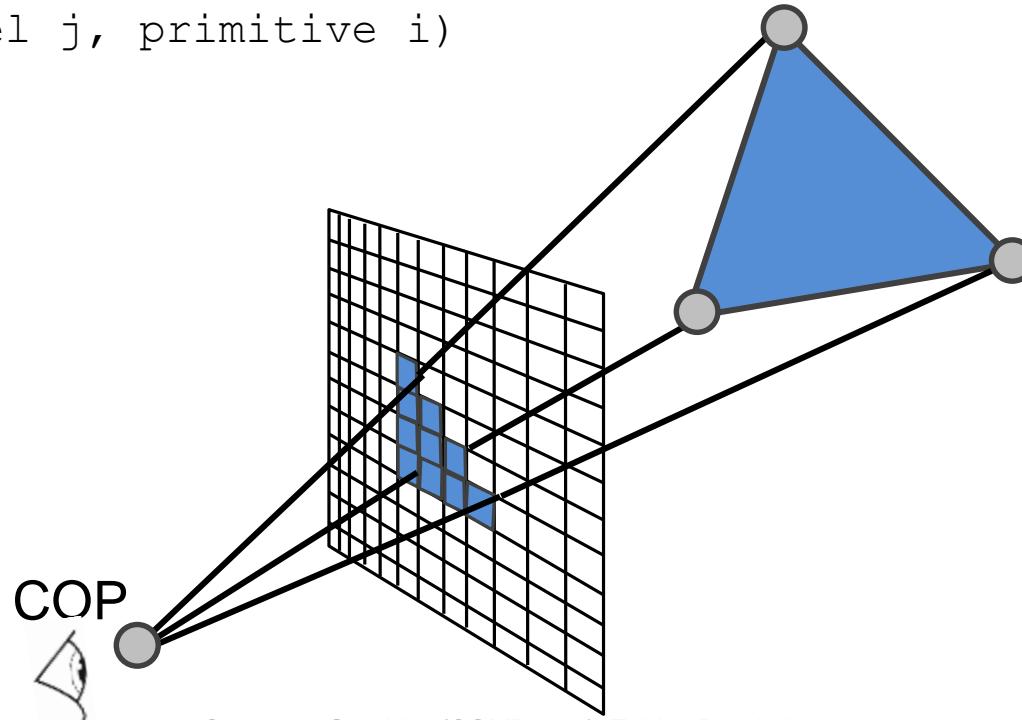
Ray-tracing

```
for all pixel rays i  
for all primitives j  
shade(intersect(pixel i, primitive j))
```

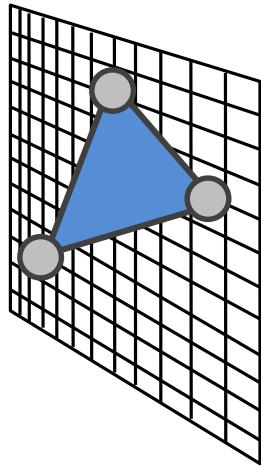


Rasterization

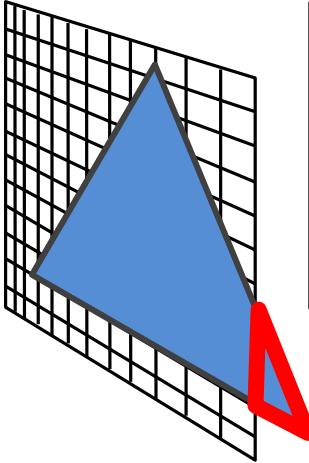
```
for all primitives i  
  for all pixels j in projection of primitive i  
    shade(pixel j, primitive i)
```



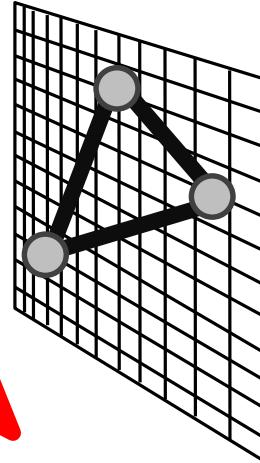
Challenges



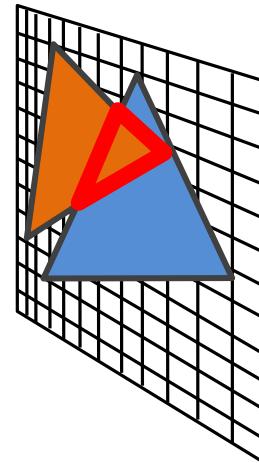
Projection



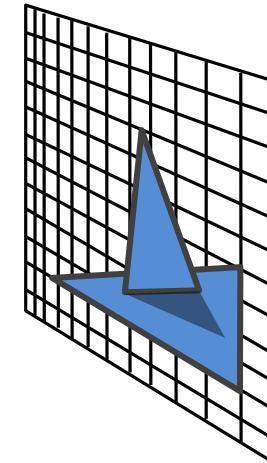
Clipping



Rasterization



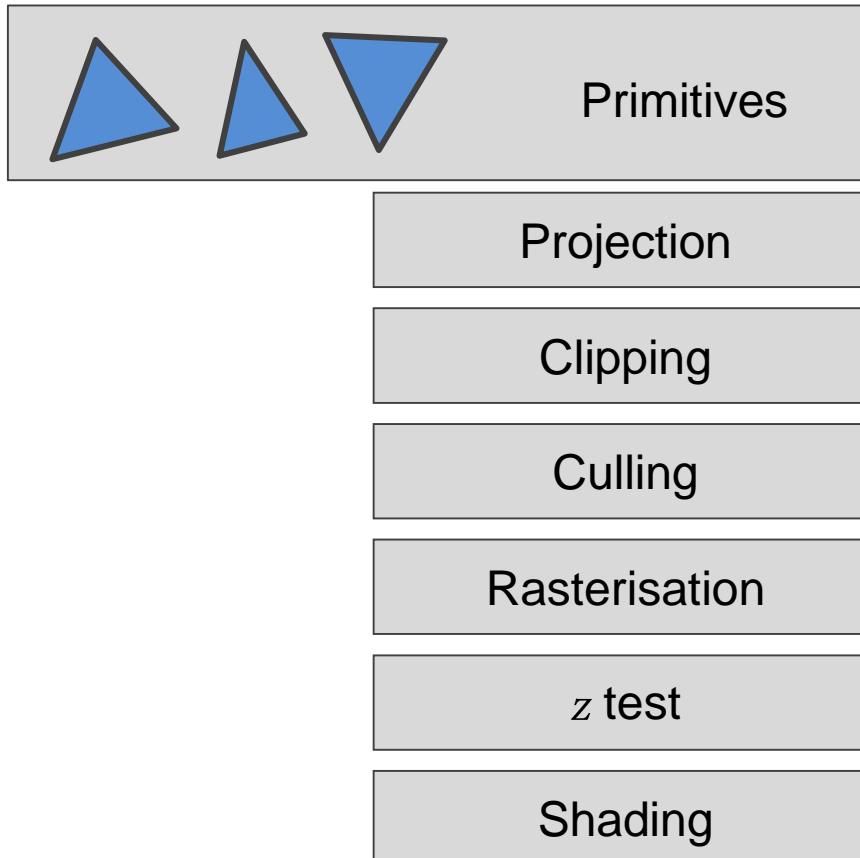
Visibility



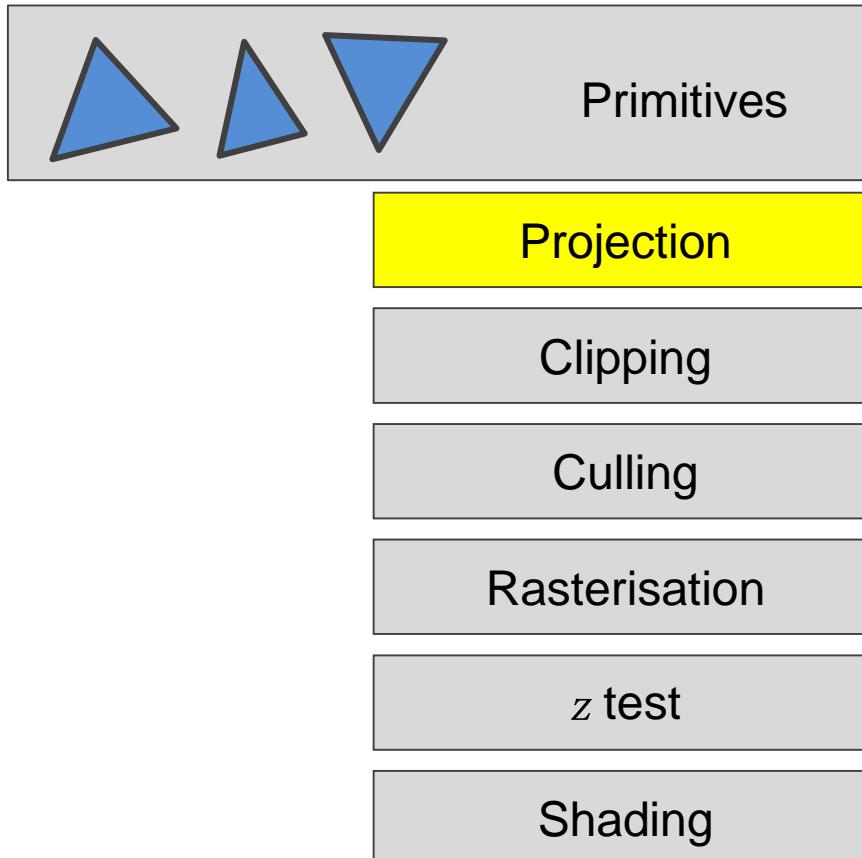
Shading

z-buffer

Pipeline

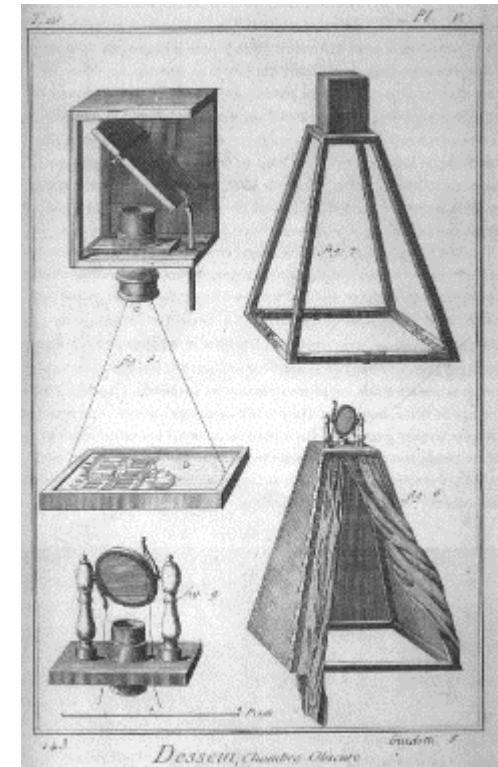
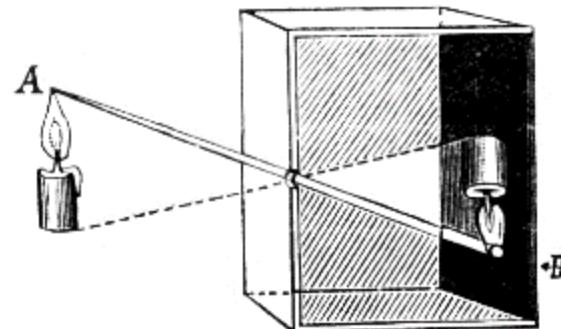


Pipeline



History of projection

- Camera Obscura
 - Earliest form of projection
 - Mo-Ti (470-390 BC)
 - Aristotle (384-322 BC)



History of Projection

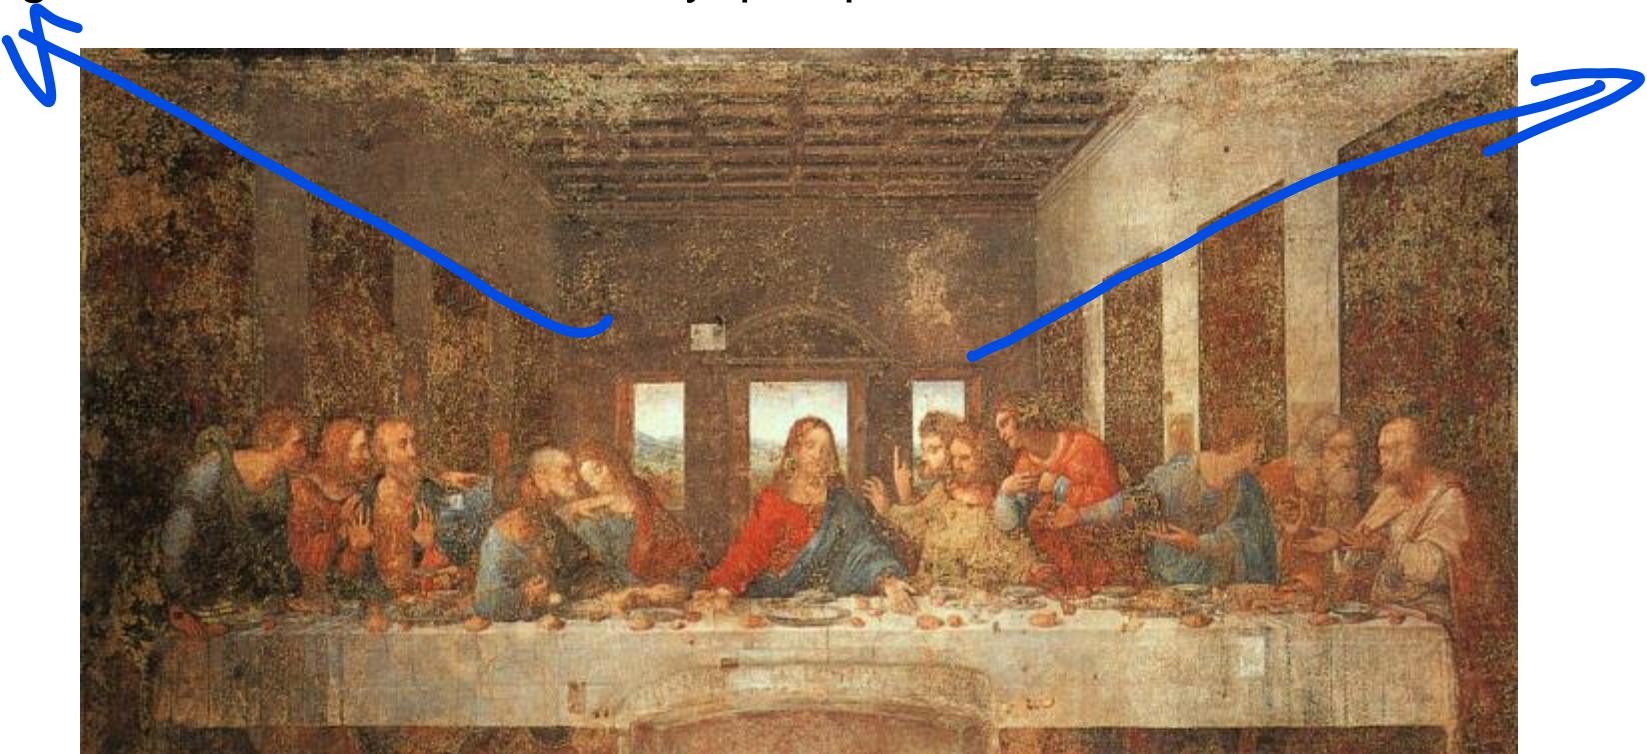
- Ancient Greeks knew the laws of perspective
- Renaissance: perspective is adopted by artists



196. From an Apulian calyx-krater. Würzburg

History of Projection

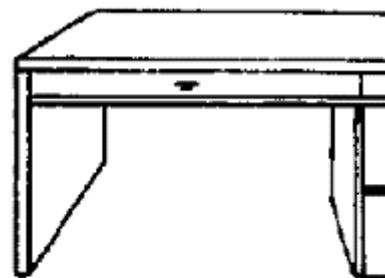
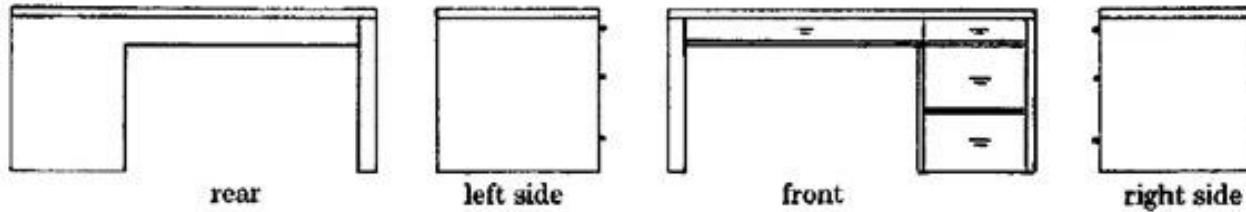
- High Renaissance, 15th century: perspective formalized



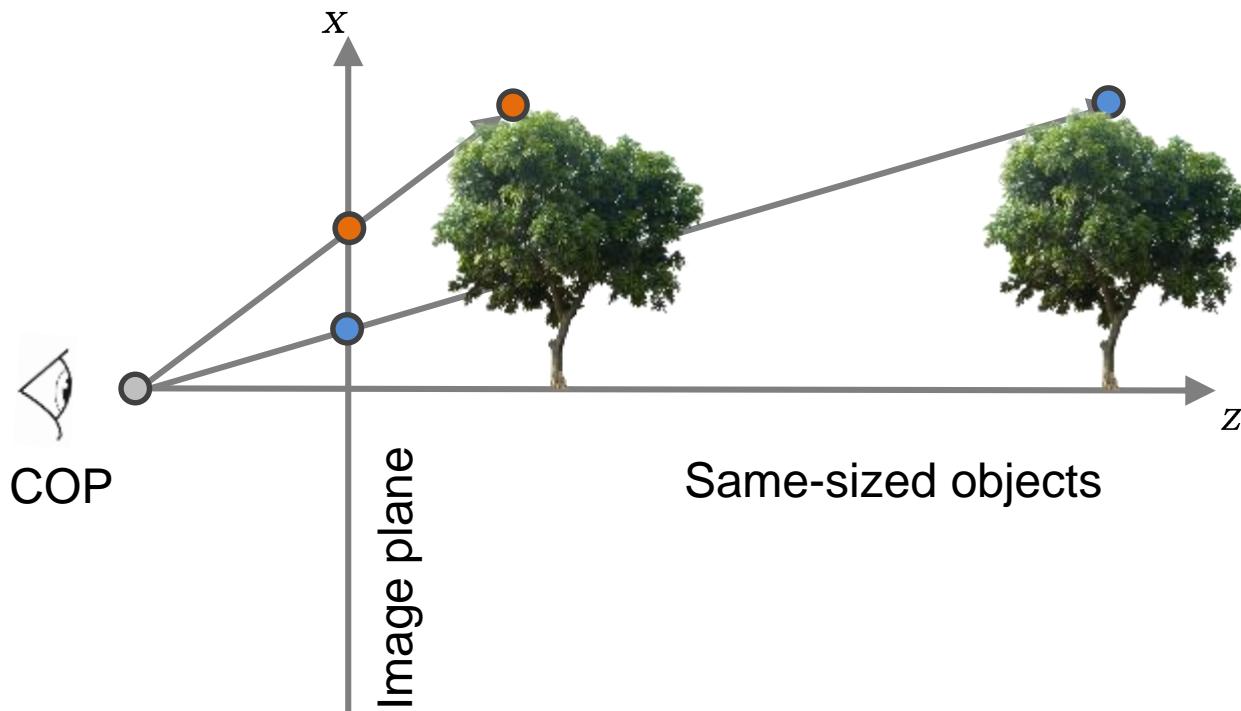
正
交

透視

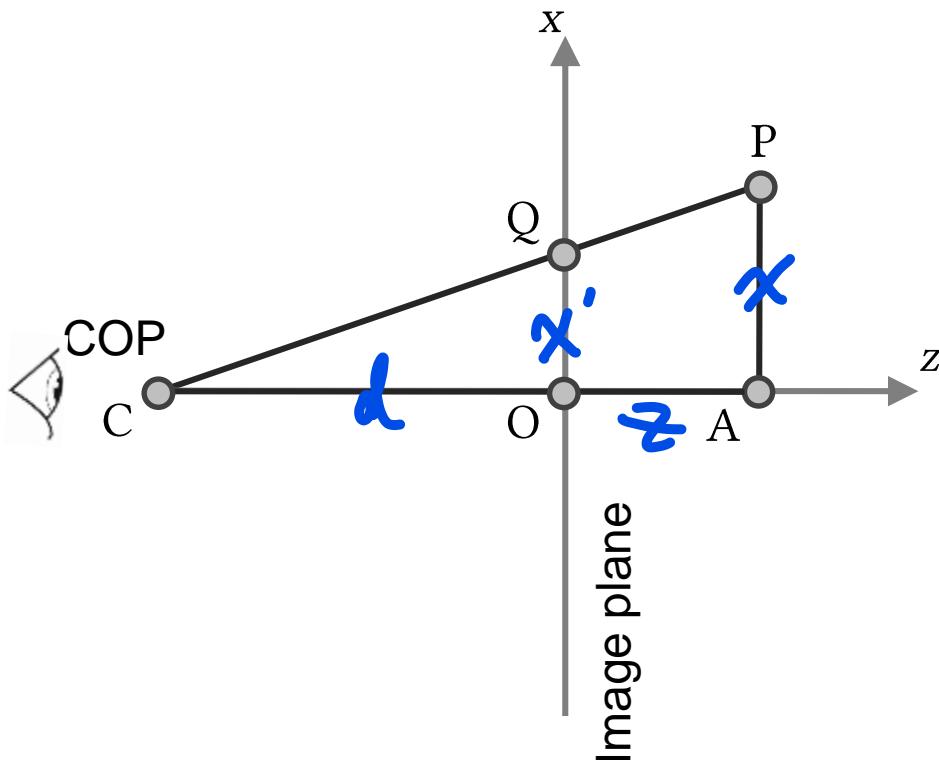
Orthographic vs. Perspective



Perspective projection



Perspective projection



We know:
 $QO/CO = PA/CA$

For x , we define
 $x' = QO$
 $x = PA$
 $d = CO$
 $z = AO$

$$x'/d = x / (d + z)$$

so

$$x' = d x / (d + z)$$

E.g., for $d = 1$:
 $x' = x / (z + 1)$

这也太特殊了

Perspective division

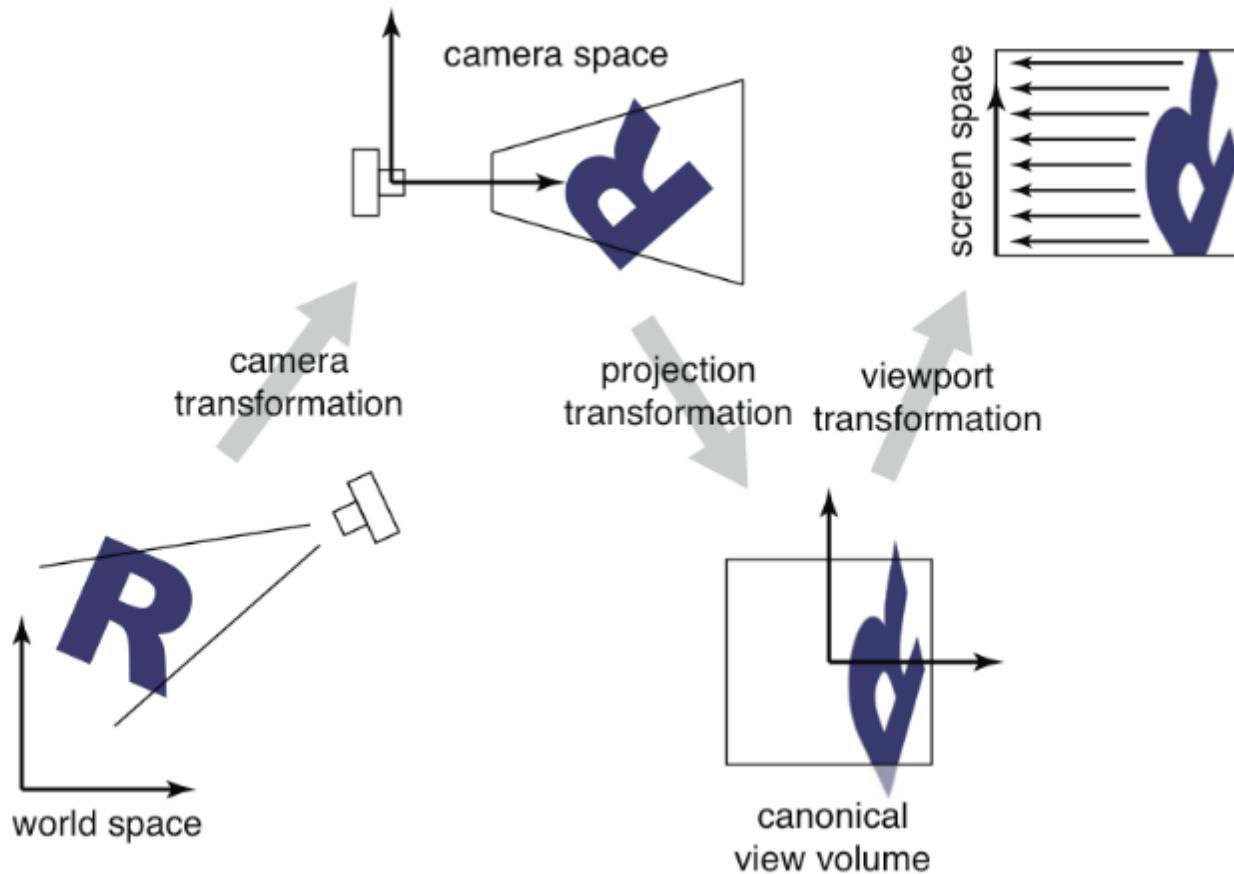
$$P = \begin{pmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 1 & 1 \end{pmatrix}$$

- Given a point $(x, y, z, 1)$
- Another good reason for 4x4 matrices
- Its transform is

$A \cdot P = \begin{pmatrix} x \\ y \\ z \\ z + 1 \end{pmatrix} = \begin{pmatrix} \frac{x}{z+1} \\ \frac{y}{z+1} \\ \frac{z}{z+1} \\ \frac{1}{z+1} \end{pmatrix}$

这样才是
homogeneous
equivalent

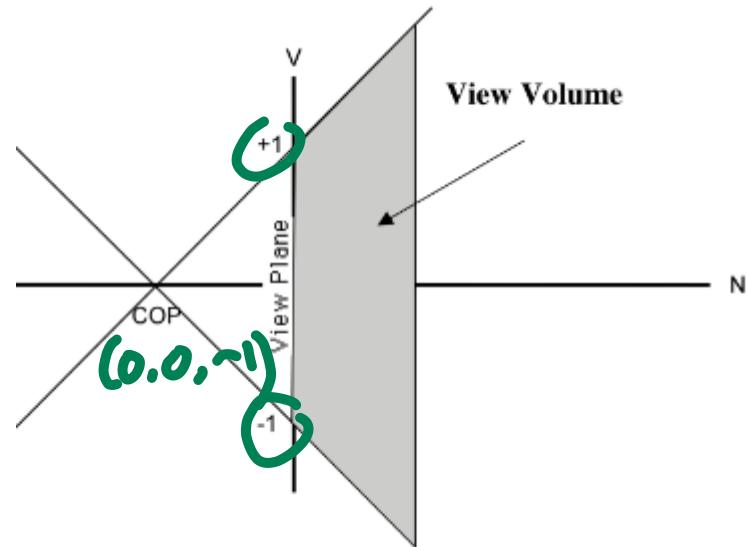
Projection Pipeline



Canonical Frames

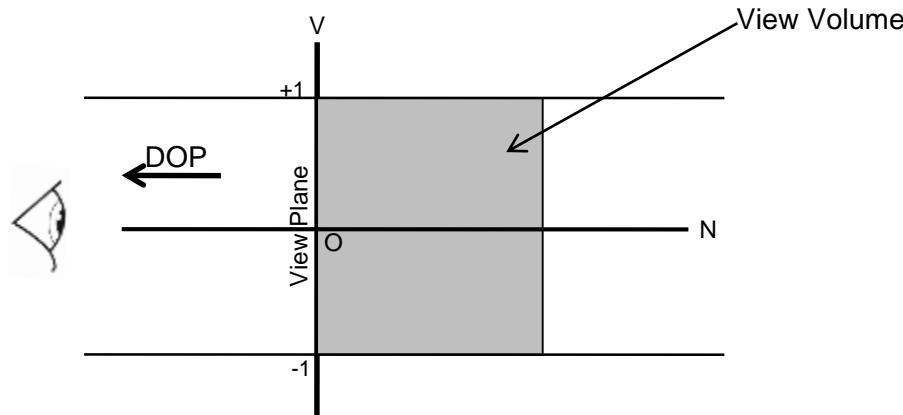
史圭框架

- We use **canonical frames** as intermediate stages from which we know how to proceed
- Canonical Frame for Perspective Projection:
 - COP at $(0, 0, -1)$
 - View plane coincident with UV plane
 - Viewplane window bounded by –
1 to +1



Canonical Frame for Parallel Projection

- Orthographic parallel projection
- Direction of projection (DOP) is $(0,0,-1)$
- View volume bounded by -1 and $+1$ on U and V
- And by 0 and 1 on the N axis
- $p' = (x, y, 0)$

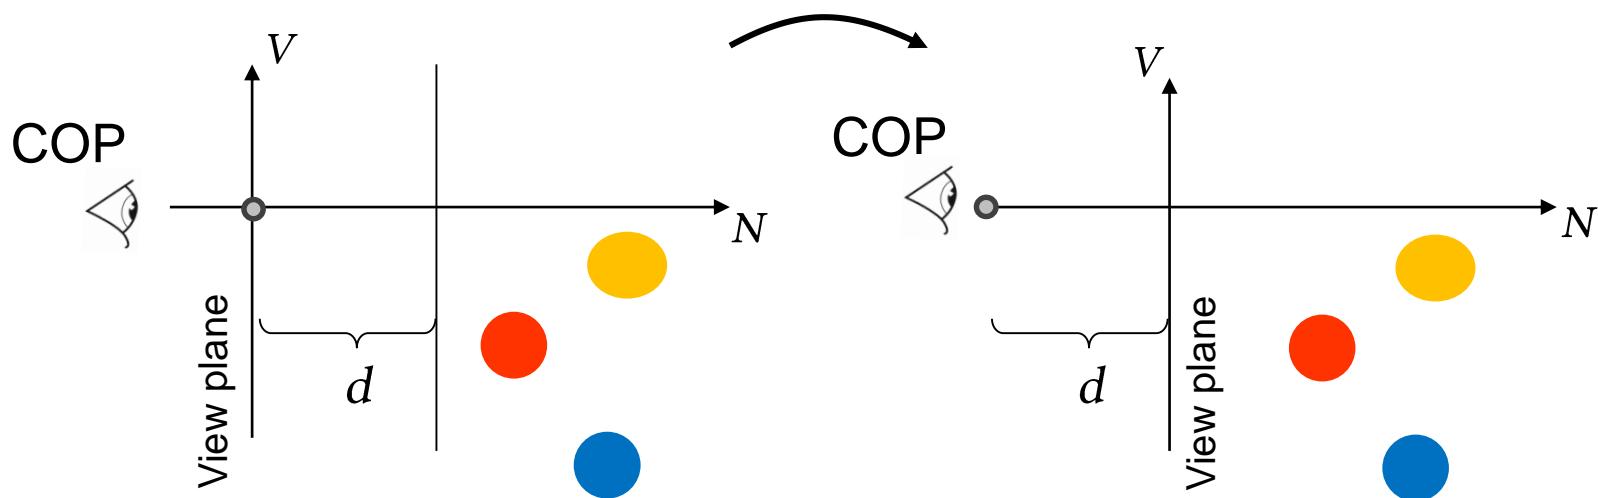


General Persp. to Canonical Persp.

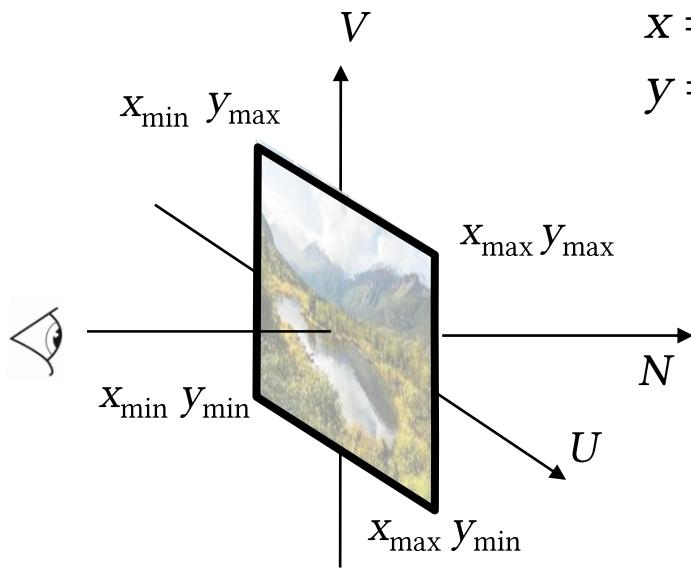
- We will apply **four** transformation matrices
- Each “corrects” one aspect of the projection
- Finally, we multiply them all together

Step 1: Move to UV plane

$$\begin{pmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & -d \\ 0 & 0 & 0 & 1 \end{pmatrix}$$



Recall the window



$$W = X_{\max} - X_{\min}$$

$$h = Y_{\max} - Y_{\min}$$

$$x = X_{\max} + X_{\min}$$

$$y = Y_{\max} + Y_{\min}$$

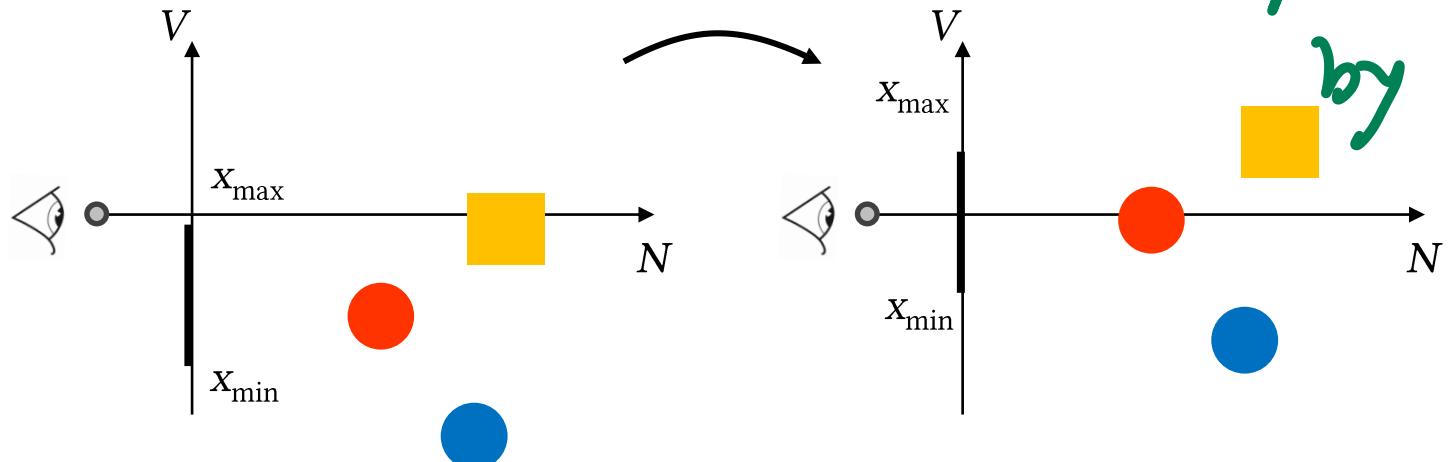
Step 2: Regular pyramid

$$\begin{pmatrix} 1 & 0 & 0 & -x \\ 0 & 1 & 0 & -y \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix}$$

← world move up by x

$$\begin{aligned} W &= X_{\max} - X_{\min} \\ h &= Y_{\max} - Y_{\min} \\ x &= X_{\max} + X_{\min} \\ y &= Y_{\max} + Y_{\min} \end{aligned}$$

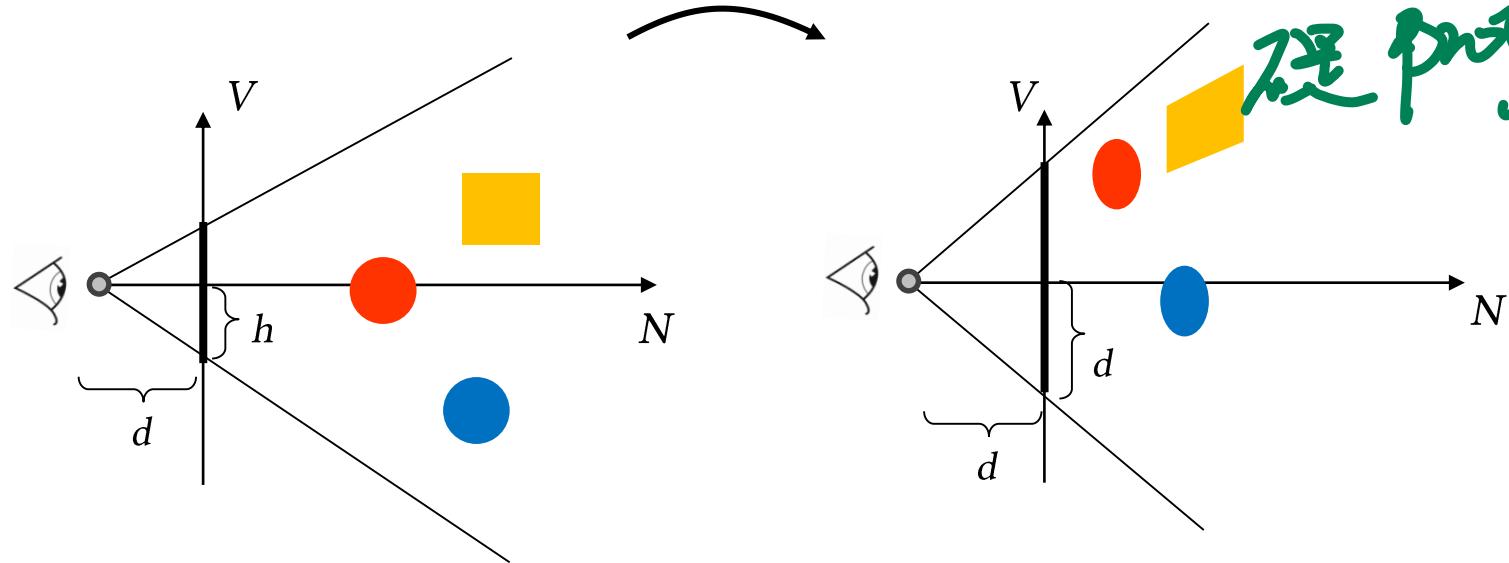
loop move down by x



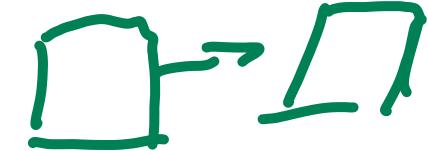
Step 3: Regular pyramid

This is not uniform scale: z did not change!

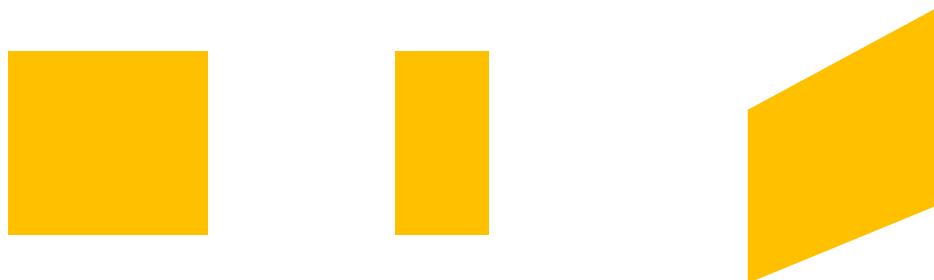
透視射影
非直角形
基礎



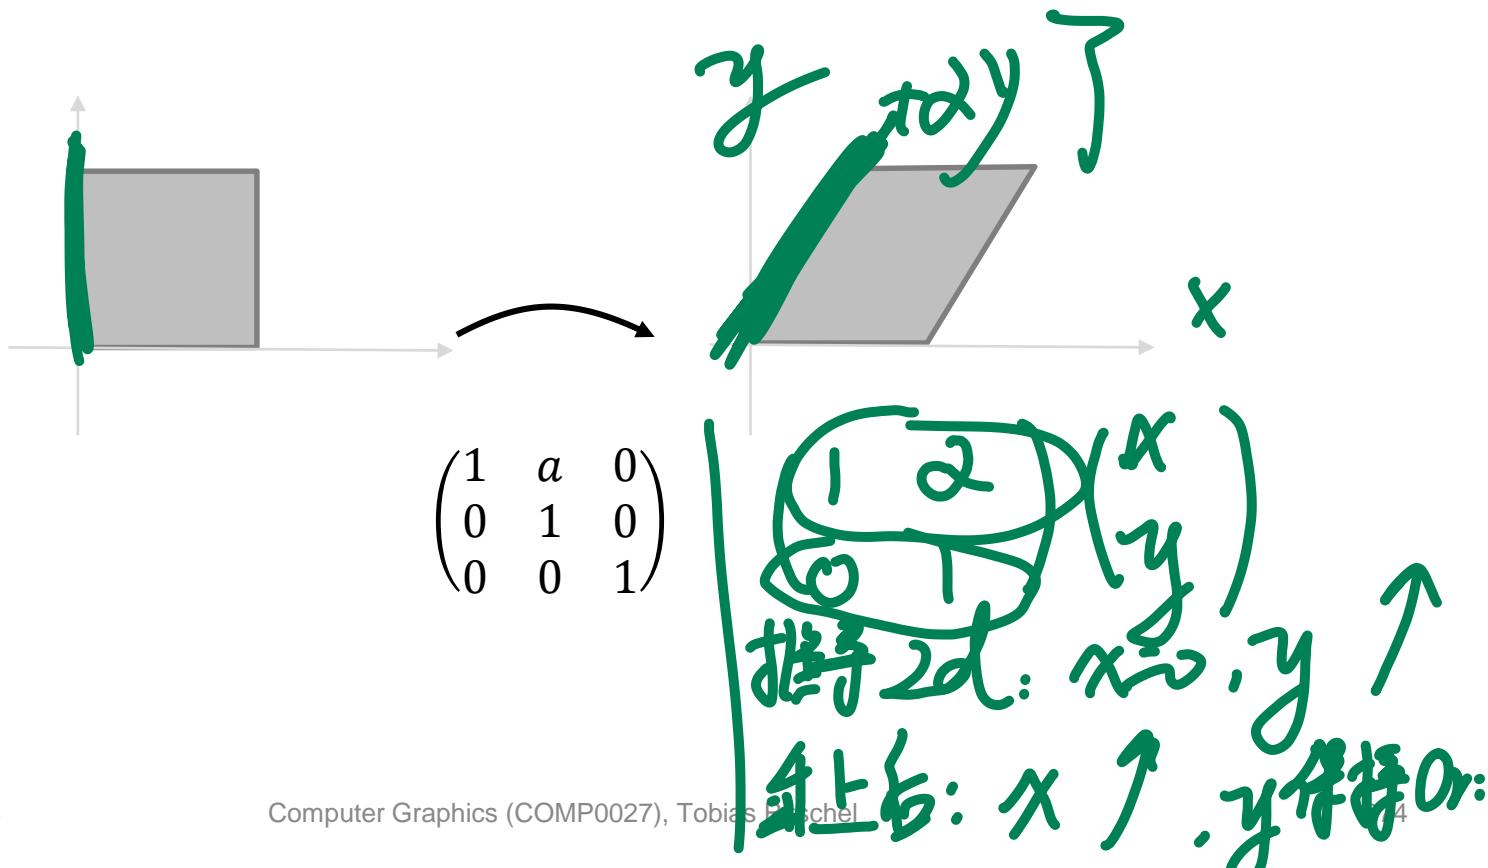
Shear



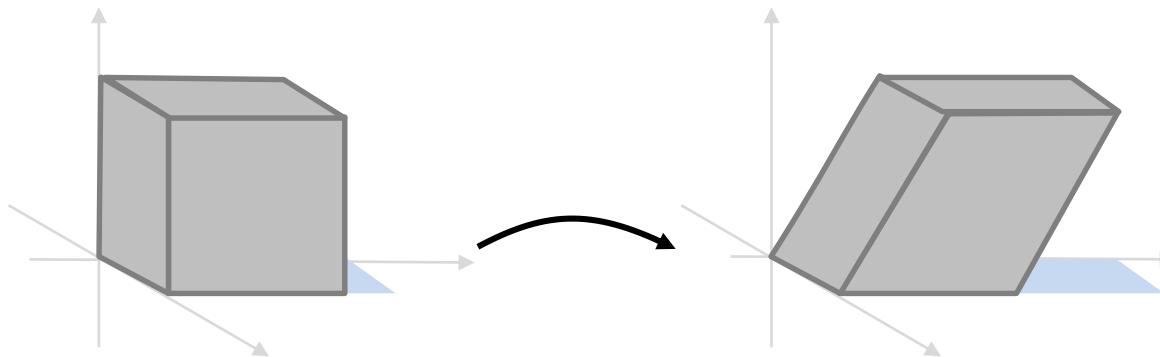
- This needs to be a **shear**
- If you want a regular pyramid to become a box ..
- A box has to become some other regular pyramid
- What is shear?



Interlude: Shear 2D

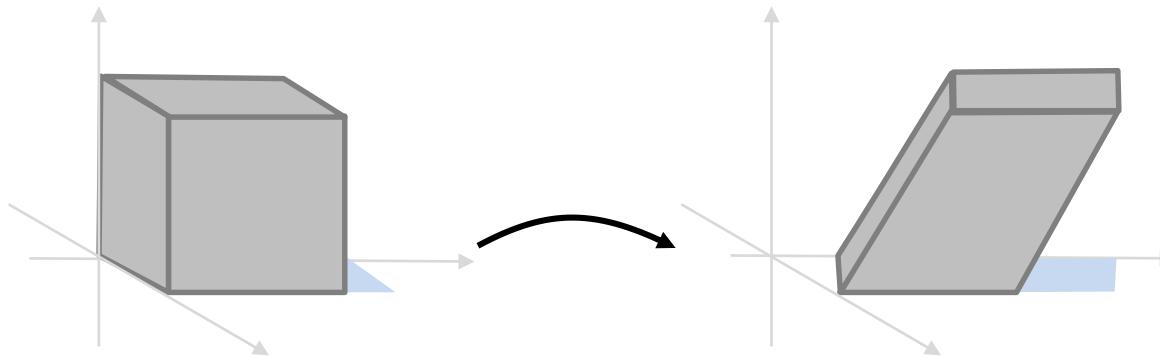


Interlude: Shear 3D



$$\begin{pmatrix} 1 & a & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix}$$

Interlude: Shear 3D

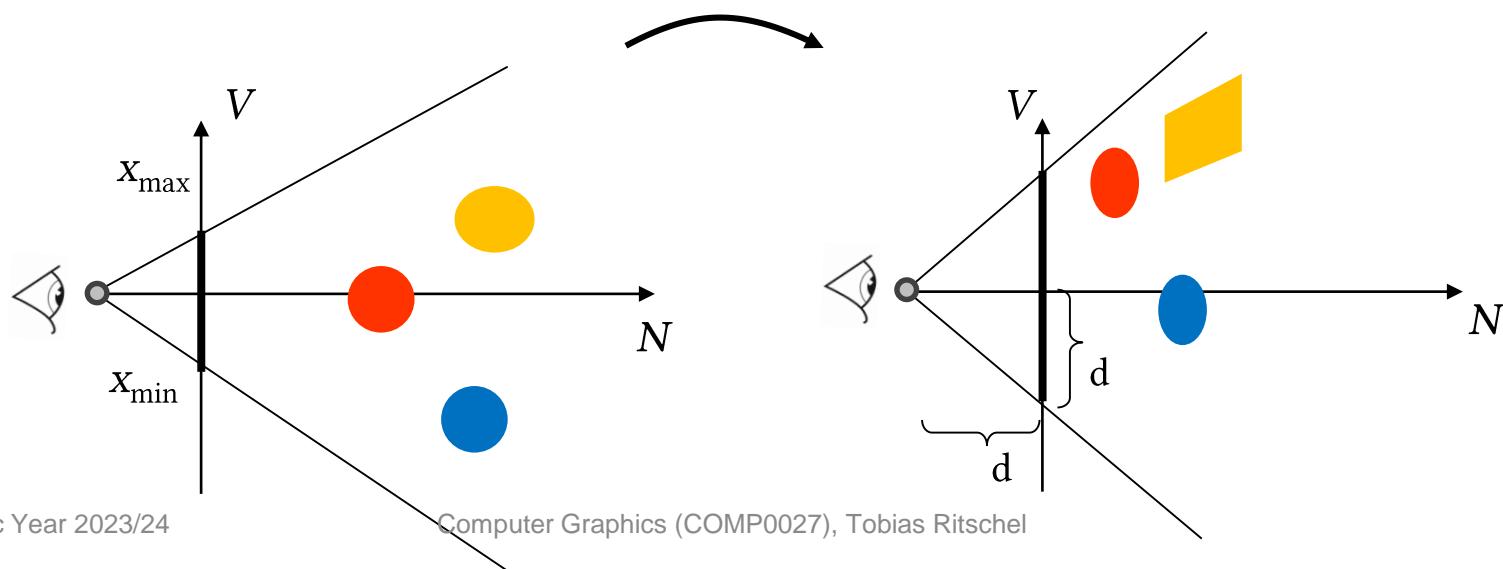


$$\begin{pmatrix} 1 & a & b & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix}$$

Step 3: Regular pyramid

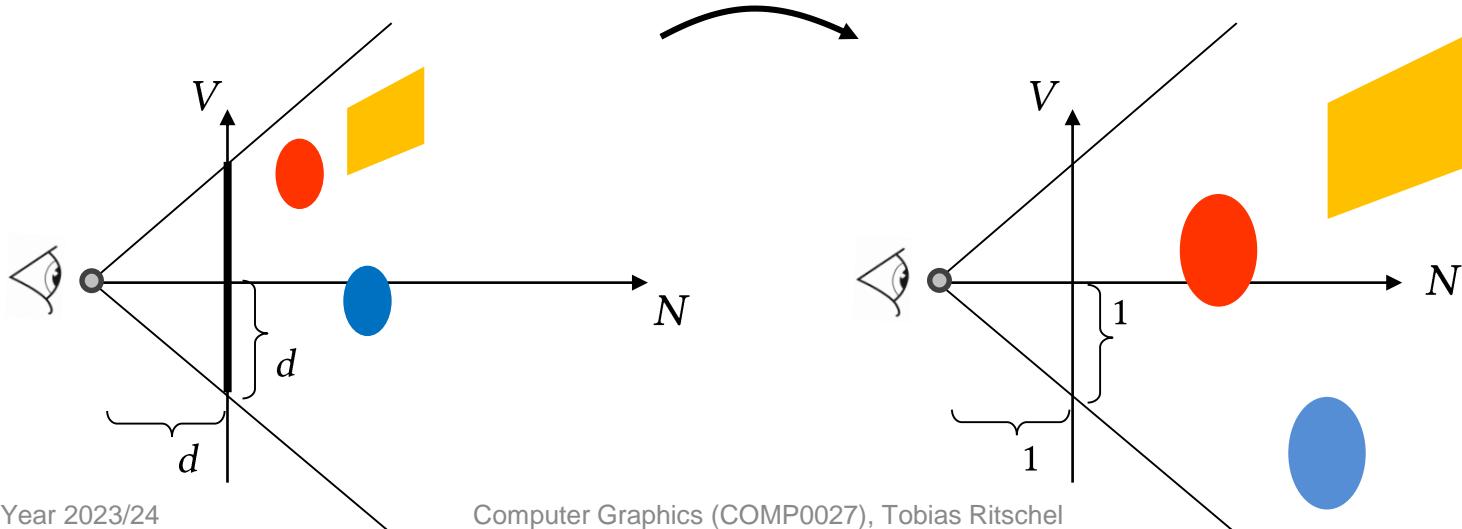
$$\begin{pmatrix} d/w & 0/w & -x/w & 0/w \\ 0/h & d/h & -y/h & 0/h \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix}$$

$$\begin{aligned} w &= X_{\max} - X_{\min} \\ h &= Y_{\max} - Y_{\min} \\ x &= X_{\max} + X_{\min} \\ y &= Y_{\max} + Y_{\min} \end{aligned}$$



Step 4: Scale by $1/d$

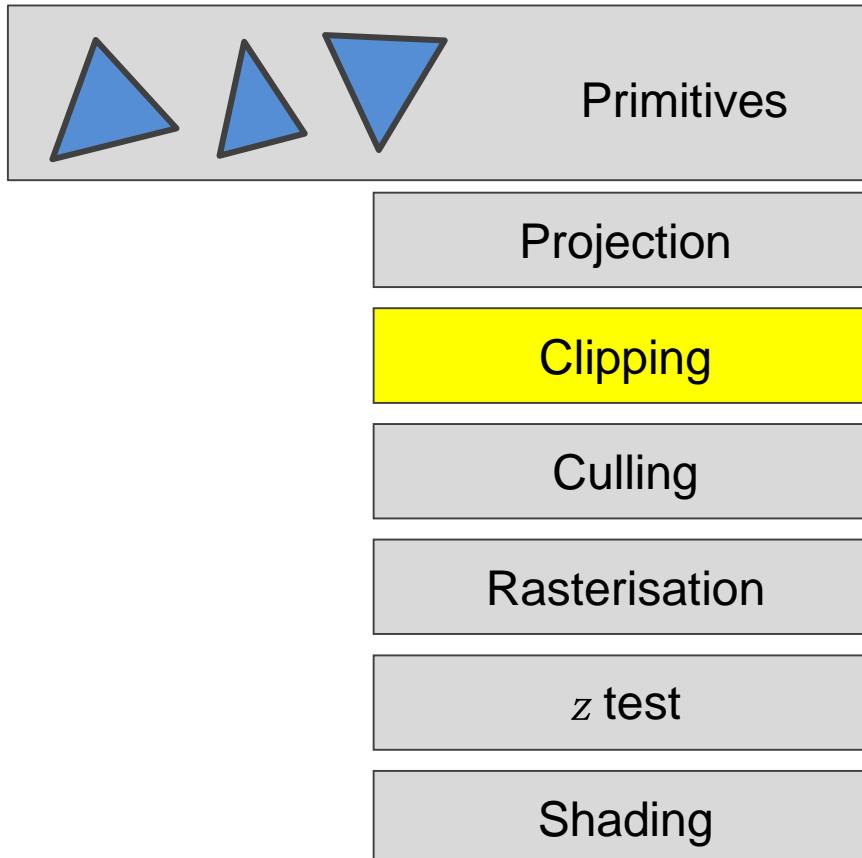
$$\begin{pmatrix} 1/d & 0 & 0 & 0 \\ 0 & 1/d & 0 & 0 \\ 0 & 0 & 1/d & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix}$$



Two simplifications

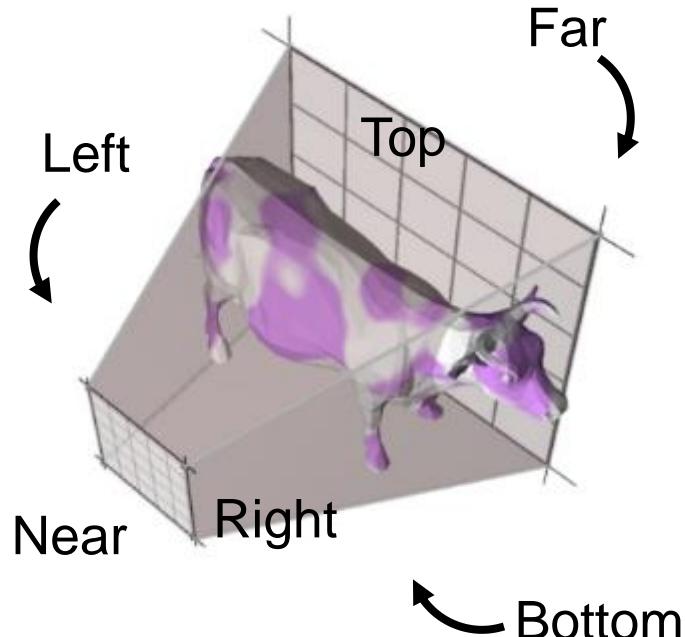
1. The COP might not fall on the z axis
 - In our model it does
 - Required in professional photography
 - Required in virtual reality (Stereo, HMDs)
2. Might want to re-scale z to fit a range
 - Will get back to this adjustment later

Pipeline



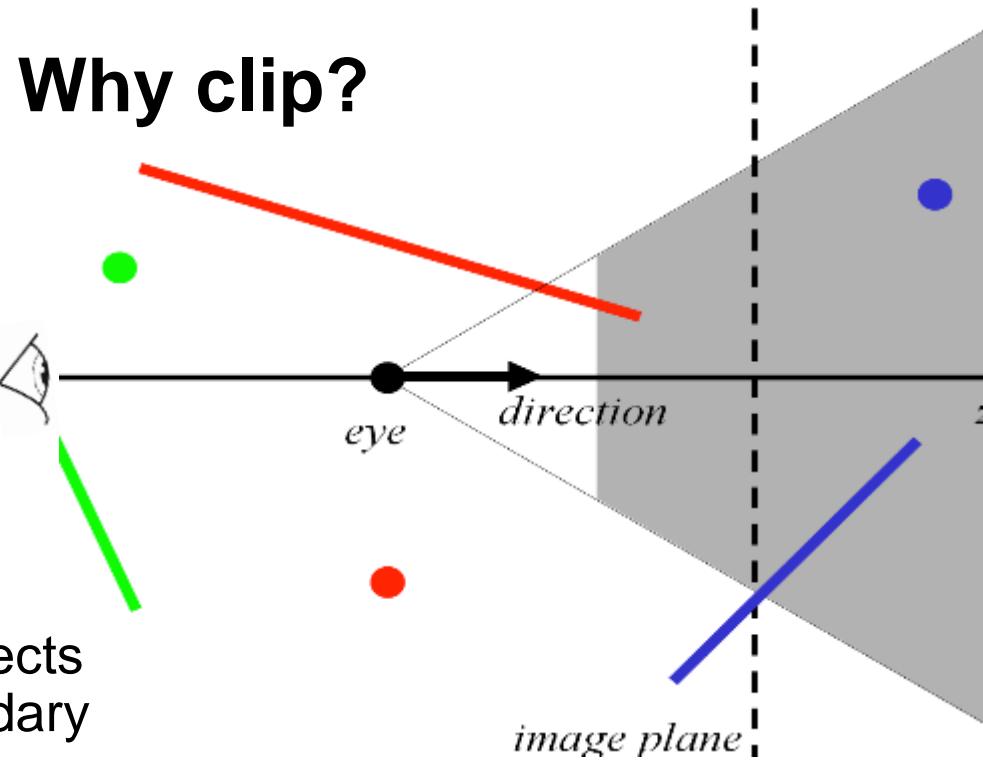
Clipping 截剪 剔除多(Cull) off

- Eliminate portions of objects outside the viewing frustum
- View frustum
 - Boundaries of the image plane projected in 3D
 - A near & far clipping plane
- User may define additional clipping planes



退化

- Avoid degeneracies
 - Don't draw stuff behind the eye
 - Avoid division by 0 and overflow
- Efficiency
 - Don't waste time on objects outside the image boundary
- Other graphics applications (often non-convex)
 - Hidden-surface removal, Shadows, Picking, Binning, CSG (Boolean) operations (2D & 3D)

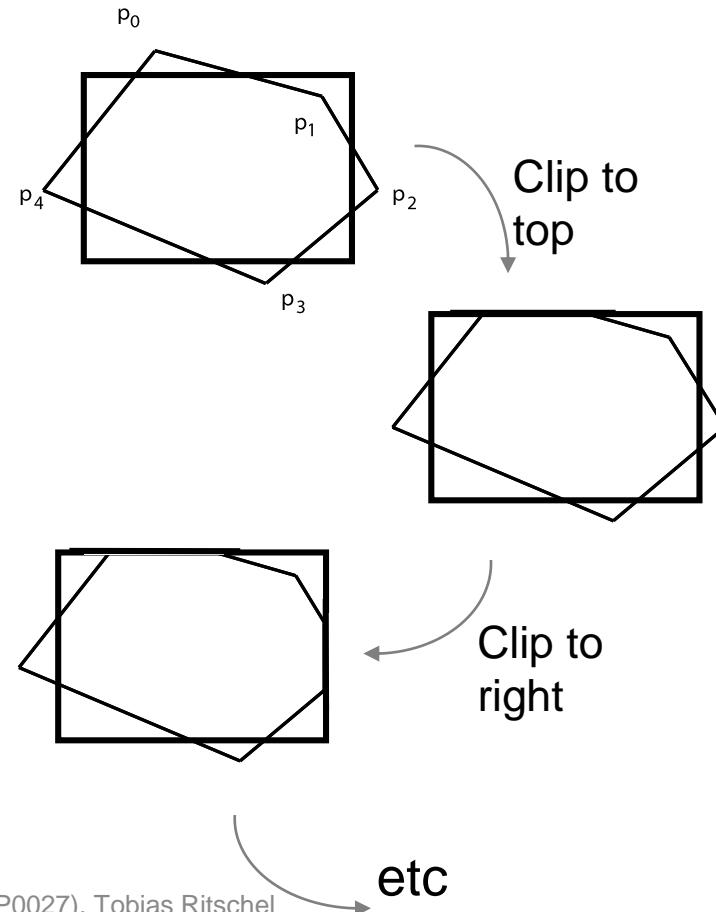


Clipping Summary

- It's the process of finding the exact part of a polygon lying inside the view volume
- To maintain consistency, clipping of a polygon should result in a polygon, not a sequence of partially unconnected lines
- We will first look at two different 2D solutions and then extend one to 3D

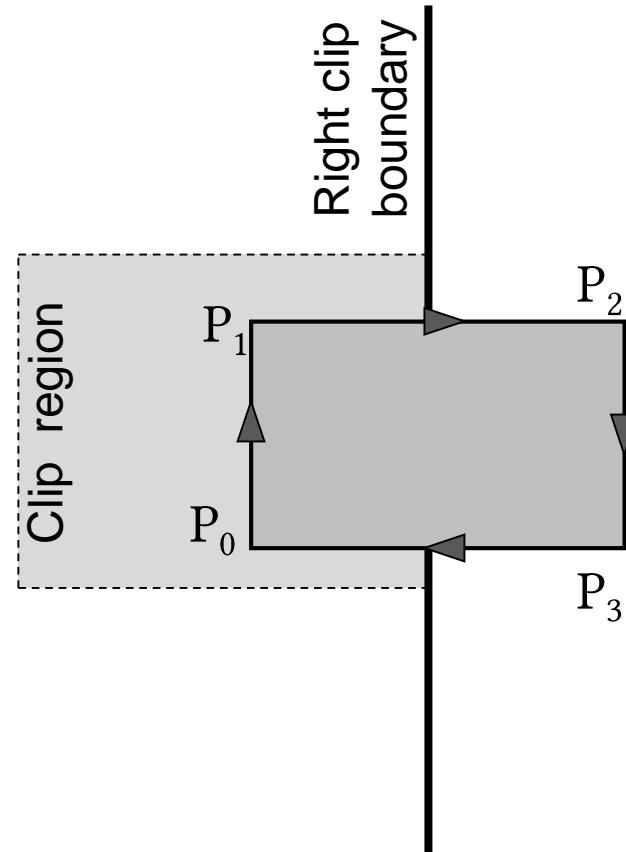
Sutherland-Hodgman Algorithm

- Clip the polygon against each boundary of the clip region successively
- Result is possibly NULL if polygon is outside
- Can be generalised to work for any polygonal clip region, not just rectangular



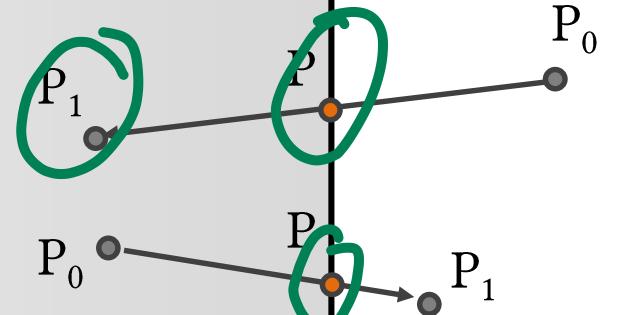
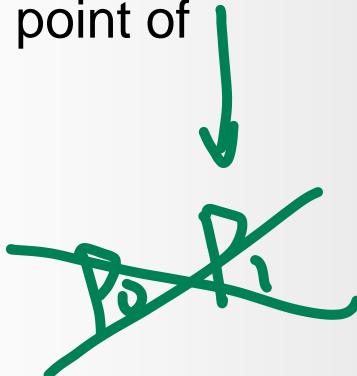
Clipping To A Region

- To find the new polygon
 - iterate through each of the polygon edges and construct a new sequence of points
 - starting with an empty sequence
 - for each edge there are 4 possible cases to consider

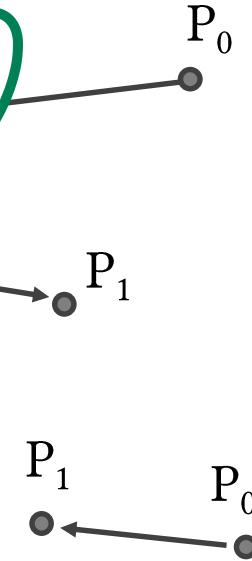


Clipping polygon edge against boundary

- Given an edge P_0, P_1 we have 4 case. Can ...
 - enter the clip region, add P and P_1
 - leave the region, **add only P**
 - be entirely outside, do nothing
 - be entirely inside, **add only P_1**
- Where P is the point of intersection



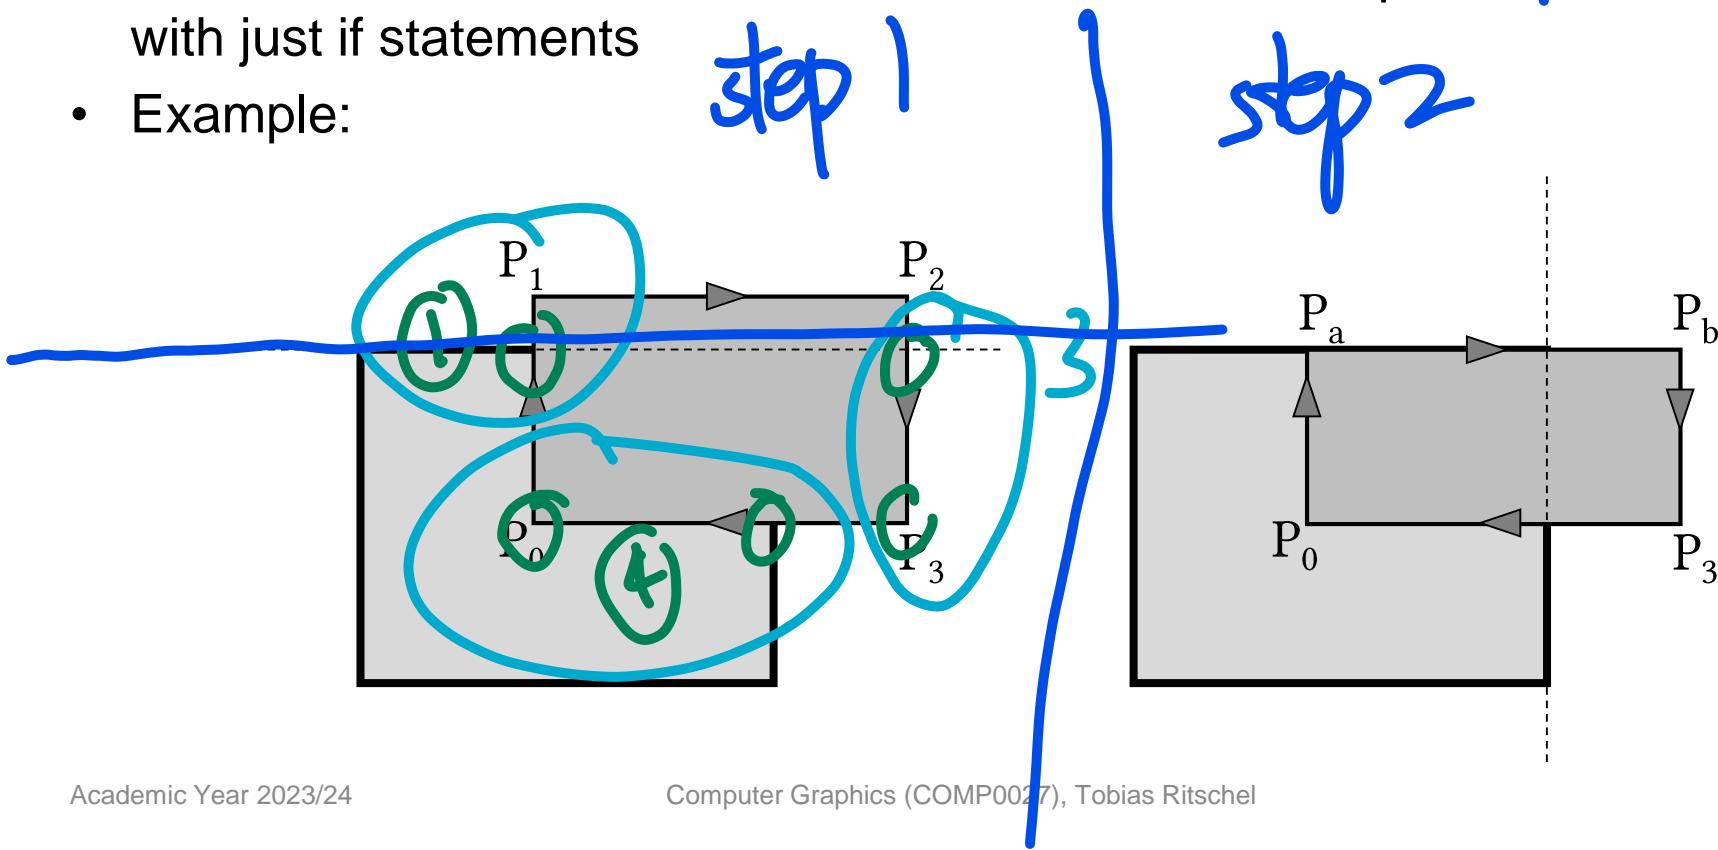
In



Out

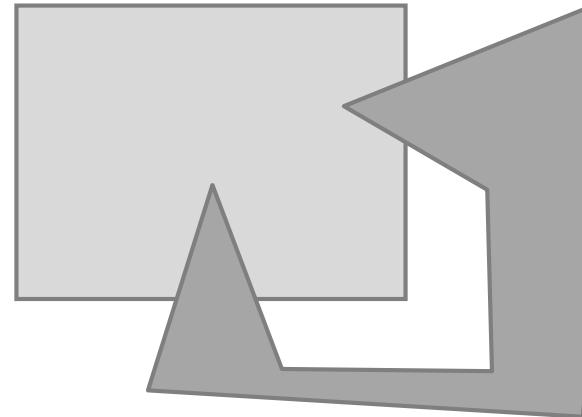
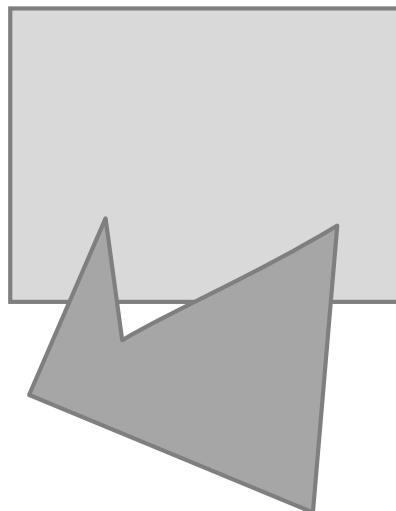
Still the Sutherland-Hodgman

- We can determine which of the 4 cases and also the point of intersection with just if statements
- Example:



Sutherland-Hodgman Problem

- Clipping is an example where convex polys make our lives easier
- Concave: Weiler-Atherton algorithm

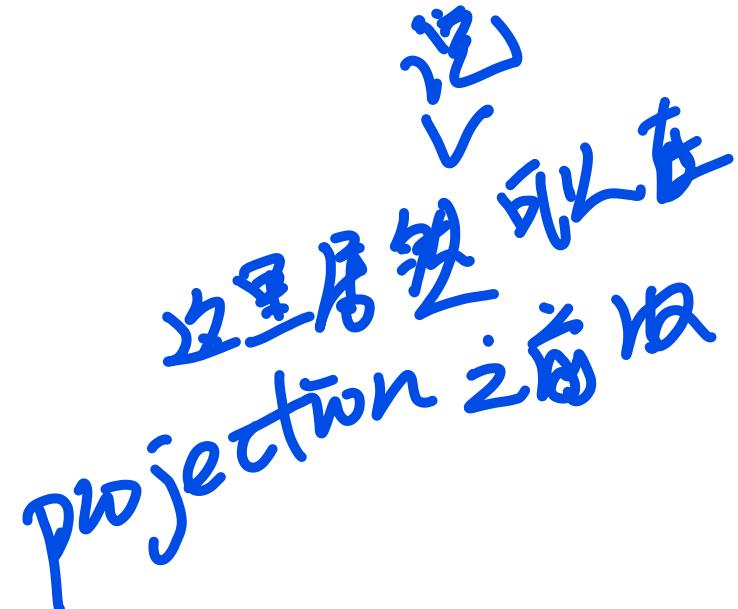


Clipping Polygons in 3D

- The Sutherland-Hodgman can easily be extended to 3D
 - Clipping boundaries are 6 planes instead of 4 lines
 - Intersection calculation is done by comparing an edge to a plane instead of edge to edge

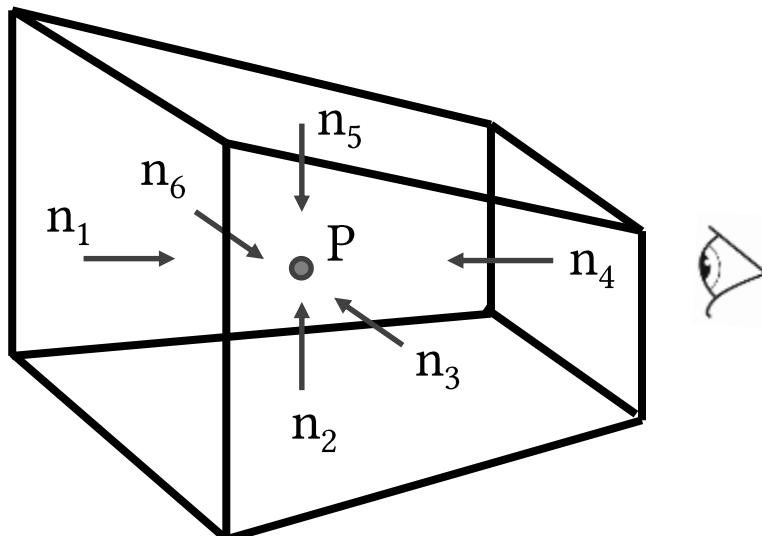
When to clip?

- Before perspective transform in 3D space
 - Use the equation of 6 planes
 - Natural, not too degenerate
- In homogeneous coordinates after perspective transform (Clip space)
 - **Before** perspective divide (4D space, weird w values)
 - Canonical, independent of camera
 - The simplest to implement in fact
- In the transformed 3D screen space after perspective division (canonical par.)
 - Problem: objects in the plane of the camera



Clipping with respect to View Frustum

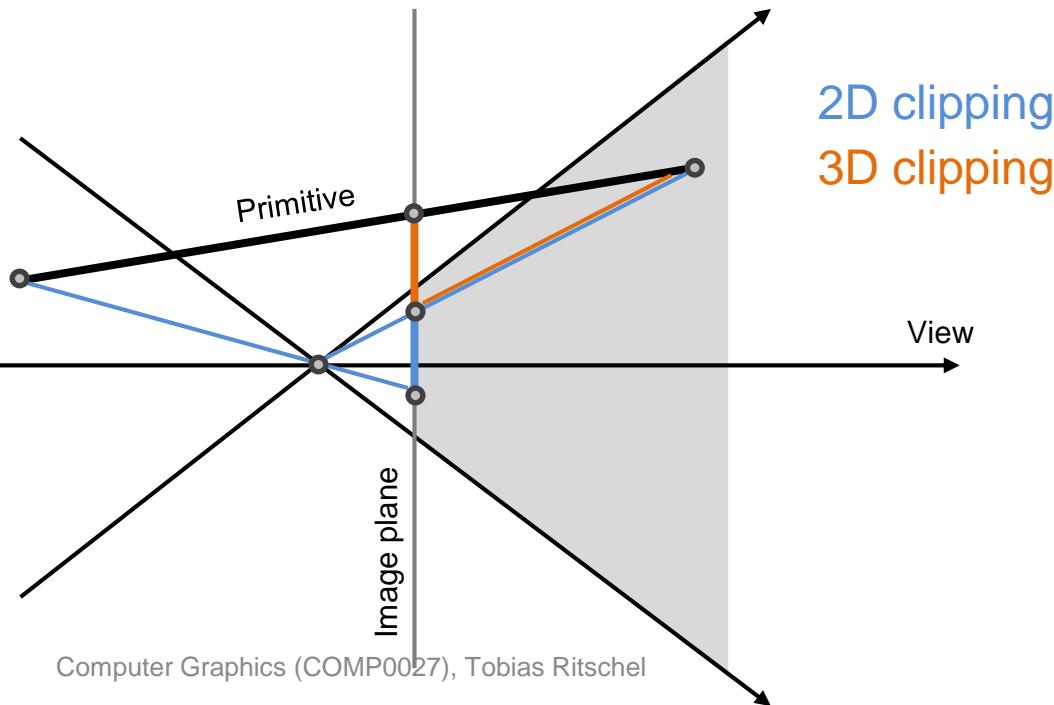
- Test against each of the 6 planes
 - Normals oriented towards the interior
- Clip / cull / reject point P if any $\langle \mathbf{n}_i, \mathbf{p} \rangle < d_i$



Clipping After Projecting

- When we have an edge that extends from the front to behind the COP, then if we clip after projection (which in effect is what e.g. the PS 1 did) we might get wrong results

投影
Projection
2D clipping
3D clipping



A note on implementation

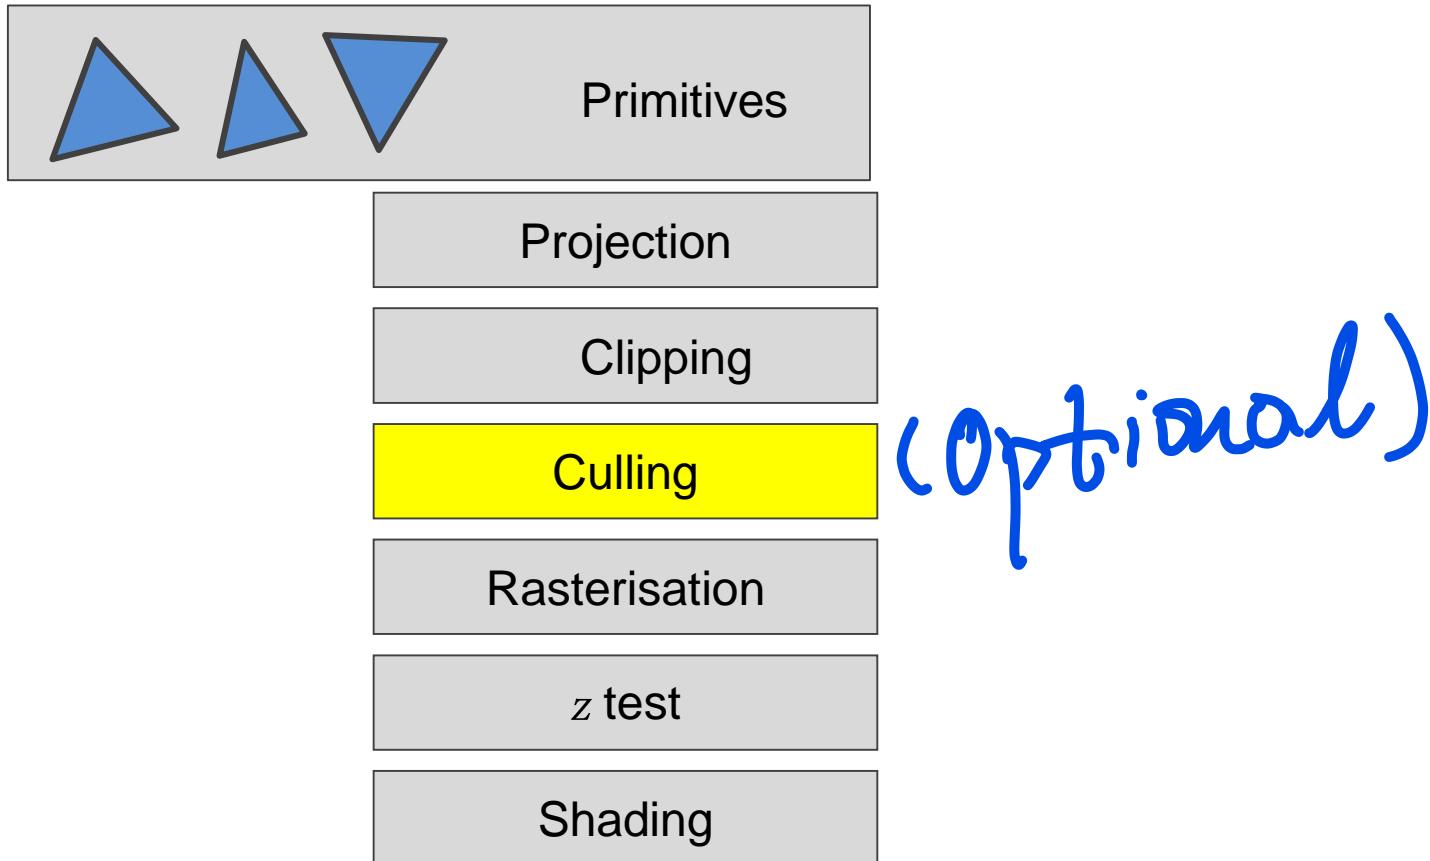
- Vanilla implementation uses dynamic data structure

```
iterate windowVertices edges {
    iterate primitiveVertices
        if(...)
            primitiveVertcies.add(new vertex)
        if(...)
            primitiveVertcies.remove (new vertex)
    } }
```

- Better, without

```
for MAX_EDGES edges {
    for MAX_CLIPPED_VERTICES {
        if(...)
            primitiveVertcies[vertexCount++] (new vertex)
        if(...)
            vertexCount++
    } }
```

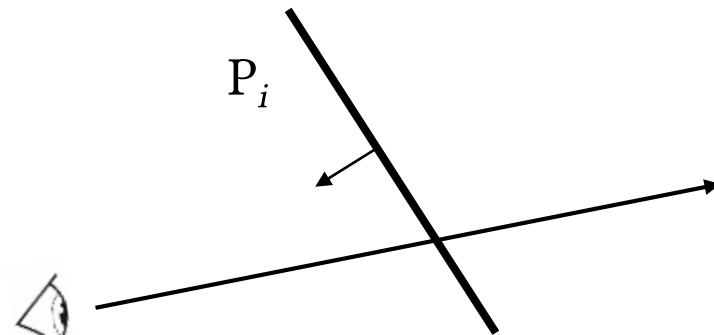
Pipeline



Idea

- Optional speed-up
- Remove primitives P_i not facing the camera
- Polygons whose normal does not face the viewpoint, are not rendered

这里也不适合
culling



cull to back face

这里也不适合
culling

Back Face Culling

- Thus if we have the plane equation

$$l(x, y, z) = ax + by + cz - d = 0$$

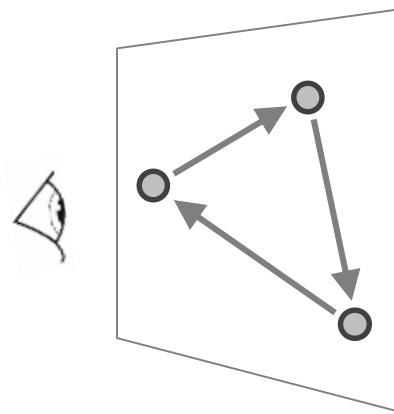
- and the COP (c_x, c_y, c_z), then

$$l(c_x, c_y, c_z) > 0$$

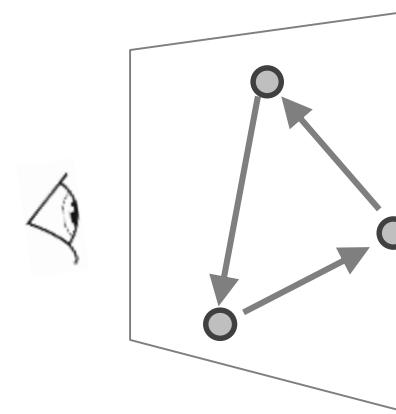
- if the COP is in front of the polygon
- Otherwise: cull!

Clockwise/counter-clockwise

Typically done without normals: All polygons that are not counter-clockwise after projection are culled



Clockwise: Culled



Counter clockwise: Not culled

Limitation

Does not ensure correct visible surfaces determination,
unless there is only a single convex object



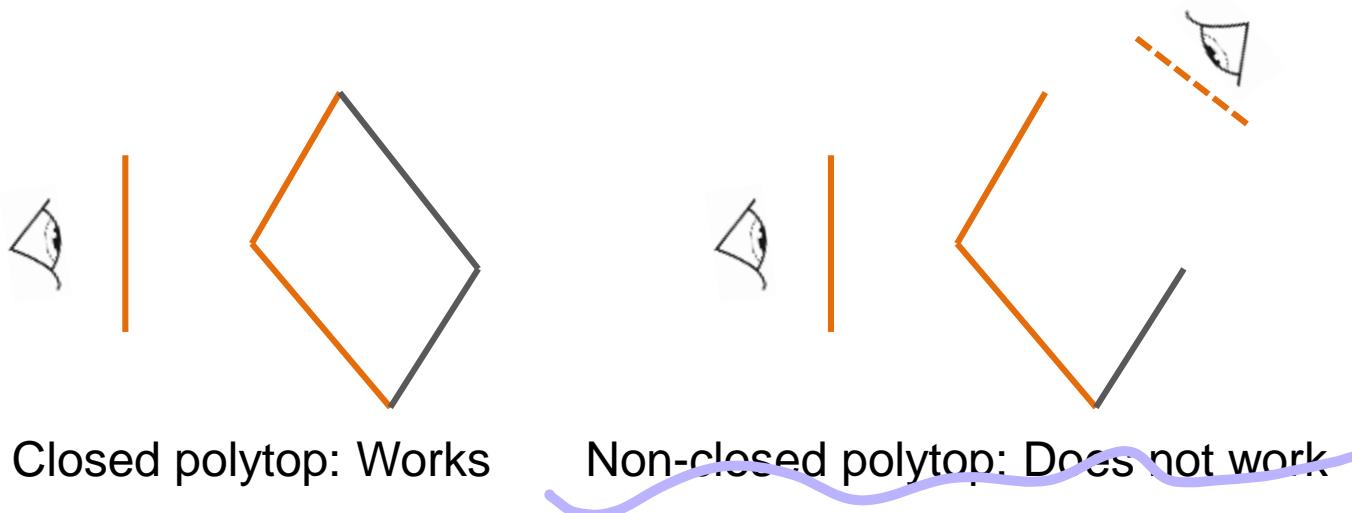
Convex: Works



Concave: Does not work

Limitation

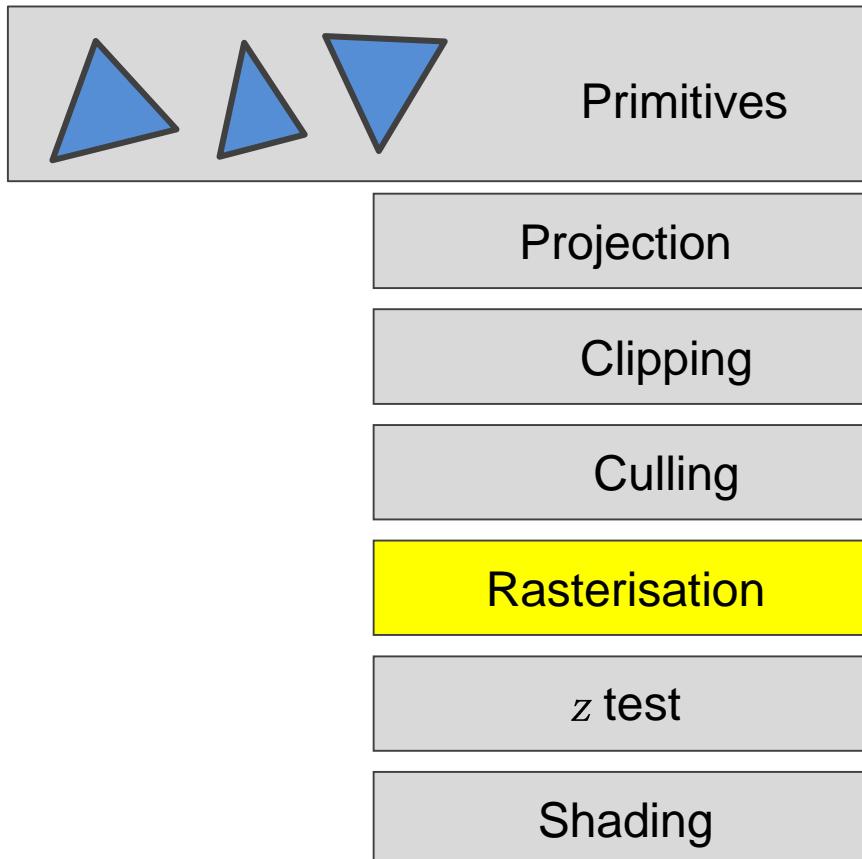
Only reliable if backsides are never required. Therefore, does not work for non-closed polytopes.



Culling recap

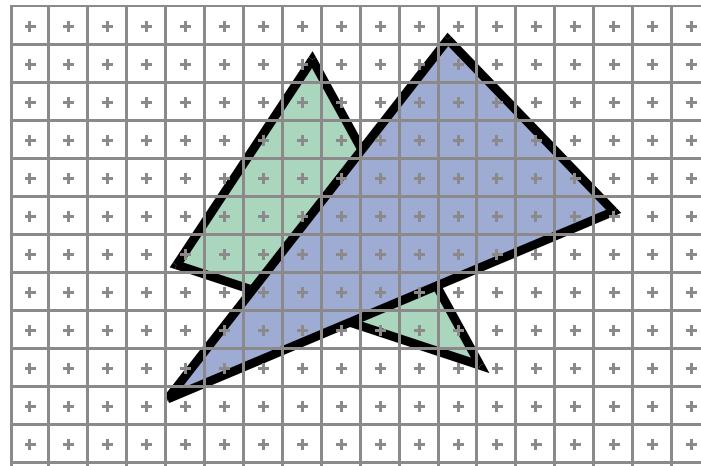
- Culling can easily & quickly remove (some!) invisible polygons from the pipeline
- Still some (partially) invisible ones remain

Pipeline



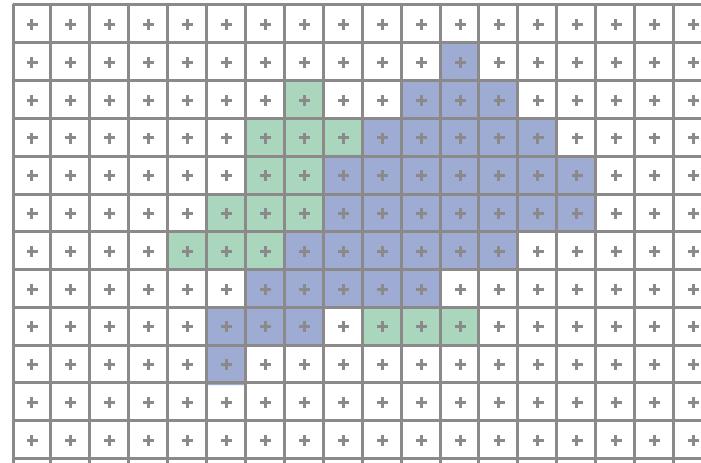
2D Scan Conversion

- Primitives are continuous; screen is discrete
 - Triangles defined by discrete set of vertices
 - But they define a continuous area on screen



2D Scan Conversion

- Solution: compute discrete approximation
- Scan Conversion (Rasterization):
algorithms for efficient generation of the samples comprising this approximation



Naïve Filling Algorithm

- Find a point inside the polygon
- Do a **flood fill**:
 - Keep a stack of points to be tested
 - When the stack is not empty
 - Pop the top point (Q)
 - Test if Q is inside or outside
 - If Inside, colour Q , push neighbours of Q if not already tested
 - If outside discard
 - Mark Q as tested

2 BFS

Critique 行評

- Horribly slow
 - Explicit in/out test at every point
 - But still very common in paint packages!
 - Only solid color
- Stack might be very deep
- Want to do this
 - Without much memory
 - In parallel

Brute Force Solution for Triangles

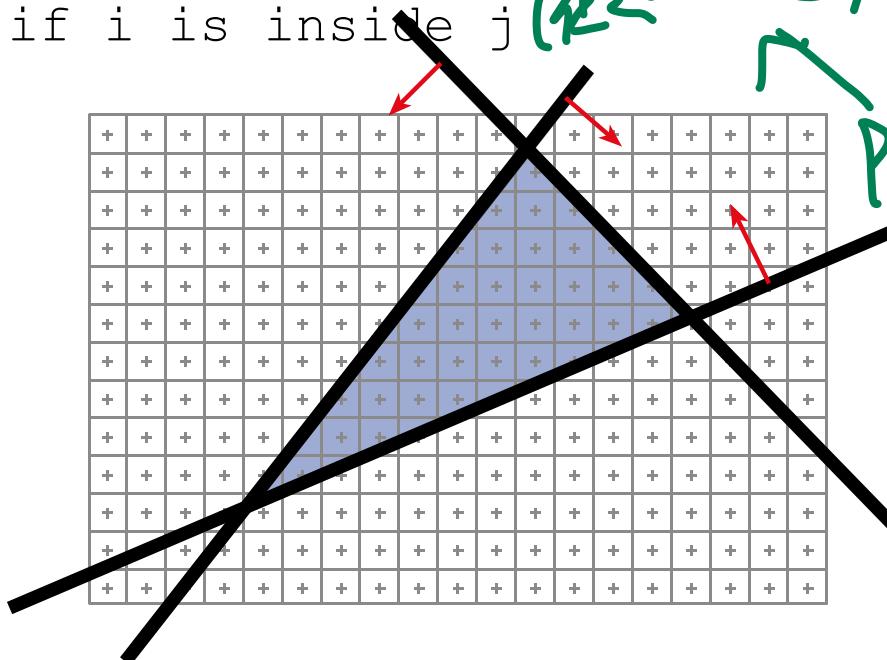
For all pixels i

For all edges j

Test if i is inside j

($E \cap \sim \infty$)

positive
half-space



Half-Space Test Reminder

- For each edge compute line equation (analogue to plane equation):

$$L_i(x, y) = a_i x + b_i y + c_i$$

- If $L_i(x, y) > 0$ point in **positive** half-space
- If $L_i(x, y) < 0$ point in **negative** half-space
- If all $L_{1,2,3}(x, y) \geq 0$ point inside triangle



Half-Space Test Reminder

- For each edge compute line equation (analogue to plane equation):

$$L_i(x, y) = a_i x + b_i y + c_i$$

- Example:

- Assuming $(2, 2)$ to $(4, 7)$,
- The slope is $dy/dx = 5/2 = 2.5$
- Hence $y - 7 = 2.5(x - 4)$,
- So $L_1(x, y) = 2.5x - y - 3.$

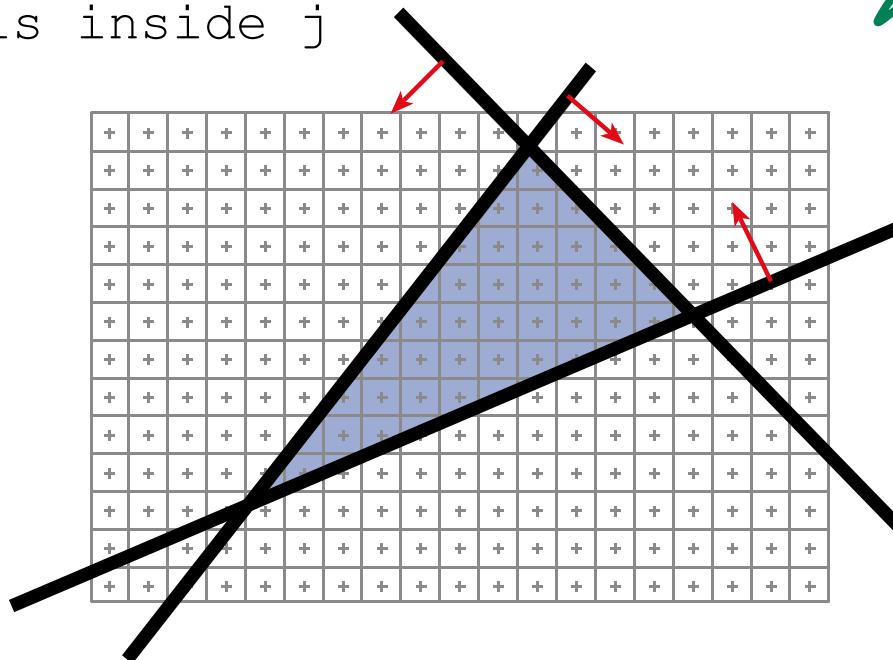
Brute Force Solution for Triangles

For all pixels i

For all edges j

Test if i is inside j

冗長
暴力

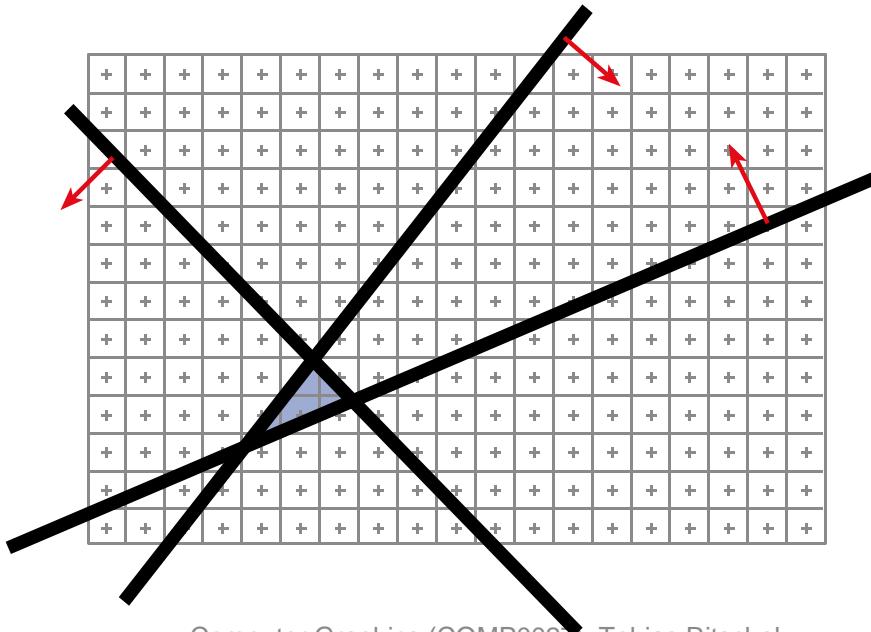


Brute Force Solution for Triangles

For all pixels i

For all edges j

Test if i is inside j

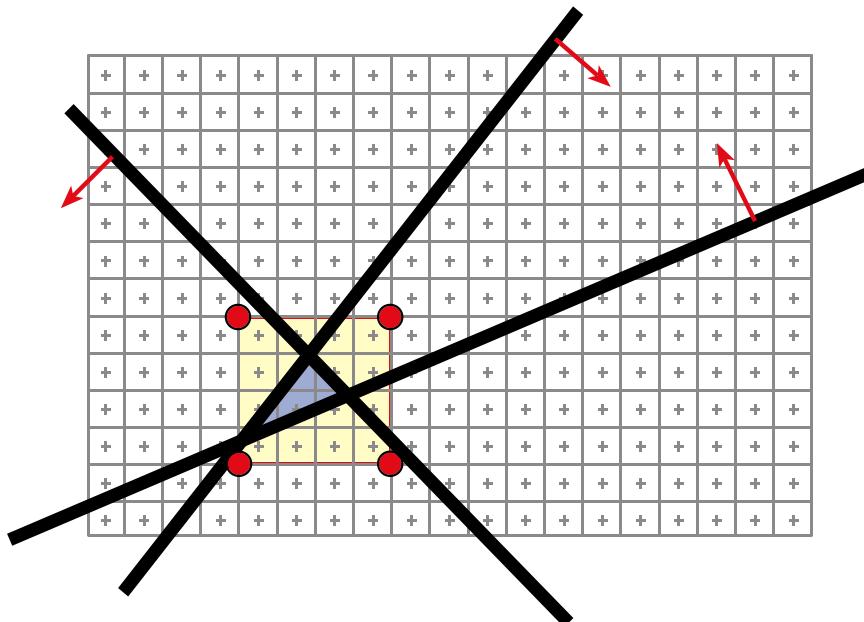


Problem?
If the triangle is small,
a lot of useless
computation!



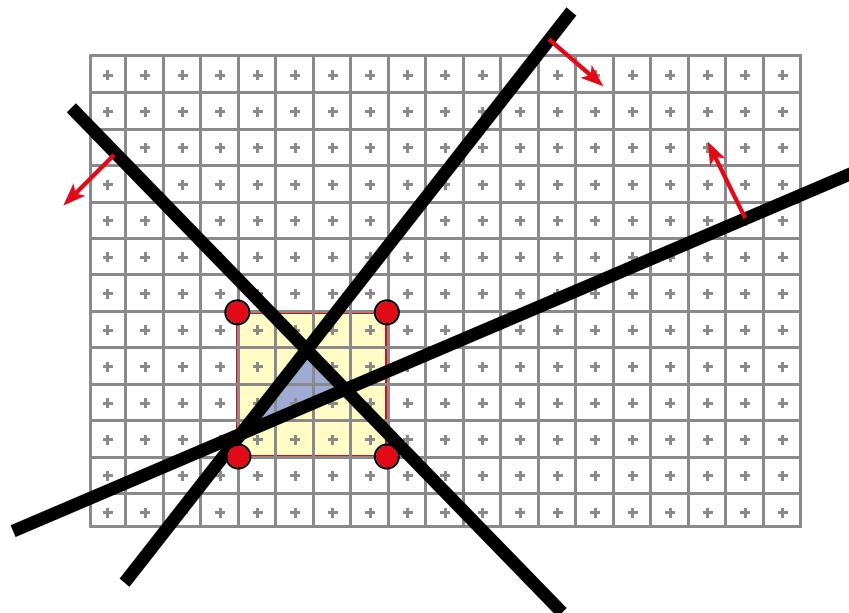
Brute Force Solution for Triangles

- Improvement: Compute only for the *screen bounding box* of the triangle
- How do we get such a bounding box?
 - $x_{\min}, x_{\max}, y_{\min}, y_{\max}$ of the triangle vertices



Rasterisation on Graphics Cards

- Triangles are usually very small
 - Setup cost are becoming more troublesome
- Brute force is tractable



Rasterisation on Graphics Cards

For every triangle

 ComputeProjection

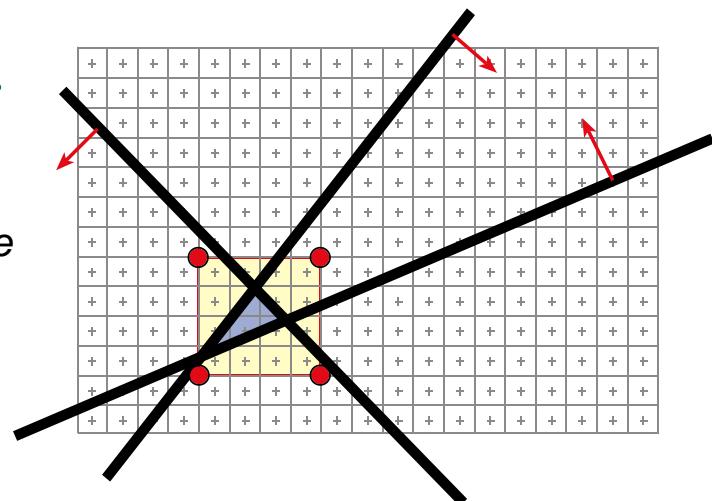
 Compute bbox, clip bbox to screen limits

 Compute line equations

 For all pixels in bbox

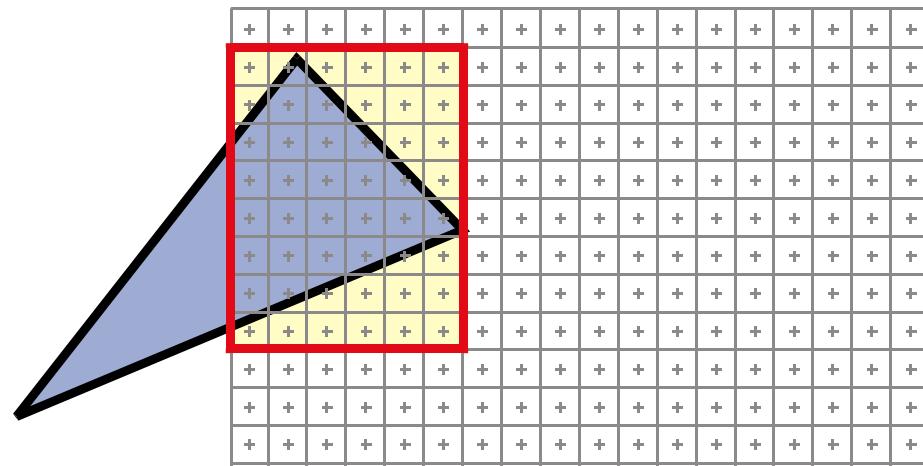
 If all line equations > 0 //Pixel [x,y] in triangle

 framebuffer[x,y] = color;



Rasterisation on Graphics Cards

Note that Bbox clipping is trivial, unlike triangle clipping

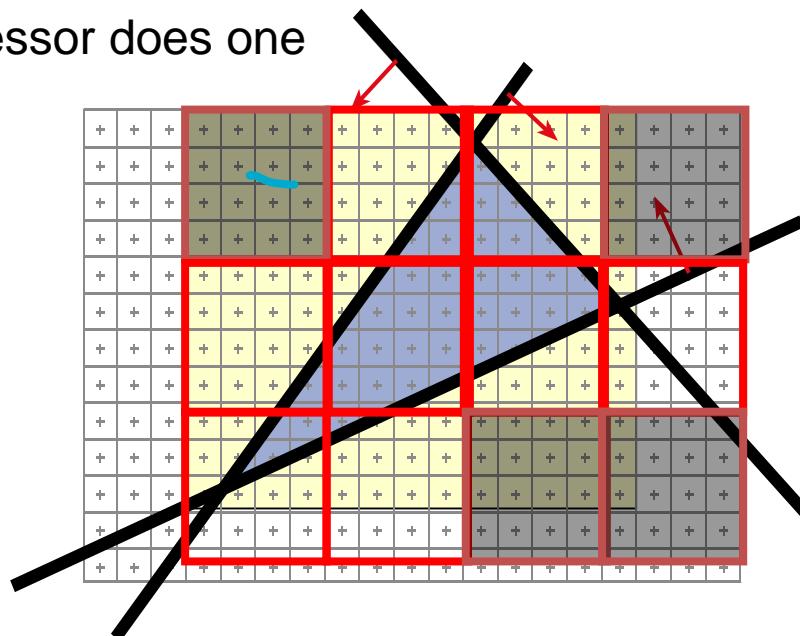
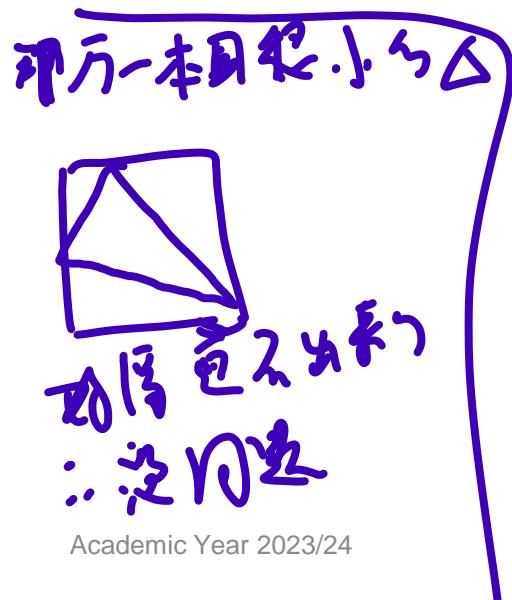


(直接按行按列分)
4x4
bbox

Tiling ✗

(没学过而已)

- Subdivide BBox into smaller tiles
 - Early rejection of tiles
 - Memory access coherence
 - Each microprocessor does one



查看 tile 与
四个角在不在
三角形里
(黑色的 tile
能可以不看)

Recap

- Moving away from ray-tracing to projection
- How to
 - Project a point
 - Clip a polygon
 - Cull a polygon
 - Fill the interior (rasterise)
- We'll spend next lectures tidying up the remaining problems!