

Computer Graphics (COMP0027) 2022/23

Path tracing

Tobias Ritschel

Today

- Solving integrals (approximately)
 - What did ray-tracing do?
 - Analytic vs. Numeric
 - Monte Carlo method
- Solving the RE using Monte Carlo
- Variance reduction
 - Good Sampling patterns
 - Importance sampling

Why is this hard to solve?

- It involves an integral with no analytic solution
- It is an integral equation, so the RHS contains the LHS in an integrand

$$L(\mathbf{x}, \omega_o) = L_e(\mathbf{x}, \omega_o) + \int f_r(\omega_i, \omega_o) L(\mathbf{y}, -\omega_i) \cos \theta_i d\omega_i$$

(Spherical) integration

Same thing!

↑
normal dot product

How ray-tracing solves the RE

- Many got suspicious about ray-tracing
- Example: Does metal have finite gloss?
 - Yes, cause otherwise I cant see highlight!
 - No, reflections are not blurry in steel balls!
- Contradiction!



How ray-tracing solves the RE

rendering
equation

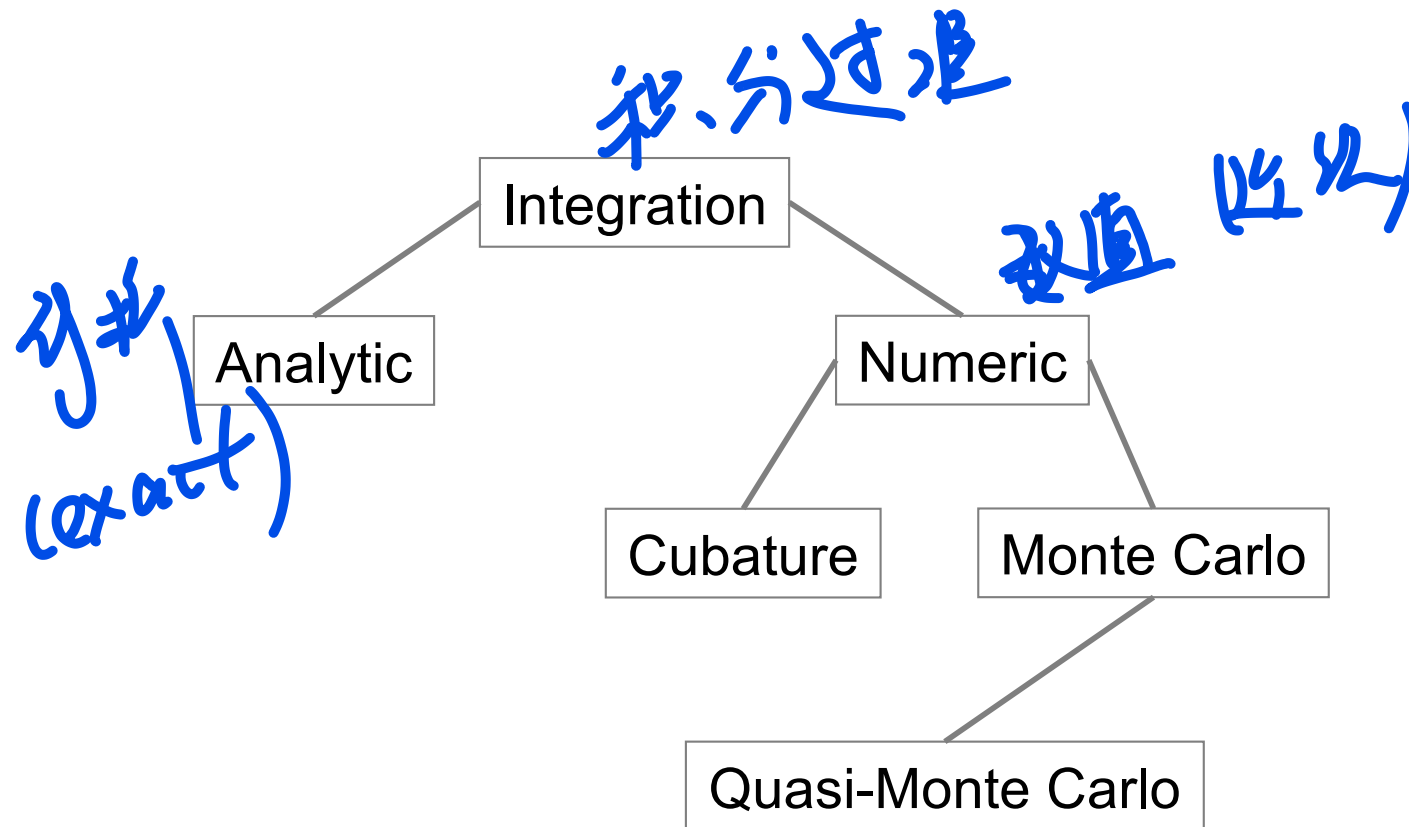
- The integral is split into a sum of two:
 - A Dirac-paths, solved by binary recusing
 - A path connecting to a point light, can evaluate without recursion

.

Solving integrals

- Analytic (accurate)
 - Given the function f , try to find F by symbolic manipulation
- Numeric (approximate)
 - Cubature
 - The Monte Carlo method
 - Both are approximate


Classification



Analytic


- Examples
 - $f(x) = x$, $F(x) = \frac{1}{2} x^2$
 - $f(x) = \sin(x)$, $F(x) = -\cos(x)$
- Difficult for a function such as we have
 - Impossible, as: The input is not even analytic (what is f for a bunny in the sun?)
 - Difficult, as: Recursion
 - Also difficult: Spherical domain

Some analytic things work




$$f(\omega) = f(\theta, \phi) = \begin{cases} c & \text{if } 2 < \theta < 3 \text{ and } -1 < \phi < 1 \\ 0 & \text{otherwise} \end{cases}$$


Forget about spheres for now



Numerical integration




Numerical integration




You should recall from high school analysis course.

Numerical integration

$1 + 1 + 1 + 1 +$
 $0.2 + 0.2 + 1 + 1 +$
 $1 + 0.2 + 0.7 + 1 +$
 $1 + 1 + 1 + 1$
 \rightarrow
 $13.3 / 16 = 0.83125$
 \sim
 0.85




Numerical integration



Numerical integration

...
 →
 $54.2 / 64 = 0.846$
 ~
 0.85



Cubature: A bit more formal

$$F(\omega) \text{ on } \Omega = \int f(\omega) = \sum f(\omega_i) / N$$

To find the integral F of a function f , decompose Ω it into N as-small-as-possible cubes, evaluate f on each (**sample**) and average.

Is this it?

- Very simple method!

- To get $L(\mathbf{x}, \omega_o)$
- Subdivide hemisphere Ω above \mathbf{x} into N strata ω_i
- Evaluate integrand, i.e., send a ray
 - Triple product of light, svBRDF and geometric term
- Average
- Done!

分成 N 个
input
第 i 遍 iteration $(i=1 \sim N)$


Problem: Curse of dimensionality

- L appears on both sides
 - For $L(\mathbf{y}, -\omega_i)$ need to solve another integral
 - OK, let's go to \mathbf{y} and also compute $L(\mathbf{y}, -\omega_i)$
 - For $L(\mathbf{z}, -\omega_i)$ need to solve another integral
 - OK, let's go to \mathbf{z} and also compute $L(\mathbf{z}, -\omega_i)$
 - For $L(\mathbf{z}_2, -\omega_i)$ need to solve another integral
 - OK, let's go to \mathbf{z}_2 and also compute $L(\mathbf{z}_2, -\omega_i)$
 - » For $L(\mathbf{z}_3, -\omega_i)$ need to solve another integral
 - » OK, let's go to \mathbf{z}_3 and also compute $L(\mathbf{z}_3, -\omega_i)$
 - For $L(\mathbf{z}_2, -\omega_i)$ need to solve another integral
 - OK, let's go to \mathbf{z}_2 and also compute $L(\mathbf{z}_2, -\omega_i)$

太多递归

Recursion

- Recursion in CS is not evil
- Worked well in ray-tracing:




For depth d and reflection/refraction
the number of rays is

$$2^d$$

So three-bounce is
 $2^3 = 8$ rays / pixel

Recursion

- Recursion in CS is not evil
- Impractical for the real RE
- Typical N is maybe 100 to 1000




For depth d and N strata
the number of rays is

$$N^d$$

So three-bounce is
 $1000^3 = 1\text{B}$ rays / pixel

Another way to imagine it



Every time we want any accurate value (orange dot) at any bounce we need to resolve exponentially many others below it to proceed

作为青蛙, 如果我没有眼睛, 怎么知道这个 color

Alternative: Random samples

$$11 * 1.0 +$$
$$3 * 0.2 +$$
$$2 * 0.7$$




$$13.0 / 16 = 0.8125$$

~


0.85

✓
取16个




Alternative: Random samples


$$\begin{aligned} &11 * 1.0 + \\ &3 * 0.2 + \\ &3 * 0.7 \\ &\rightarrow \\ &13.7 / 16 = 0.8562 \\ &\sim \\ &0.85 \end{aligned}$$



What do we get from random?



Another way to imagine it



As we just need an approximate value, we can proceed
with any point without looking at all others

Monte Carlo: A bit more formal

$$F(\omega) \text{ on } \Omega = \int f(\omega) = \sum f(\omega_i) / N$$


To find the integral F of a function f , place as many samples N onto Ω , evaluate f on each and average.

This is it

- Still very simple method!
 - To get $L(\mathbf{x}, \omega_o)$
 - N times
 - random directions $\omega_{i,0}$ above \mathbf{x} , $\omega_{i,1}$ above \mathbf{y} , etc.
 - Evaluate integrand, i.e. send a ray
$$f = L(\mathbf{y}, -\omega_i) f_r(\mathbf{x}, \omega_i, \omega_o) \cos(\theta)$$
 - Triple product of light, svBRDF and geometric term
 - Done!

How to pick random directions?

- How to pick a random ray in 3D?
- `normalize(vec3(frand(), frand(), frand()))?`
- Clumps on axis and diagonals
- Bias in result!



When to stop?

- Multiple options
- Popular:
 - After a fixed depth
 - When contribution falls below a threshold

Progressiveness



1 Sample


Progressiveness

space itself as
发之海




2 Samples

Progressiveness



10 Samples

Progressiveness



1000 Samples

Progressiveness

- After N samples we get an image
- After $2N$ samples, we get an even better one
- This is called **progressive**
- Very useful for previews

Random has another reason

- **Aliasing** is the second reason for random
- Consider this integrand, not even recursive:

Cubature with 1 sample = 0.0



Cubature with 2 samples = 0.0



Cubature with 5 samples = 0.0



Ground truth = 0.5

Random has another reasons



- **Aliasing** is the second reason for random
- Consider this integrand, not even recursive:

Monte Carlo with 1 sample = 0.0



Monte Carlo with 2 samples = 0.5



Monte Carlo with 5 samples = 0.4




Ground truth = 0.5

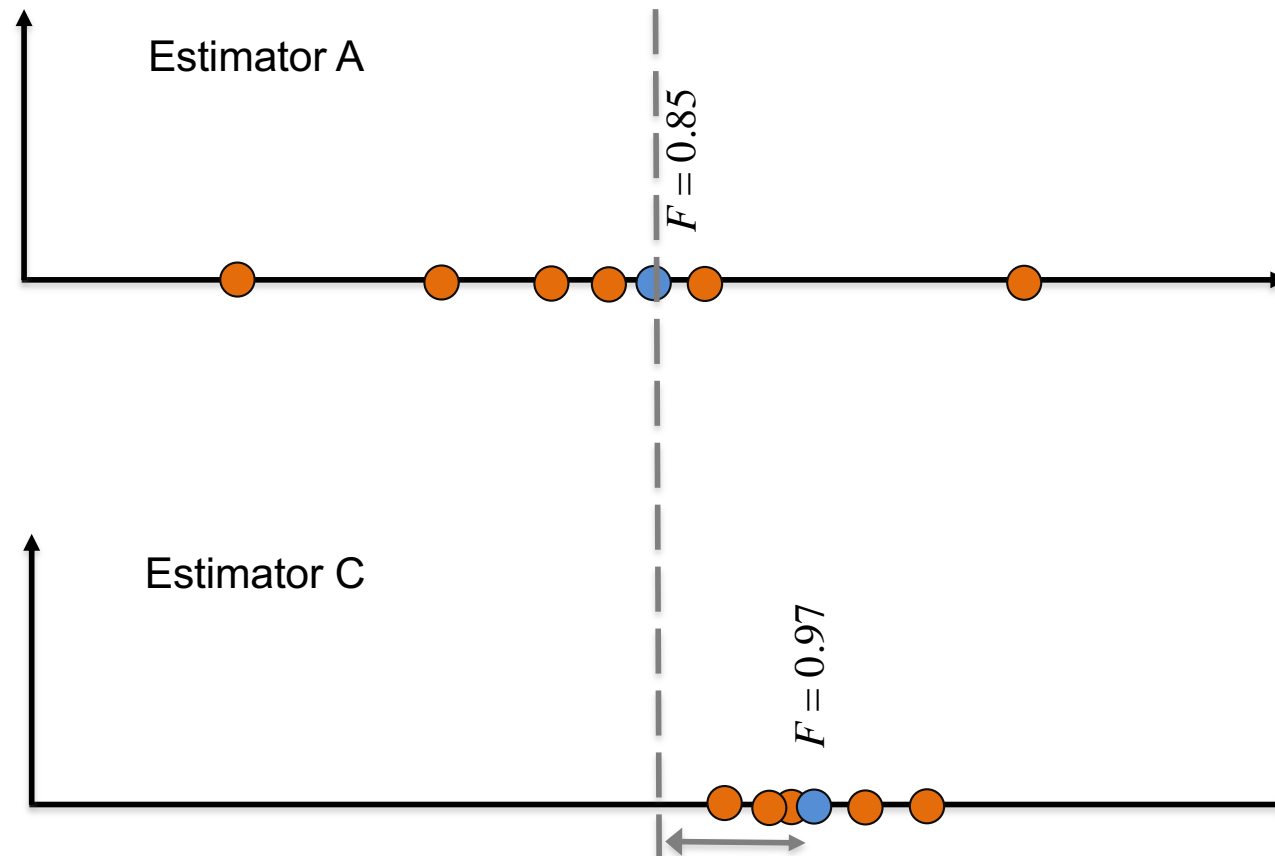
Estimator/Variance/Bias

- We get a new value for every random seed
- We map these into an **estimate**
- This is the value of the integral
- If there is a deviation, we call it **bias**
- Around this exists a distribution of values
- This distribution has a **variance**

Variance of an estimator



Bias of an estimator




Desiderata variance

- Also it has no bias, i.e., the expected value is really the solution of the integrand
- OpenGL and CW1 ray-tracing: All biased
- A good estimator has a low **variance**
- Whenever we render, we will get a value close to the true value
- To this end, we do **variance reduction**

Variance reduction


- Next-event estimation
- Good sample patterns
 - Jittered
 - Quasi-Monte Carlo
 - Blue noise
- Importance Sampling

Path tracing - High hopes




We hope for this ...

Path tracing – High hopes




We hope for this ...

Path tracing - Reality




But what we get is

Path Tracing - reality




But what we get is

Next-event estimation



But what we get is

Problem: Double-accounting



There are now **two ways** to hit the light.
Simple solution: Simply only take emission from NEE.

Computer Graphics (COMP0027), Tobias Ritschel

next event estimation


Next event estimation

- Two simple changes
 - Add a random ray in the direction to the light
 - Remove adding in emission L_e on all other paths
- Best for **small** light sources
- Result:
 - Will never miss direct light at any point
 - Still have all benefits of MC
- Glorified Whitted-style CW 1 ray-tracing




Endpoint choice

- Simple for flat area lights




Endpoint choice

- Hard for e.g., spheres



Uniform random can go bad

16
→
 $16 / 16 = 1$
~
0.85
(not so cool)




The ideal sample pattern

- Two contradicting goals:
 1. Maybe not be regular
 2. But always cover the domain uniformly, not only in the limit



Jittering *(random & not random at the same time)*

- Back to the future:
 - Do cubature first
 - Then jitter every sample inside its cell
- Suffers from the curse of dimensionality
- Prevents aliasing
- Applicable if dimensionality is low
 - Example: Area light sampling

Regular (recap)




Jittered



① in cell random
② 每个 cell 一个


Multi-Jittered




(Better than jitter)
Halton

- A way to place samples
 - Somewhat uniform
 - Without structure
 - In high dimensions
- A typical work-horse solution for rendering
- Defined on the unit hypercube


Radical inverse base $n = 2$



Radical inverse base $n = 3$





Halton in 2,3 i.e. $\pi(0), \pi(1)$



not random


Random vs. Halton



Quasi-Monte Carlo sampling

- To get an n -dimensional Halton pattern
- Build the radical inverse in the co-prime basis $\pi(0, \dots, n-1)$
- Build tuples in the order they occur in each sequence
- Improvement: **Hammersley** (regular 1st dim.)
- De-correlation: **Cranely patterson** rotation

Randomized Quasi-Monte Carlo



Random


Low-discrepancy

Structured artefacts





ca 250 Hammersley samples. No NEE.

Randomized Quasi-Monte Carlo



Random



Randomize Low-discrepancy

*cranny
rotation
rotate
ish*

Regularity


- The L2 error of the images is the same or (much) lower than random
- However, quite suspicious visually
- Easy fix:
 - Combine regular and random
 - Shift the regular pattern by a random offset
 - Use torroidal shifting
i.e., come in left when going out right



Poisson disk / Blue noise


- Patterns with a maximal minimal distance between all points is called **Poisson disk**
- Another way is to see the spectrum of the distribution of distances: It is blue, i.e. no small minimal distances

Poisson disk




Note: The smallest circle of all circles around each point is quite large

Blue noise



Receptor distribution on the macaque retina prevents aliasing

Importance sampling



Put more samples where the integrand is high, as here the errors have the largest effect.

Importance sampling

- Placing the samples non-uniformly will introduce bias
- Fortunately, random uniform is just a special case of a more general estimator formulation we will see next
 - Before we took ω_i uniform, so $p(\omega_i) = 1 / |\Omega|$
 - Any other p will work as well
 - Ideal $p \sim f$

MC with importance sampling

$$F(\omega) \text{ on } \Omega = \int f(\omega) = 1/N \sum f(\omega_i) / p(\omega_i)$$

To find the integral F of a function f , place as many samples N onto Ω according to a distribution p , evaluate f and p on each and divide.

What can be p ?


- Recall the integrand is
$$f = L(\mathbf{y}, -\boldsymbol{\omega}_i) f_r(\mathbf{x}, \boldsymbol{\omega}_i, \boldsymbol{\omega}_o) \cos(\theta)$$
- We could sample for
 - Light L : Hard, integral equation itself.
 - BRDF f_r : Not so hard, done analytically
 - Geometric term $\cos(\theta)$: Even easier analytically
 - Products of all of the above: Even harder than any alone, but doable

Example: IS for direct light



Same amount of rays ©U Virginia

Anti-aliasing



Recap

- Rendering is solving an integral equation
- Analytic and some numeric methods no-go
- Monte Carlo is the method of choice
- Suffers from noise (variance)
- Need to use variance reduction methods
 - Next event-estimation
 - Low-discrepancy Sample patterns
 - Importance sampling