

COMP0026: Image Processing

Image Filtering



COMP0026: Image Processing

Image Filtering



```
I = imread('moon.tif');
h = fspecial('unsharp');
I2 = imfilter(I,h);
imshow(I), title('Original Image')
figure, imshow(I2), title('Filtered Image')
```

COMP0026: Image Processing

Image Filtering

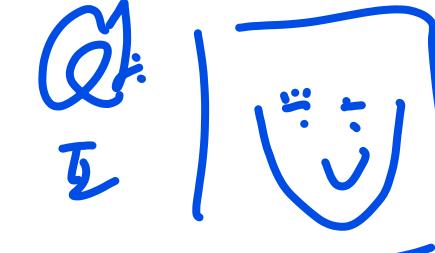


```
I = imread('moon.tif');
h = fspecial('unsharp');
I2 = imfilter(I,h);
imshow(I), title('Original Image')
figure, imshow(I2), title('Filtered Image')
```



Lectures will be Recorded

Notes: (Introduction):



P_i

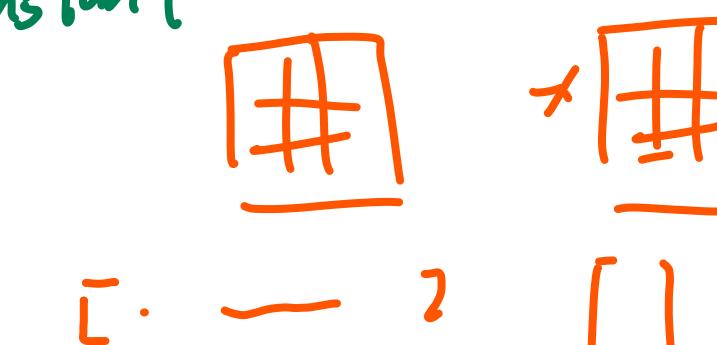


template

$$A. \text{ SSD} = \sum_{x,y} |I(x,y) - T(x,y)|^2$$

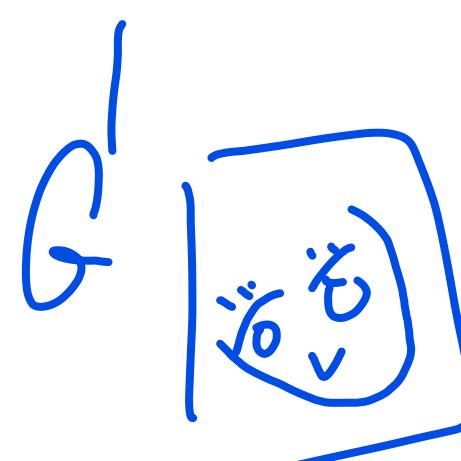
$$= I^2 + T^2 - 2 \int I(x,y)T(x,y)$$

constant correlation (convolution)

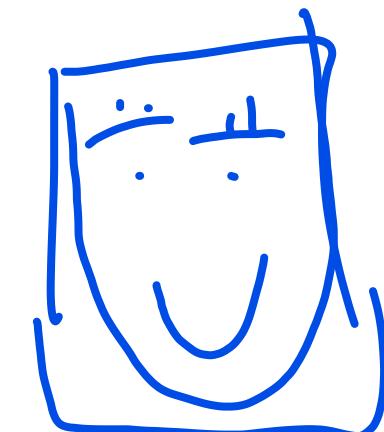


\rightarrow

invariance



↓+rotate



↓scale

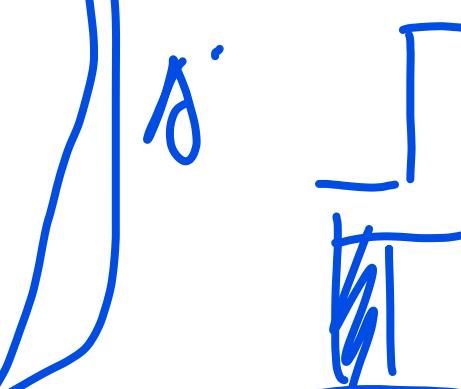


Intensity

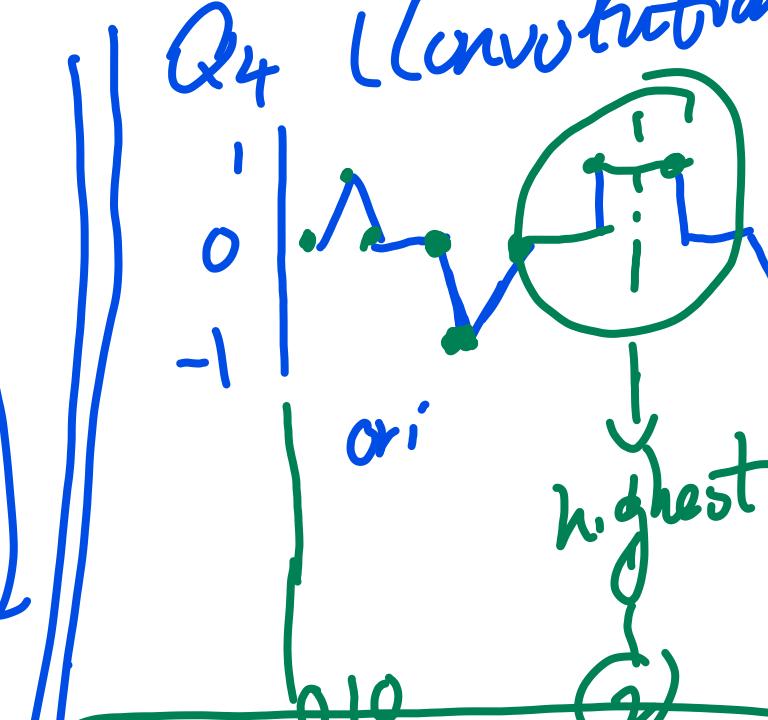
Q2: smooth pic

$$B: \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} \Rightarrow \begin{bmatrix} 1 & 0 & 1 \\ 0 & 2 & 0 \\ 1 & 0 & 1 \end{bmatrix}$$

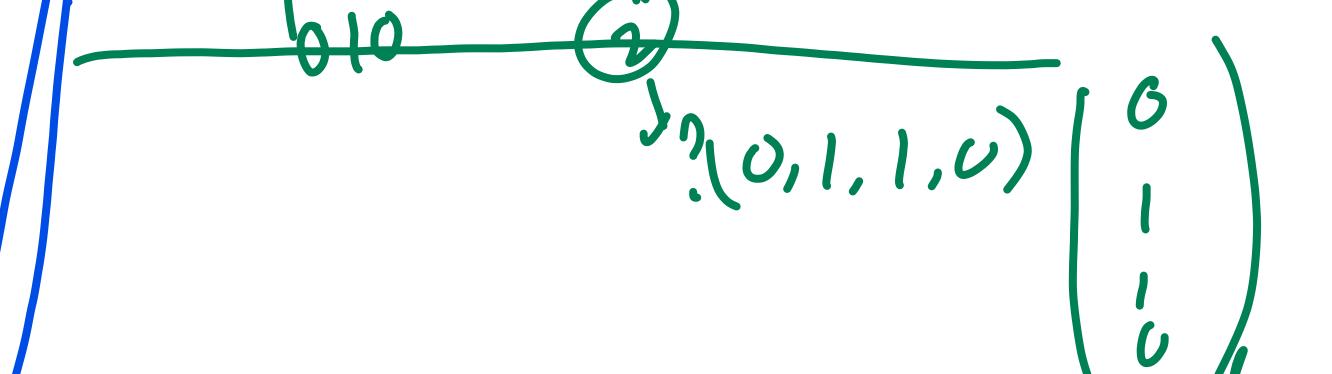
Q3: find edge.



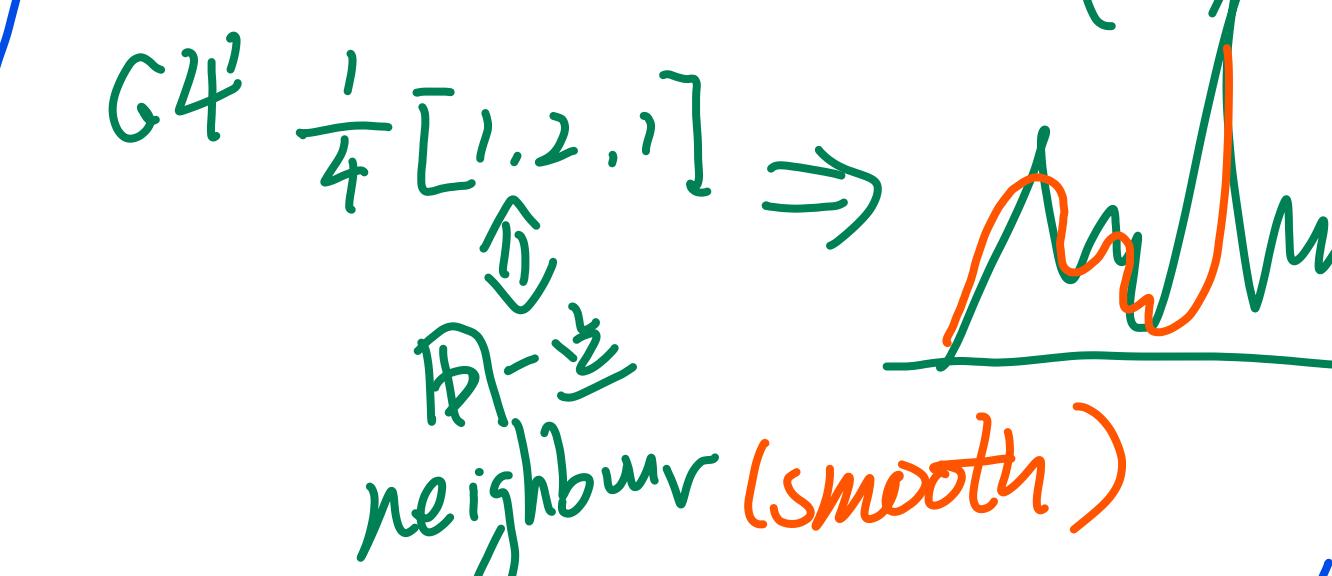
Q4: (convolution in 1D)



filter.



$$G4' \frac{1}{4} \begin{bmatrix} 1 & 2 & 1 \end{bmatrix} \Rightarrow$$



O5: spatial
(得立叶滤)

frequency

smooth (derivative)

Dif (1st, 2nd)

Sharpening

Image Filtering

input image → filter function → output image

Image Filtering

input image → filter function → output image

Functions

linear/non-linear

Neighborhoods

local/non-local

Goals

- Low-level data (e.g., pixel) processing operations
- Smoothing and noise reduction
- Feature extraction and enhancement

Non-local Filtering

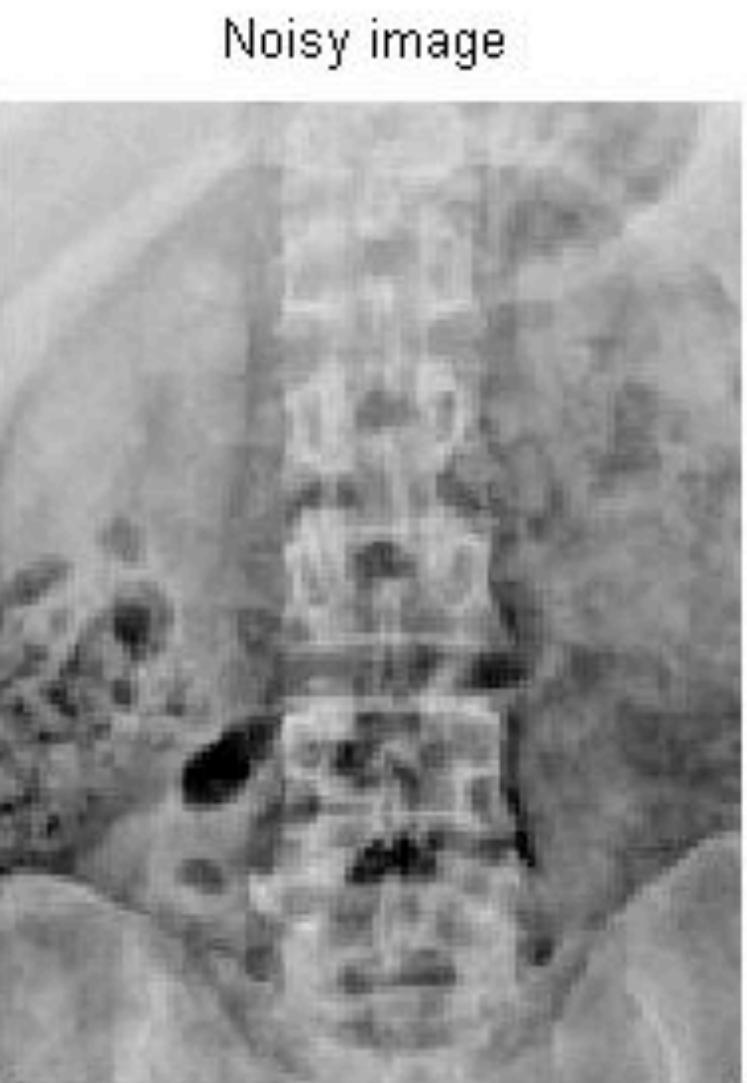
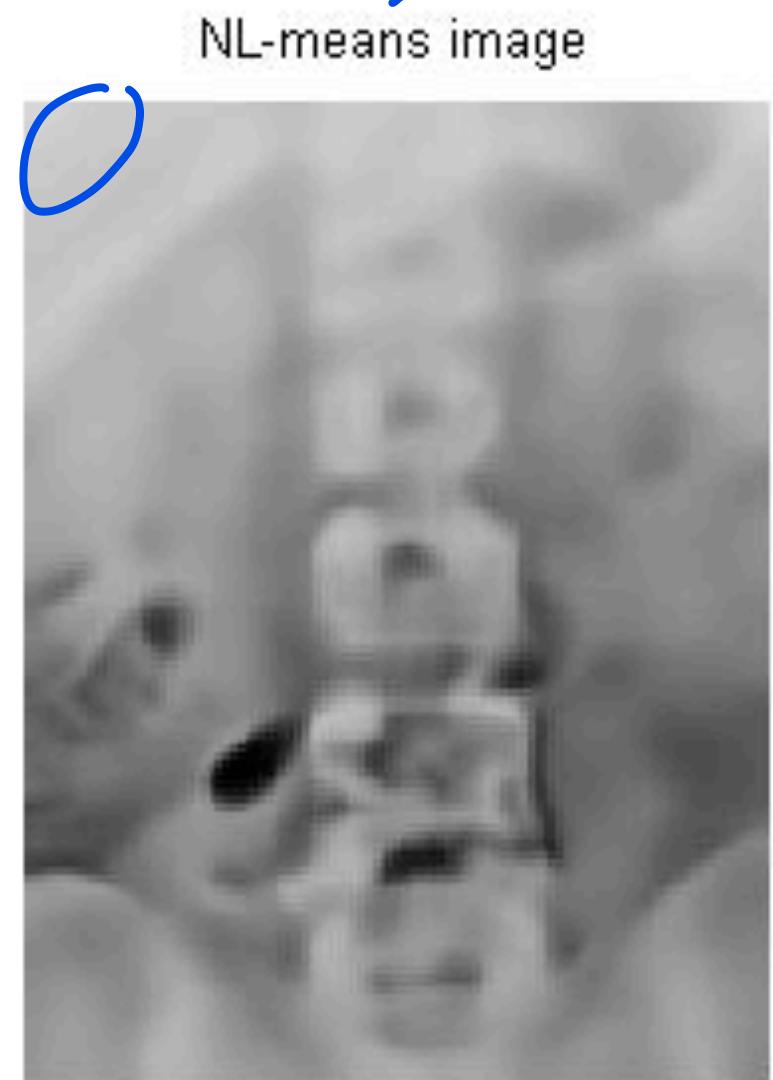


Image Filtering

Non-local Filtering



Noisy image



NL-means image



correlation
or convolution
by kernel

how similar

look all the pixels

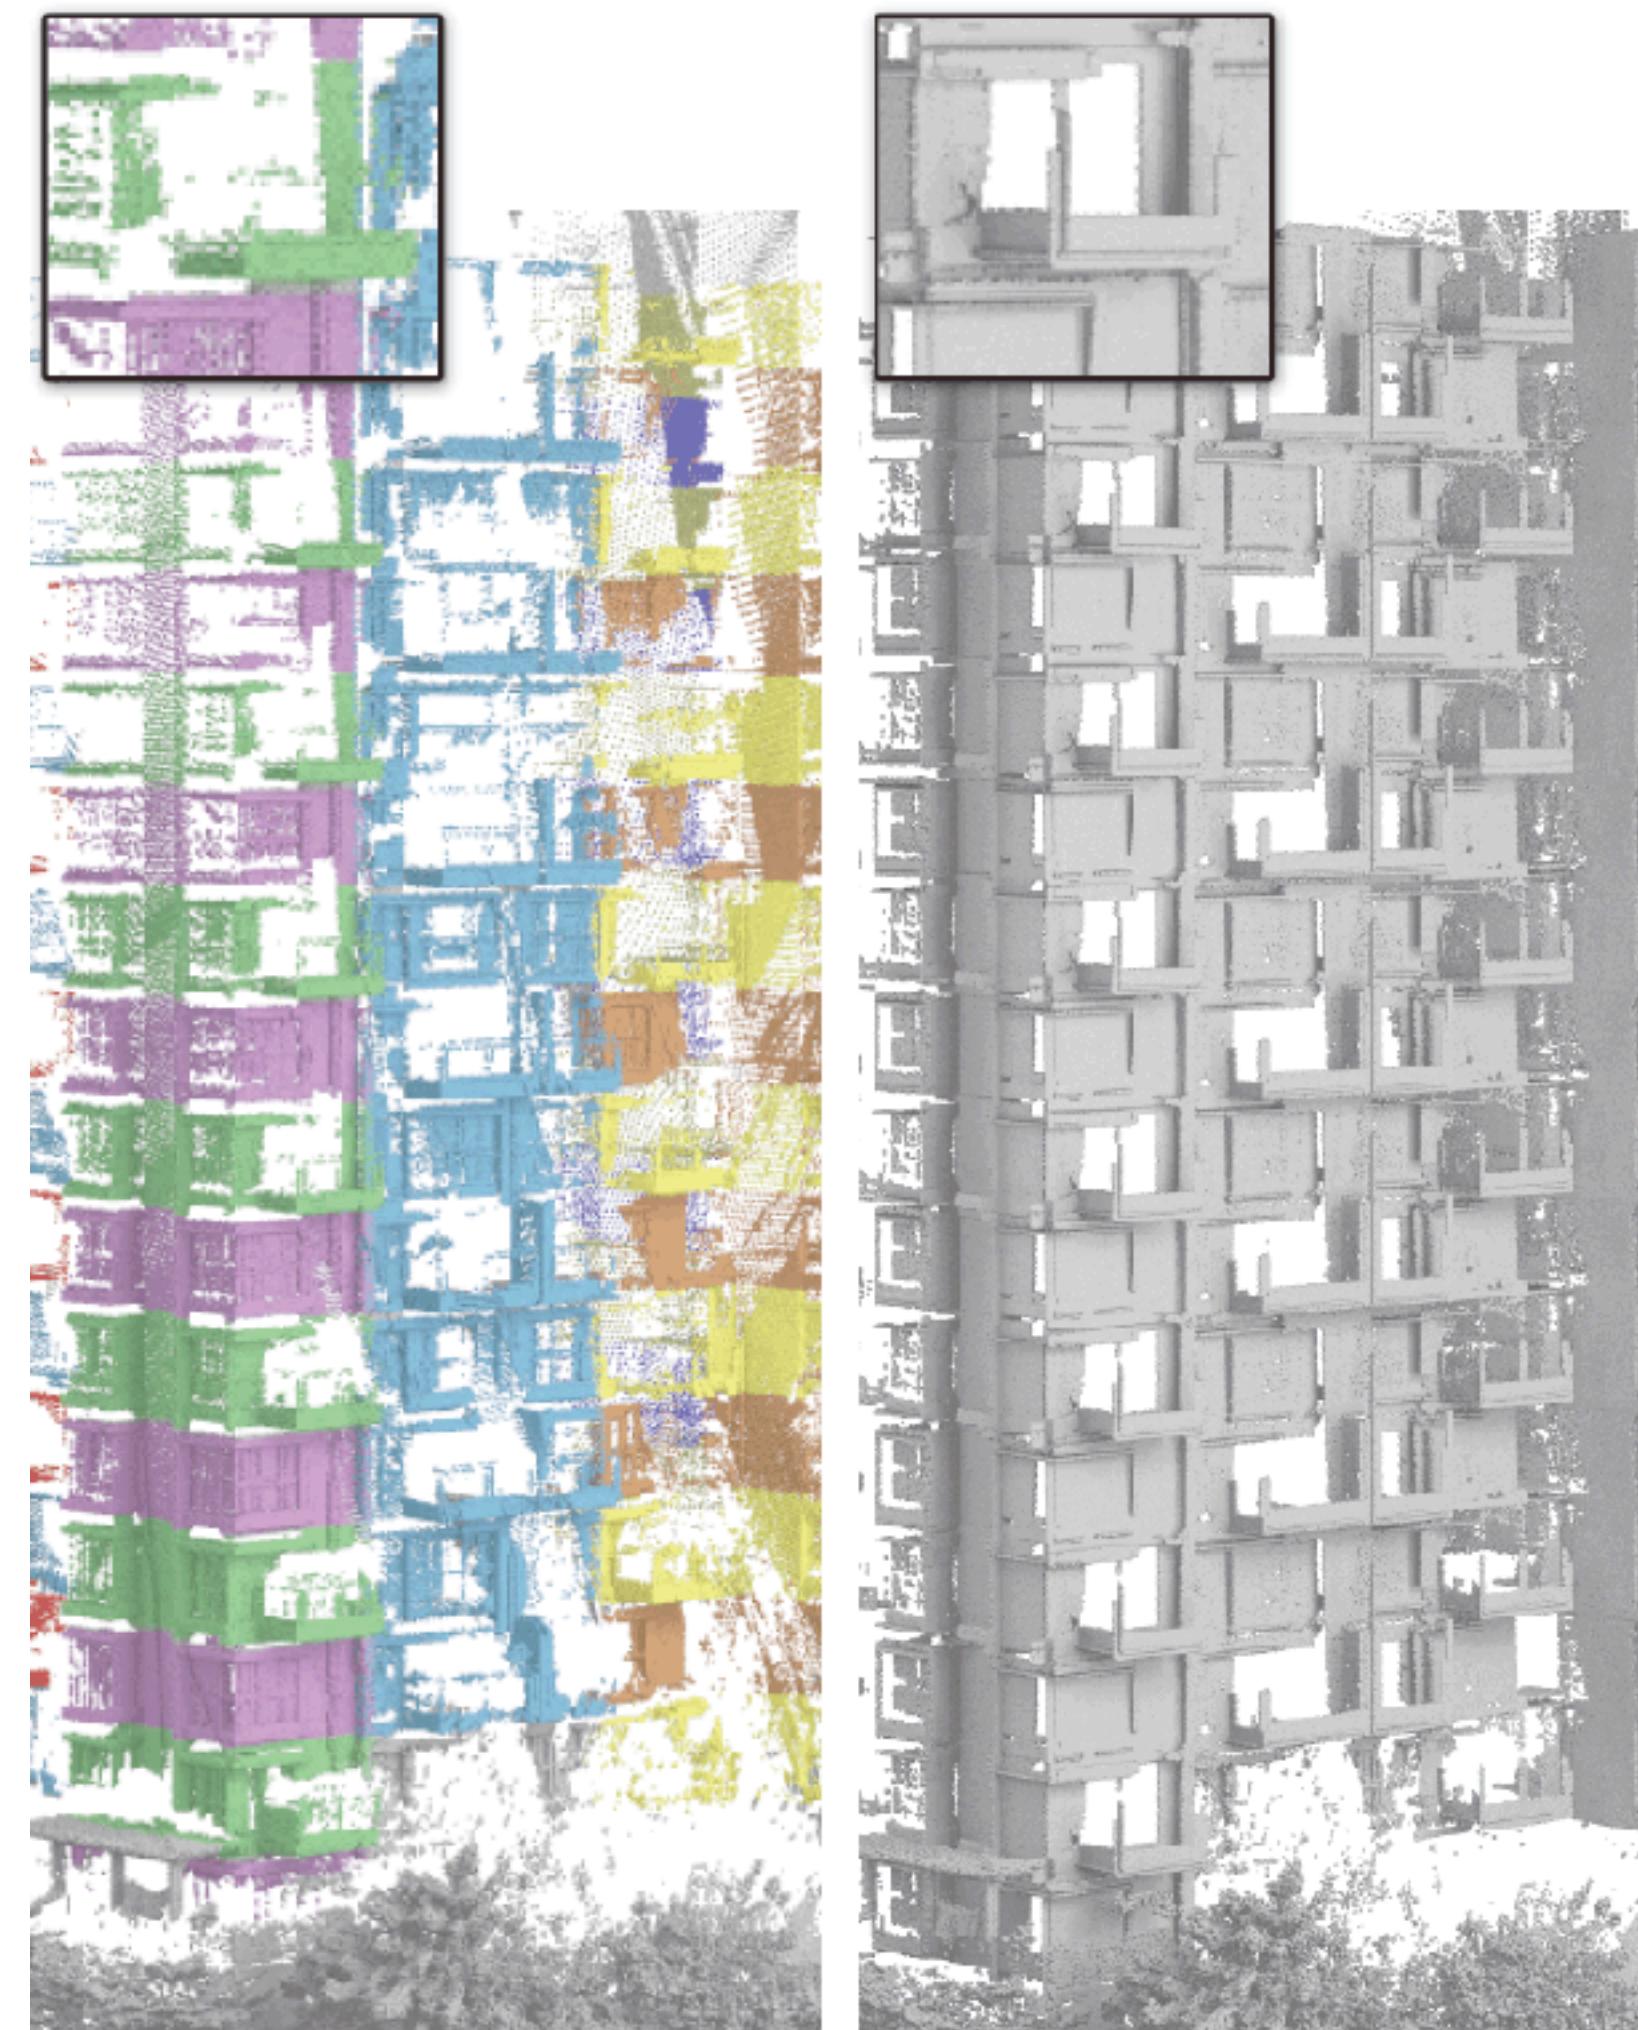
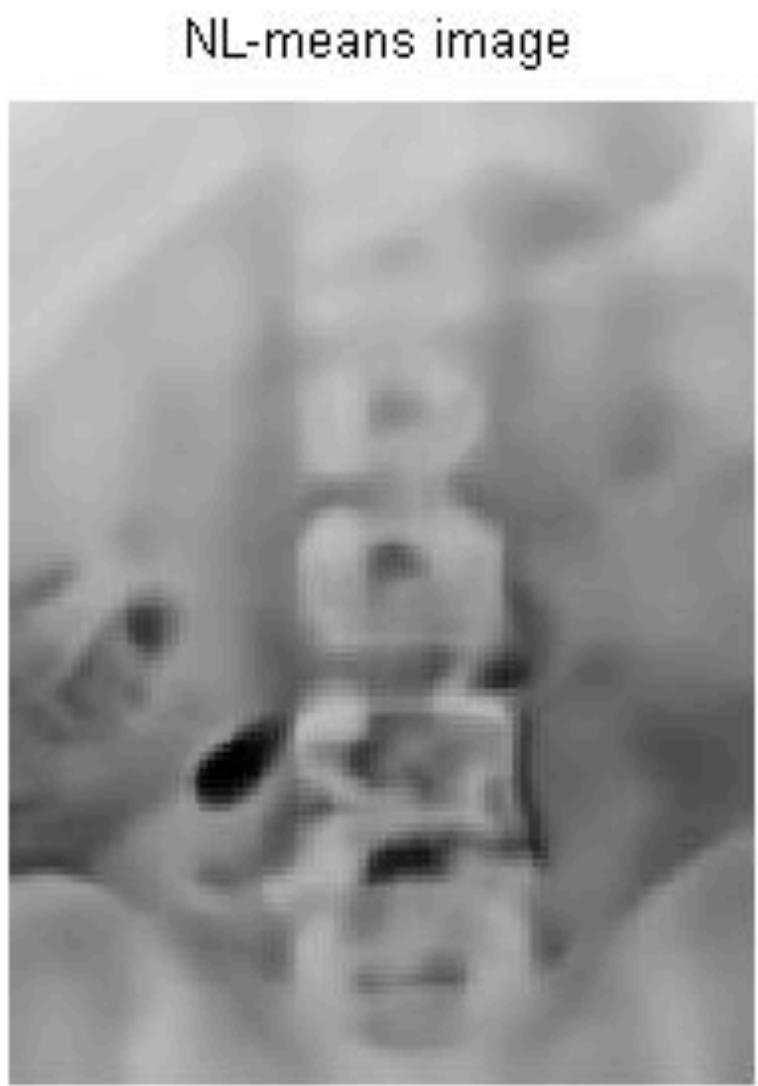
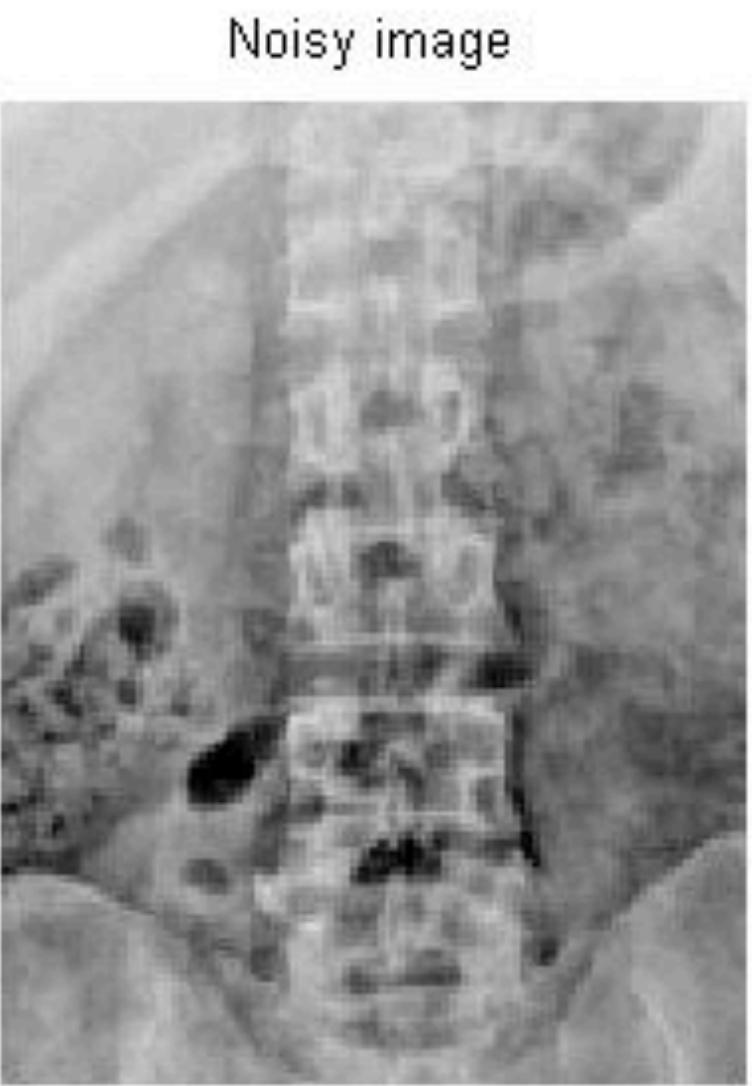
所有窗口 (像素块)

在全国找相似窗口

加权 & denoise

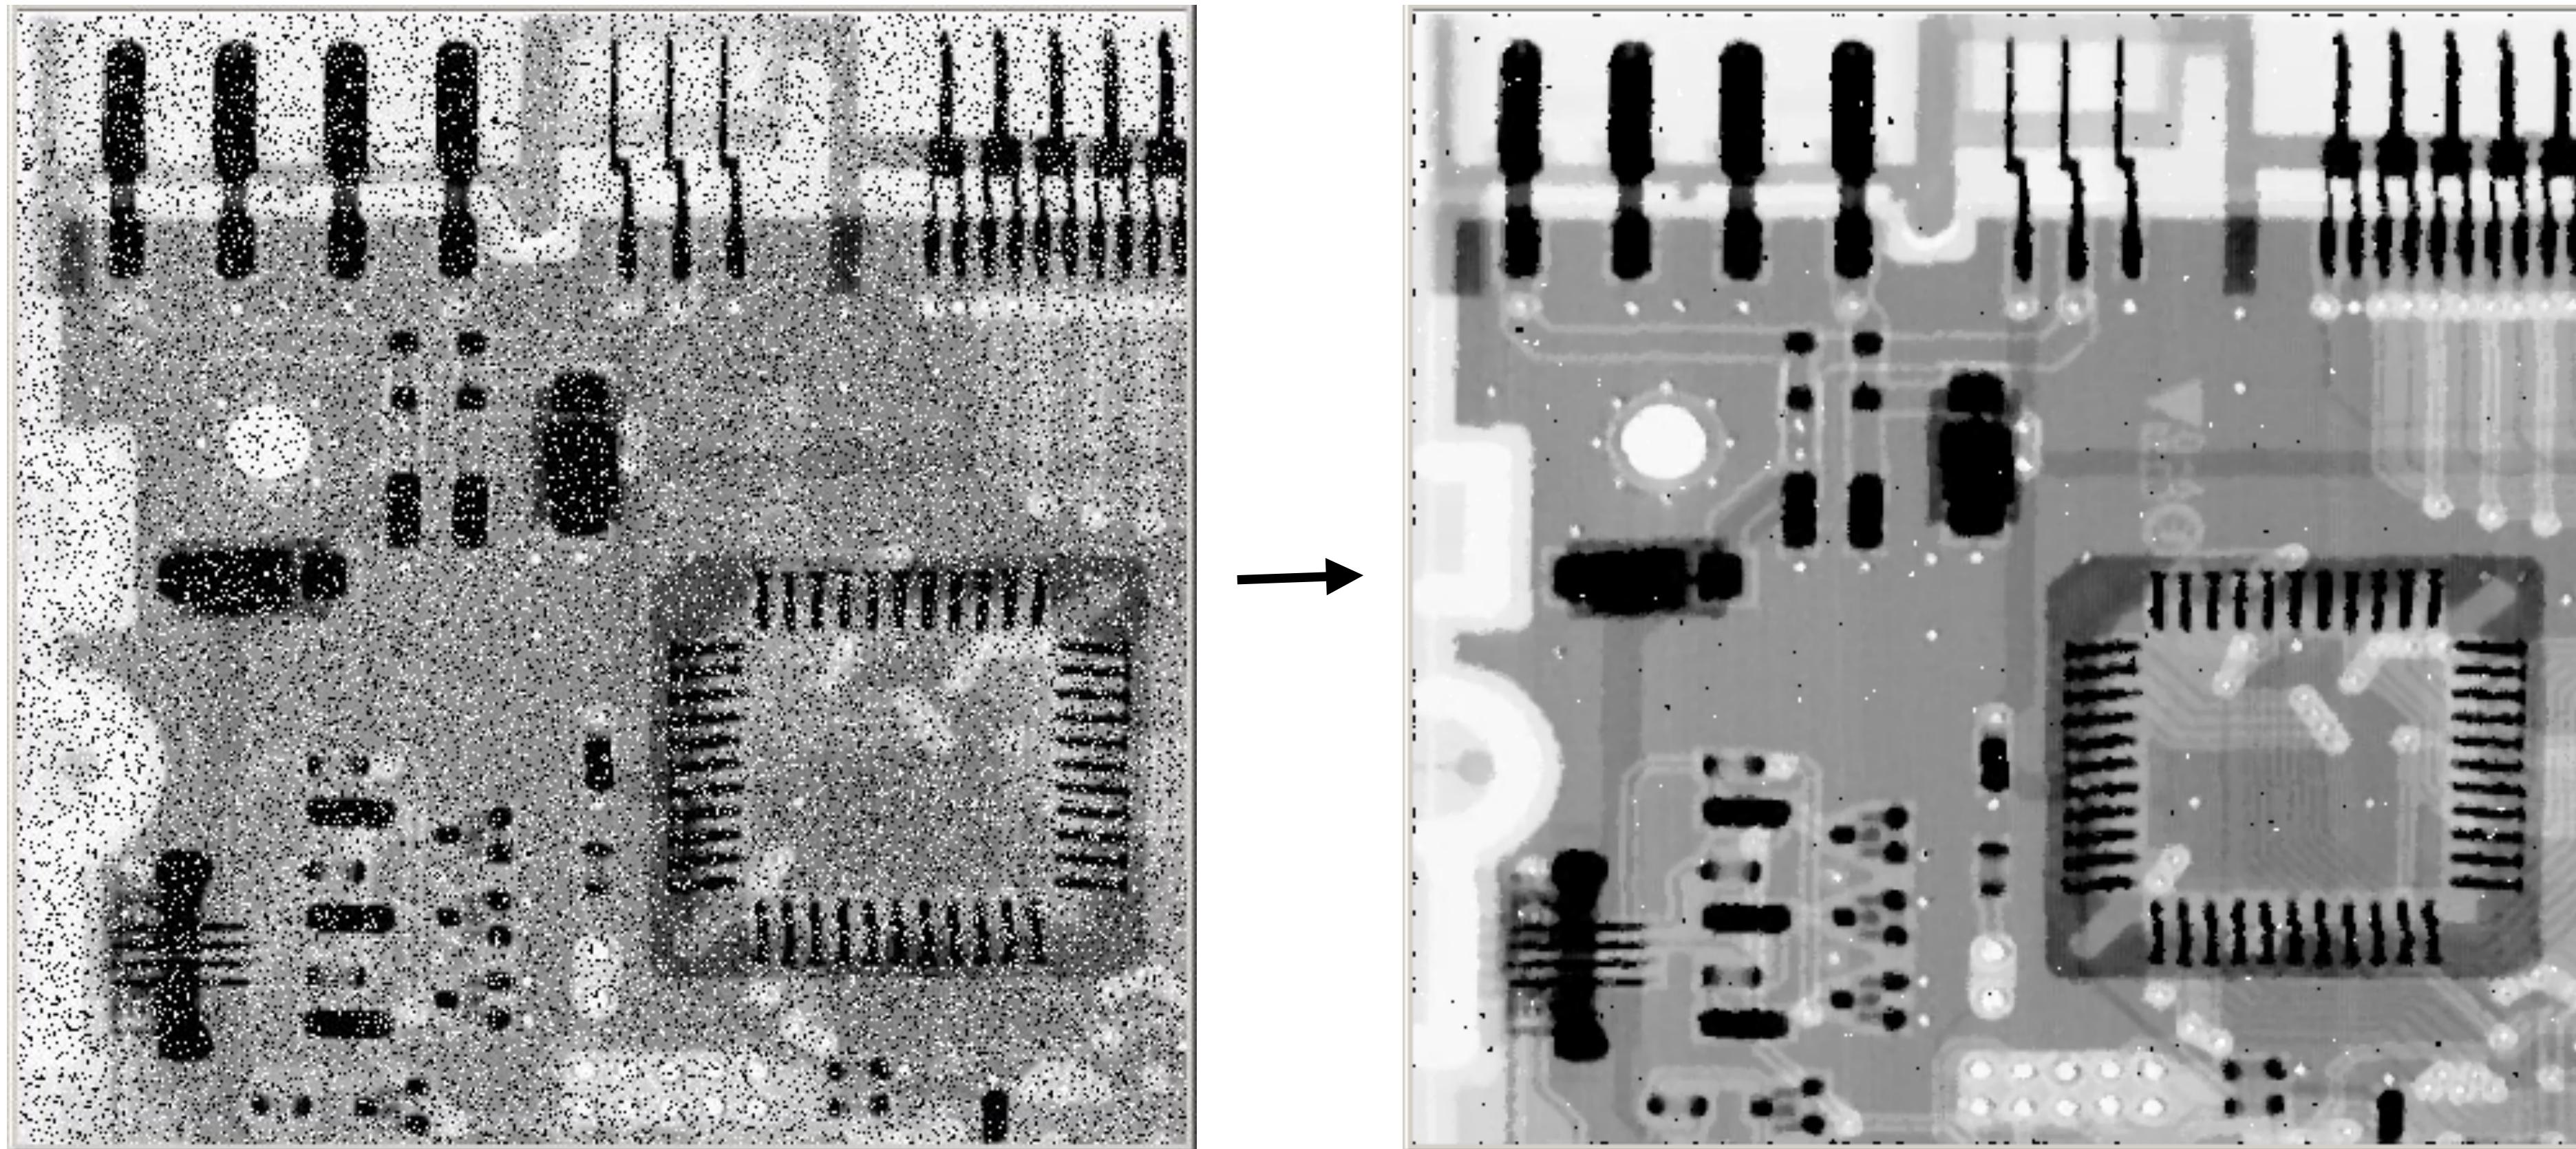
根据相似度
决定

Non-local Filtering



Example: Remove “Noise”

(if you know how to model that noise specifically!)



Input image from DIP by Gonzalez+Woods, 2008

median filter

Outline

- **Linear filtering**
- **Convolution operations**
 - Smoothing
 - Low-pass and high-pass filters
 - Sharpen
- **Filtering in the Fourier domain**
- **Rank filters**

Linear Mappings

- Equivalently, L is *linear* if

$$L(I_1 + I_2) = L(I_1) + L(I_2)$$

$$L(\alpha I) = \alpha L(I)$$

(i.e., L preserves linear combinations)

$$L(\alpha I_1 + \beta I_2) = \alpha L(I_1) + \beta L(I_2)$$

Convolutional Filters

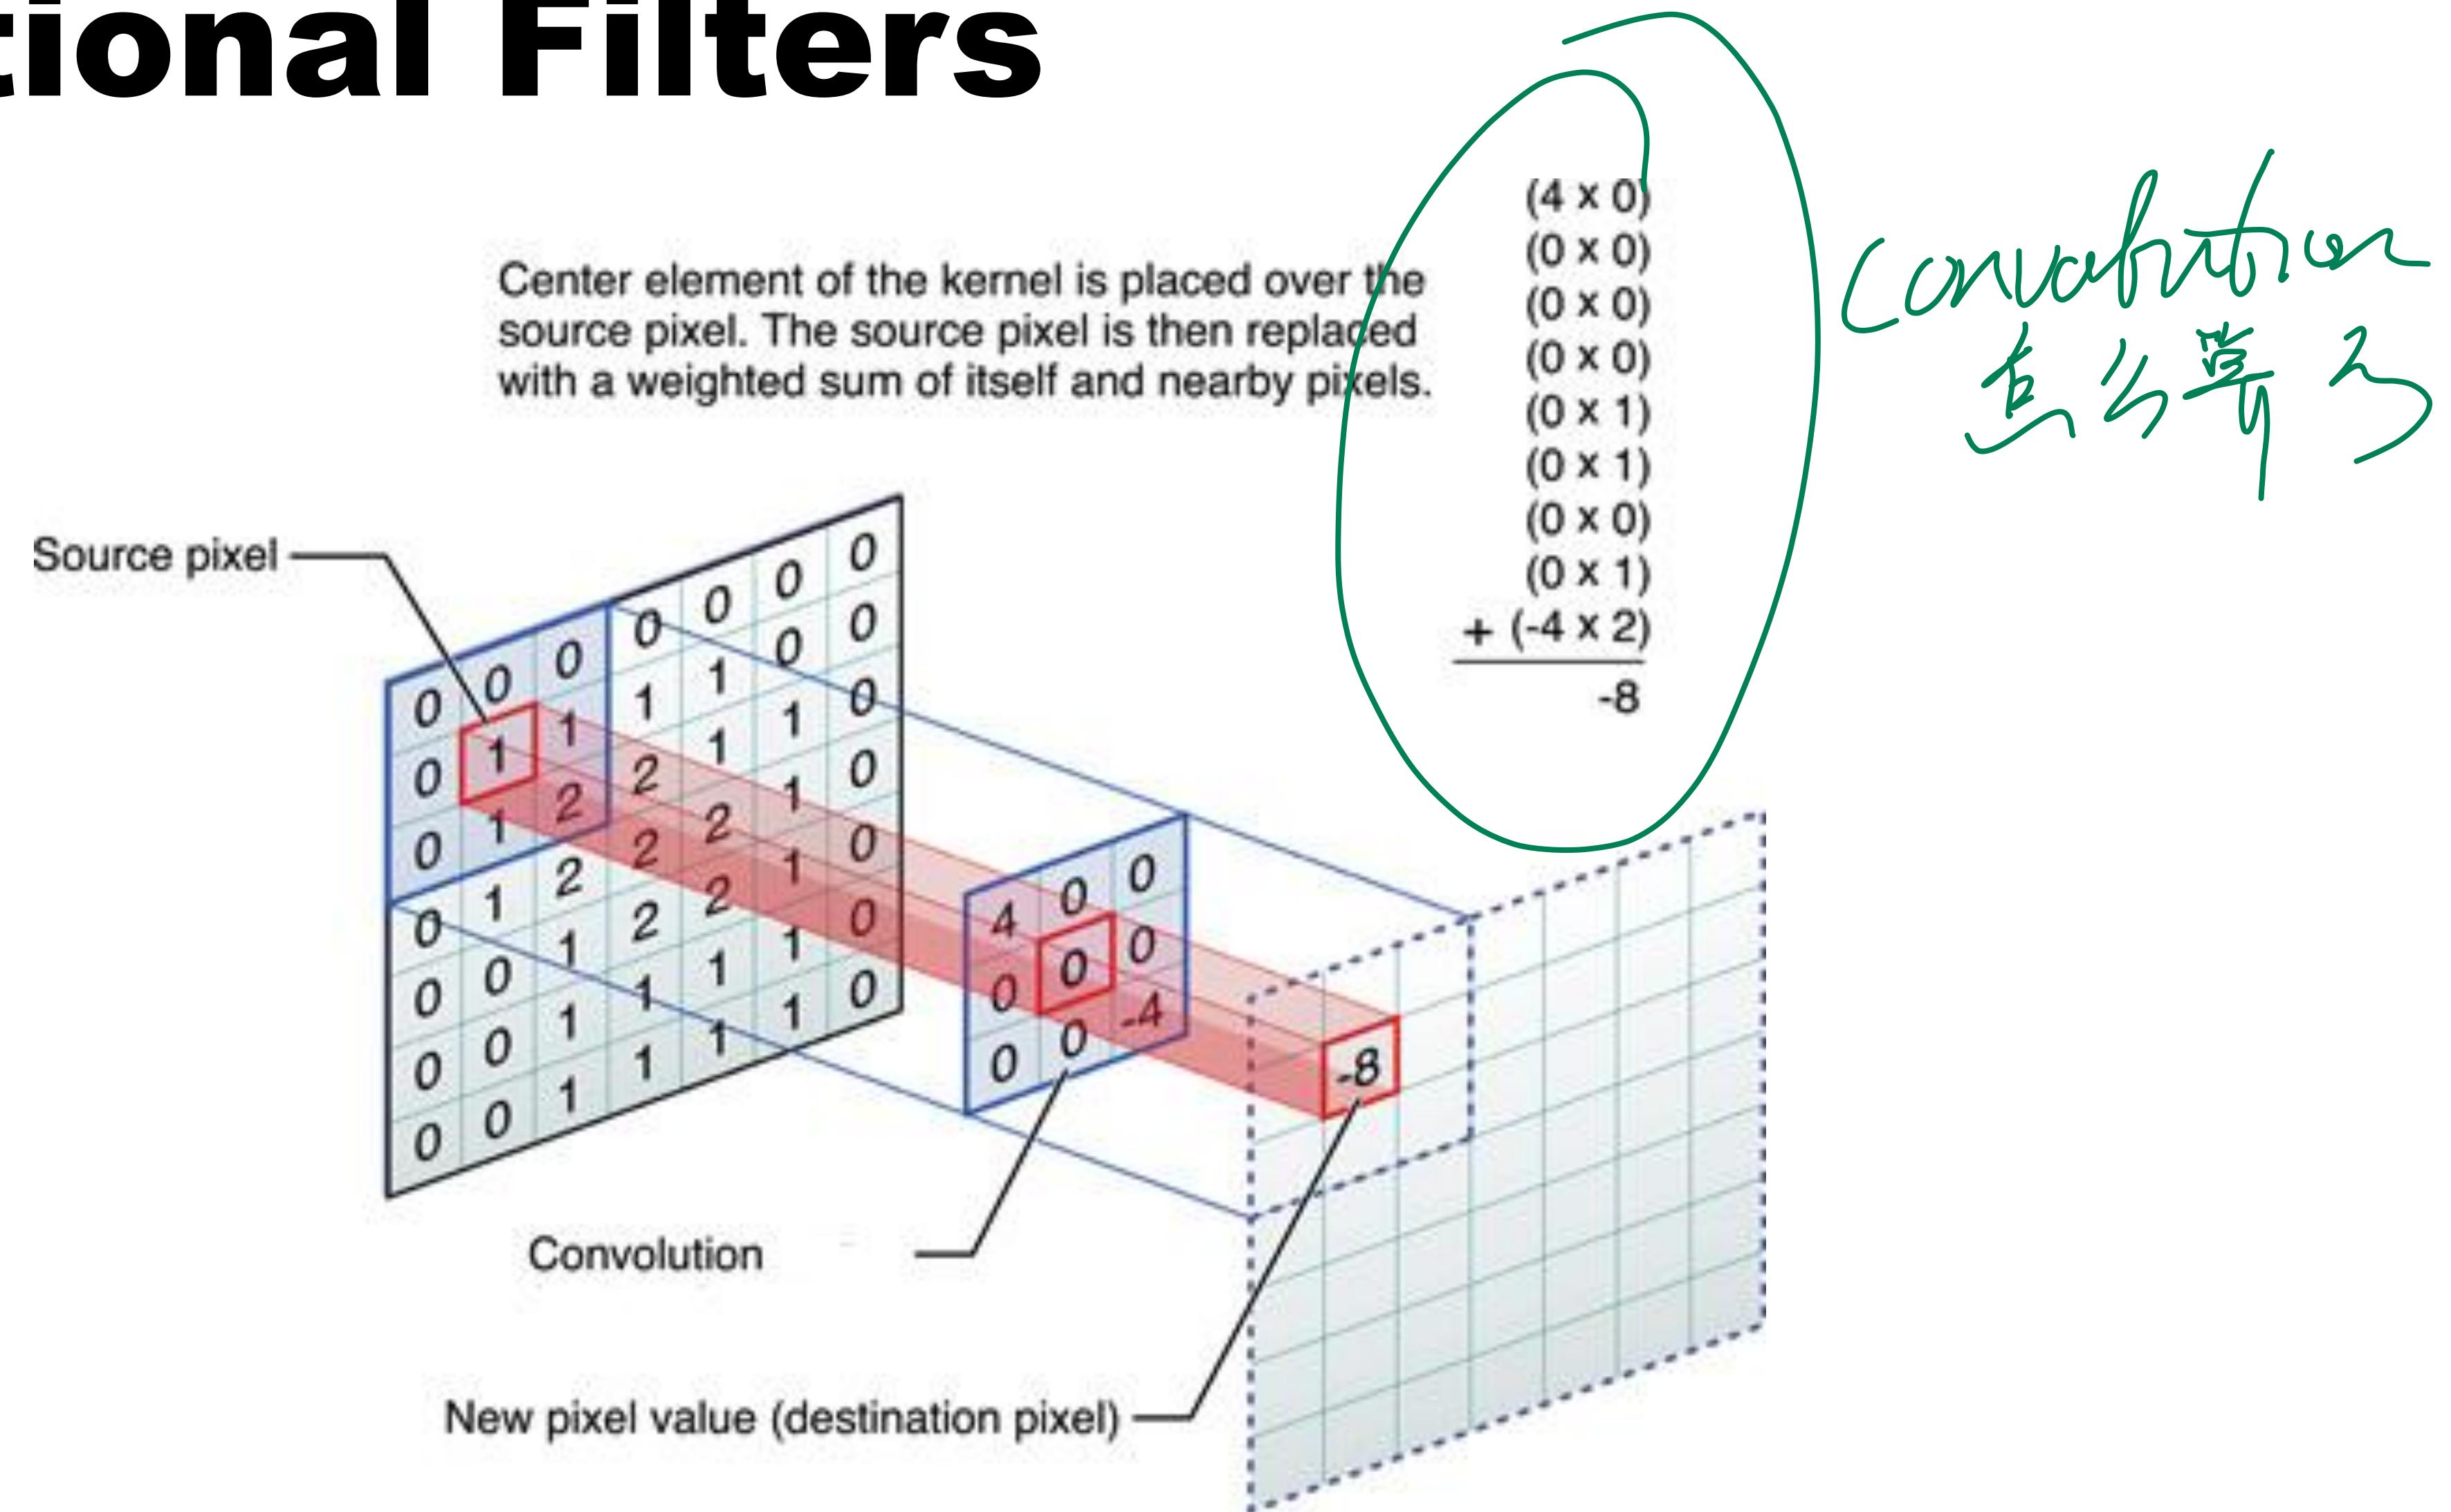
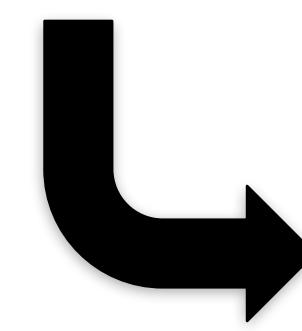
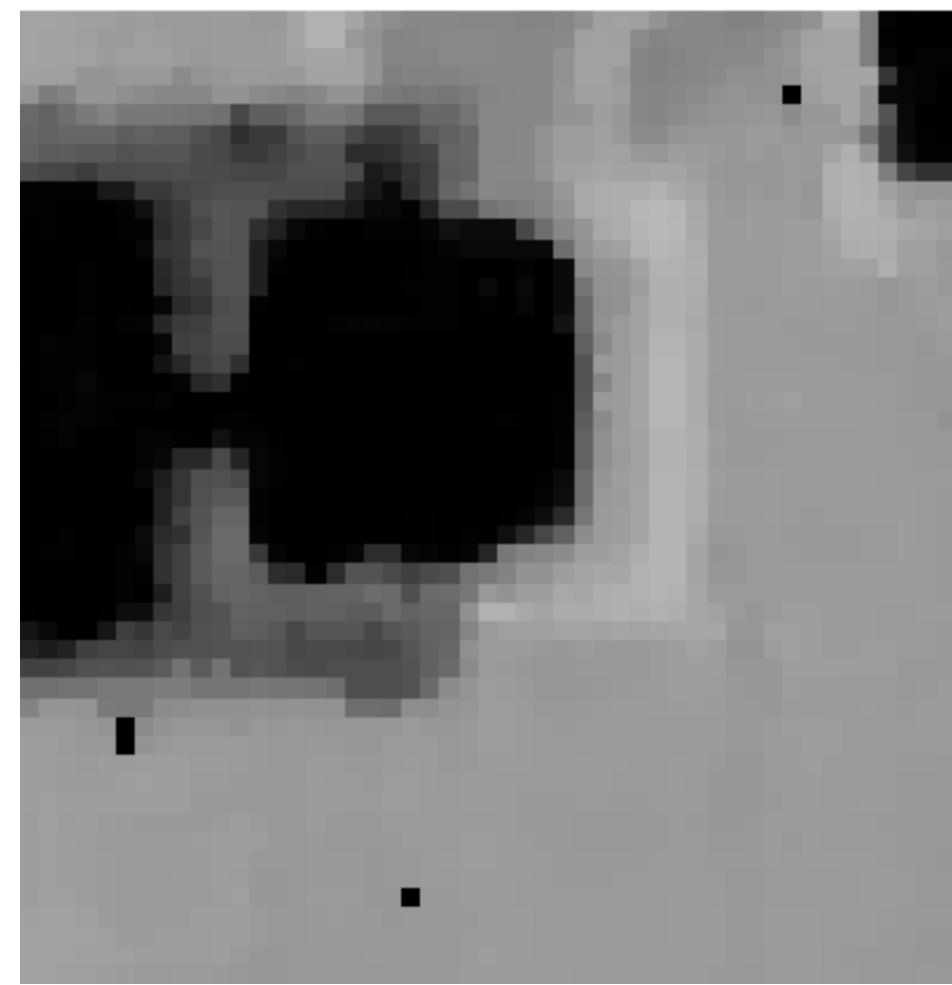
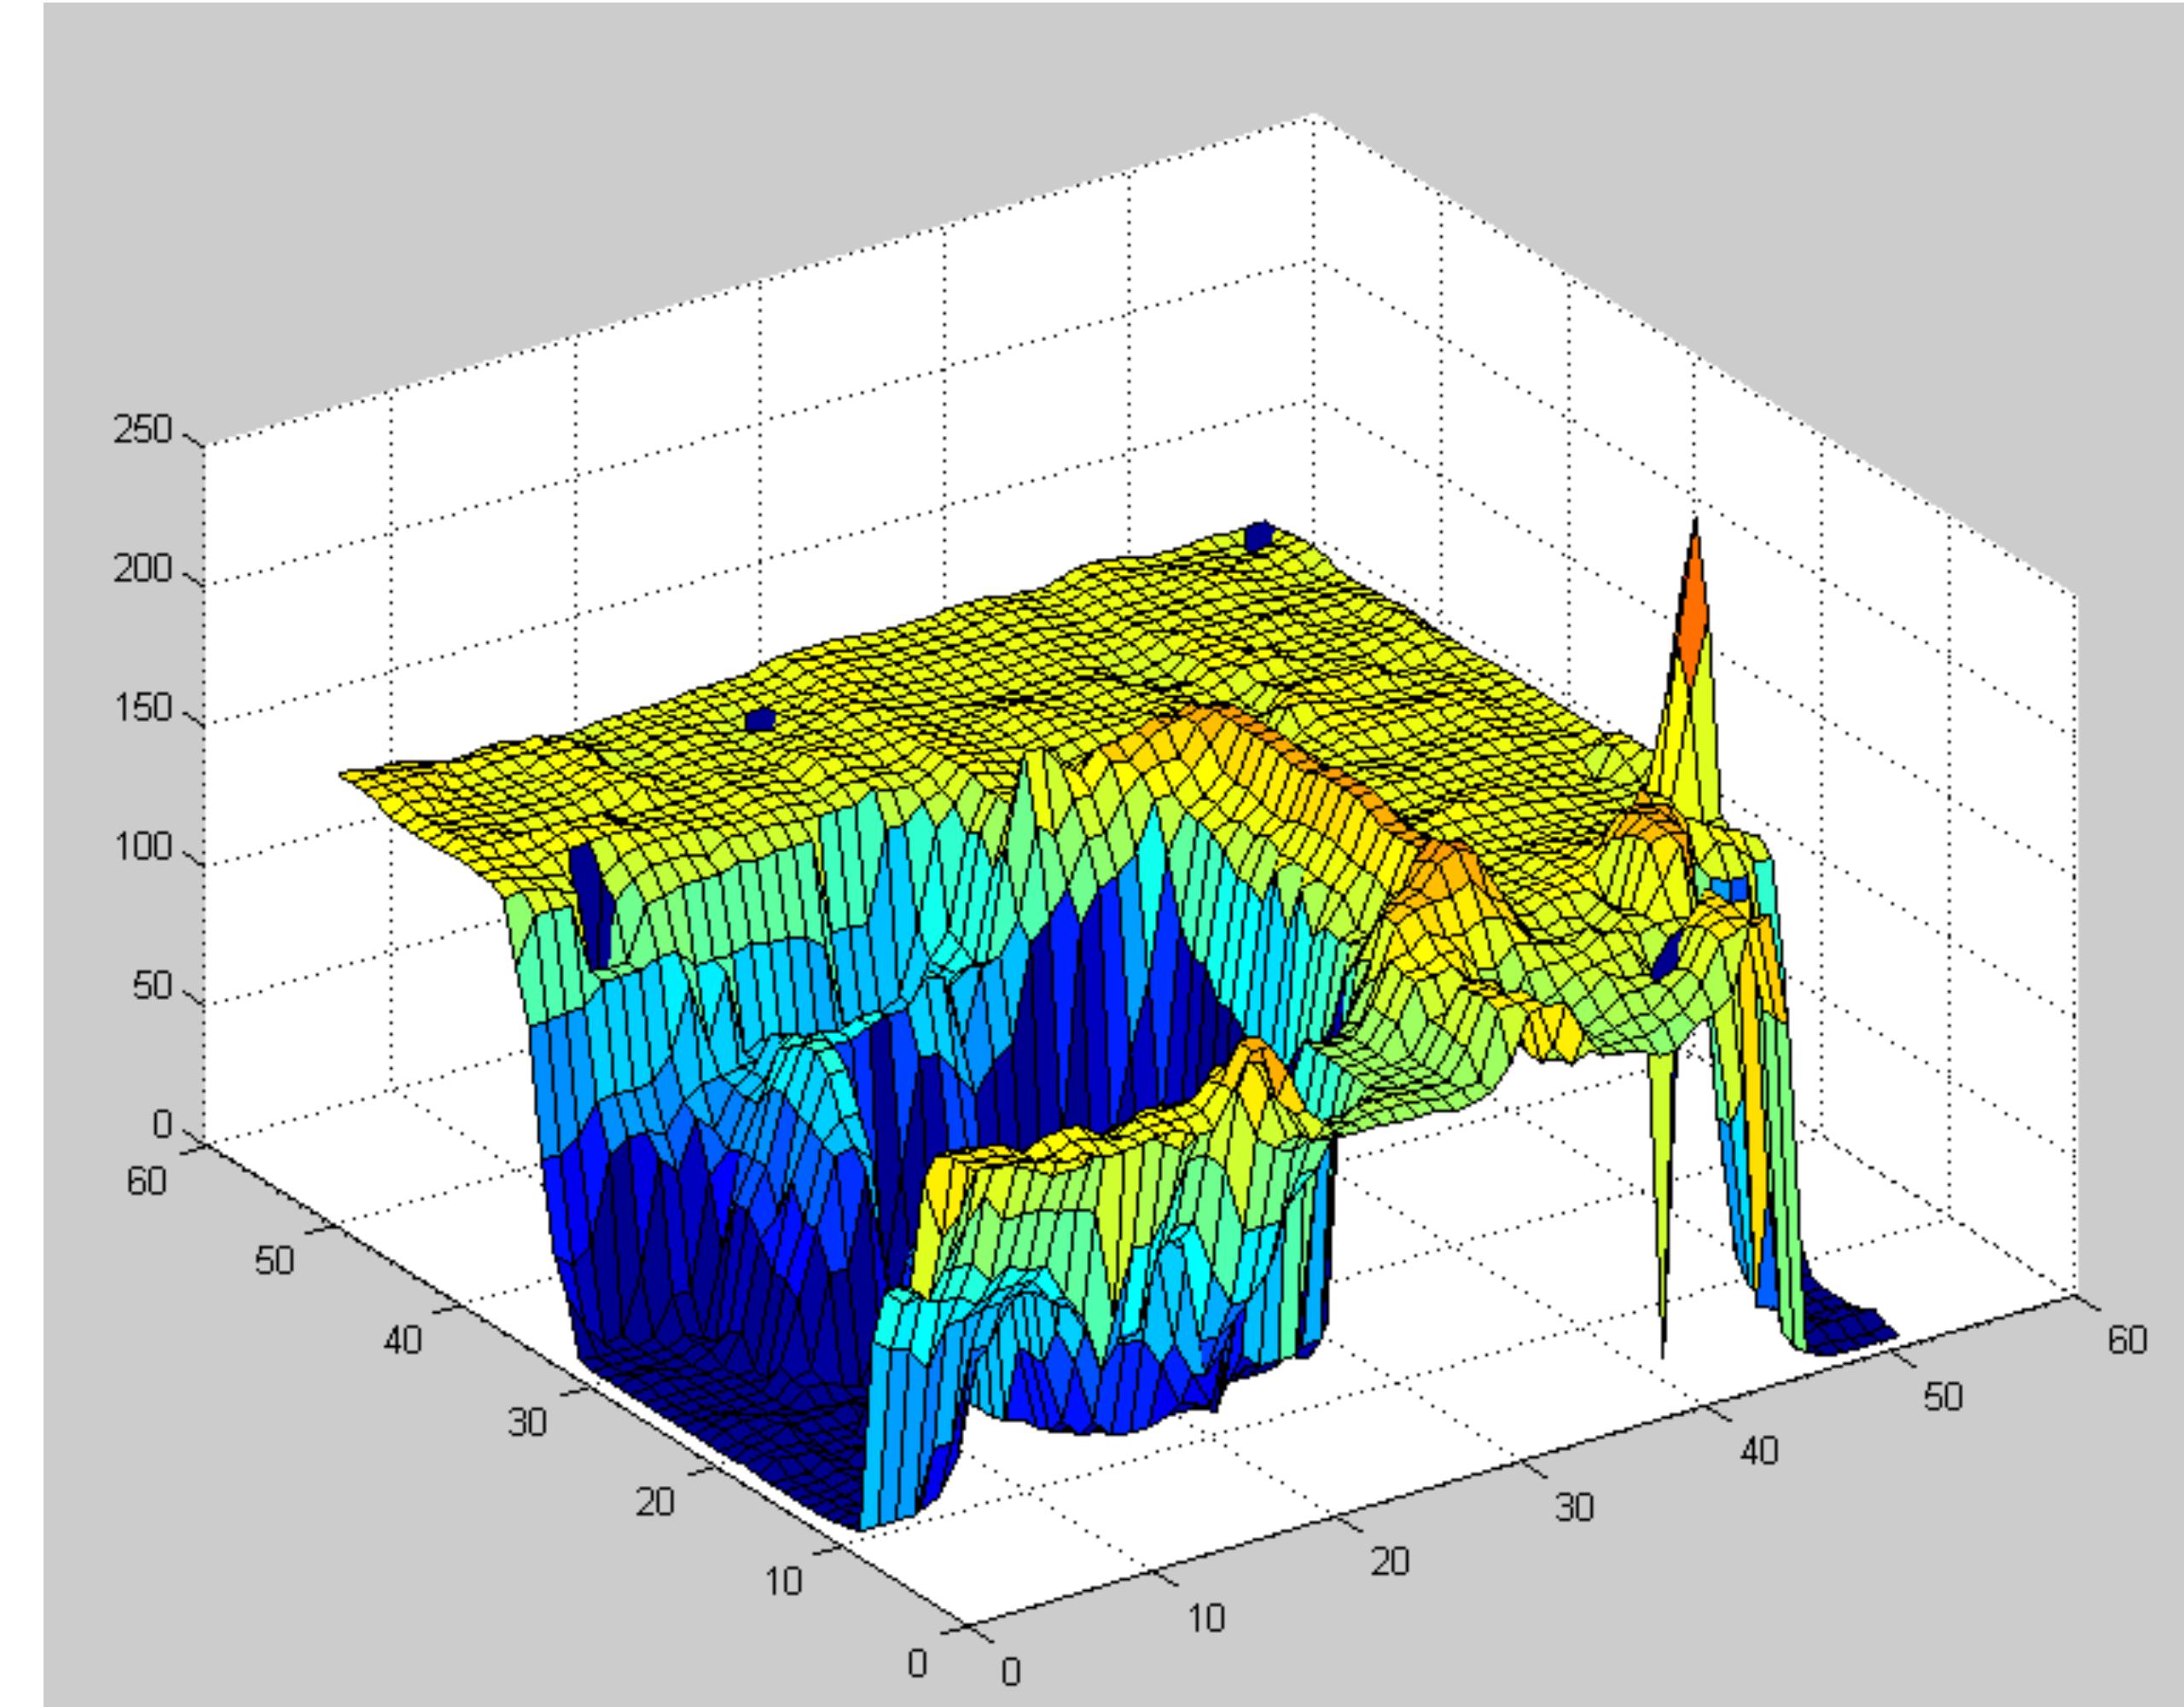


Image ~ Heightfield



$z = I$'s intensity



Filter: Let Intensity Through Selectively

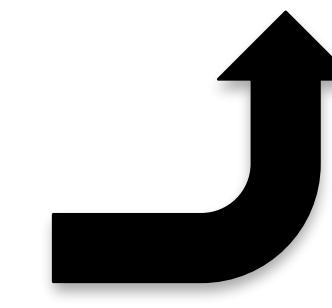
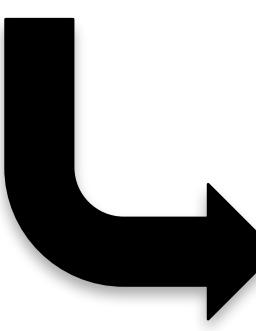
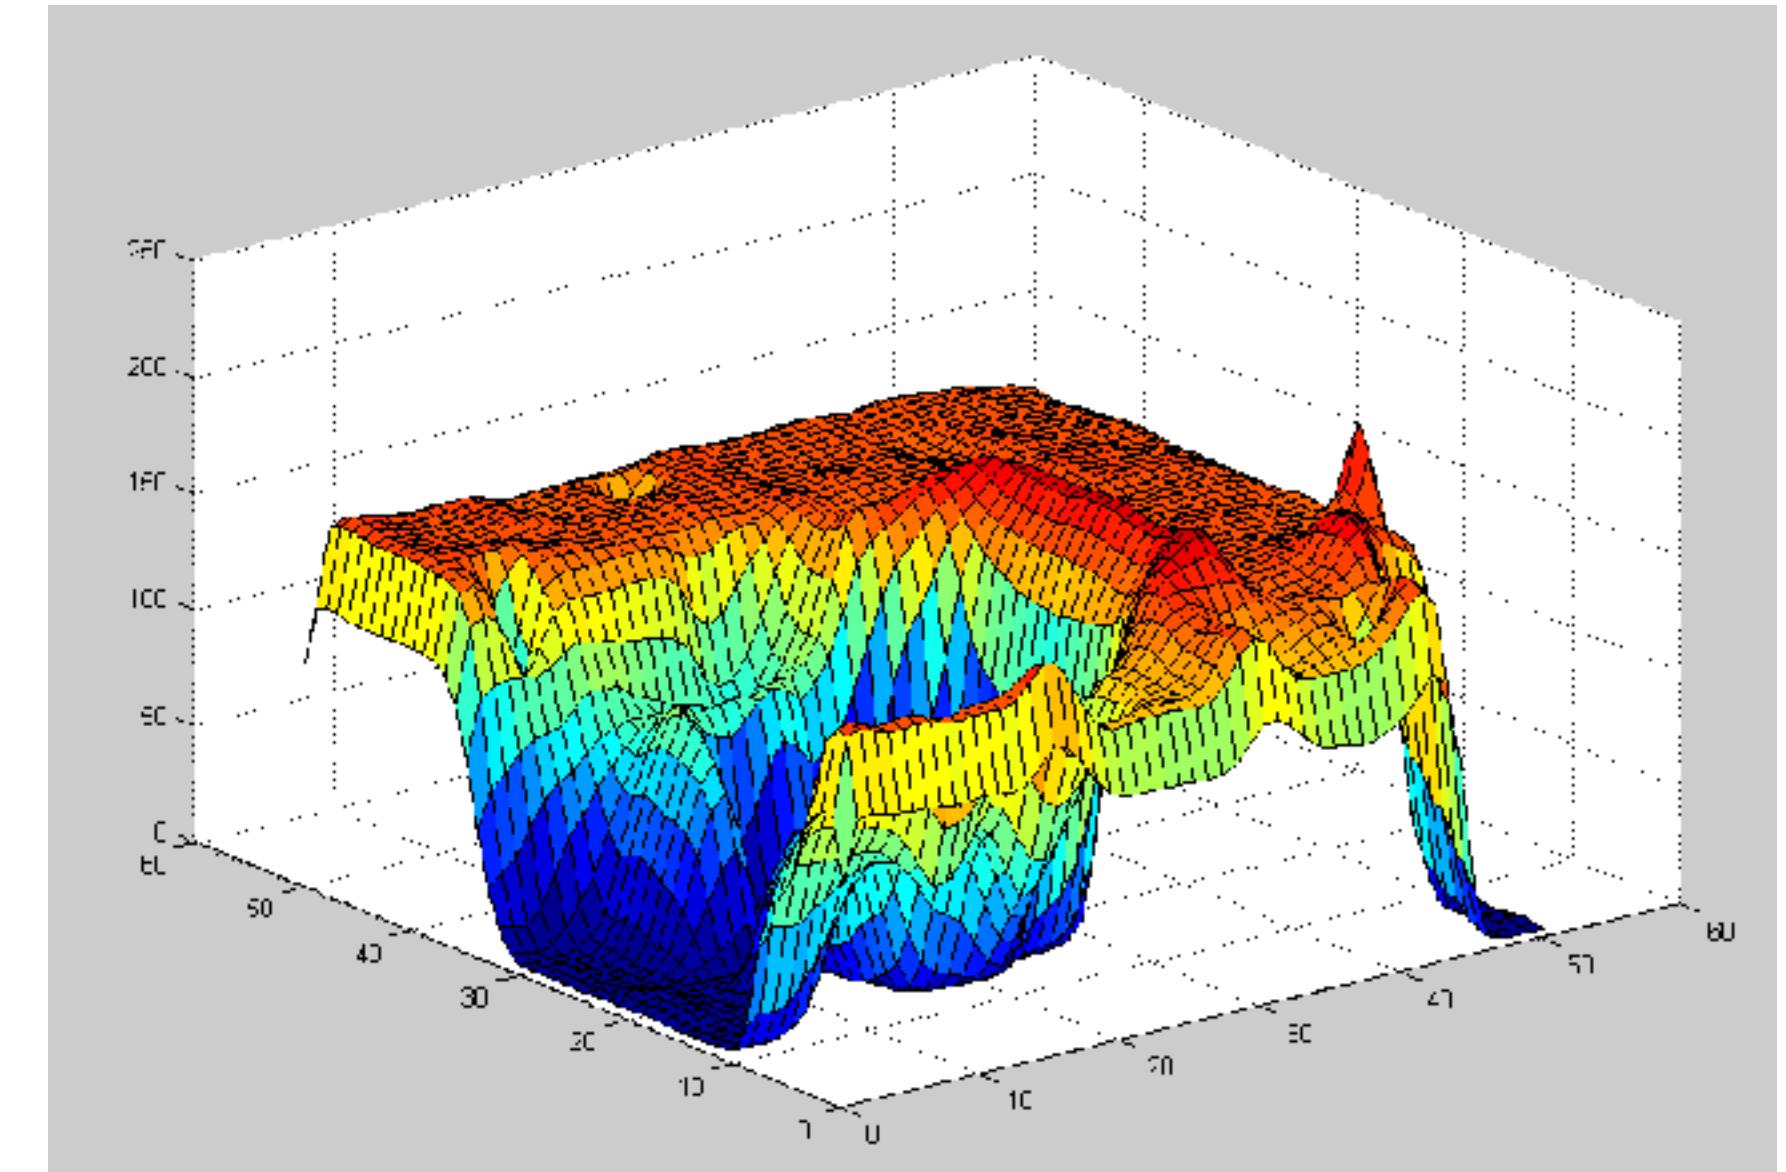
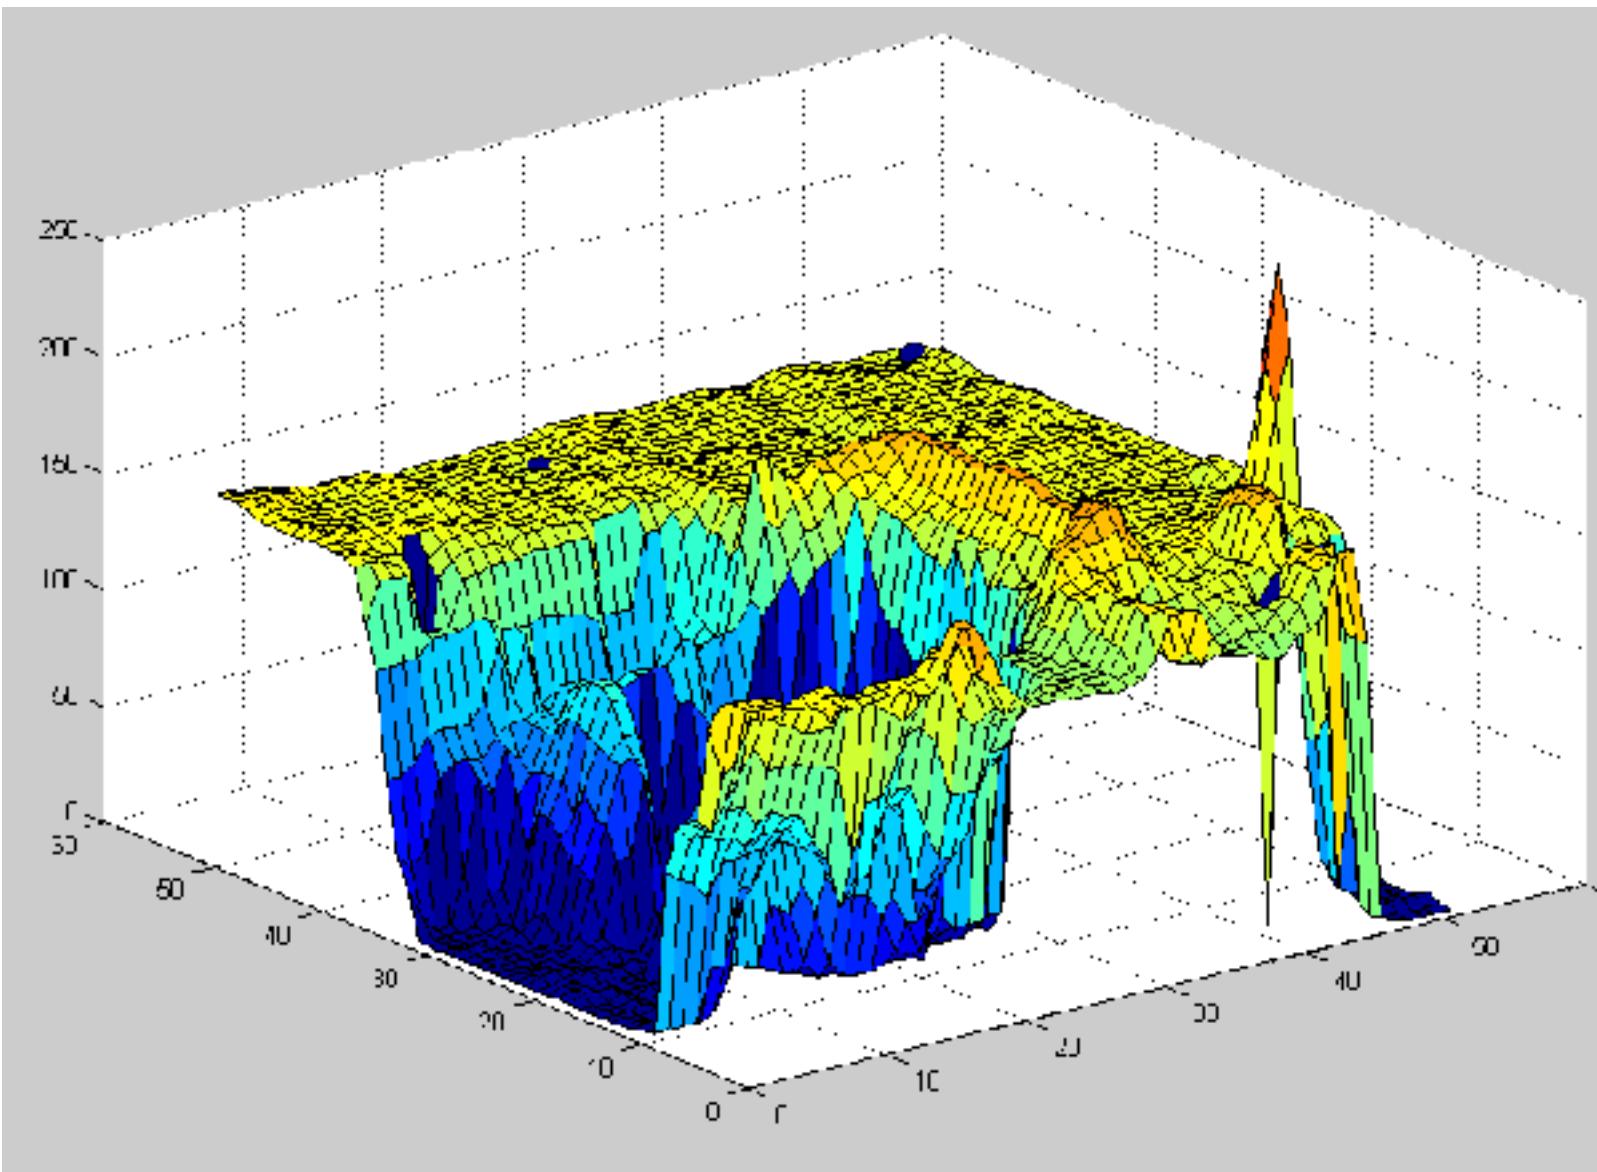
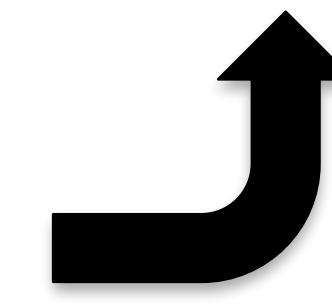
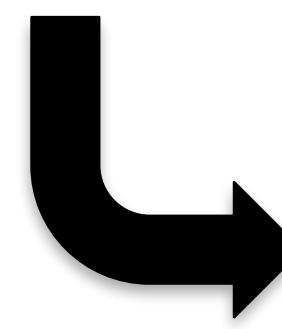
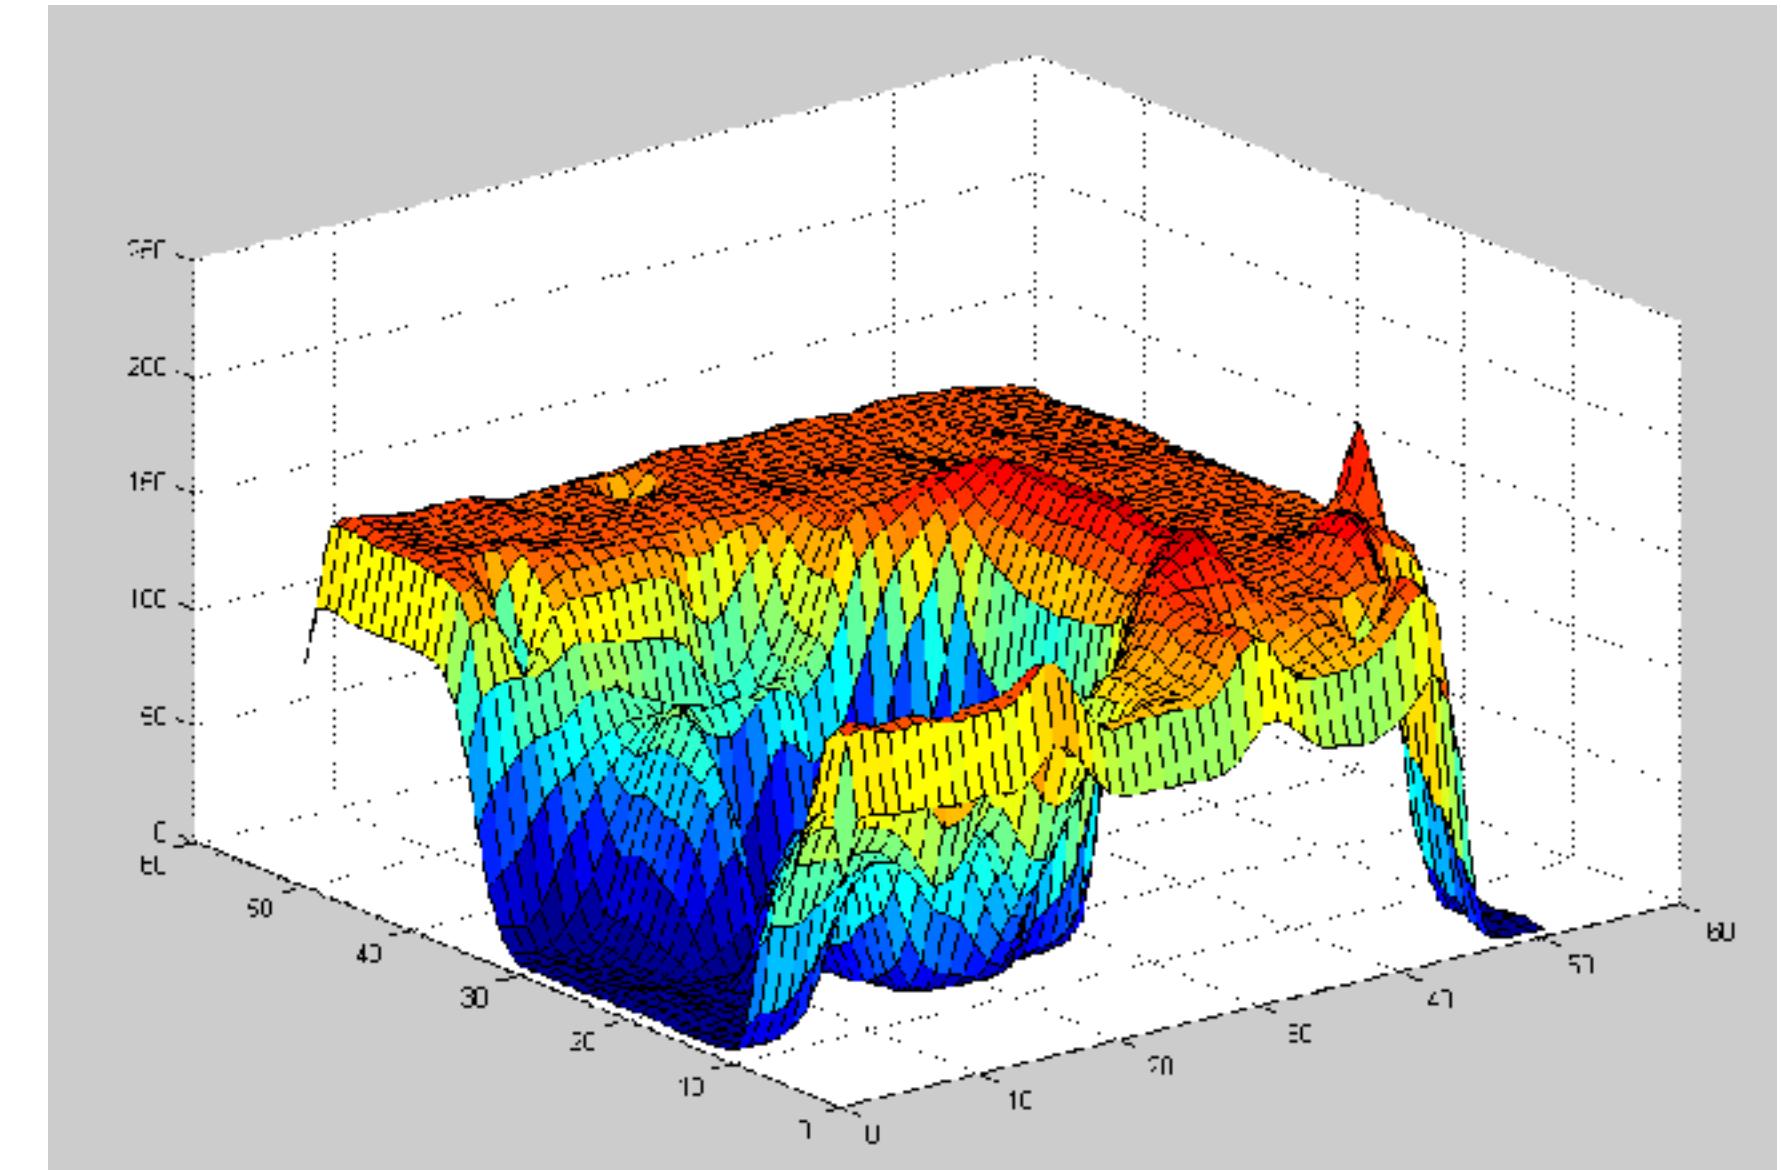
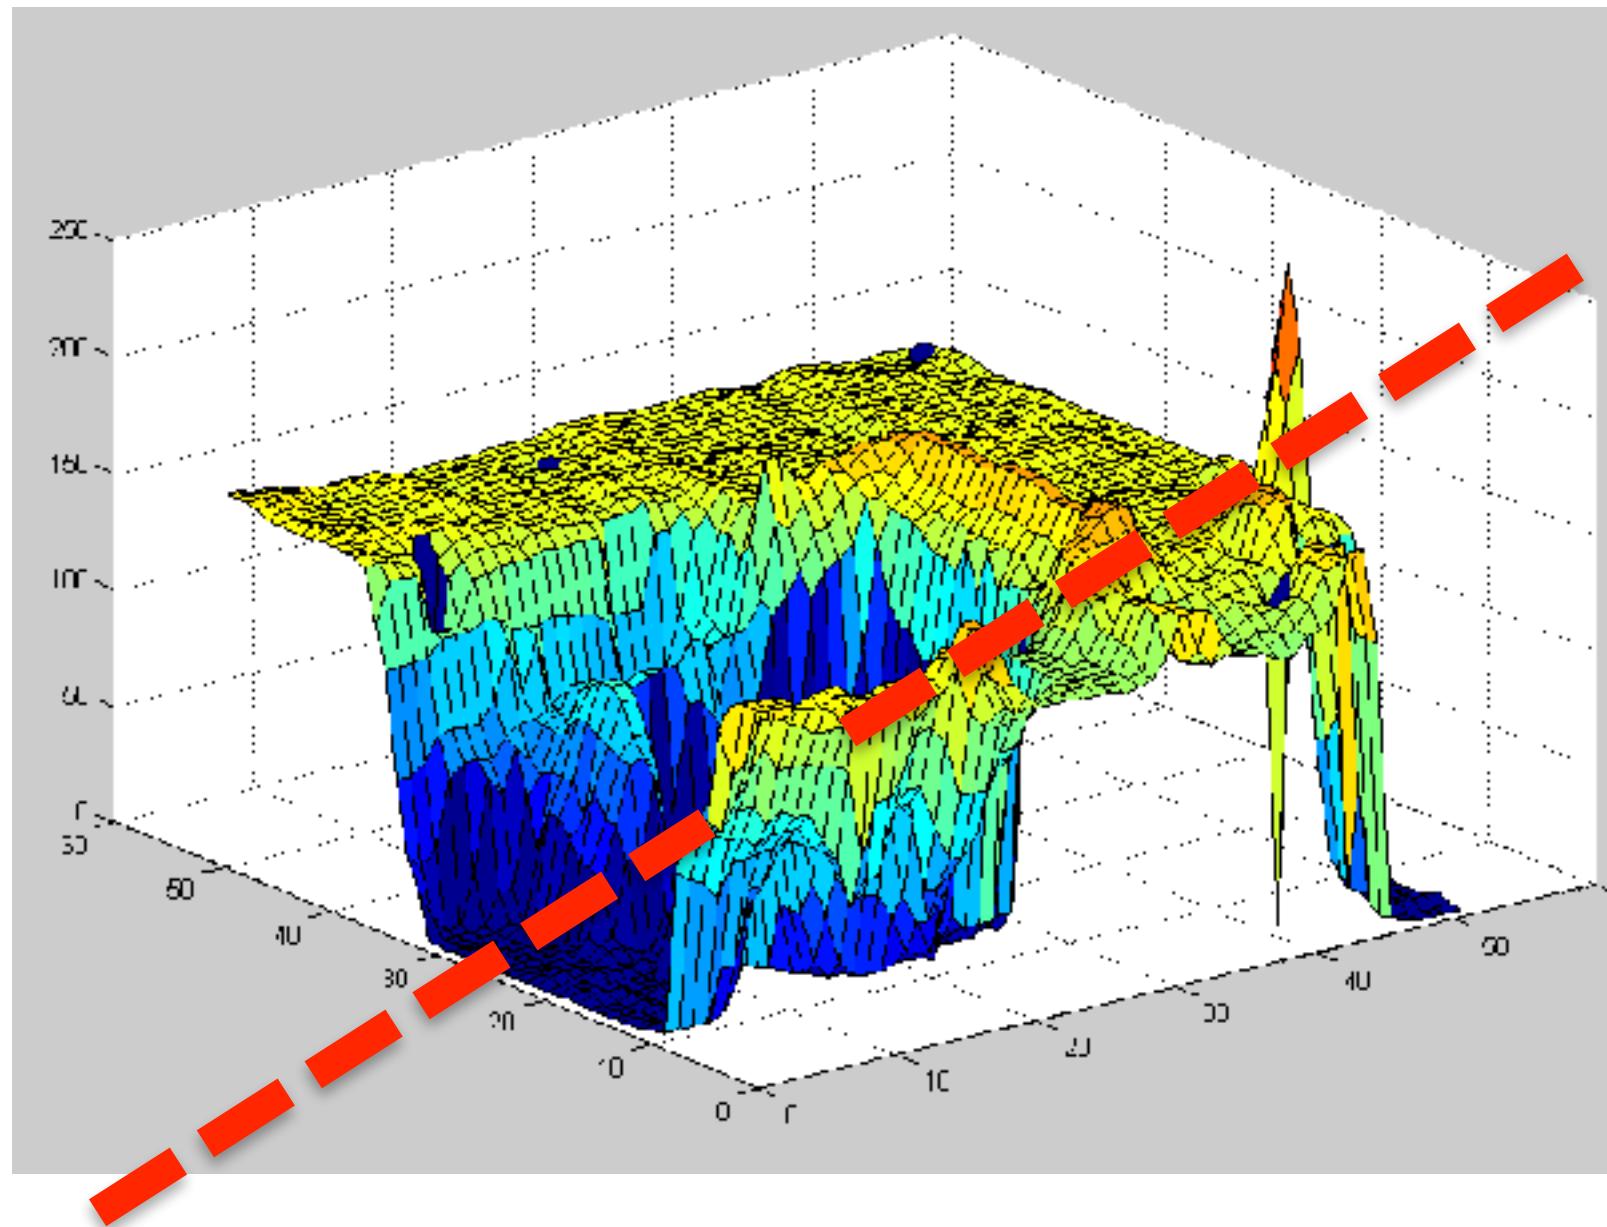


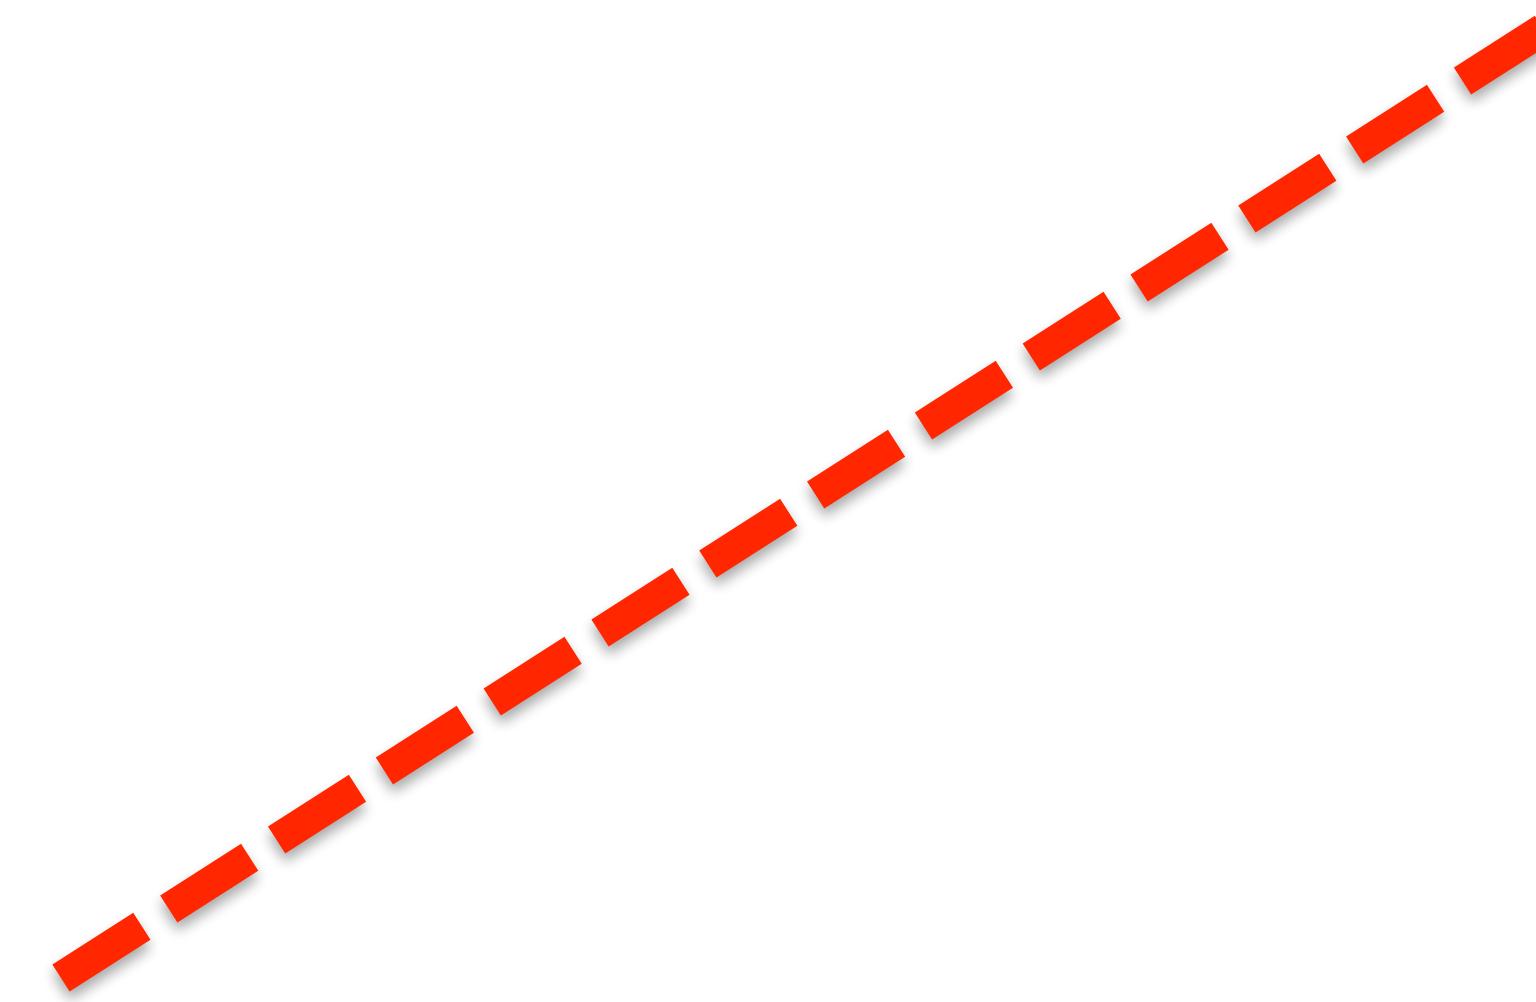
Image Filtering

let the intensity through
conditionally

Filter: Let Intensity Through Selectively



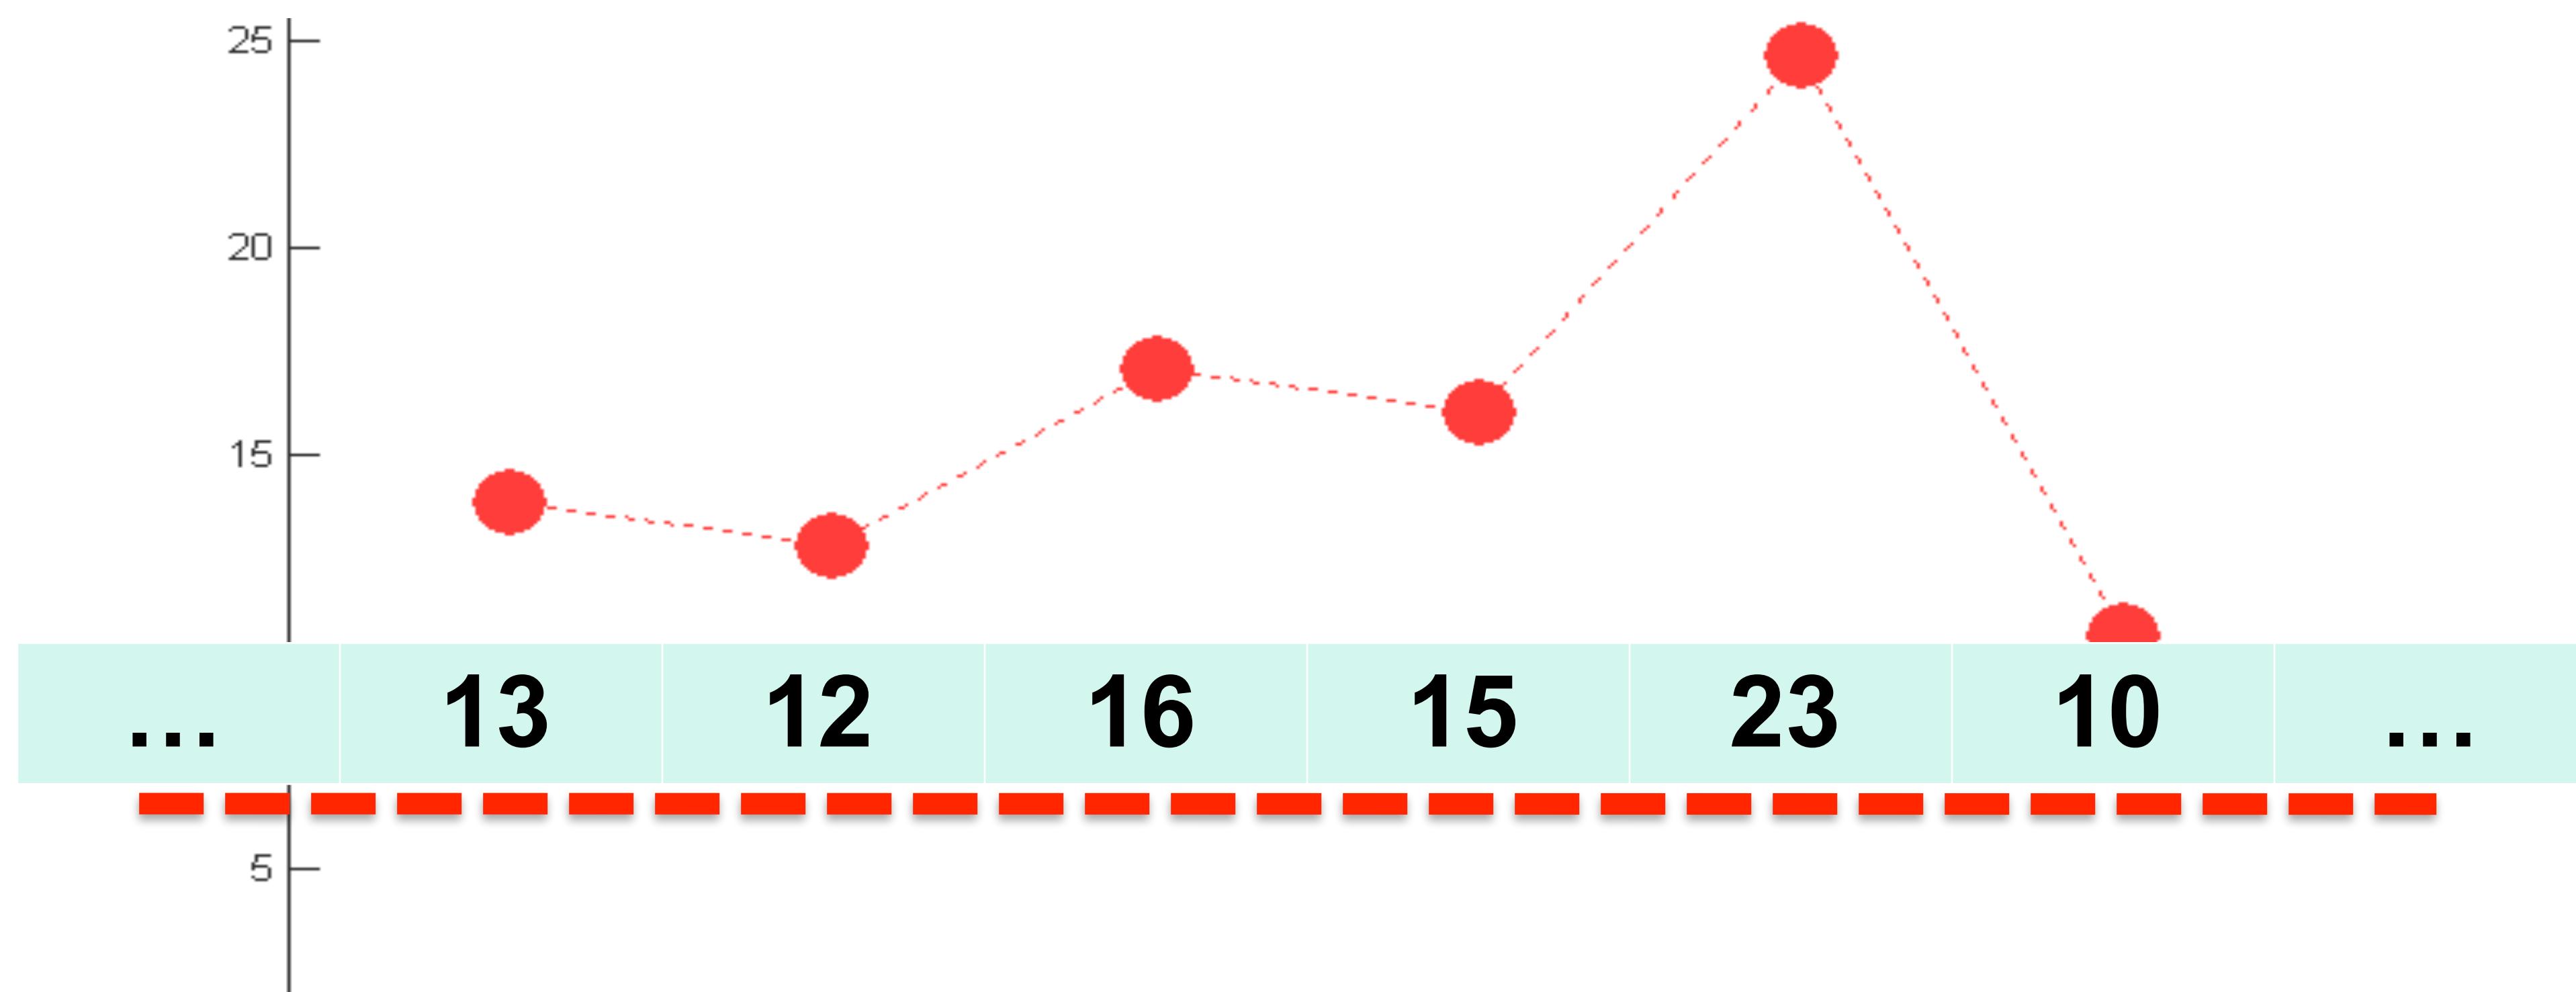
Filter: Let Intensity Through Selectively



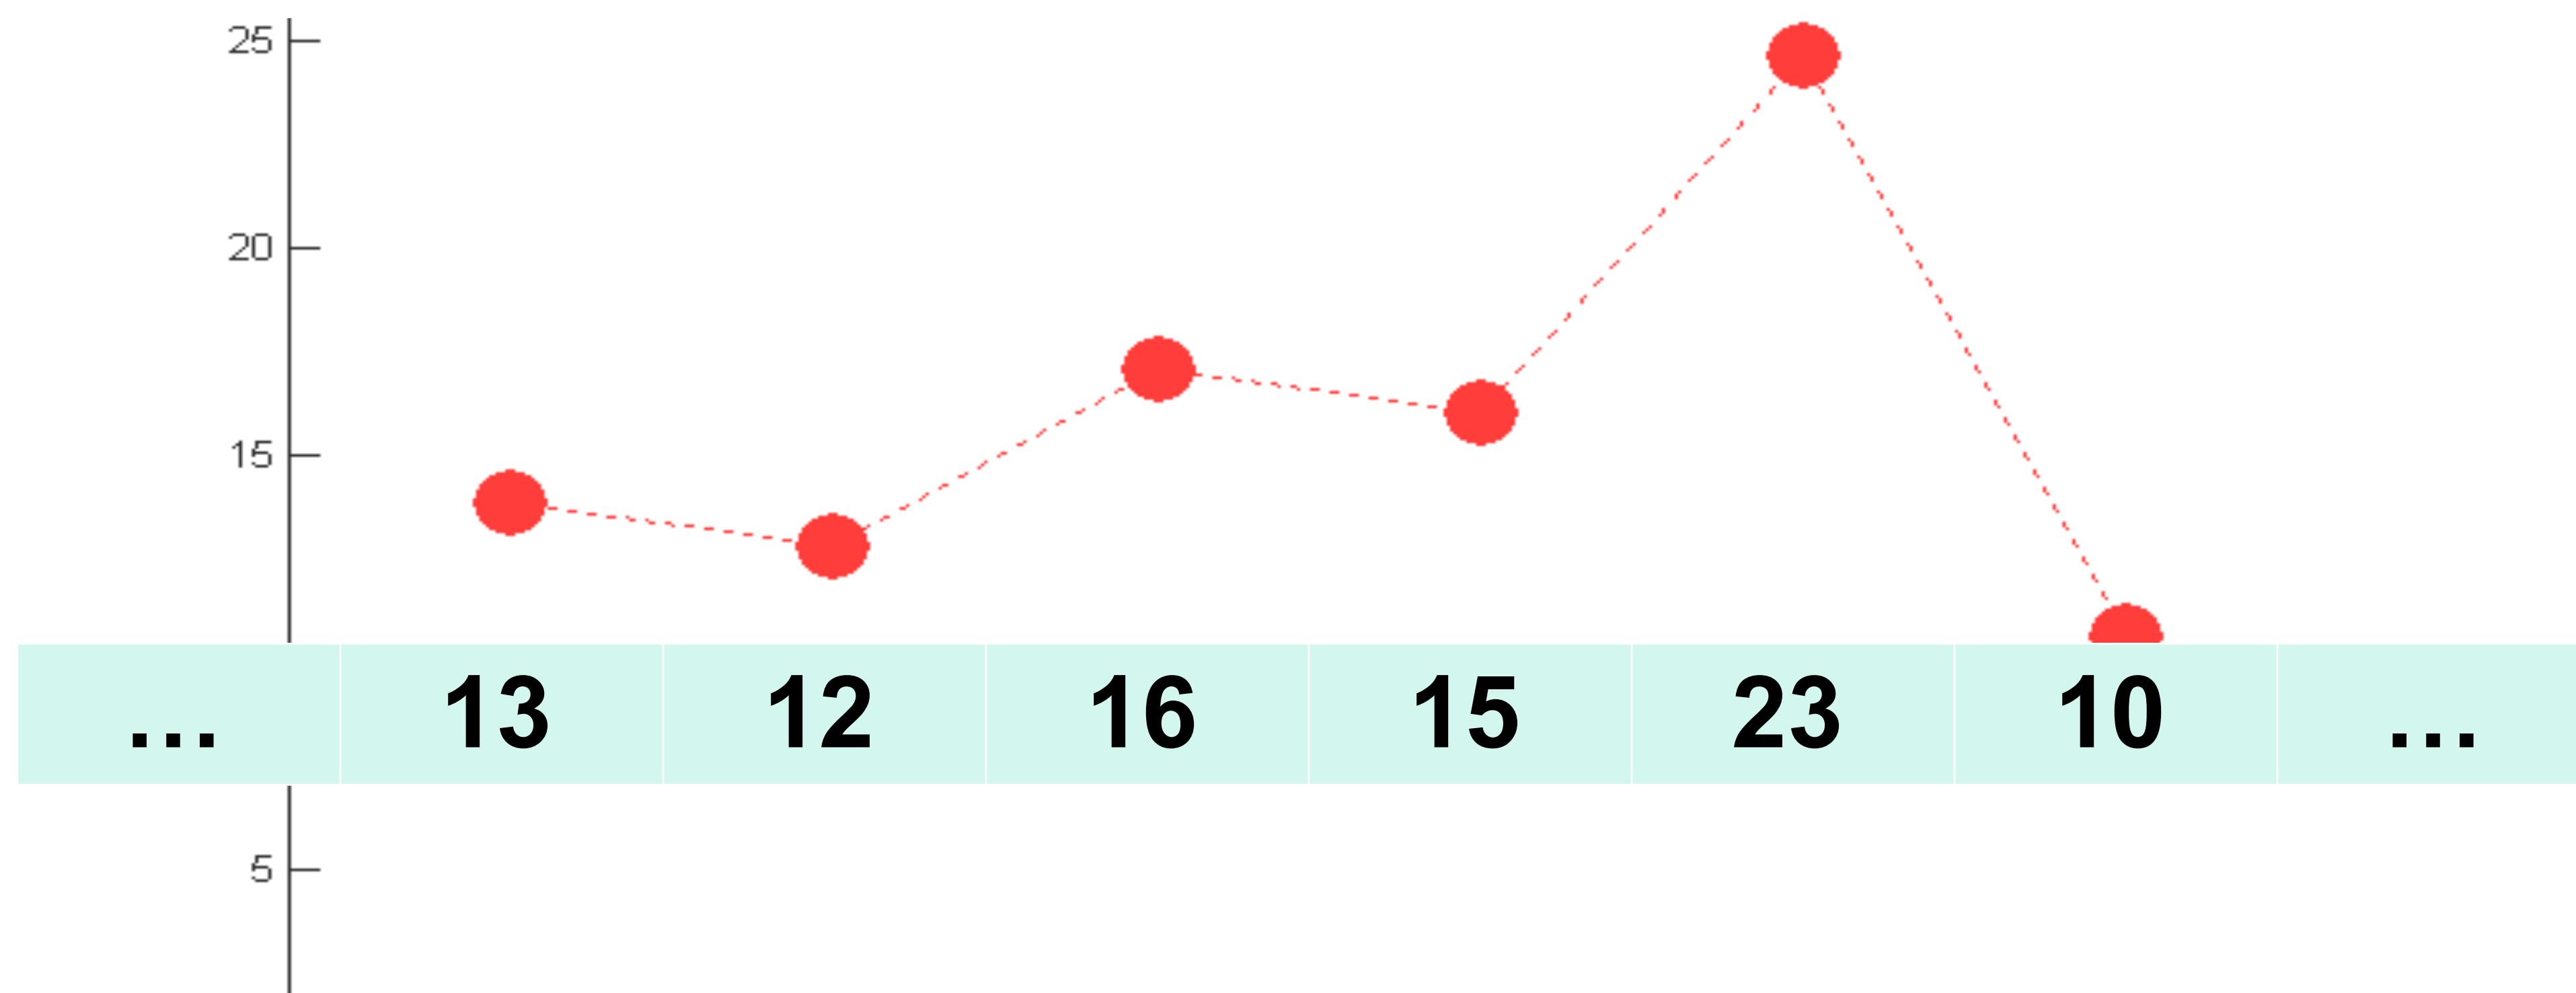
Filter: Let Intensity Through Selectively



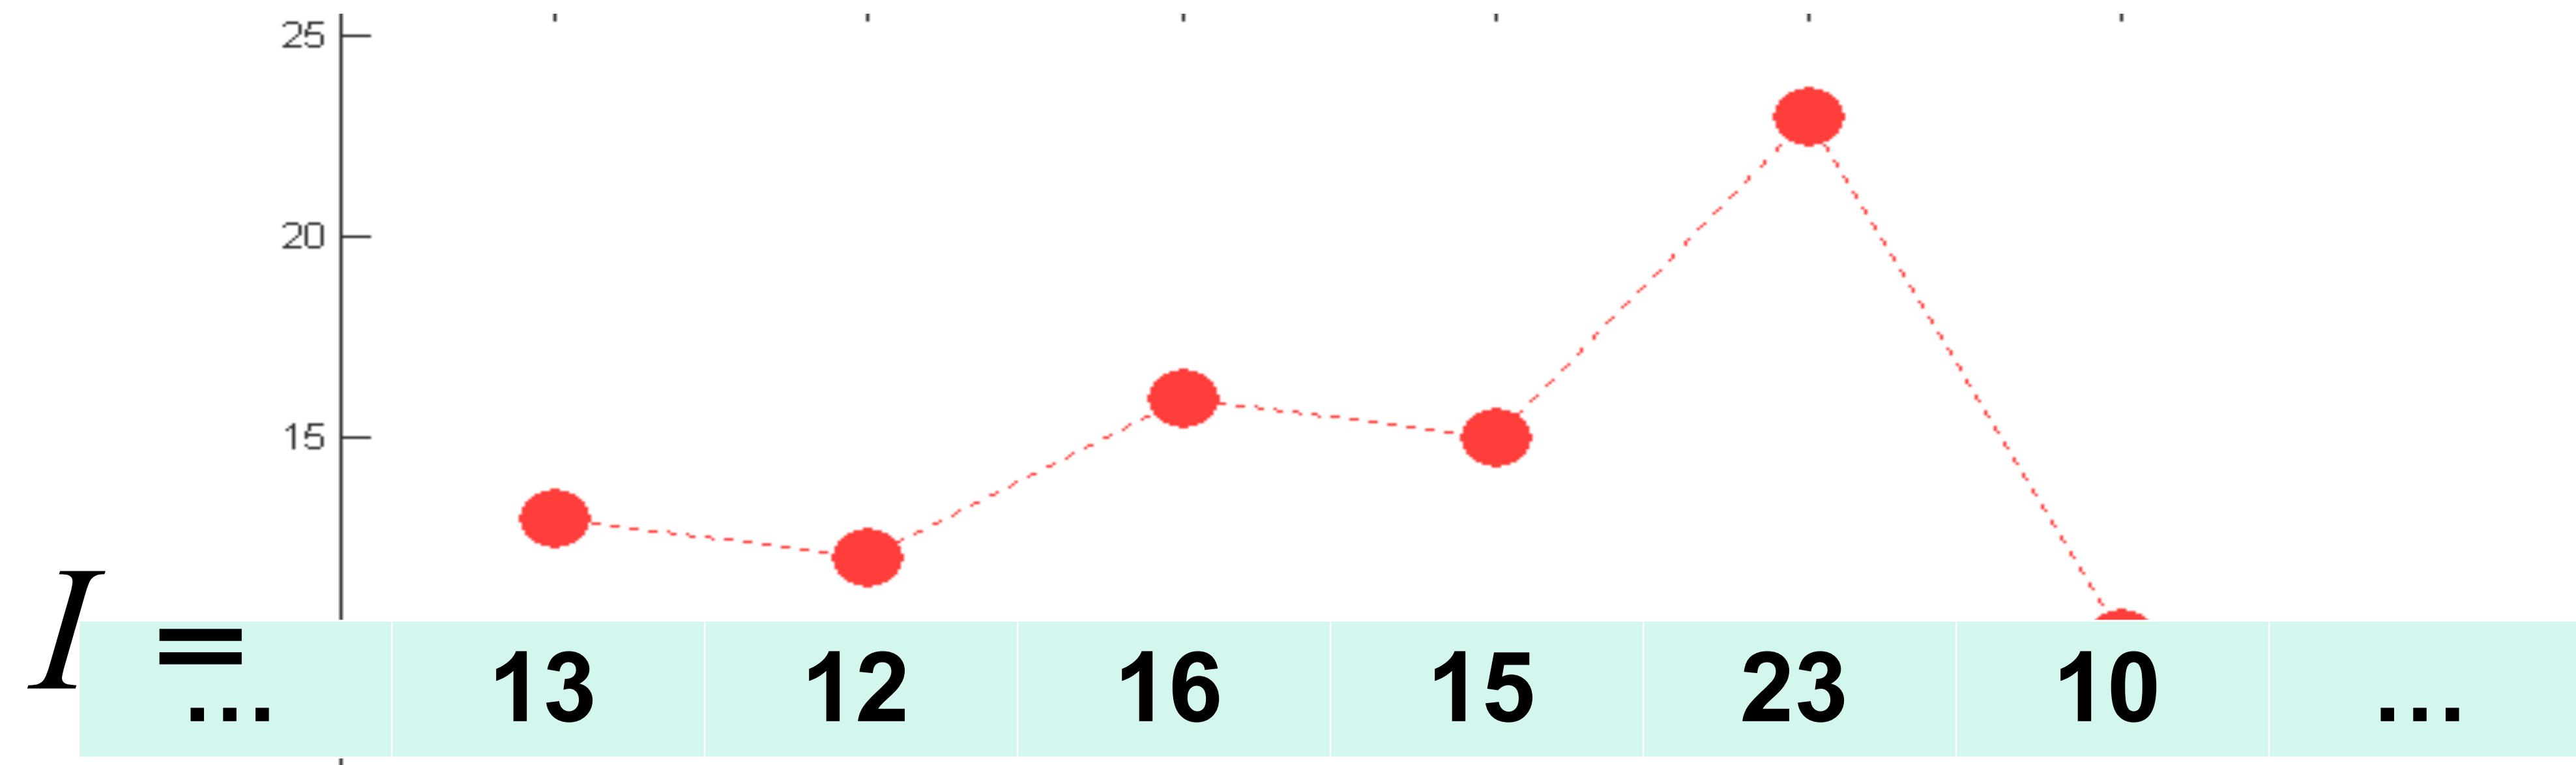
A Row of Pixel Intensities



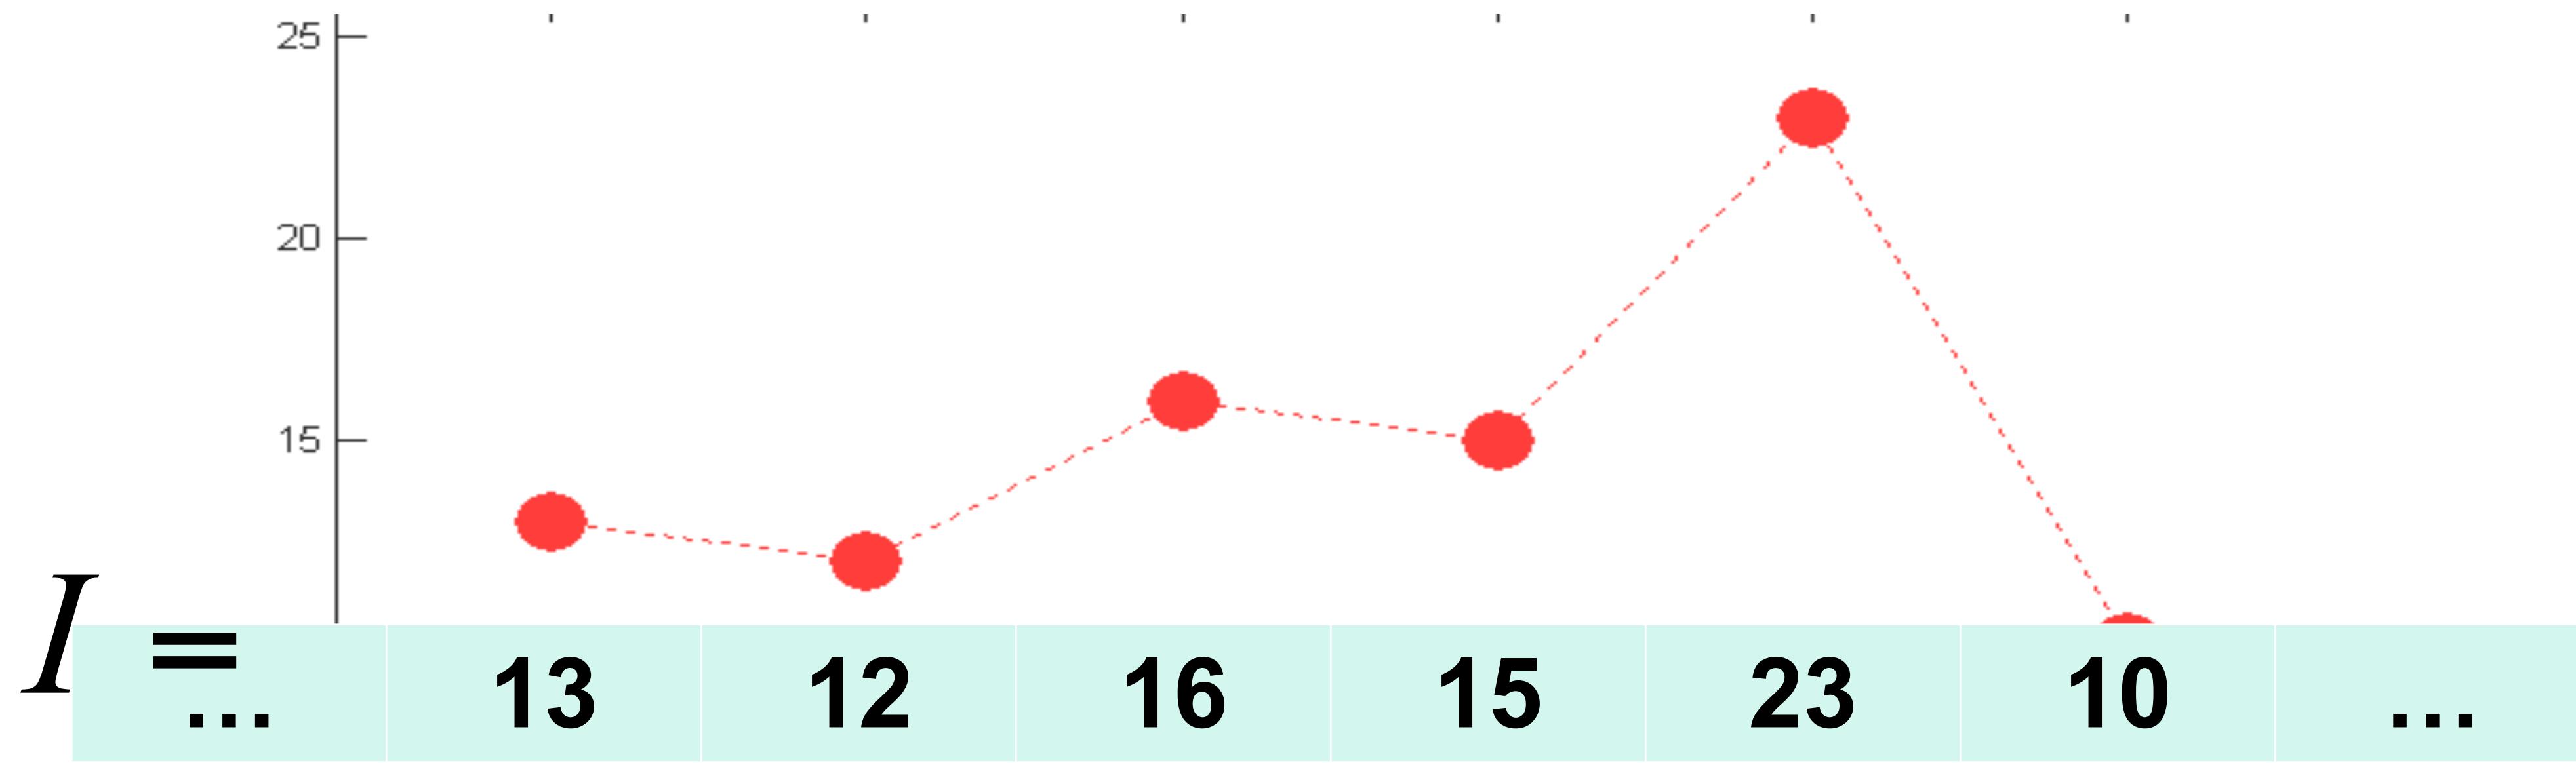
A Row of Pixel Intensities



Linear Operations: Weighted Sum

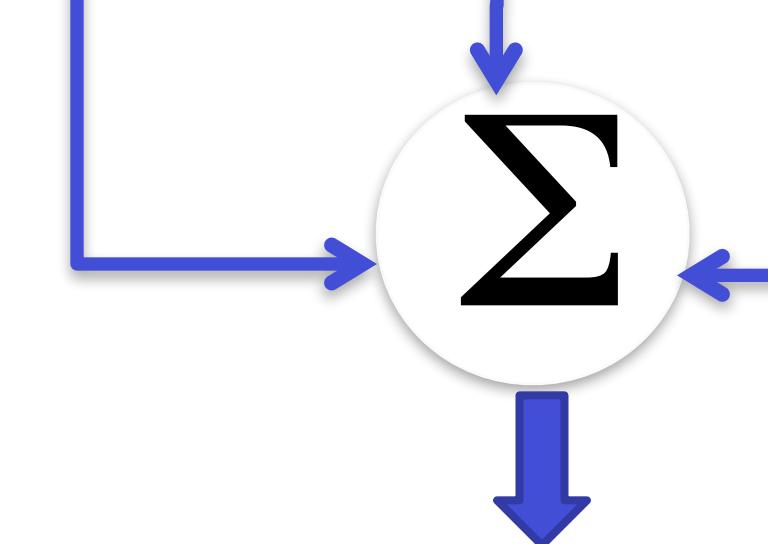


$$I' = \dots \quad ?? \quad ?? \quad ?? \quad ?? \quad \dots$$



$K =$

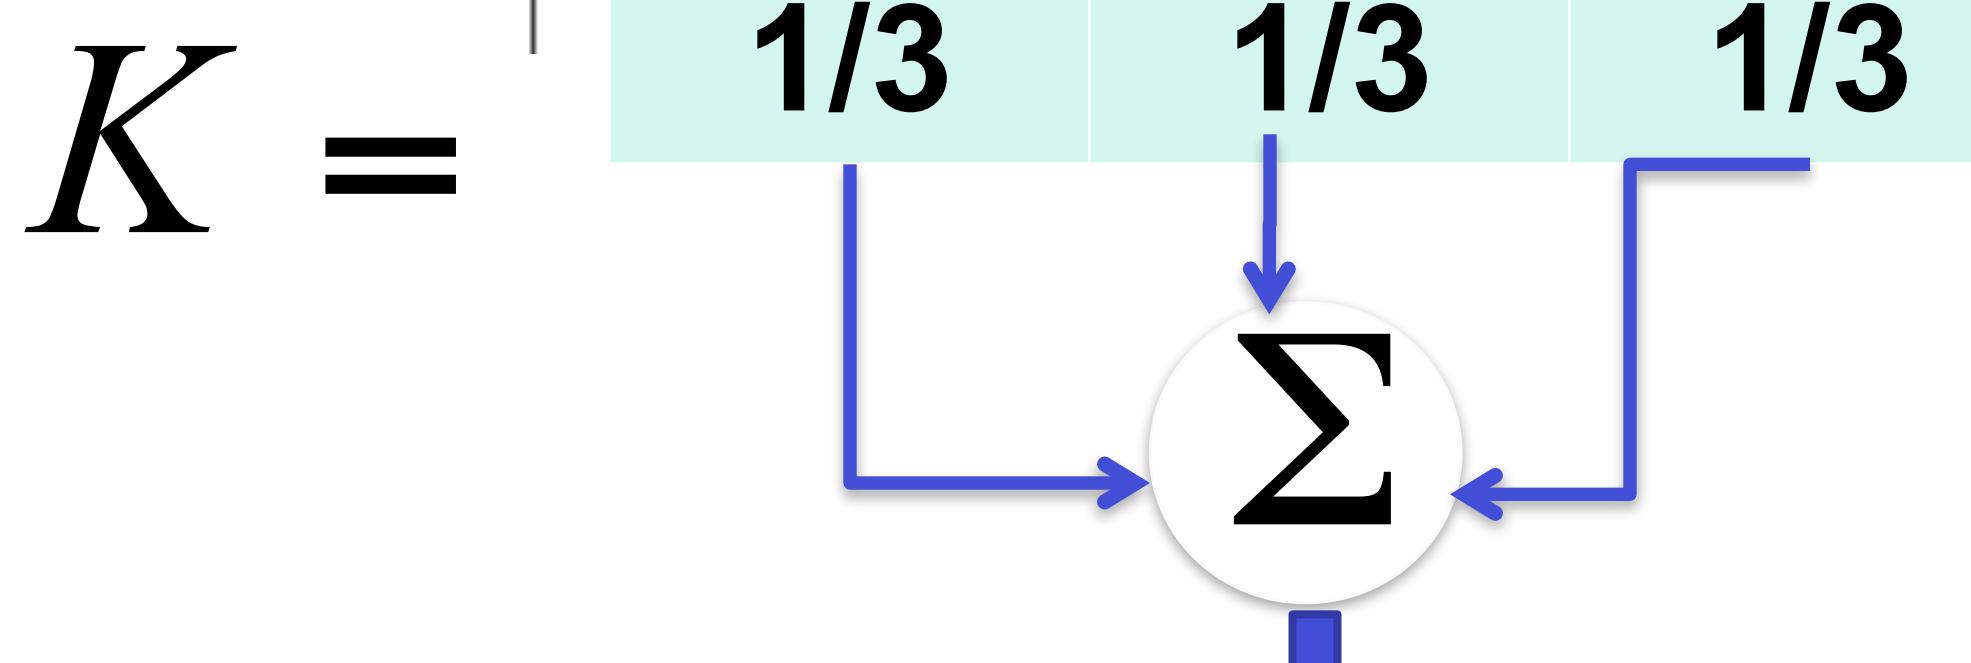
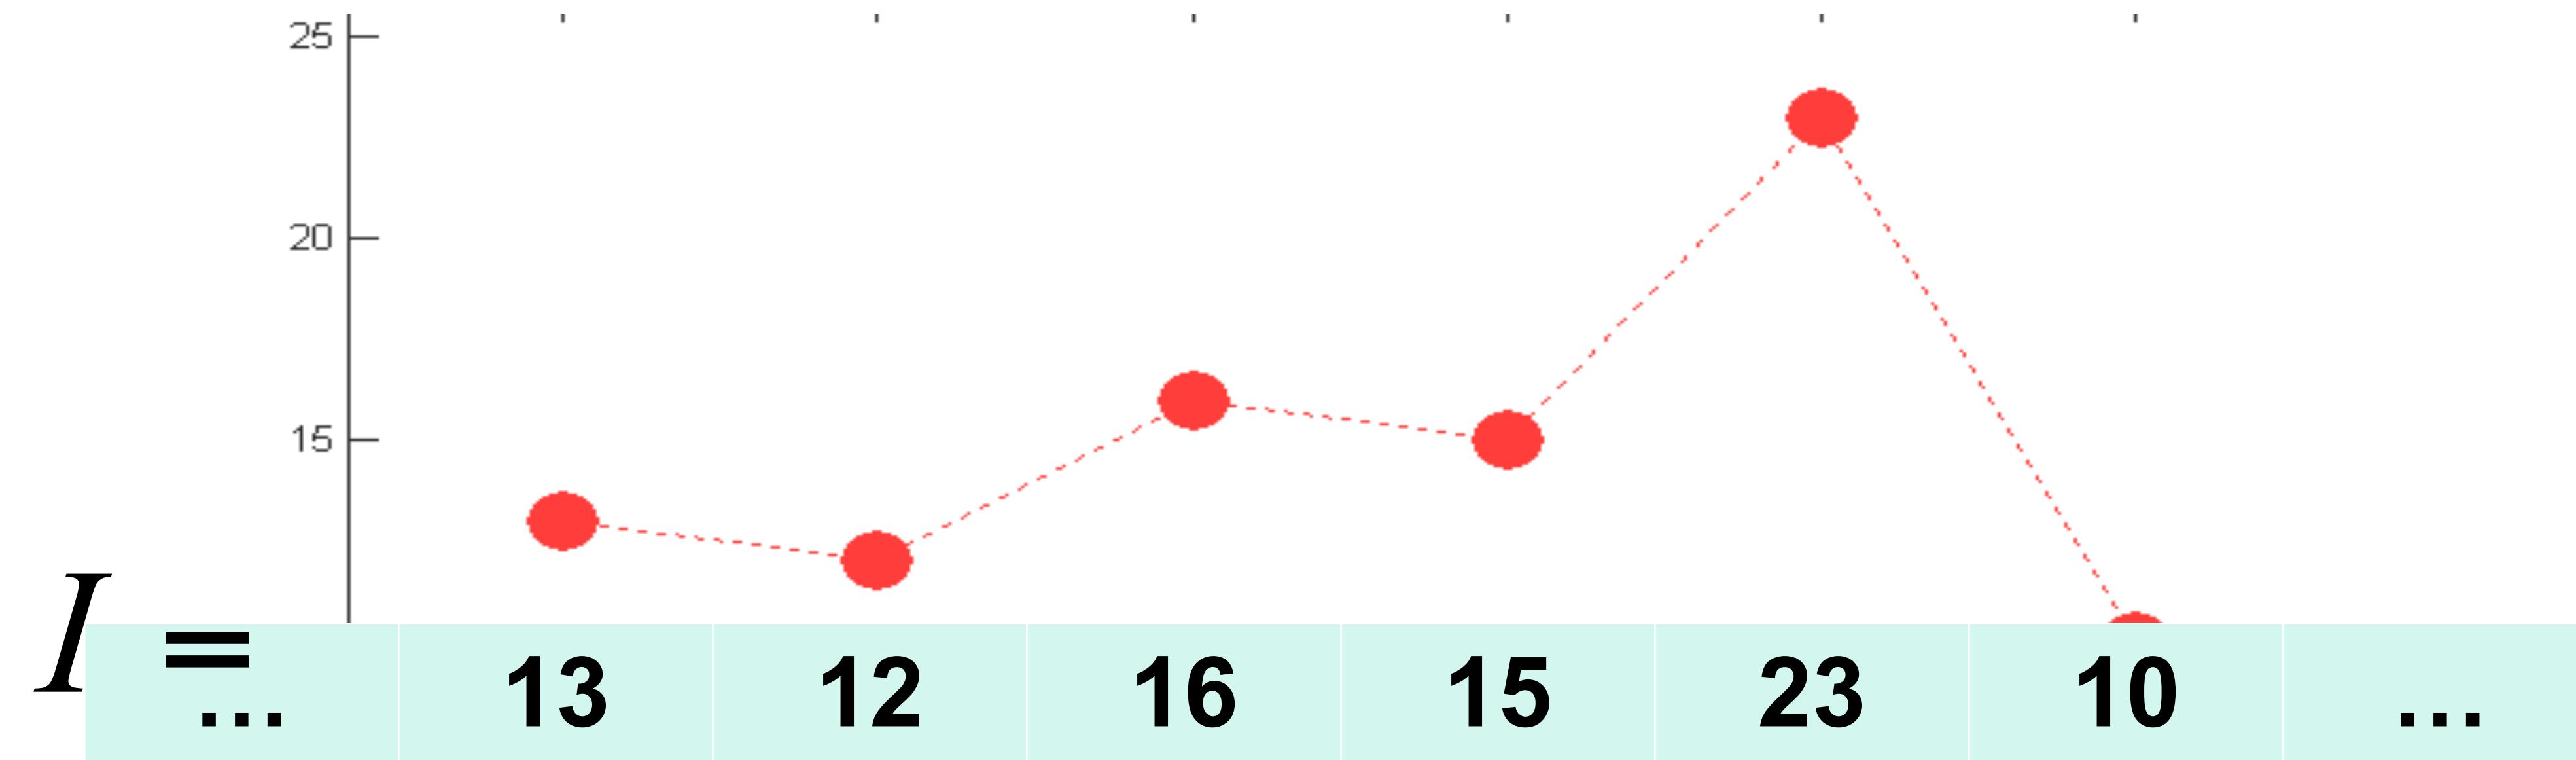
$1/3$	$1/3$	$1/3$
-------	-------	-------

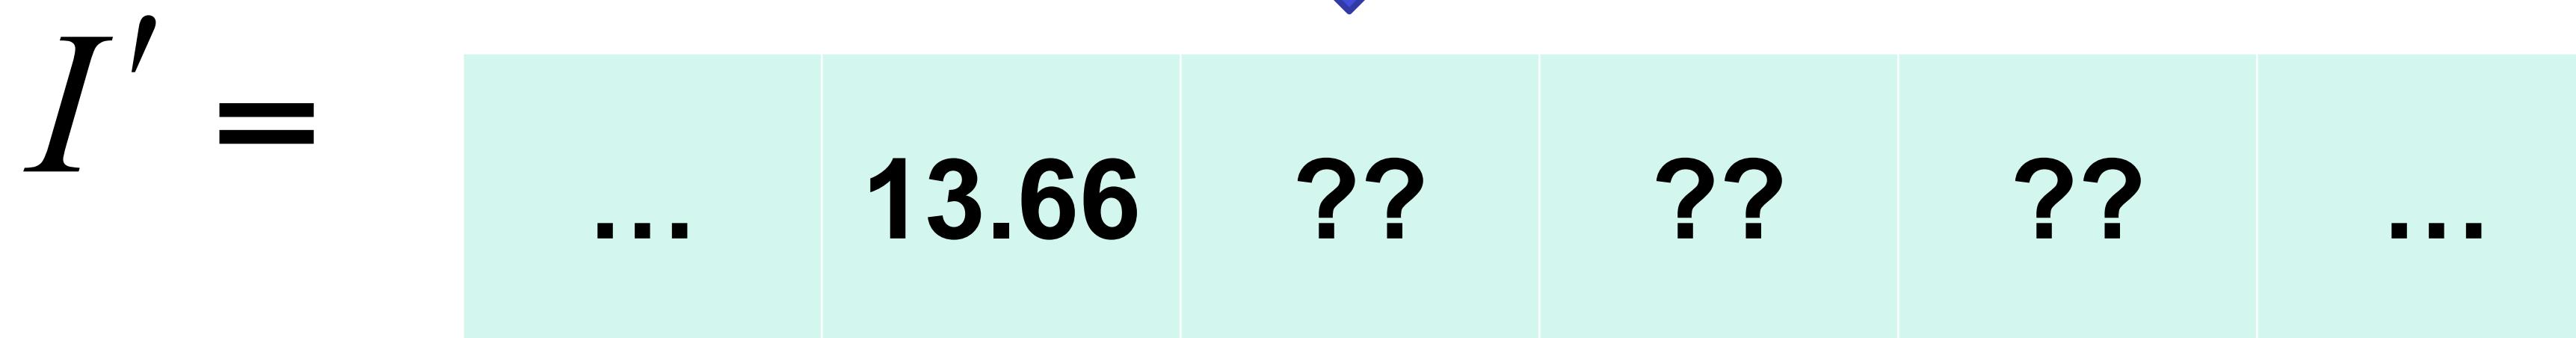
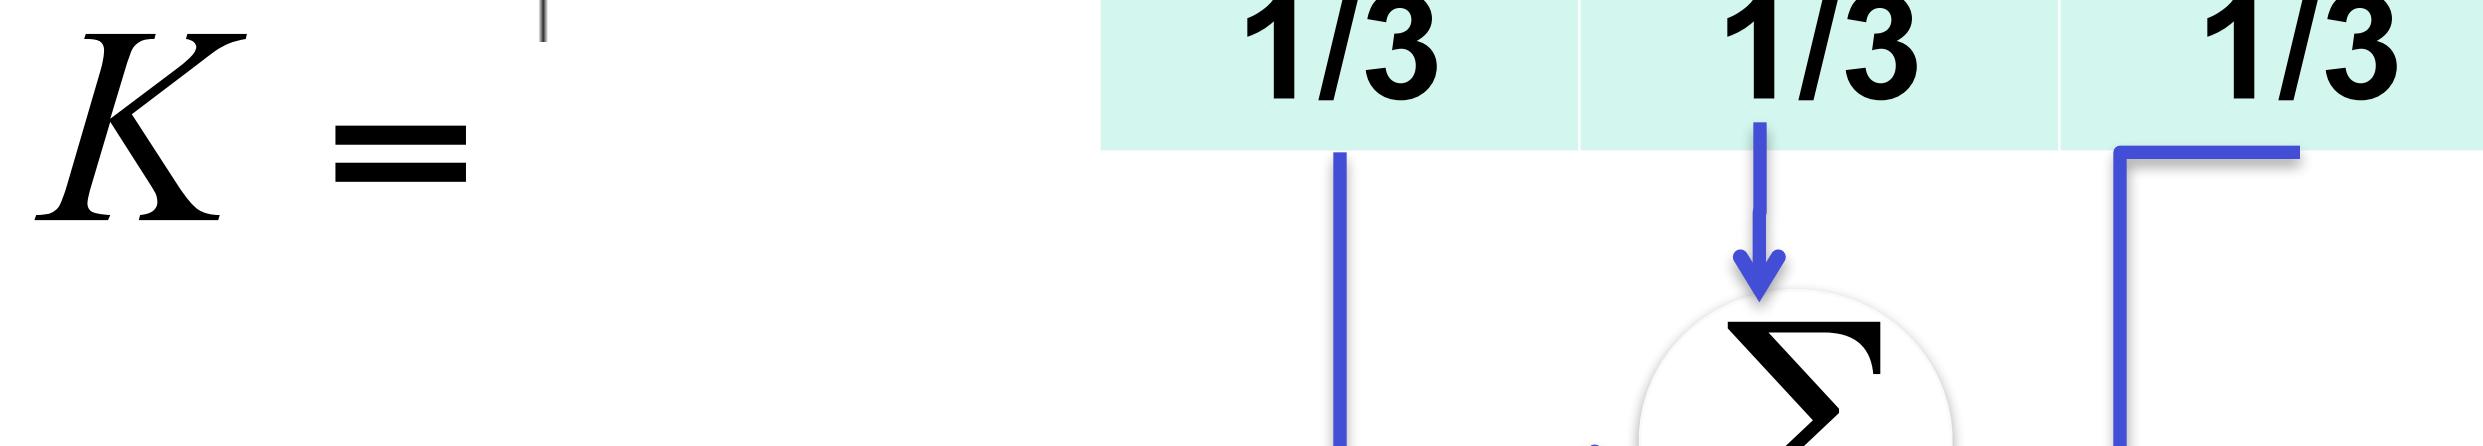
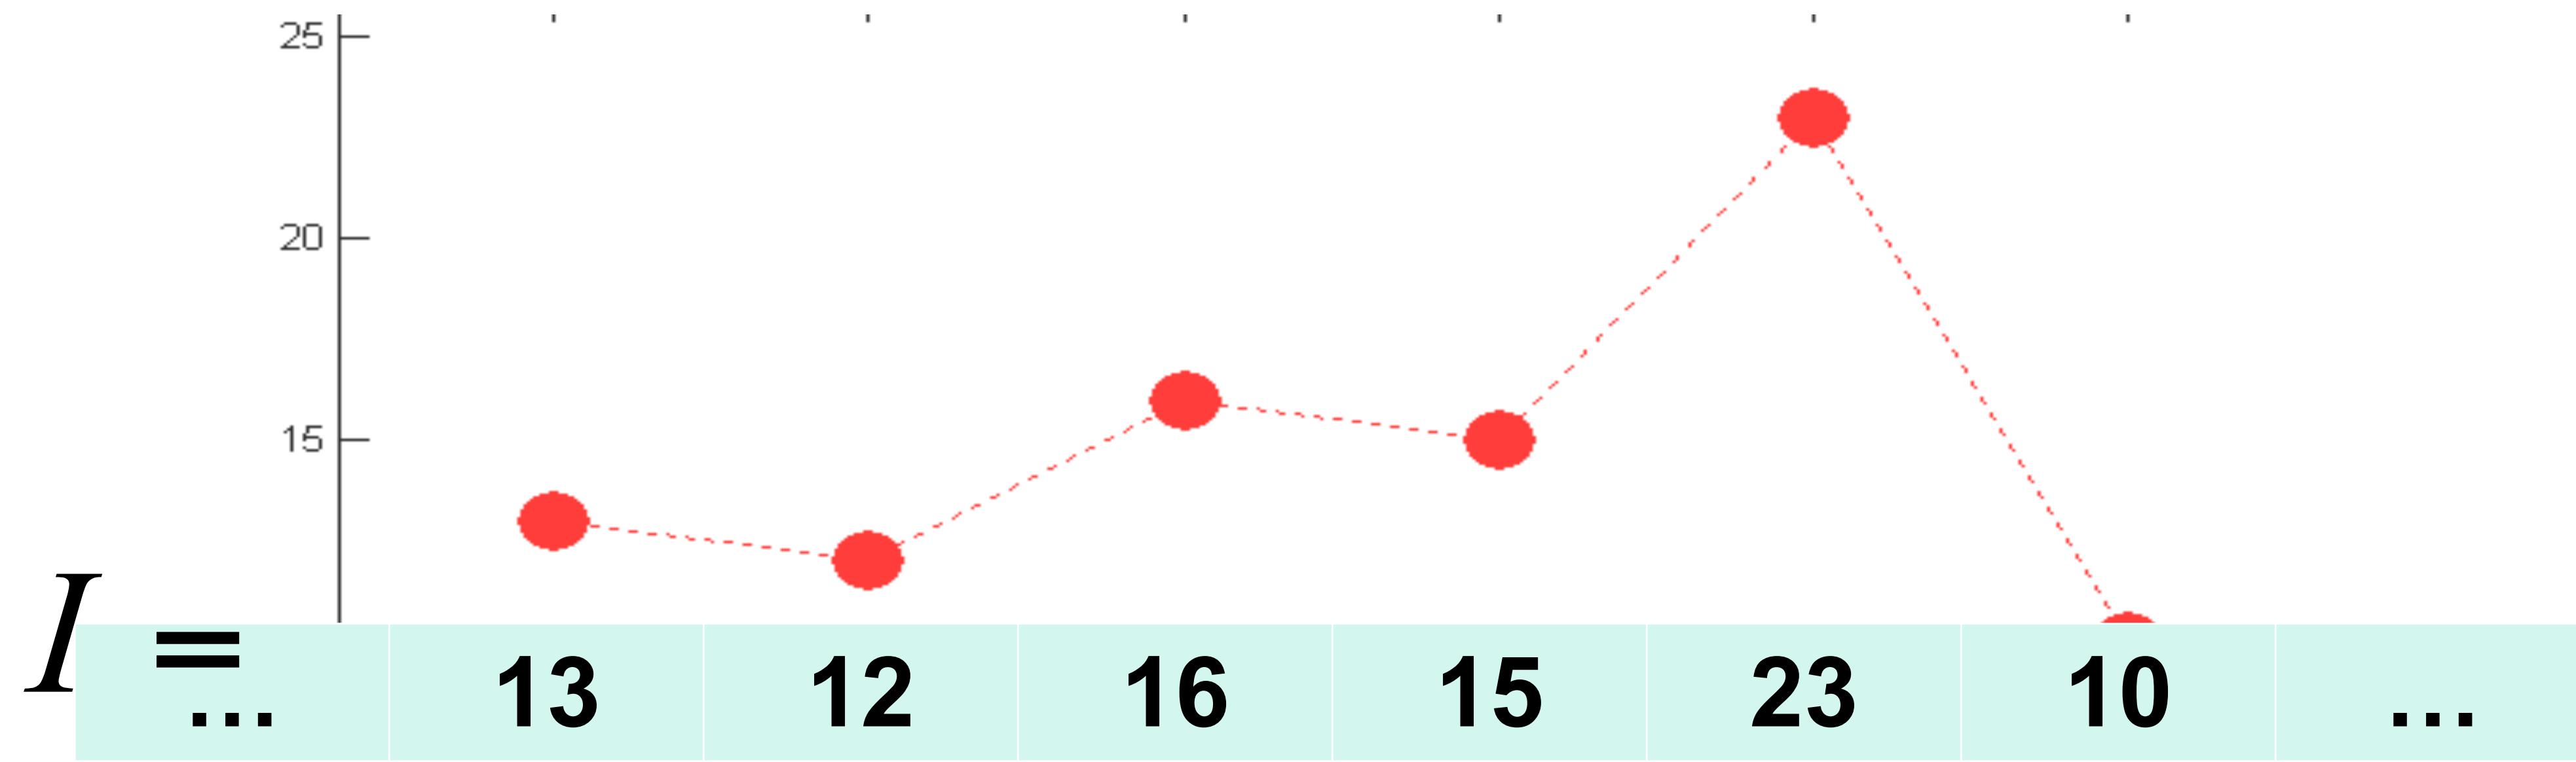


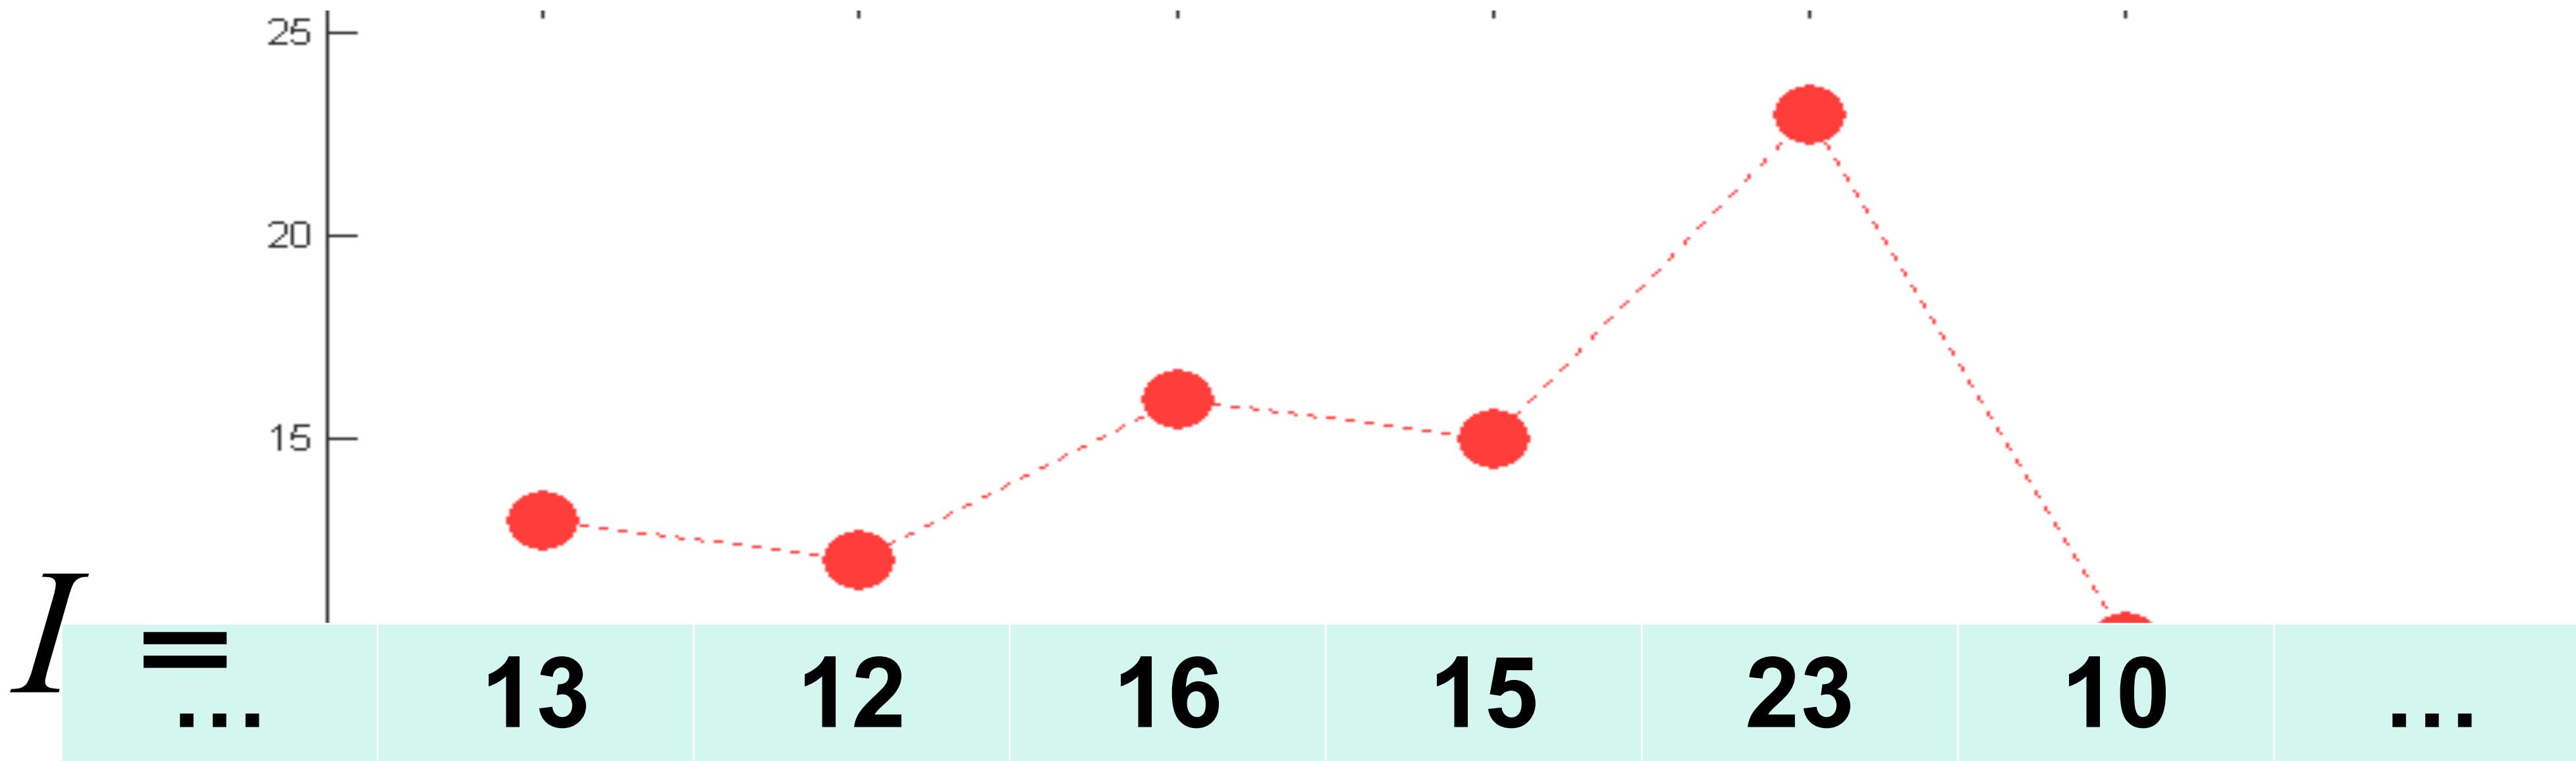
Neighborhood: $N(x) = x \pm 1$
 Operation: Average

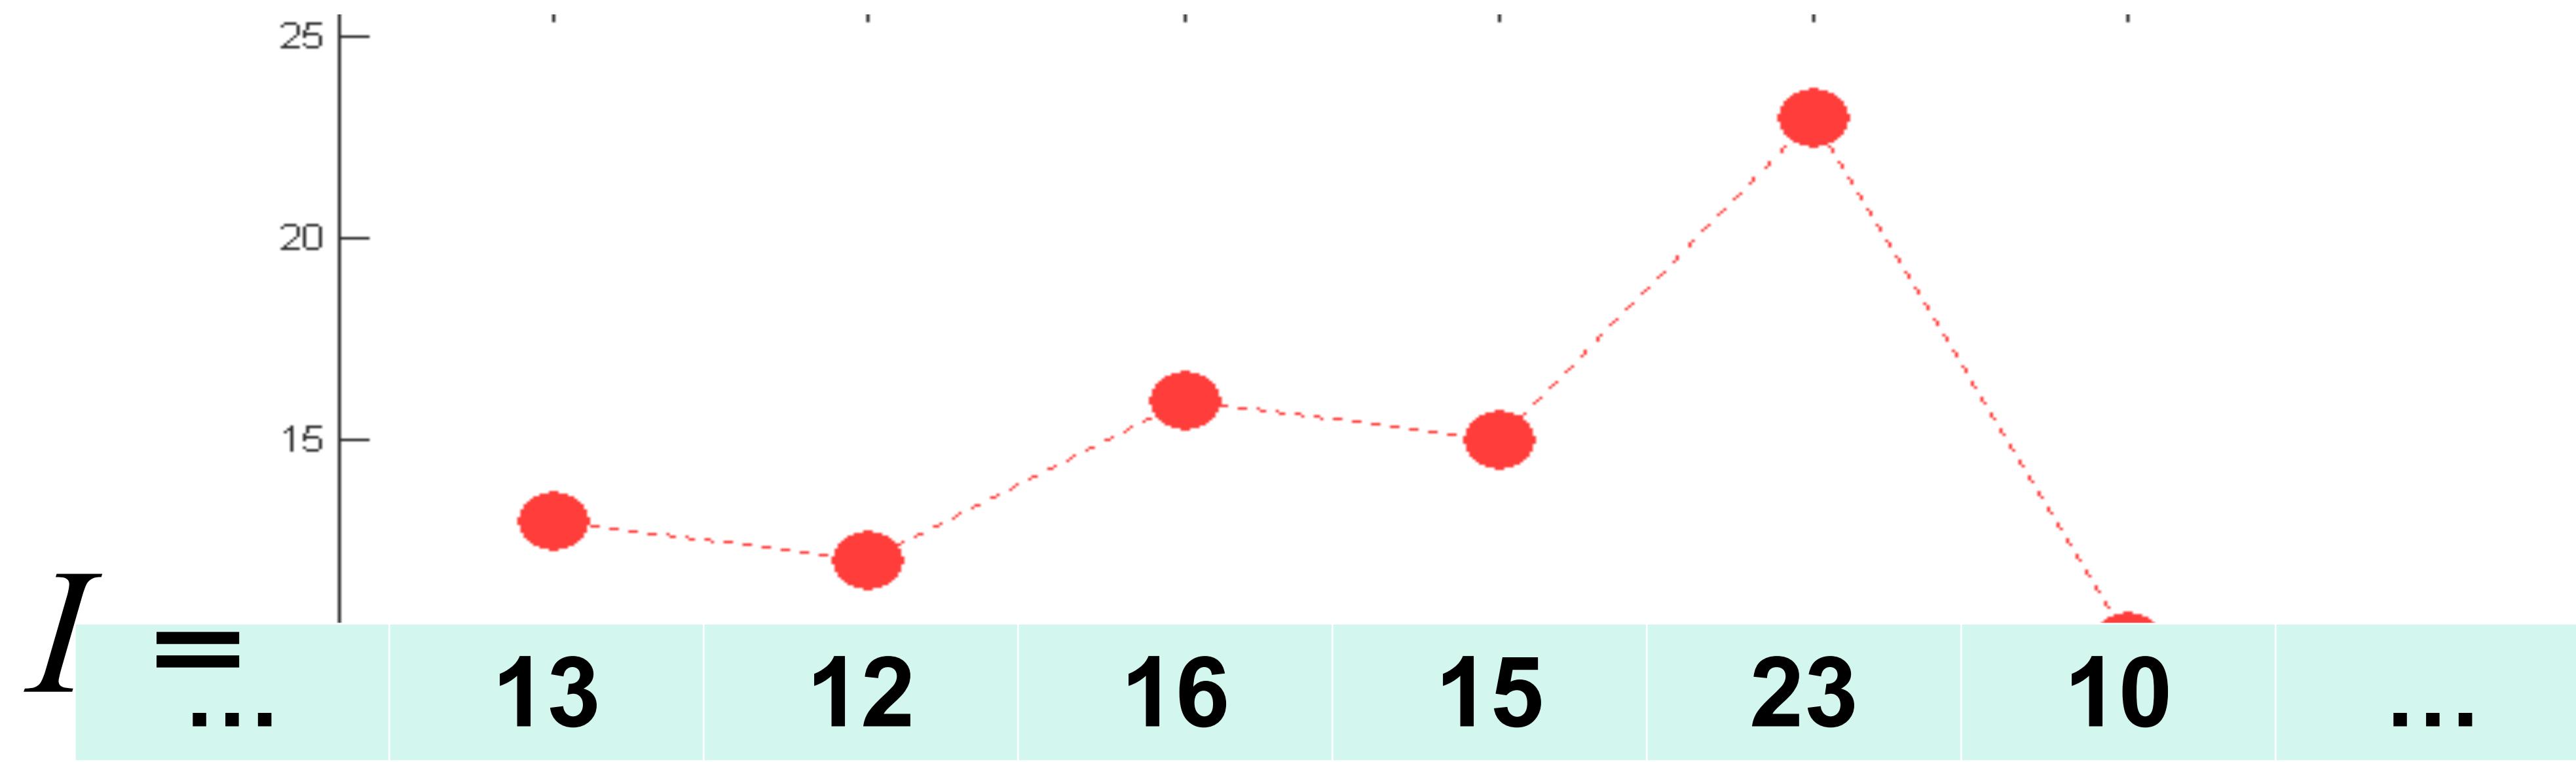
$I' =$

\dots	$??$	$??$	$??$	$??$	\dots
---------	------	------	------	------	---------









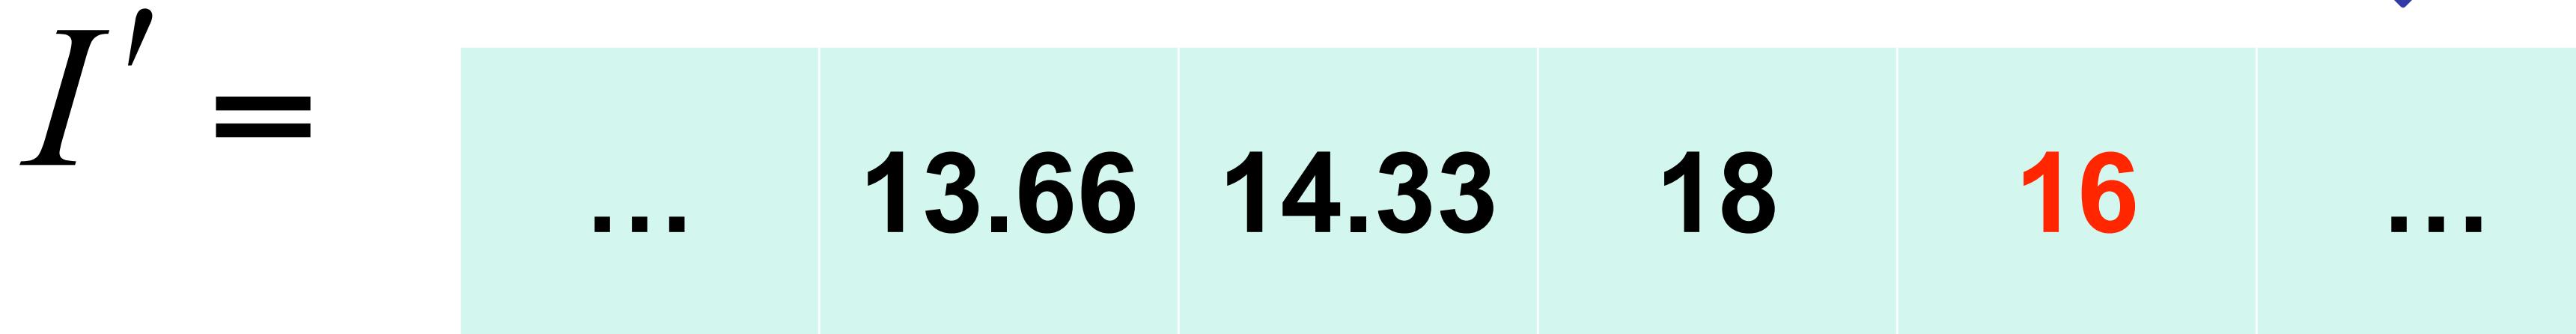
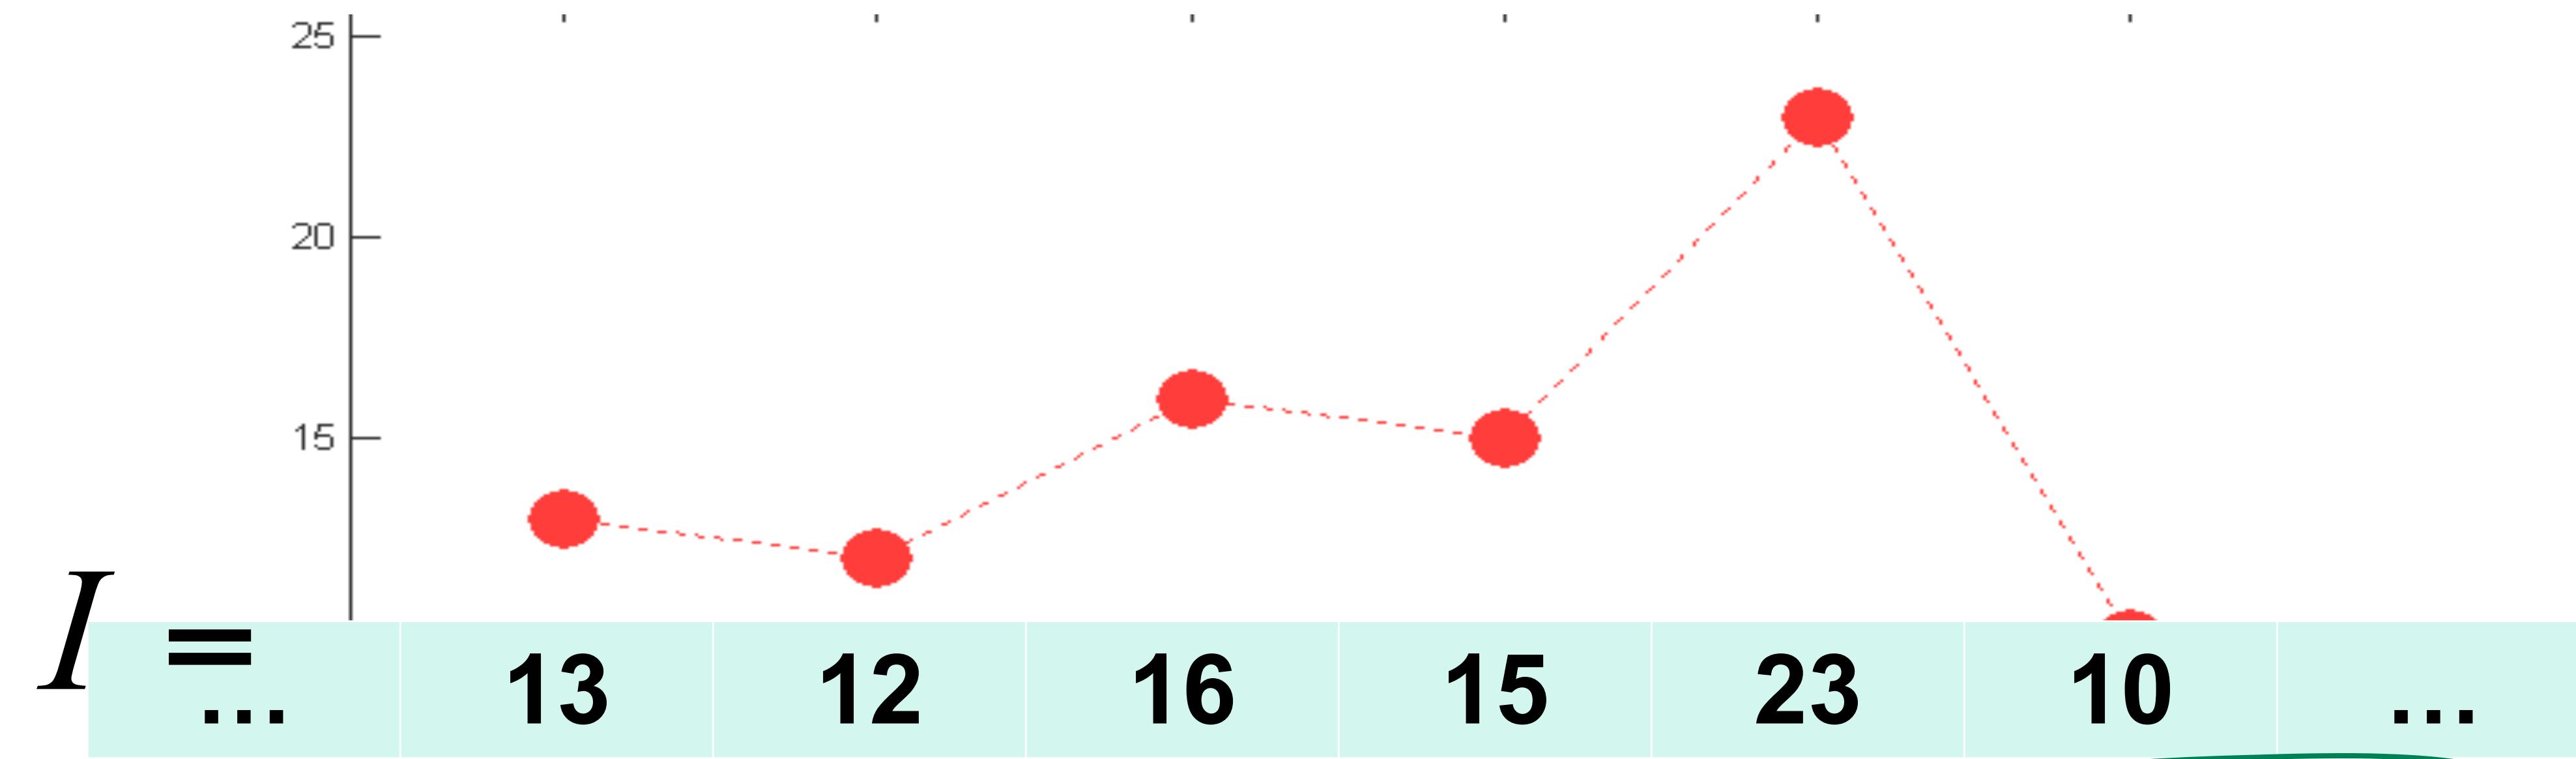
$K =$

* * *

1/3 1/3 1/3 ...

$I' =$

... 13.66 14.33 18 ?? ...



Correlation

- Linear operation in 1D:

$$I'(x) = \sum_{i=-a}^{a} K(i)I(x+i)$$

- I is the input image; I' is the output of the operation
- K is the *kernel* of the operation; many choices!
- $N(x)$ is a neighbourhood of size $(2a+1)$

Correlation

- Linear operation in 1D:

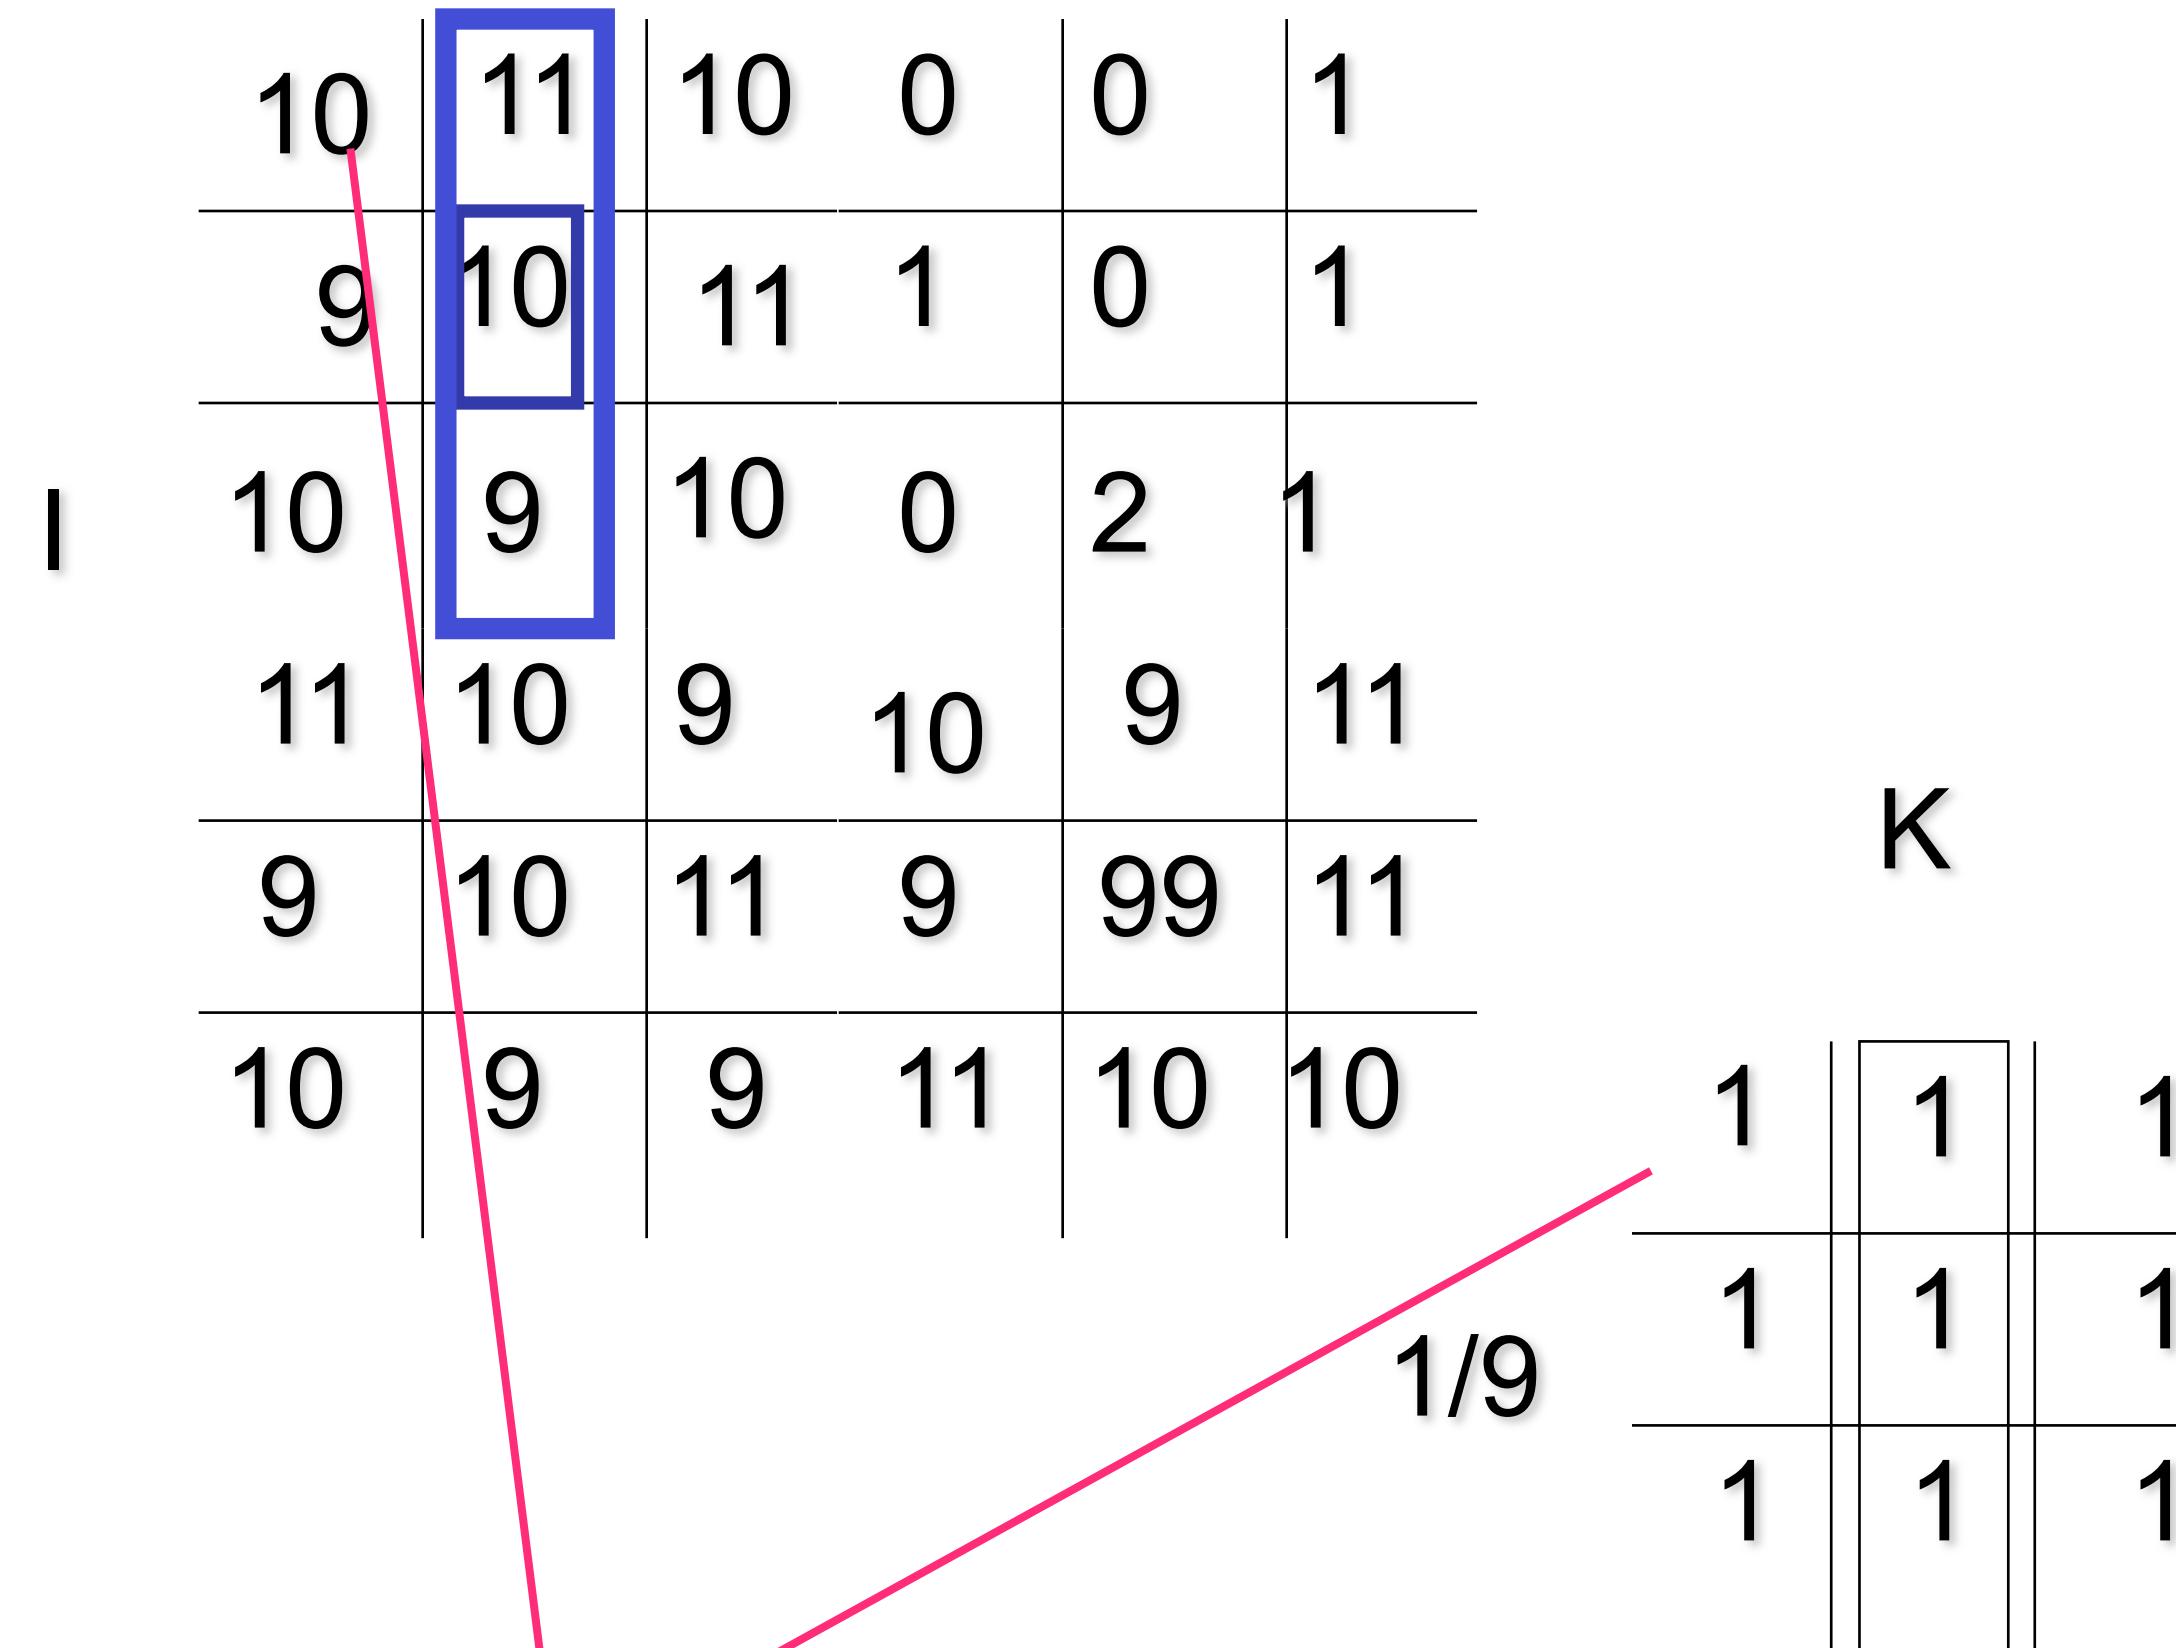
$$I'(x) = \sum_{i=-a}^a K(i)I(x+i)$$


- Linear operation in 2D:

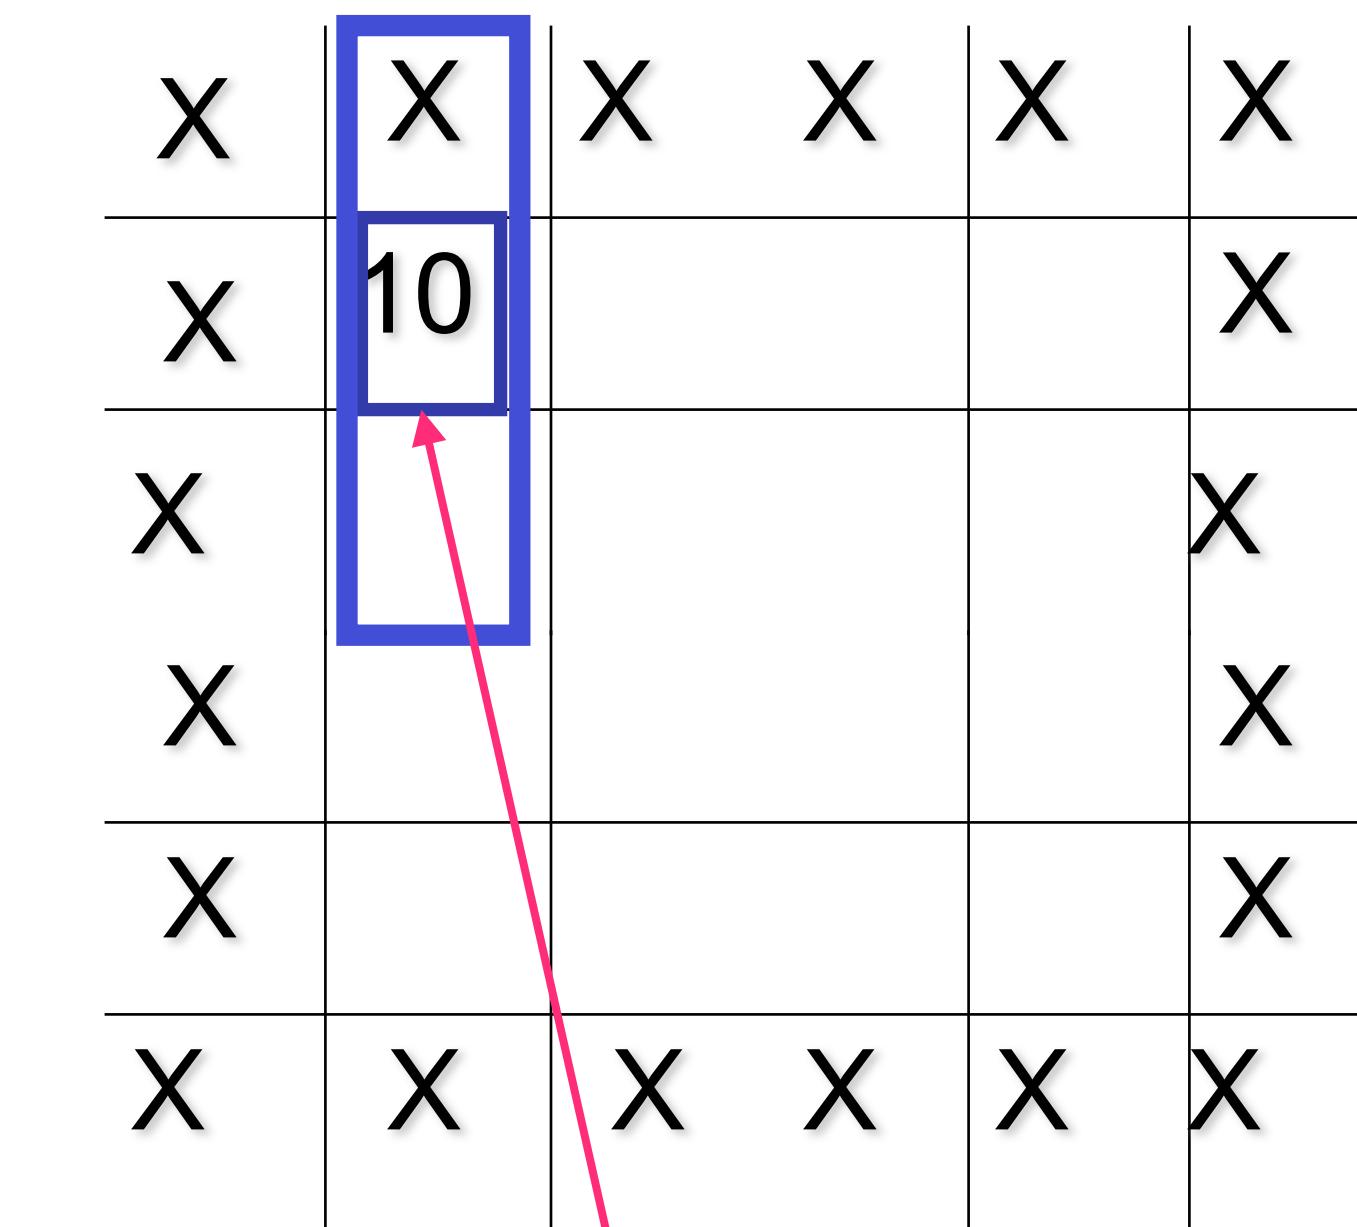
$$I'(x, y) = \sum_{i=-a}^a \sum_{j=-b}^b K(i, j)I(x+i, y+j)$$


- I is the input image; I' is the output of the operation
- K is the *kernel* of the operation; many choices!
- $N(x,y)$ is a neighbourhood of size $(2a+1) \times (2b+1)$

box filter :-

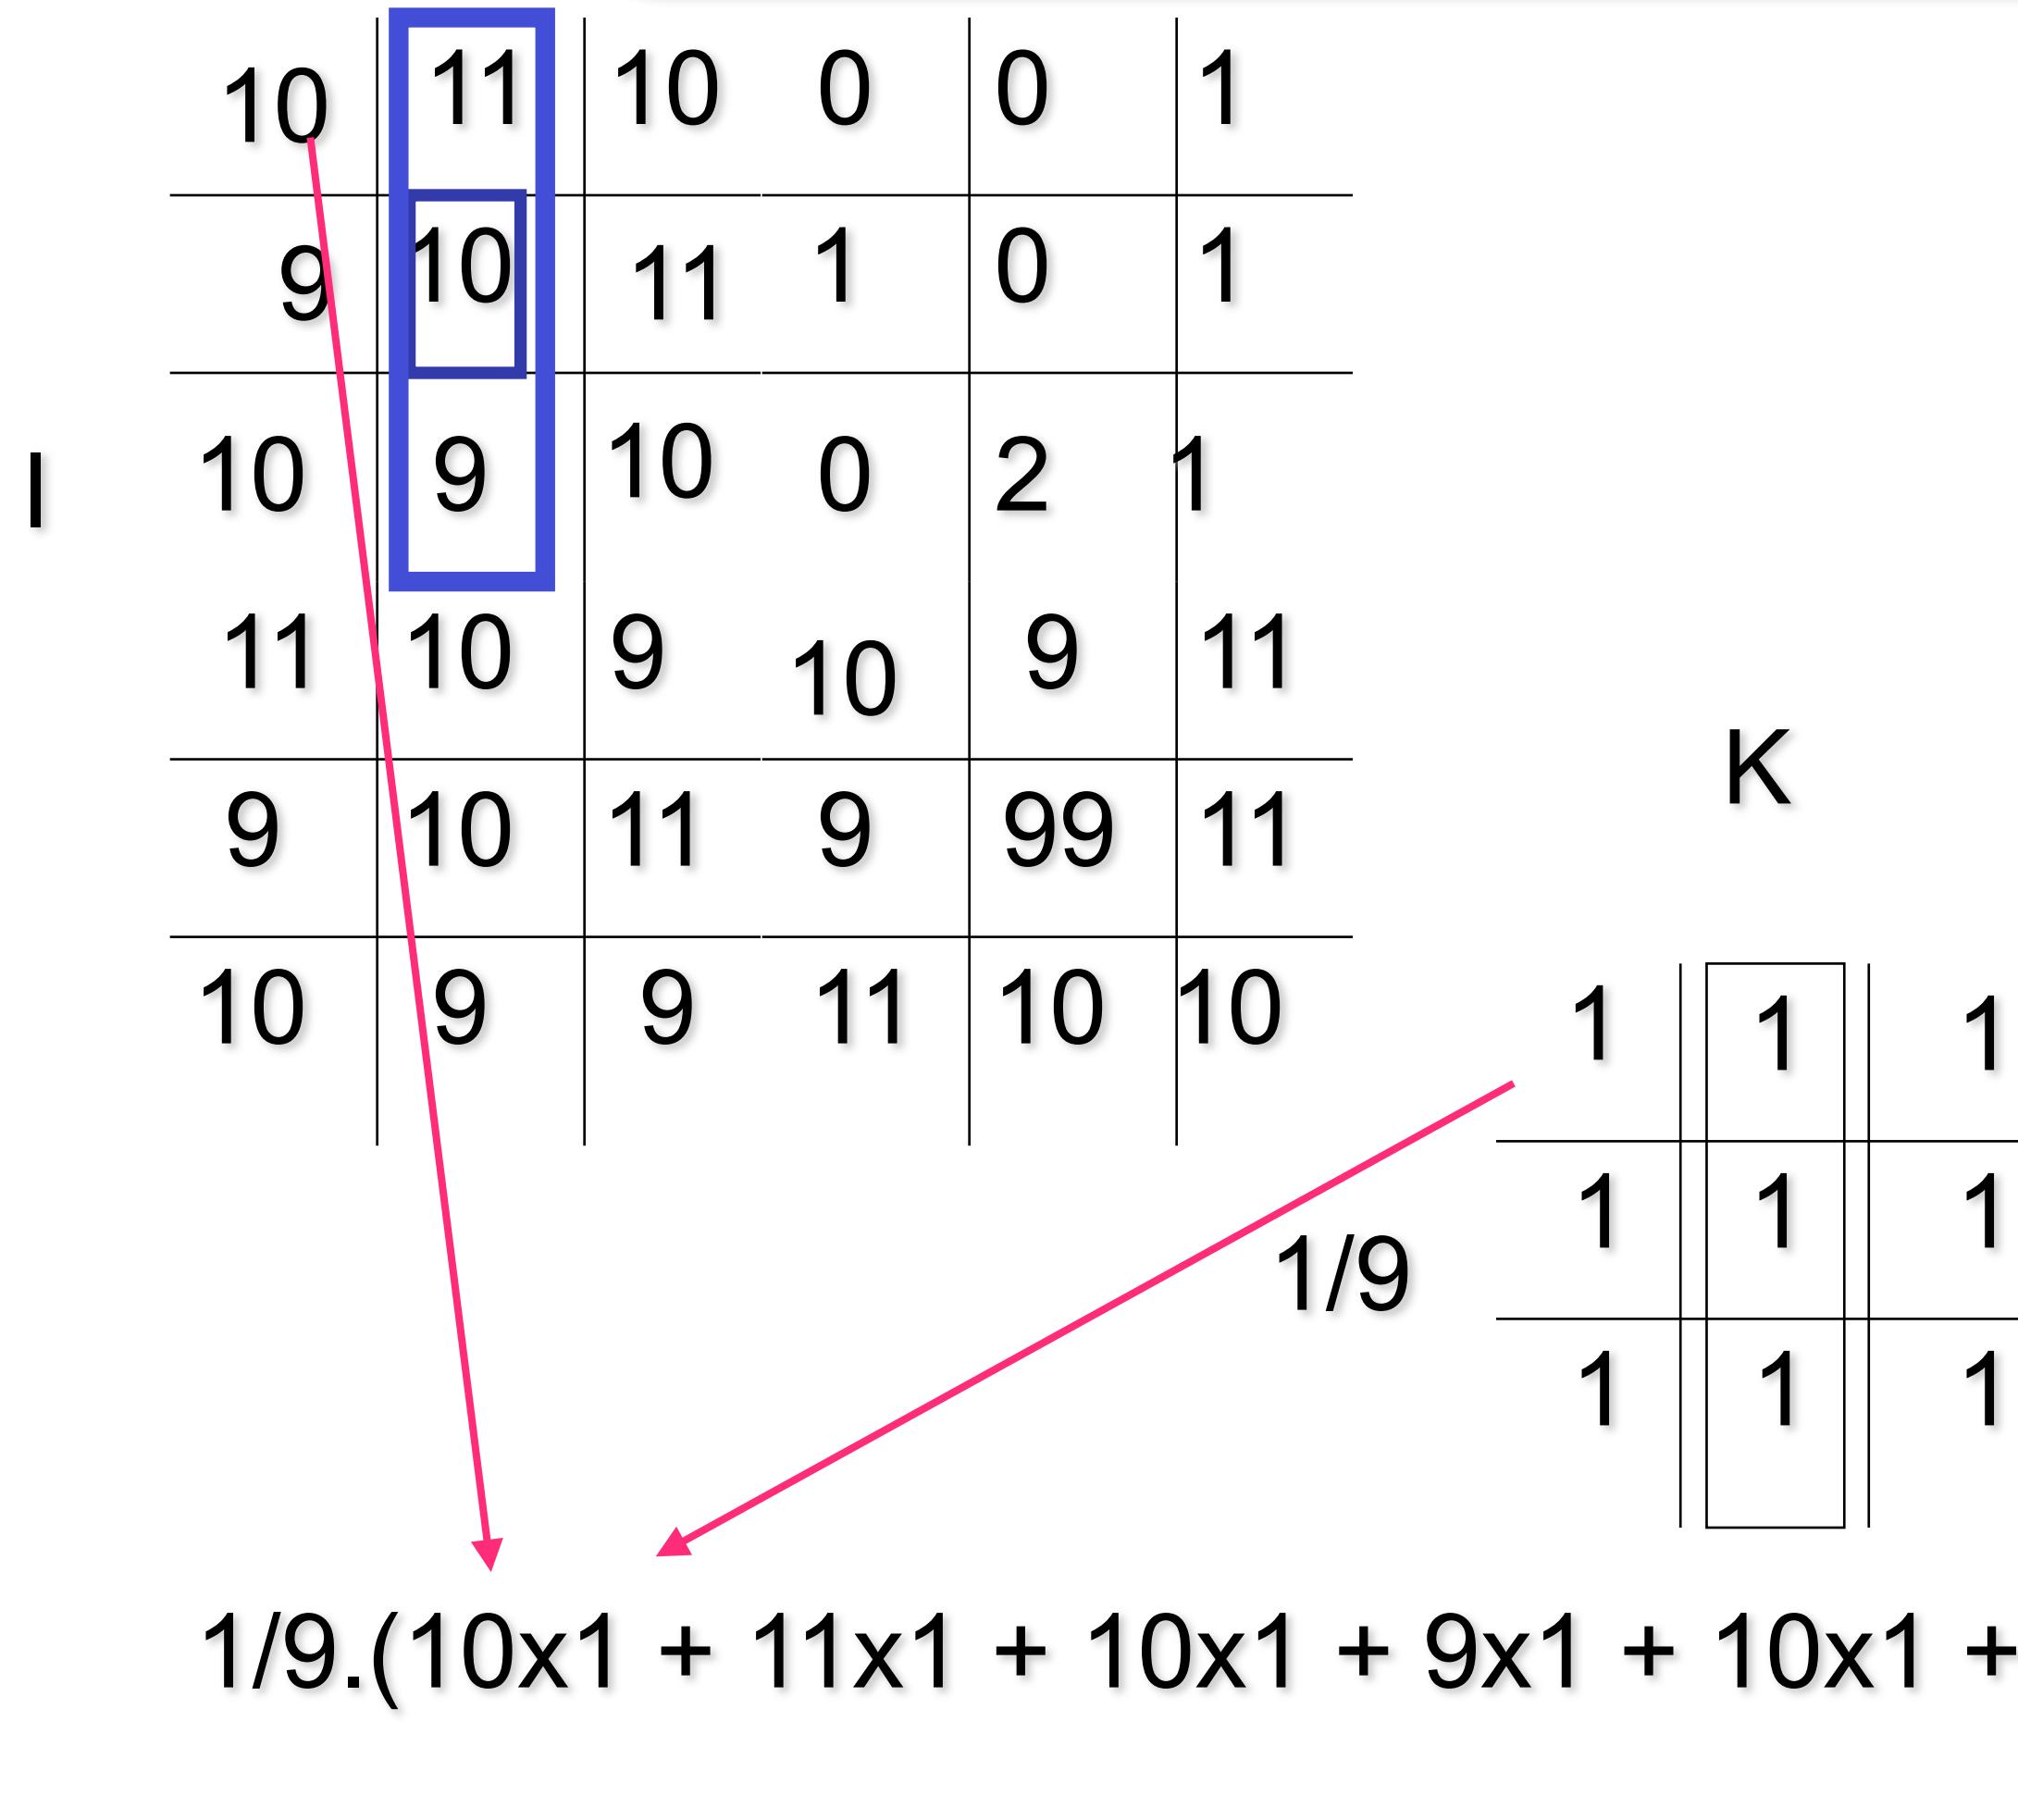


$$1/9.(10 \times 1 + 11 \times 1 + 10 \times 1 + 9 \times 1 + 10 \times 1 + 11 \times 1 + 10 \times 1 + 9 \times 1 + 10 \times 1) = \\ 1/9.(90) = 10$$



$$1/9.(90) = 10$$

$$I'(x, y) = \sum_{i=-a}^a \sum_{j=-b}^b K(i, j) I(x + i, y + j)$$



Correlation with a Unit-Impulse Image

$$I = \begin{array}{cccccccc} 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \end{array}$$

$$K = \begin{array}{ccc} 5 & 6 & 7 \end{array}$$

$$I' = \begin{array}{cccccccccccccccccc} - & - & - & - & - & - & - & - & - & - & - & - & - & - & - & - & - & - & - & - \end{array}$$

convolution 有
Image Filtering

Correlation with a Unit-Impulse Image

$$I = \begin{array}{cccccccc} 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \end{array}$$

$$K = \begin{array}{ccc} 5 & 6 & 7 \end{array}$$

$$I' = \begin{array}{cccccccc} 0 & 0 & 7 & 6 & 5 & 0 & 0 & 0 \\ \underline{-} & \underline{-} \end{array}$$

Correlation reveals copy of K , rotated by 180°

Correlation with a Unit-Impulse Image

$$I = \begin{array}{cccccccc} 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \end{array}$$

$$K = \begin{array}{ccc} 5 & 6 & 7 \end{array}$$

$$I' = \begin{array}{cccccccc} 0 & 0 & 7 & 6 & 5 & 0 & 0 & 0 \\ \underline{-} & \underline{-} \end{array}$$

Correlation reveals copy of K , rotated by 180°

Would be nice to get result that is **NOT** rotated!

Correlation/Convolution

$$I'(x, y) = \sum_{i=-a}^{a} \sum_{j=-b}^{b} K(i, j) I(x + i, y + j)$$

Correlation/Convolution

$$I'(x, y) = \sum_{i=-a}^{a} \sum_{j=-b}^{b} K(i, j)I(x + i, y + j)$$

Convolution

Correlation/Convolution

$$I'(x, y) = \sum_{i=-a}^a \sum_{j=-b}^b K(i, j) I(x + i, y + j)$$

Convolution

$$\begin{aligned} I'(x, y) &= \sum_{i=-a}^a \sum_{j=-b}^b K(i, j) I(x - i, y - j) \\ &= \sum_{i=-a}^a \sum_{j=-b}^b K(-i, -j) I(x + i, y + j) \end{aligned}$$

Correlation/Convolution

$$I'(x, y) = \sum_{i=-a}^a \sum_{j=-b}^b K(i, j) I(x + i, y + j)$$

Convolution

*if I is image
 K is kernel*

$$\begin{aligned} I'(x, y) &= \sum_{i=-a}^a \sum_{j=-b}^b K(i, j) I(x - i, y - j) \\ &= \sum_{i=-a}^a \sum_{j=-b}^b K(-i, -j) I(x + i, y + j) \end{aligned}$$

slipping the kernel

$$K(i, j) = K(-i, -j) \Leftrightarrow \text{convolution} \equiv \text{correlation}$$

Correlation/Convolution

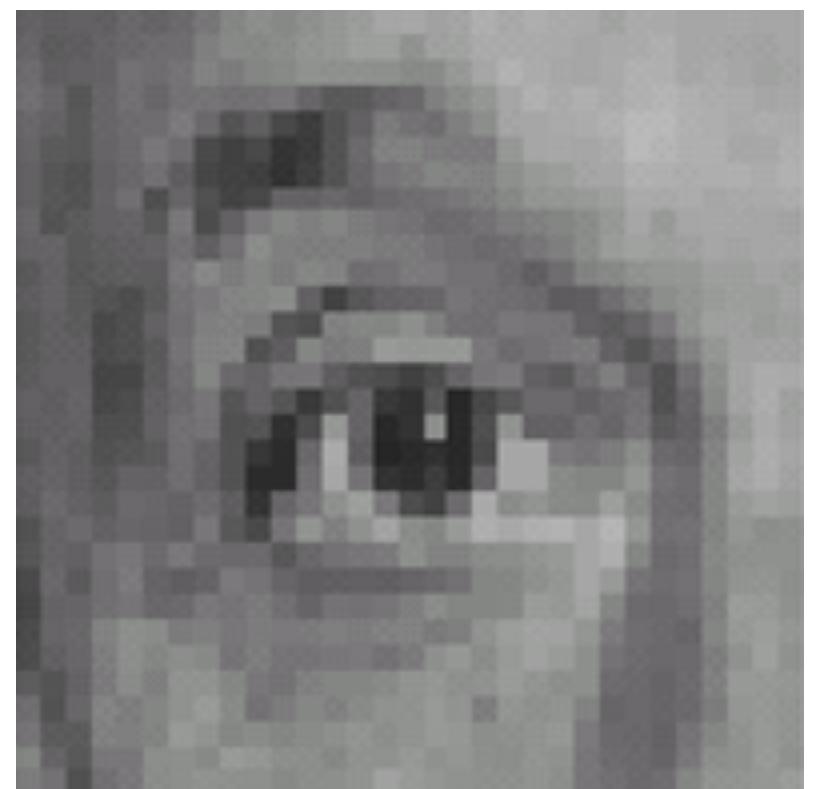
$$I'(x, y) = \sum_{i=-a}^a \sum_{j=-b}^b K(i, j) I(x + i, y + j)$$

Convolution

$$I'(x, y) = \sum_{i=-a}^a \sum_{j=-b}^b K(i, j) I(x - i, y - j)$$

$$I' = K * I$$

Linear Filtering (warm-up)



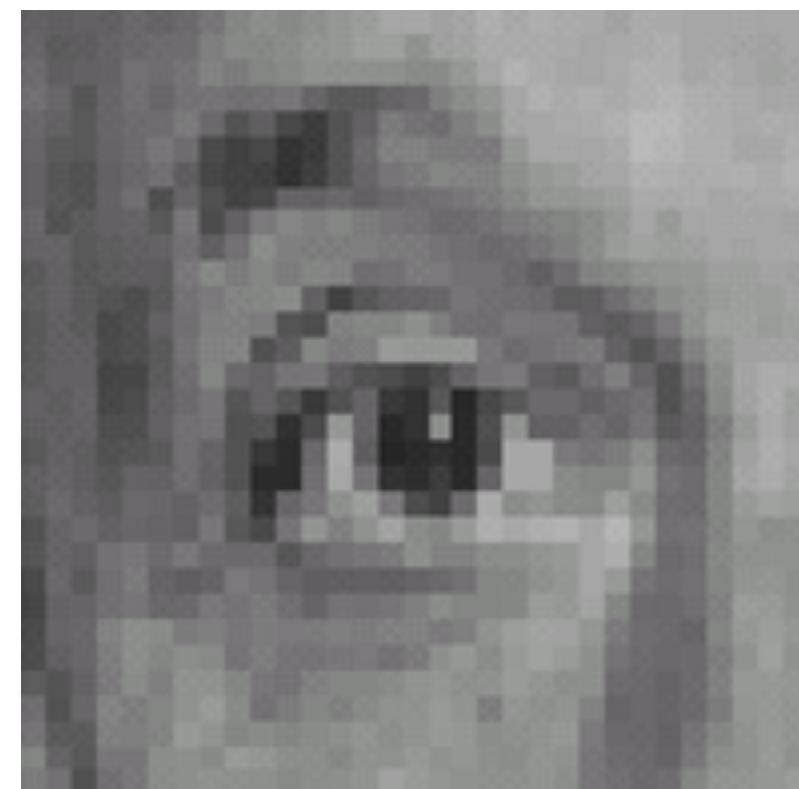
Original

0	0	0
0	1	0
0	0	0

?

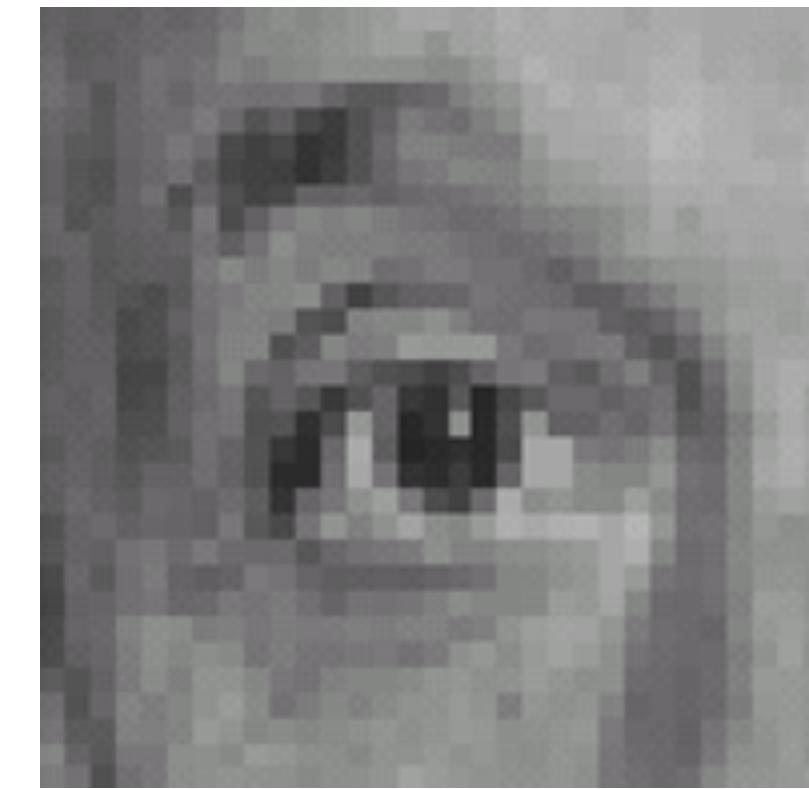
Slide credit: D.A. Forsyth

Linear Filtering (warm-up)



Original

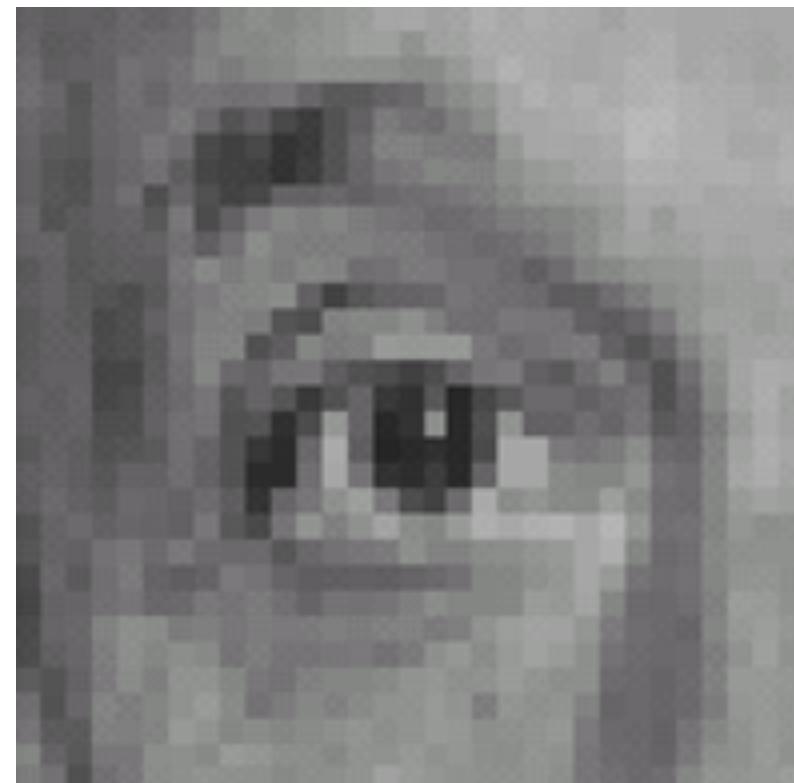
0	0	0
0	1	0
0	0	0



Filtered
(no change)

Linear Filtering (Convolution)

$$I'(x, y) = \sum_{i=-a}^a \sum_{j=-b}^b K(i, j) I(x - i, y - j)$$



Original

0	0	0
1	0	0
0	0	0

?

Linear Filtering (Convolution)

$$I'(x, y) = \sum_{i=-a}^a \sum_{j=-b}^b K(i, j) I(x - i, y - j)$$



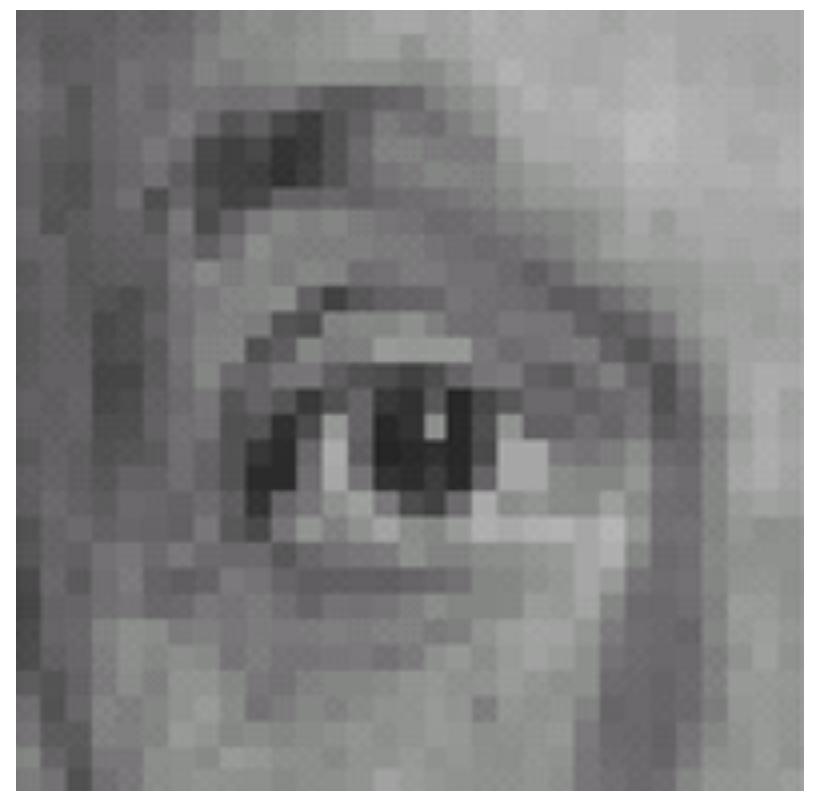
Original

0	0	0
1	0	0
0	0	0



Shifted left
By 1 pixel

Linear Filtering



Original

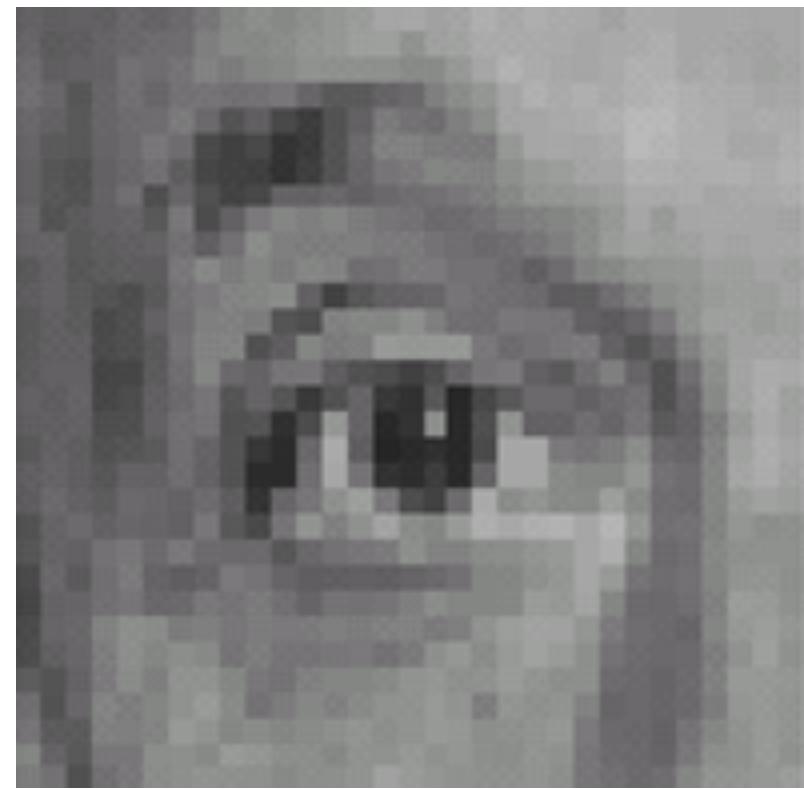
$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

?

Image Filtering

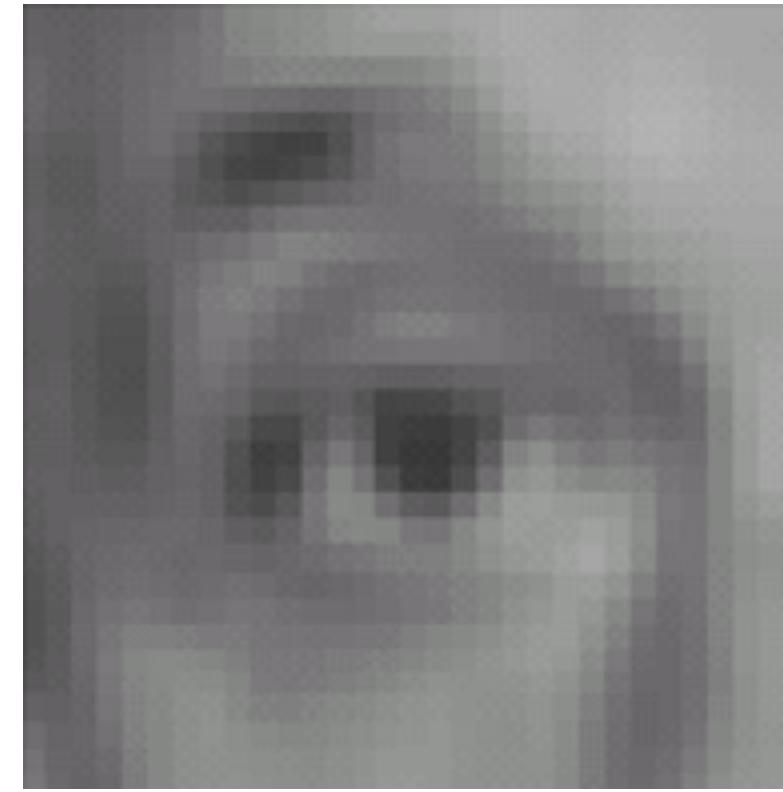
Linear Filtering

BOX FILTER



Original

$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

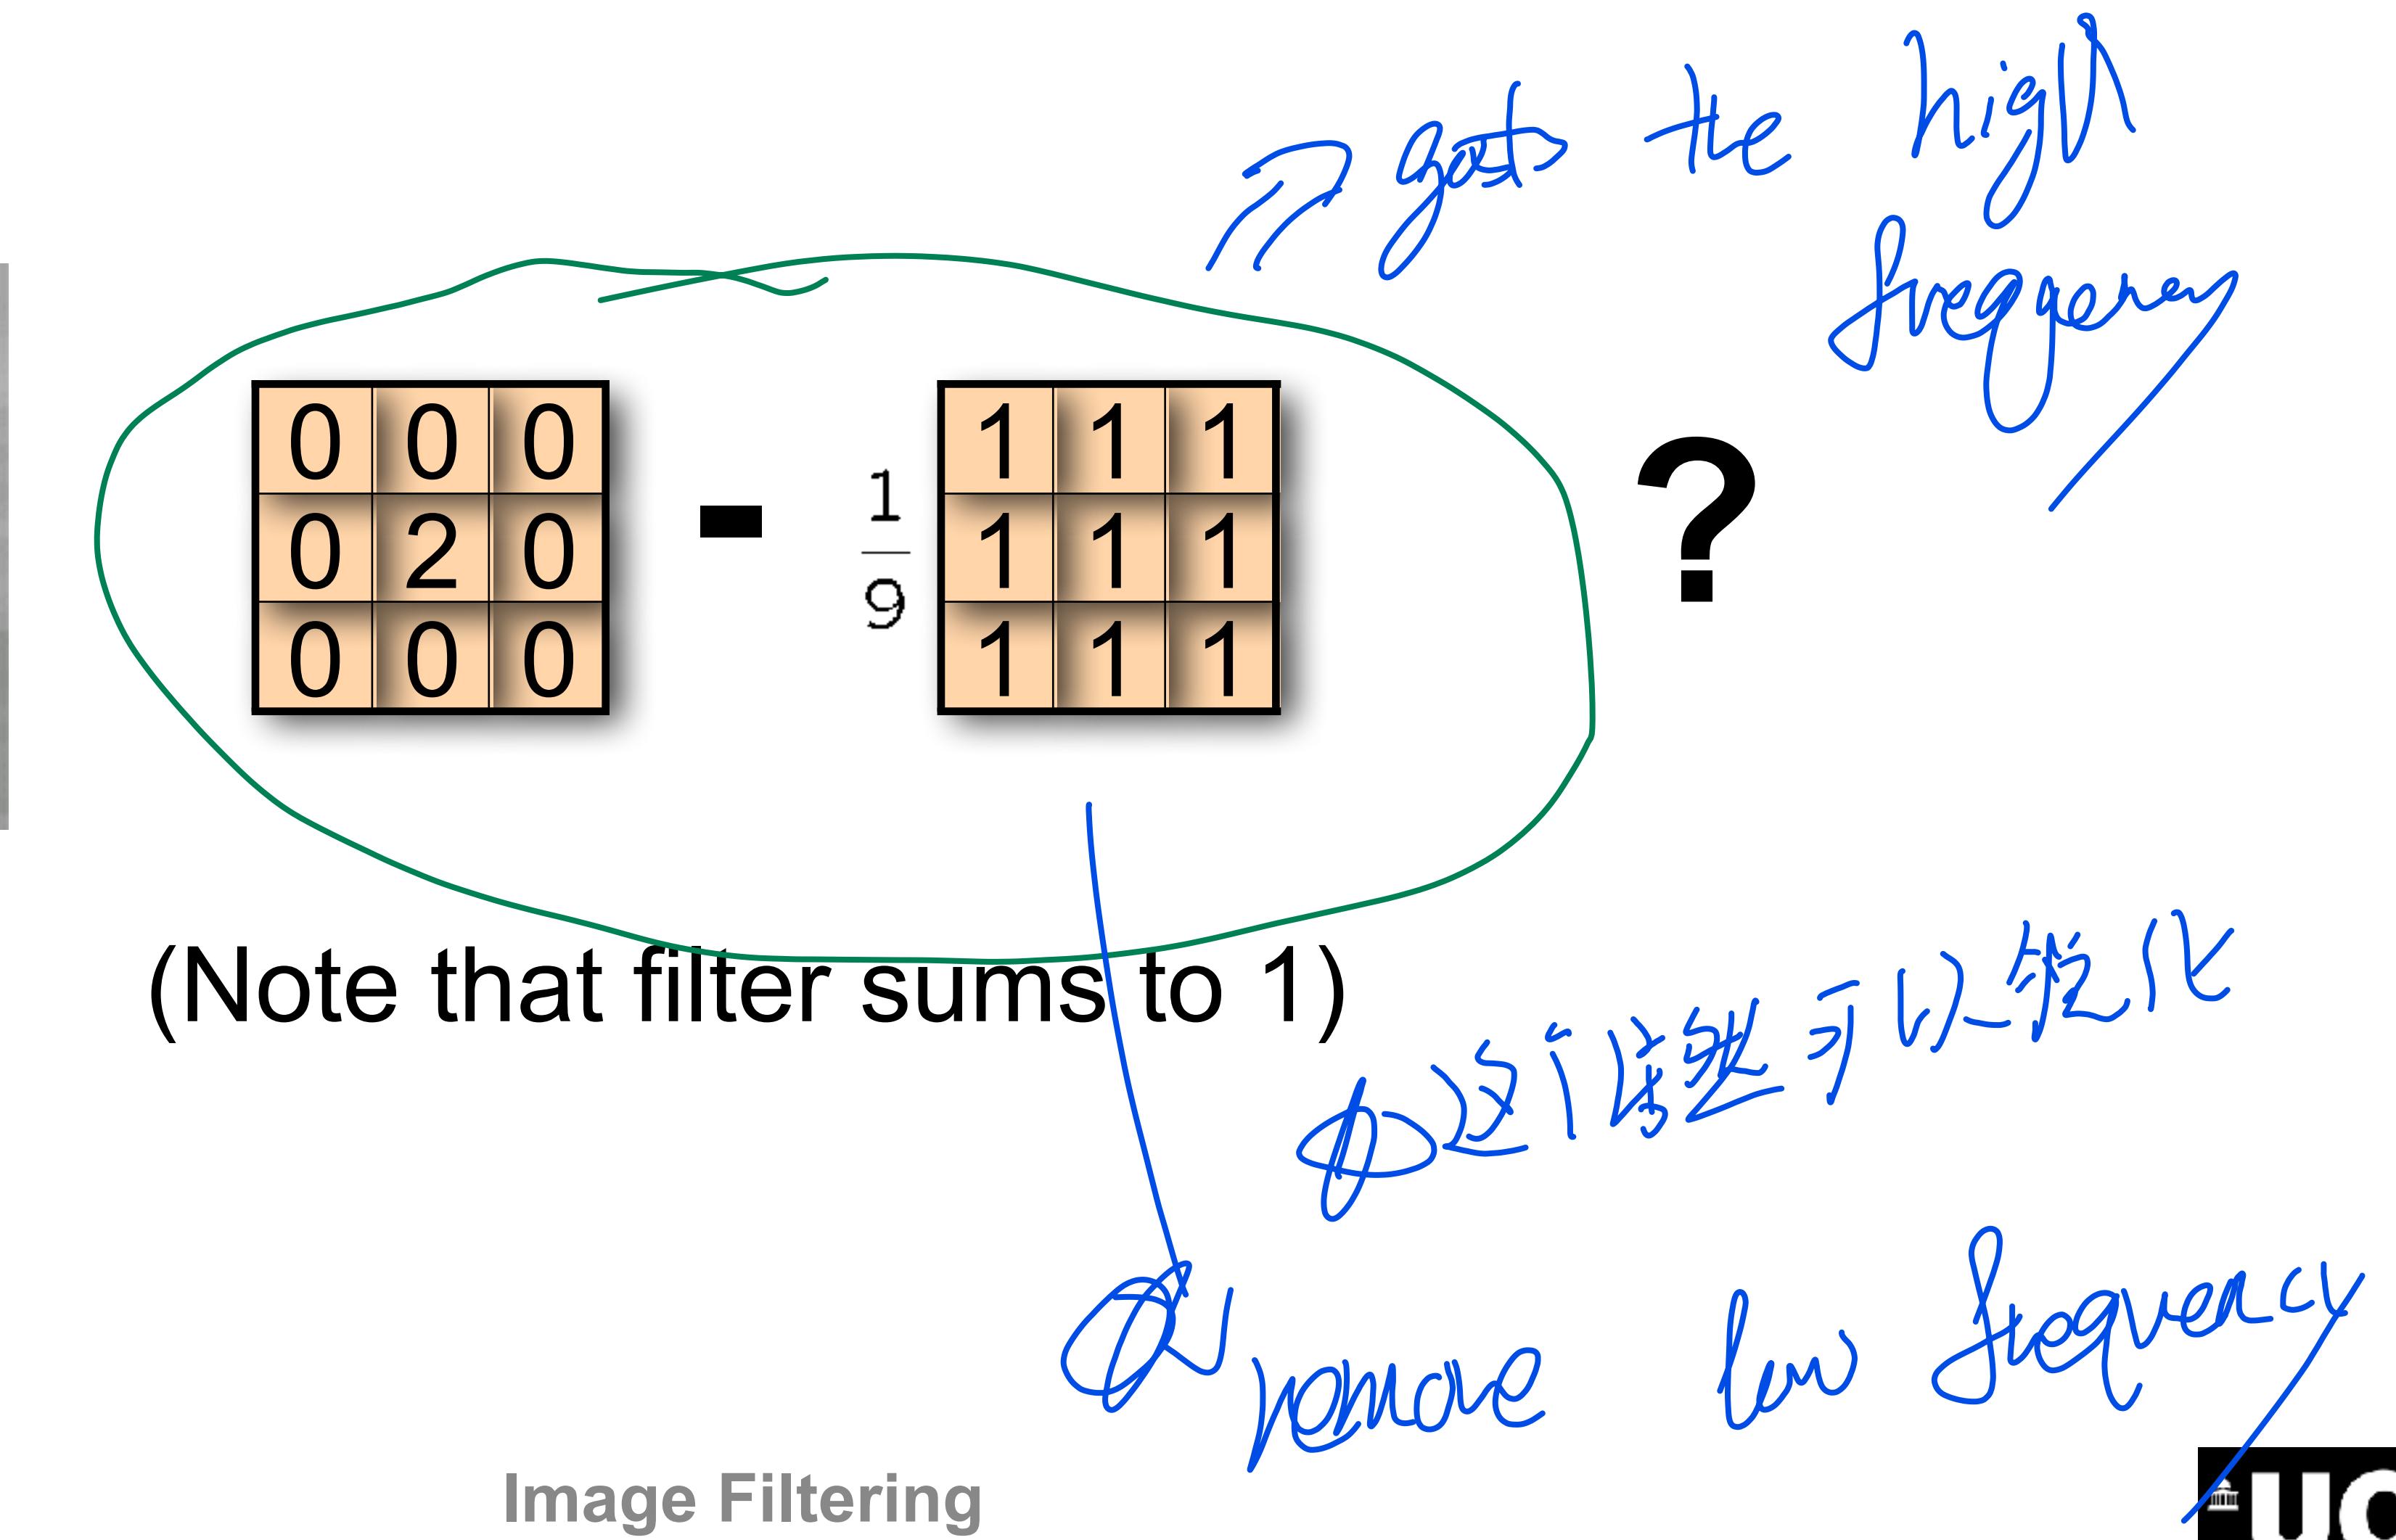


Blur (with a
box filter)

Linear Filtering



Original



Linear Filtering



Original

$$\begin{matrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{matrix}$$

-

$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

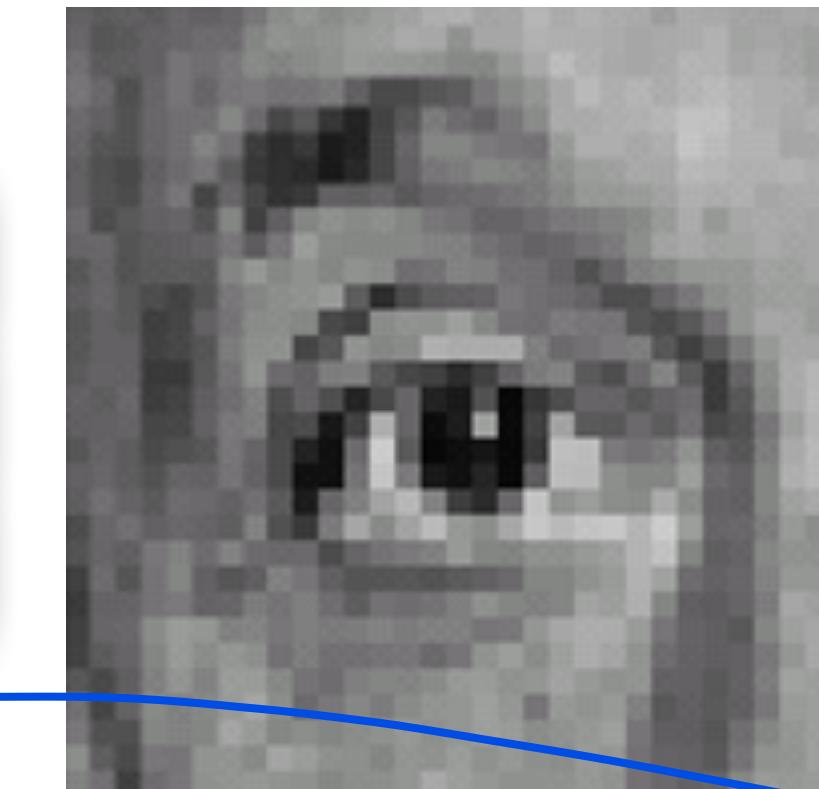


Image Filtering

Sharpening filter

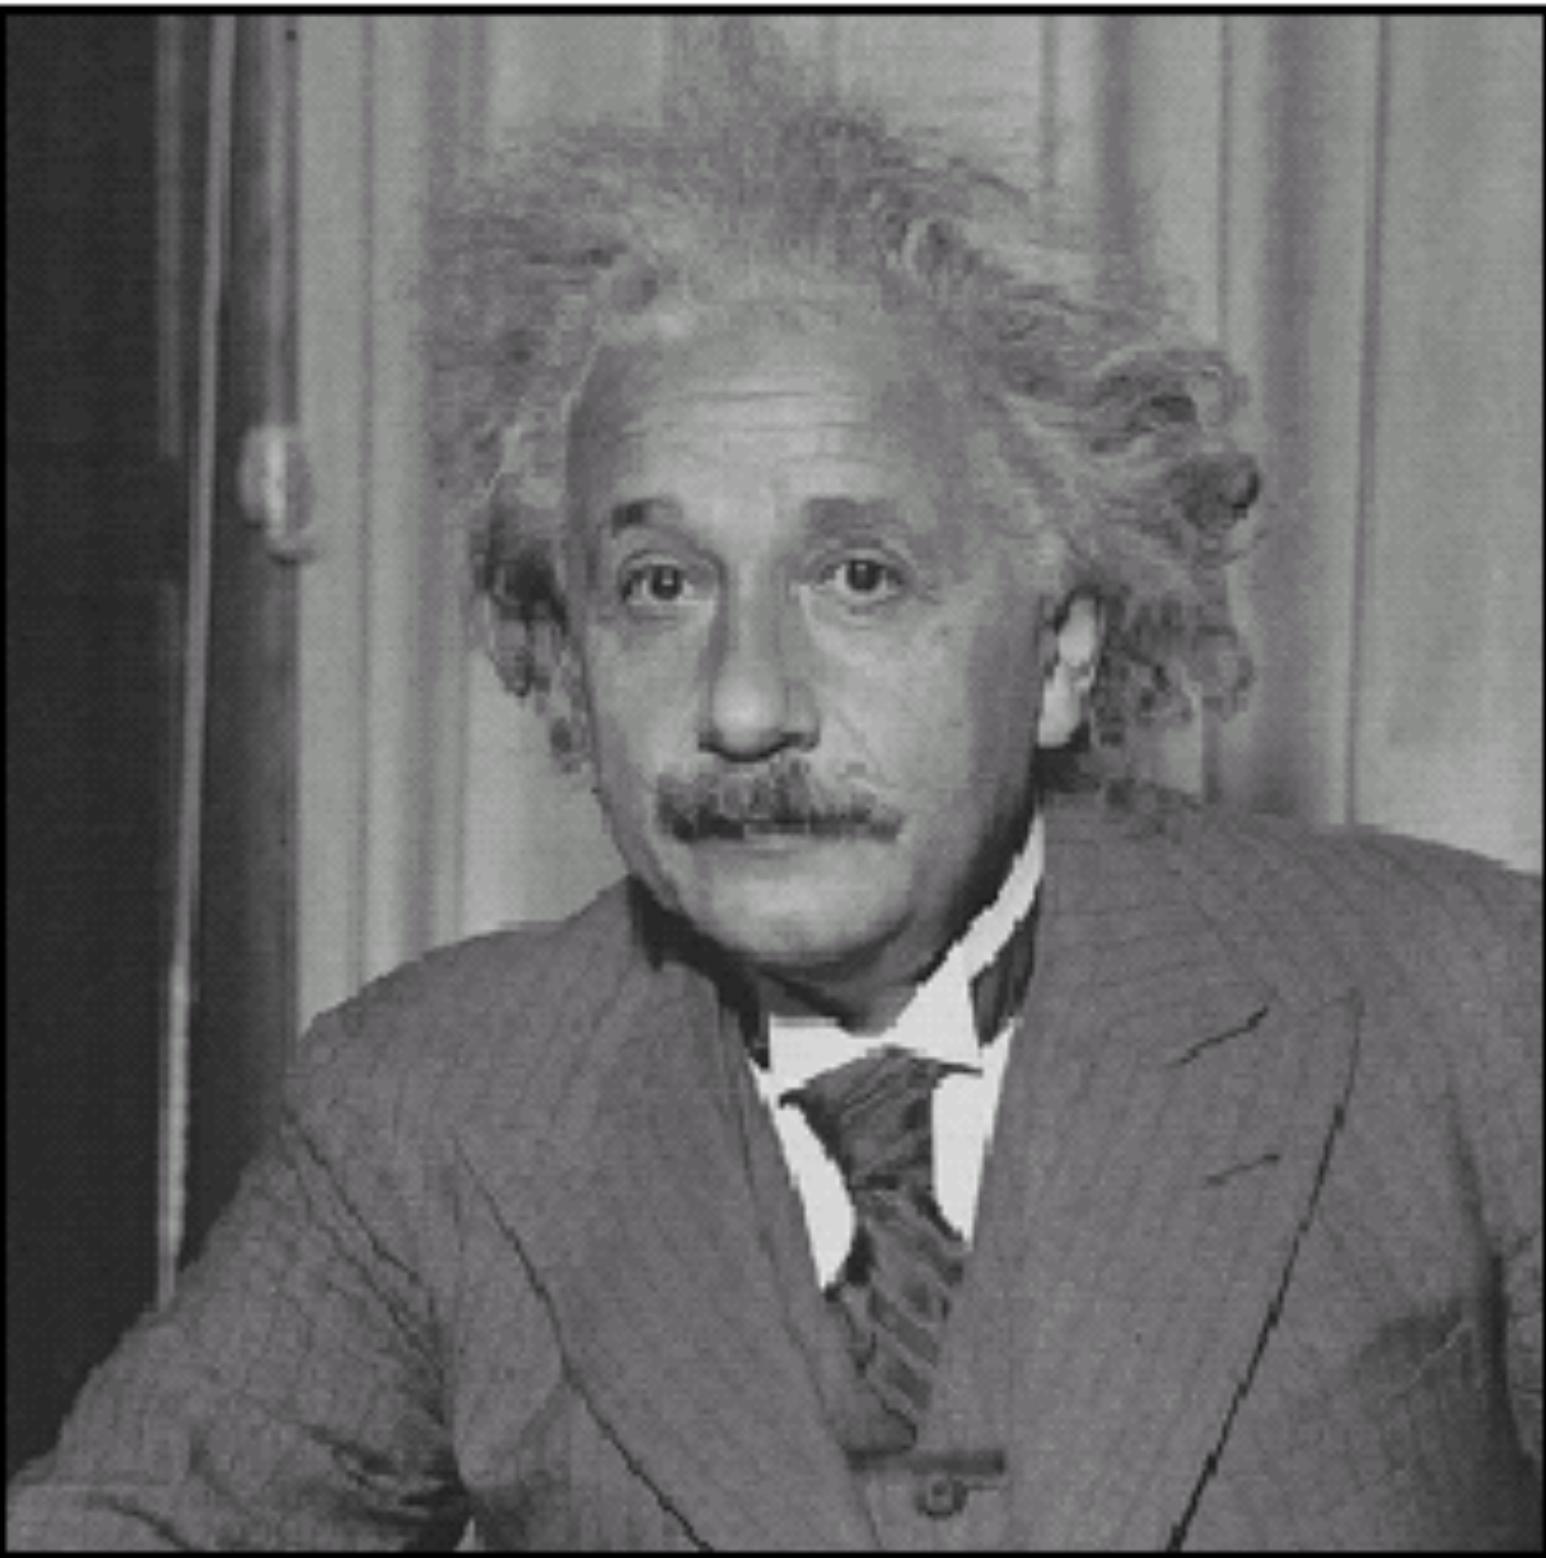
- Amplifies differences with local average
- Also known as Laplacian

COMP0026: Image Processing

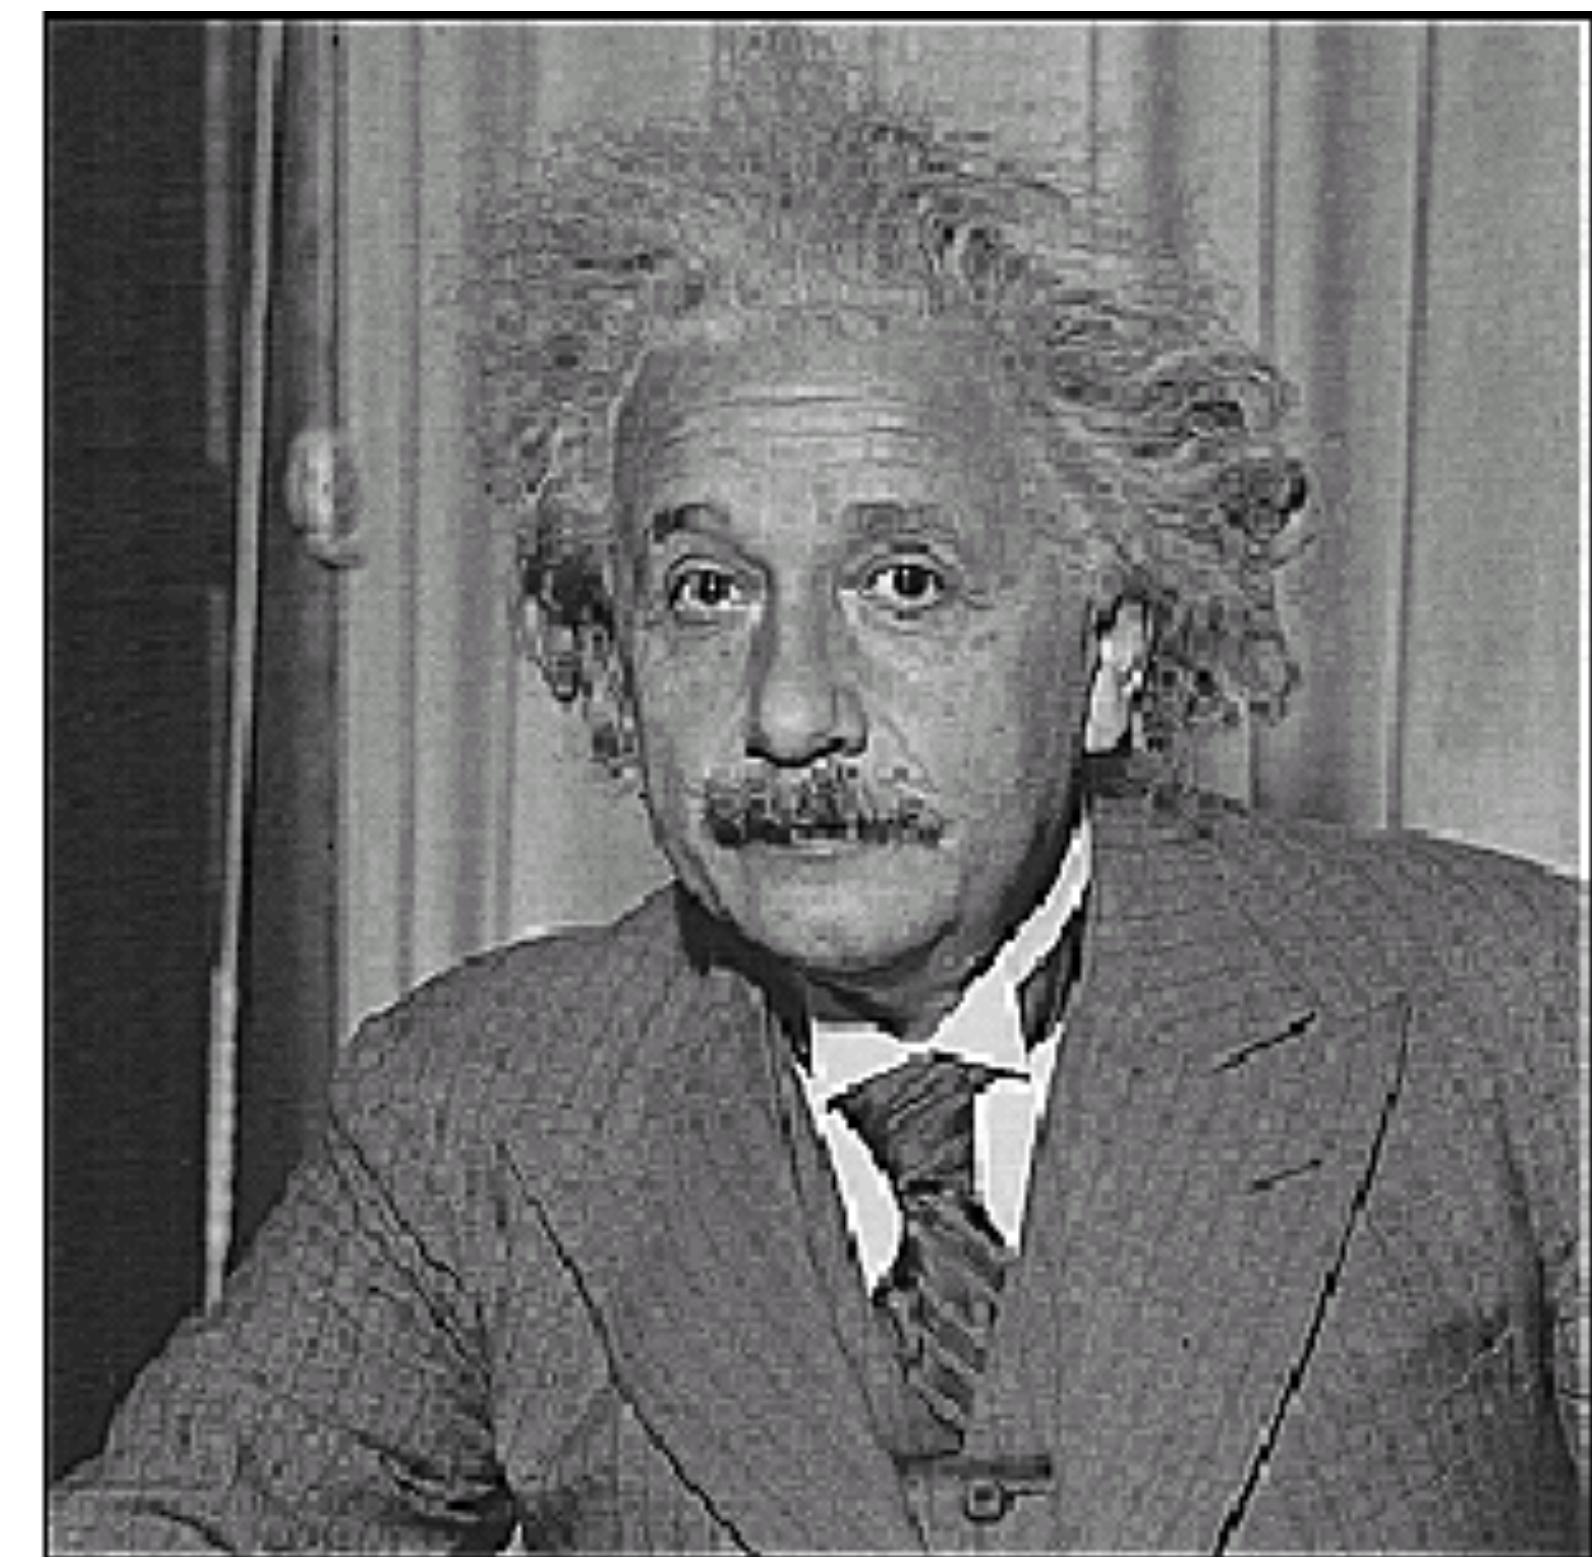
UCL

• 拉普拉斯算子

Sharpening



before



after

Example

```
K=ones(9,9);
```

```
I2=conv2(I,K);
```



Example

```
K=ones(9,9);
```

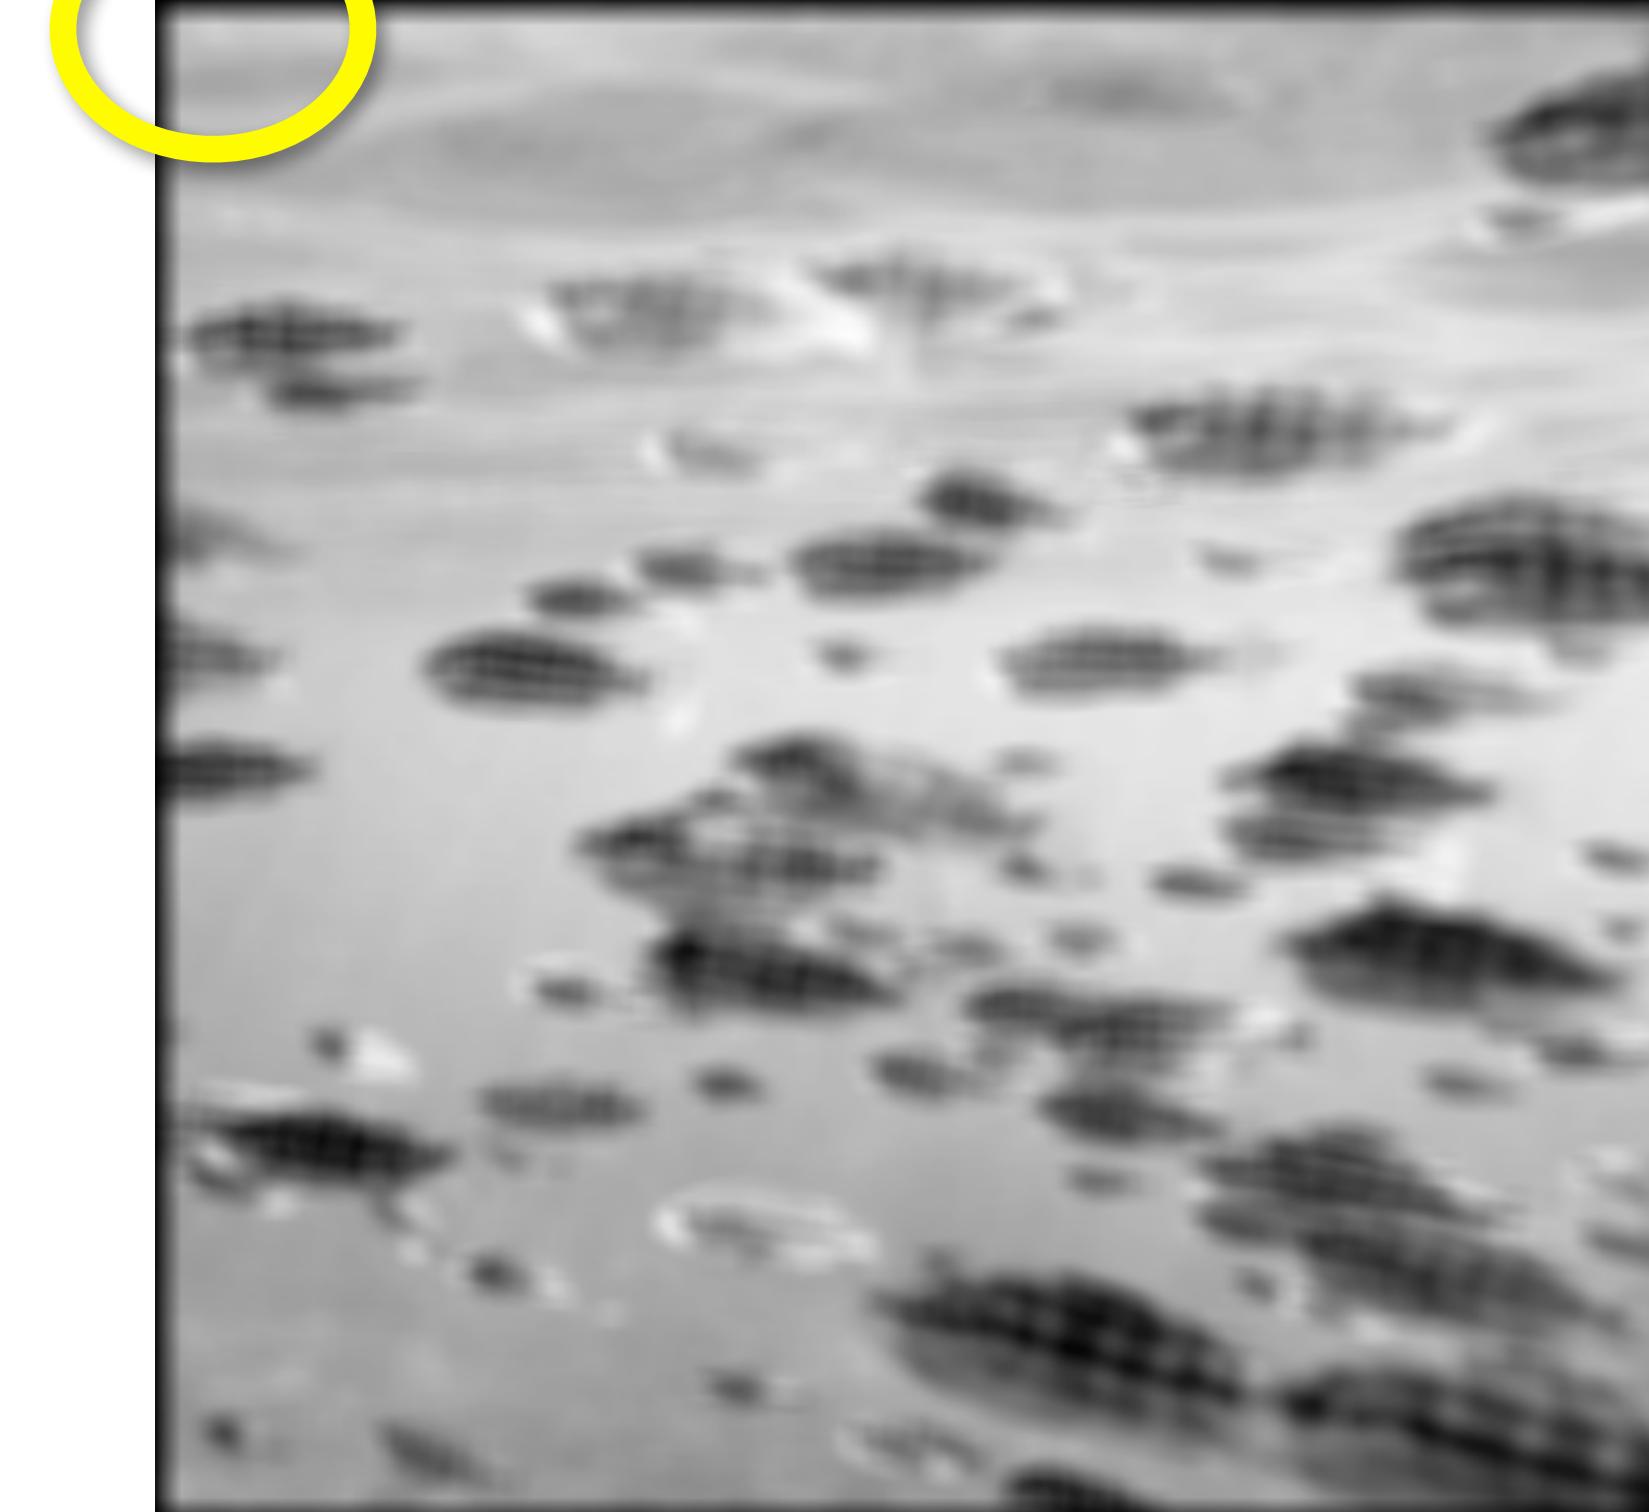
```
I2=conv2(I,K);
```



Example

```
K=ones(9,9);
```

```
I2=conv2(I,K);
```



border ~~issue~~

‘Boring’ Details

- When filter window falls off the edge of the image?
 - need to extrapolate
 - methods:
 - clip filter (black)
 - wrap around
 - copy edge
 - reflect across edge
 - vary filter near edge



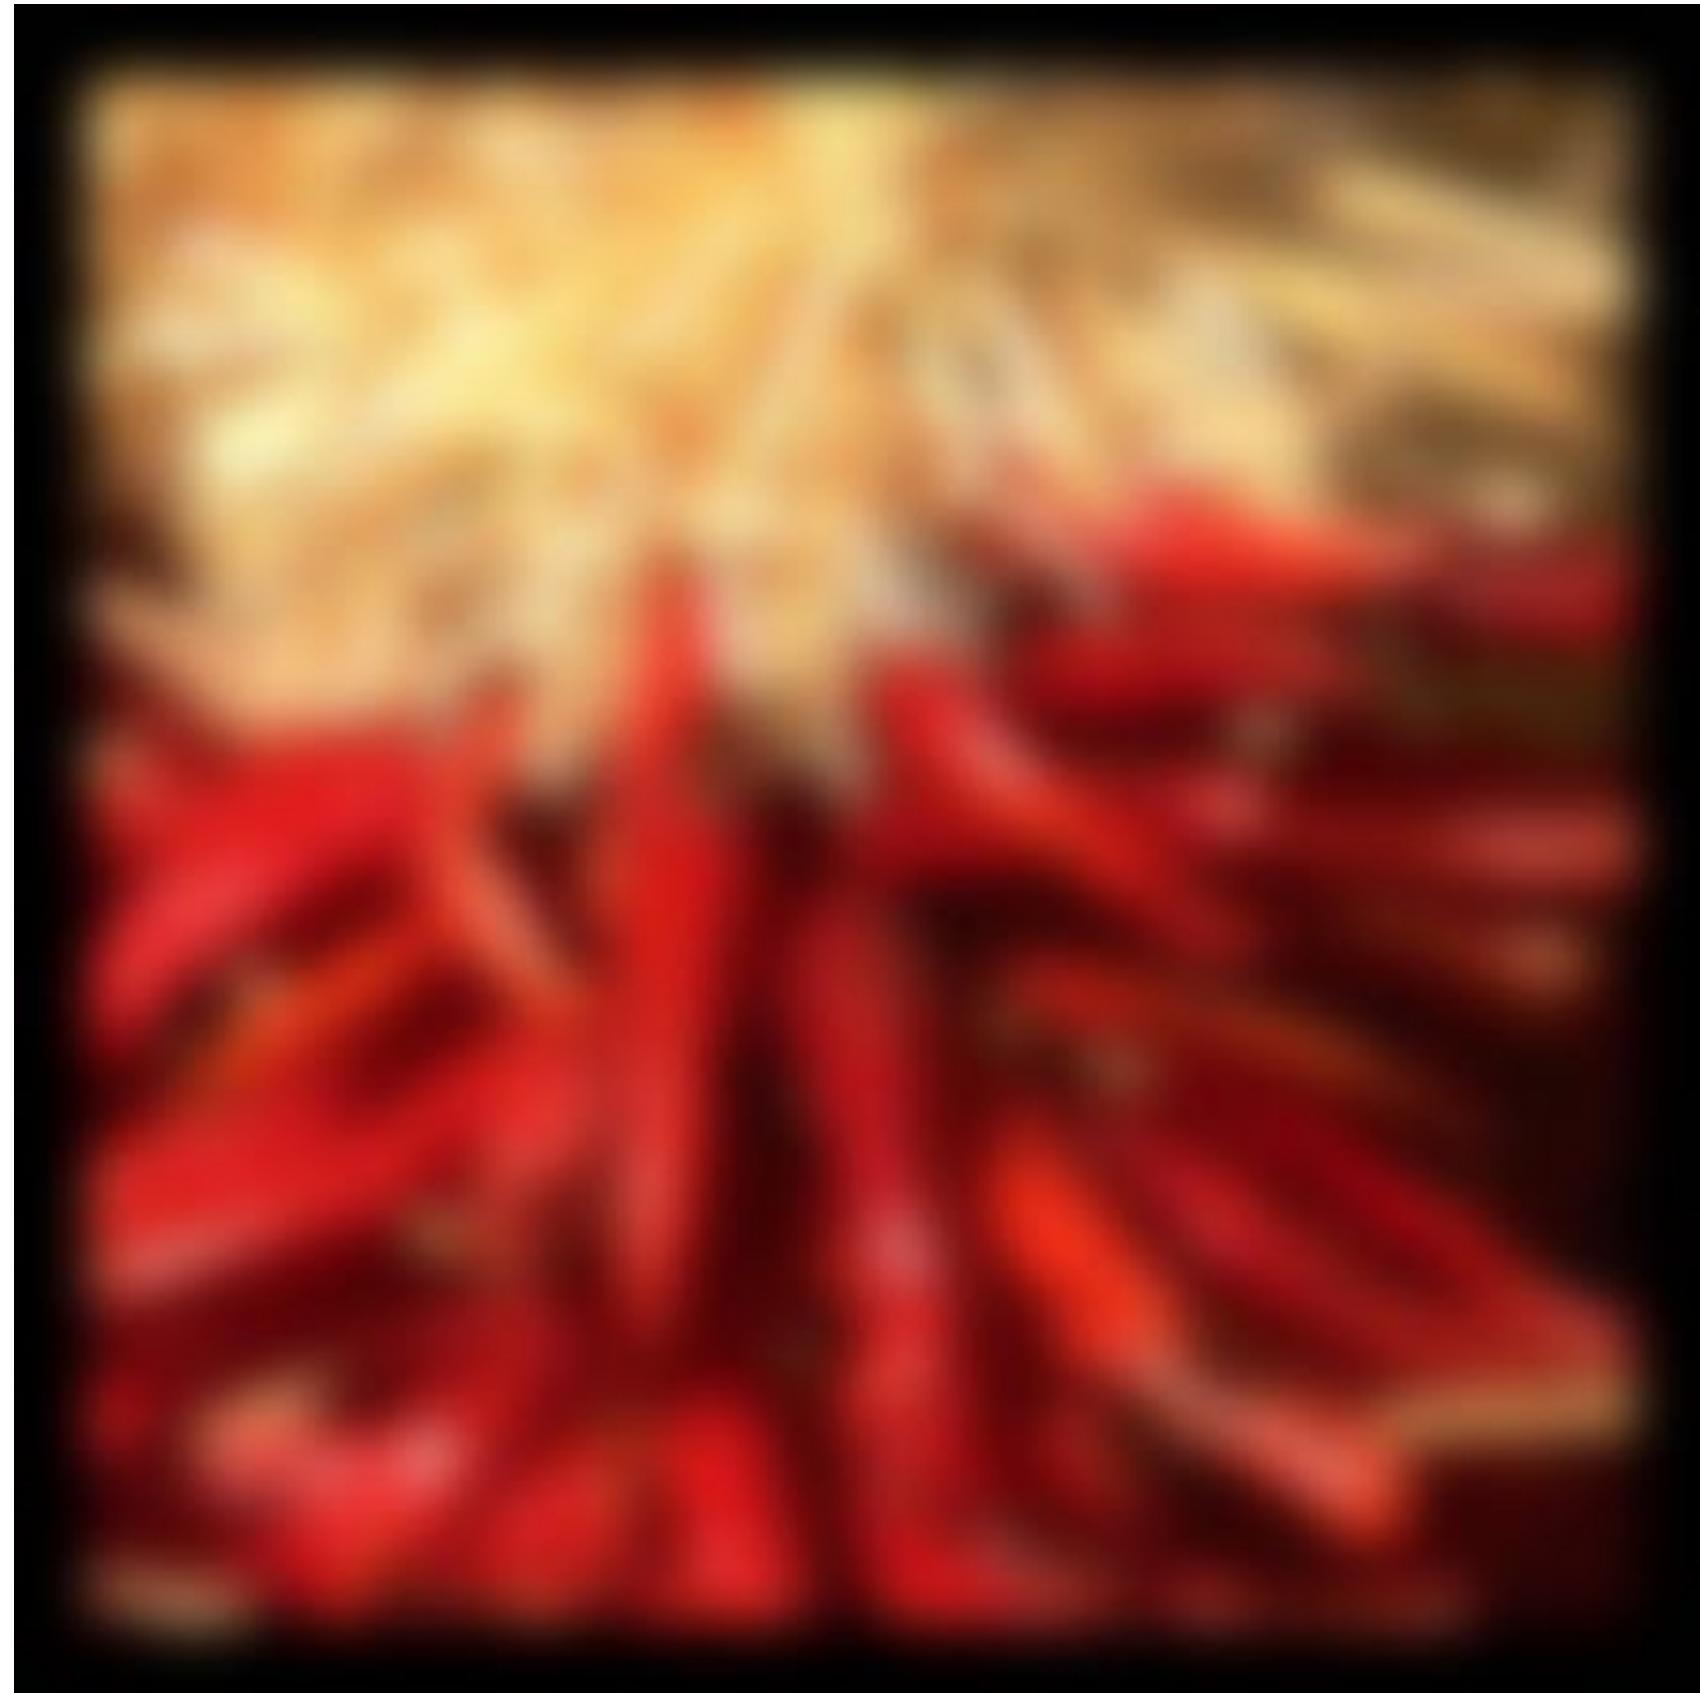
‘Boring’ Details

- When filter window falls off the edge of the image?
 - need to extrapolate
 - methods:
 - clip filter (black)
 - wrap around
 - copy edge
 - reflect across edge
 - vary filter near edge



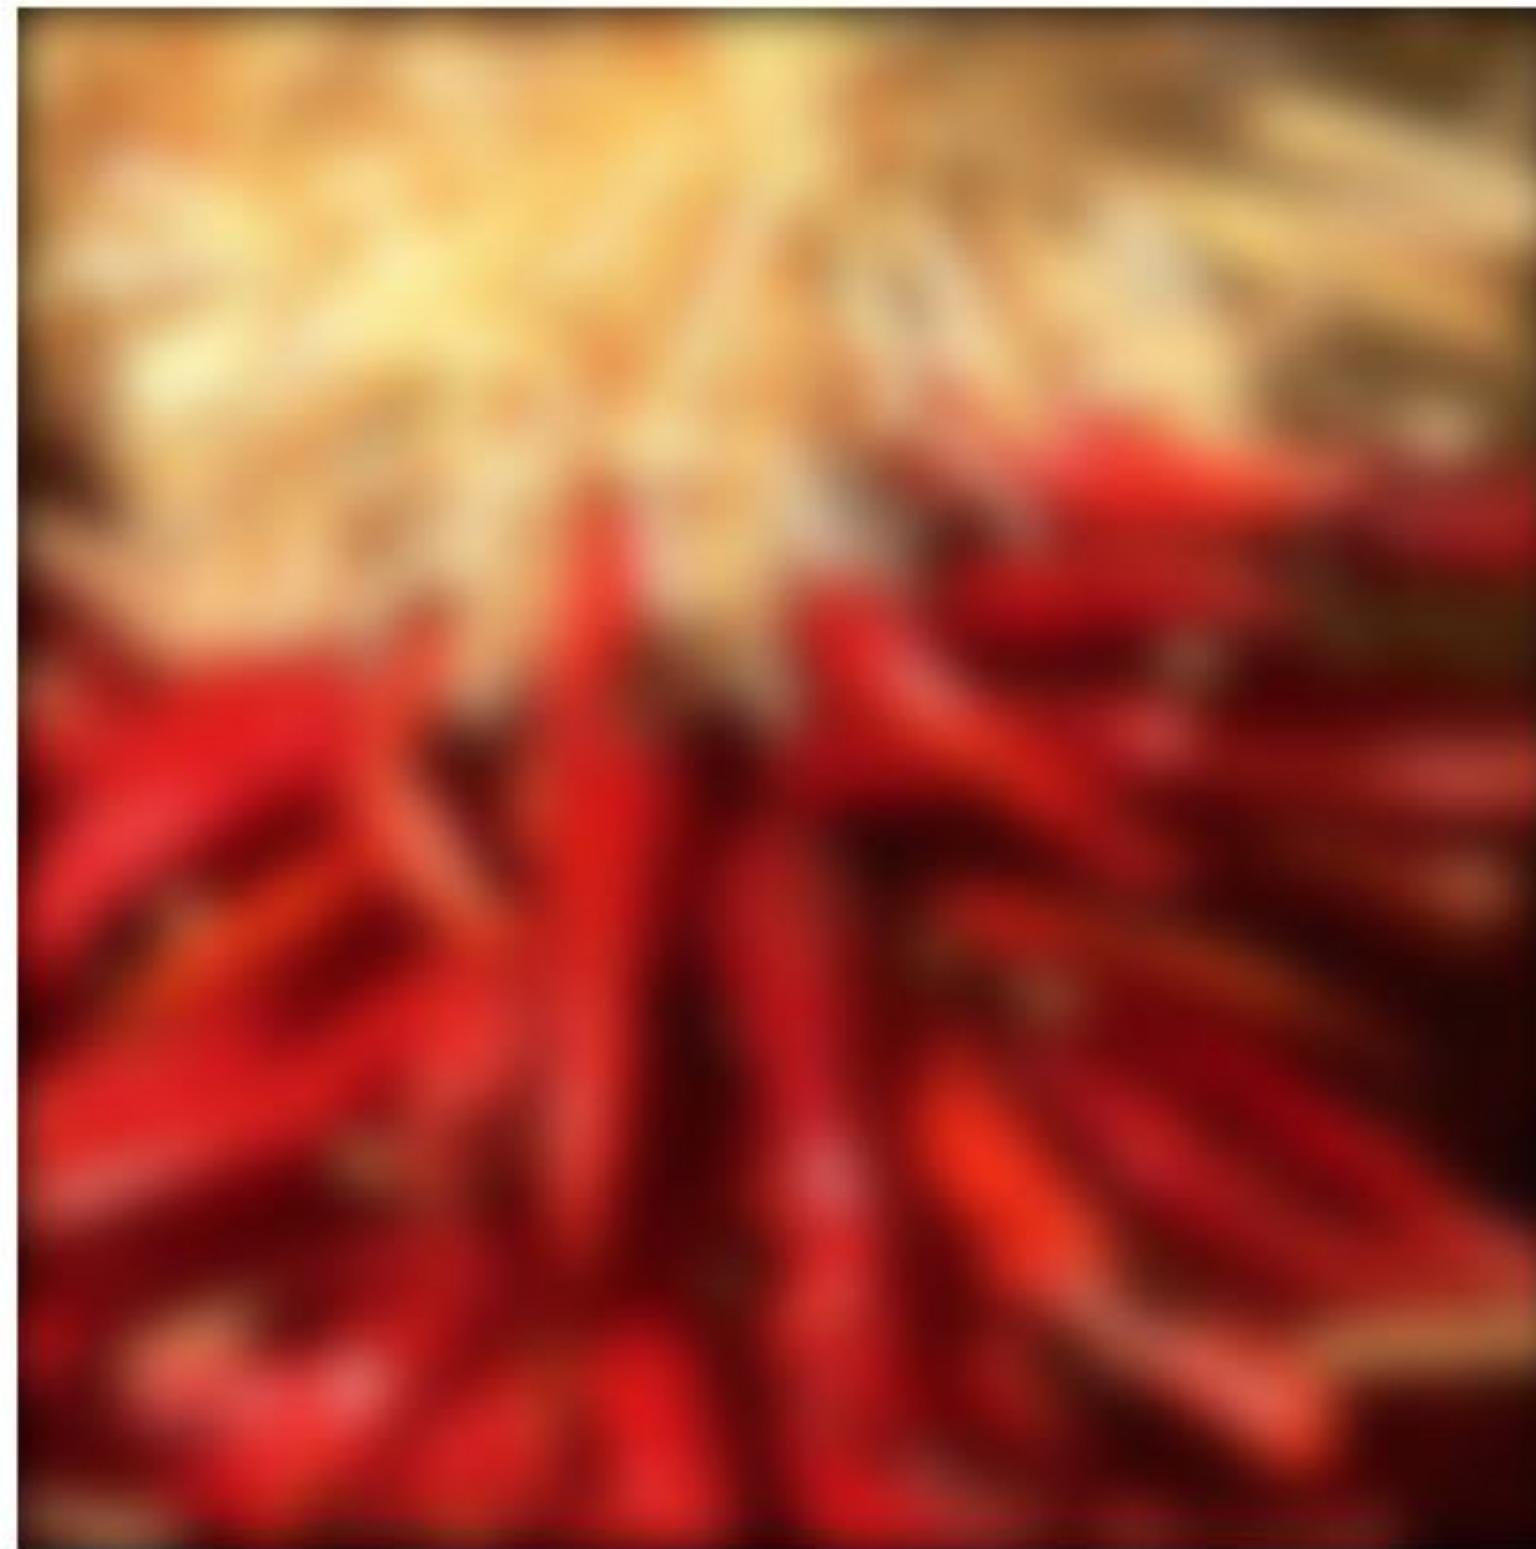
‘Boring’ Details

- When filter window falls off the edge of the image?
 - need to extrapolate
 - methods:
 - clip filter (black)
 - wrap around
 - copy edge
 - reflect across edge
 - vary filter near edge



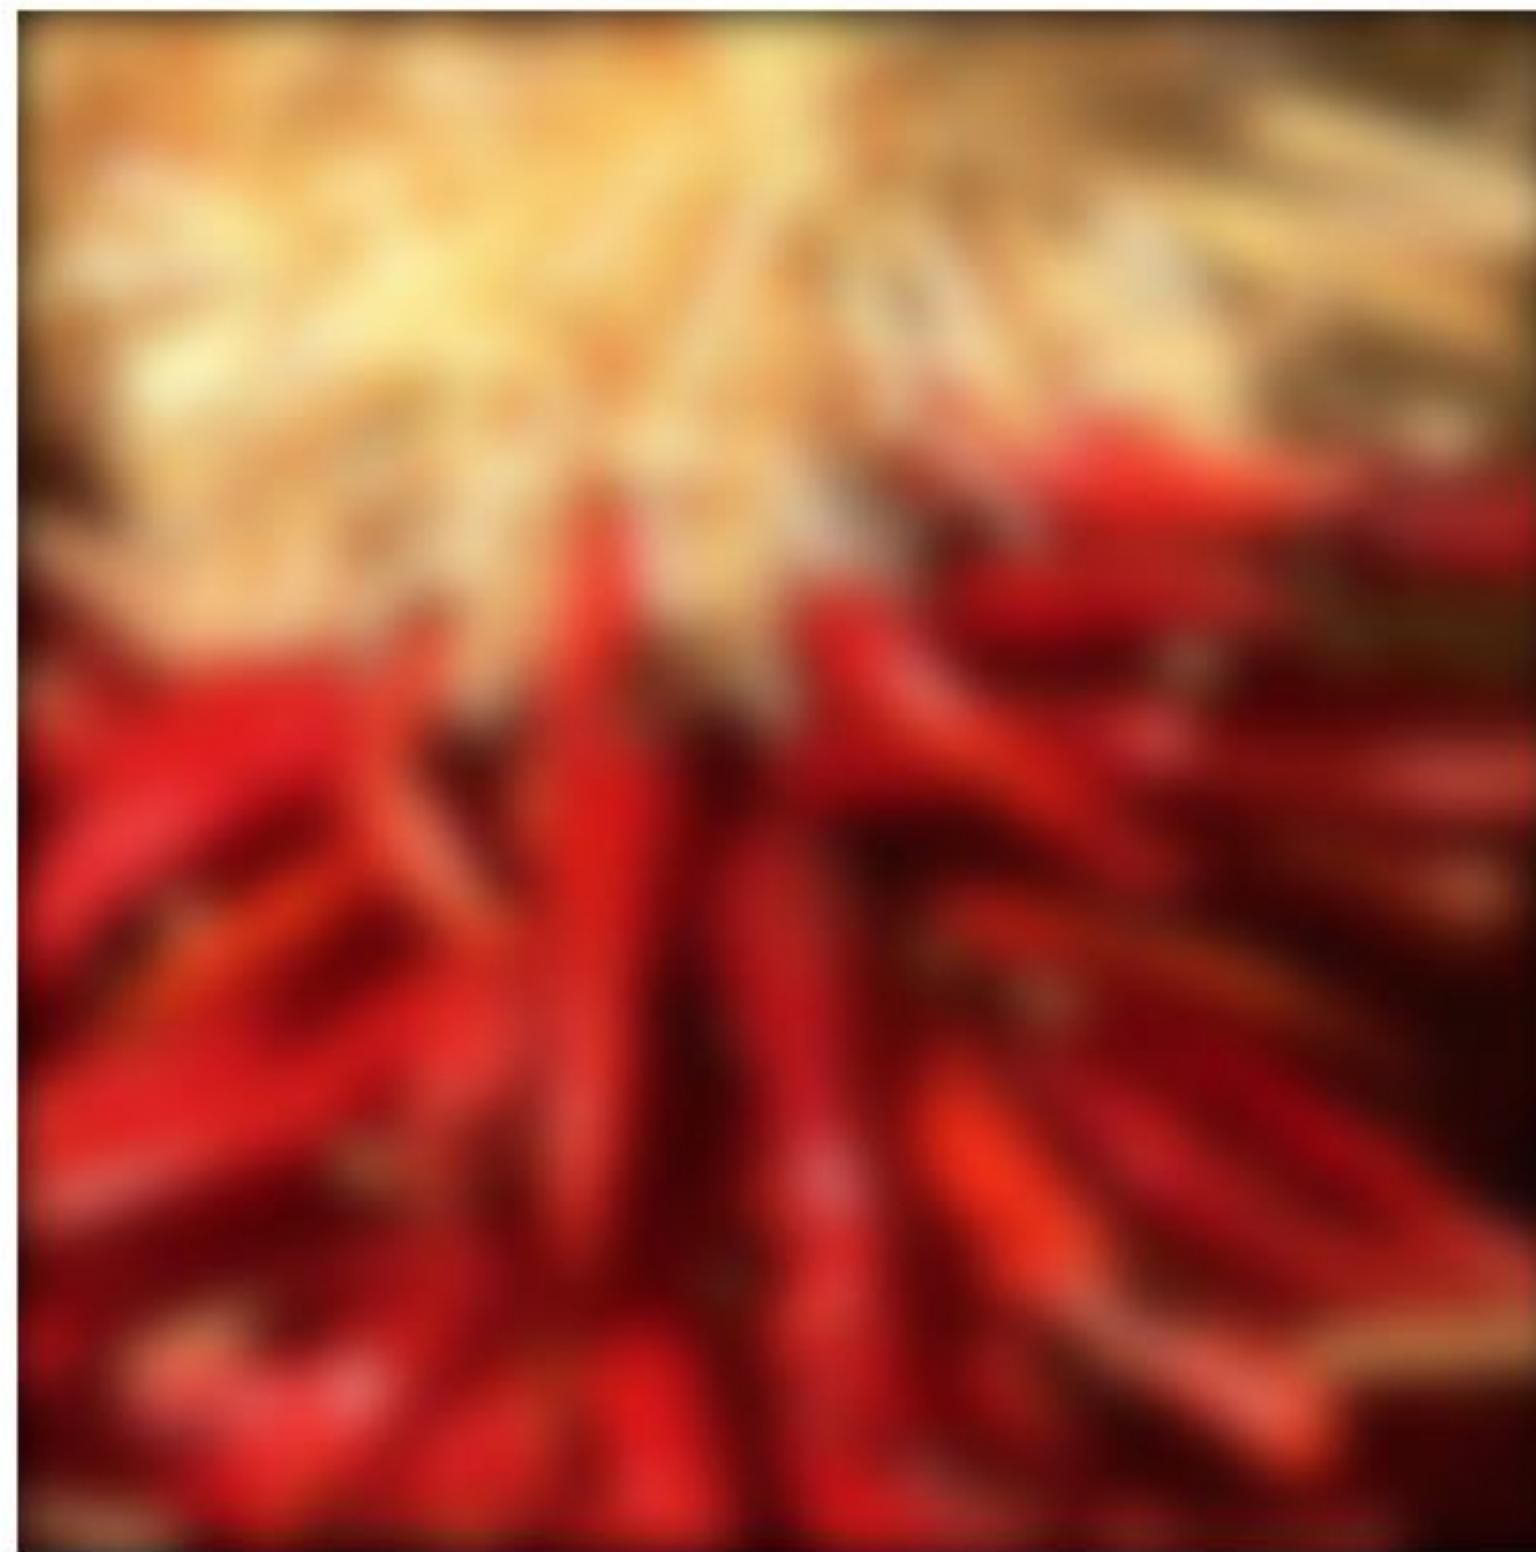
‘Boring’ Details

- When filter window falls off the edge of the image?
 - need to extrapolate
 - methods:
 - clip filter (black)
 - wrap around
 - copy edge
 - reflect across edge
 - vary filter near edge



‘Boring’ Details

- When filter window falls off the edge of the image?
 - need to extrapolate
 - methods:
 - clip filter (black)
 - wrap around
 - copy edge
 - reflect across edge
 - vary filter near edge



'Boring' Details

- When filter window falls off the edge of the image?

- need to extrapolate

- methods:

- clip filter (black)

- wrap around

- copy edge

- reflect across edge

- vary filter near edge

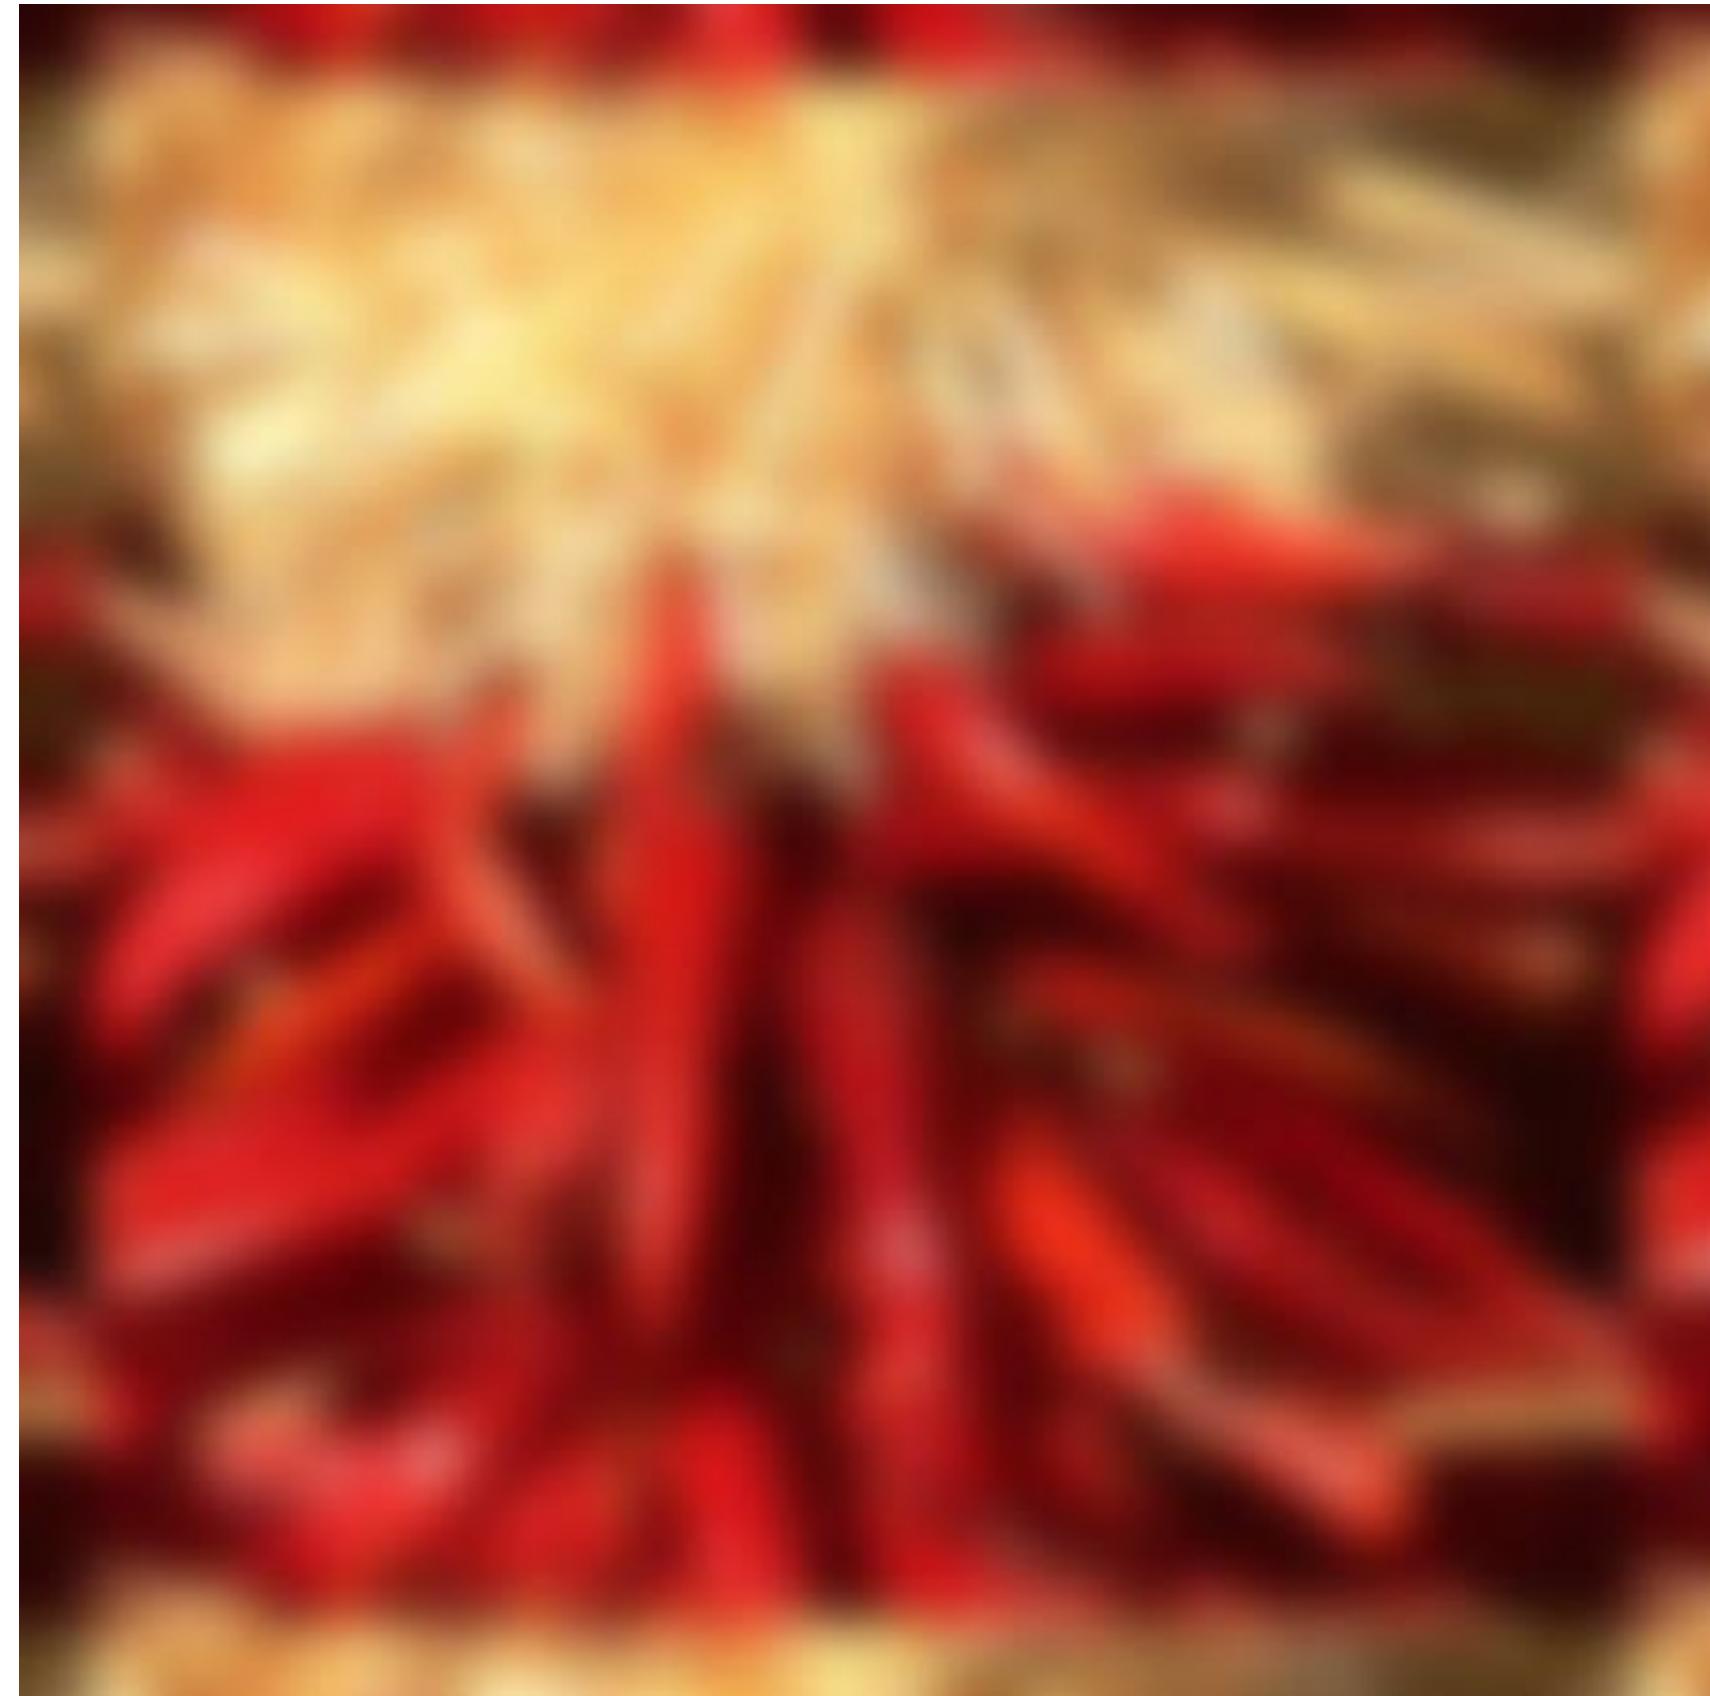
In this case,

not a very good strategy



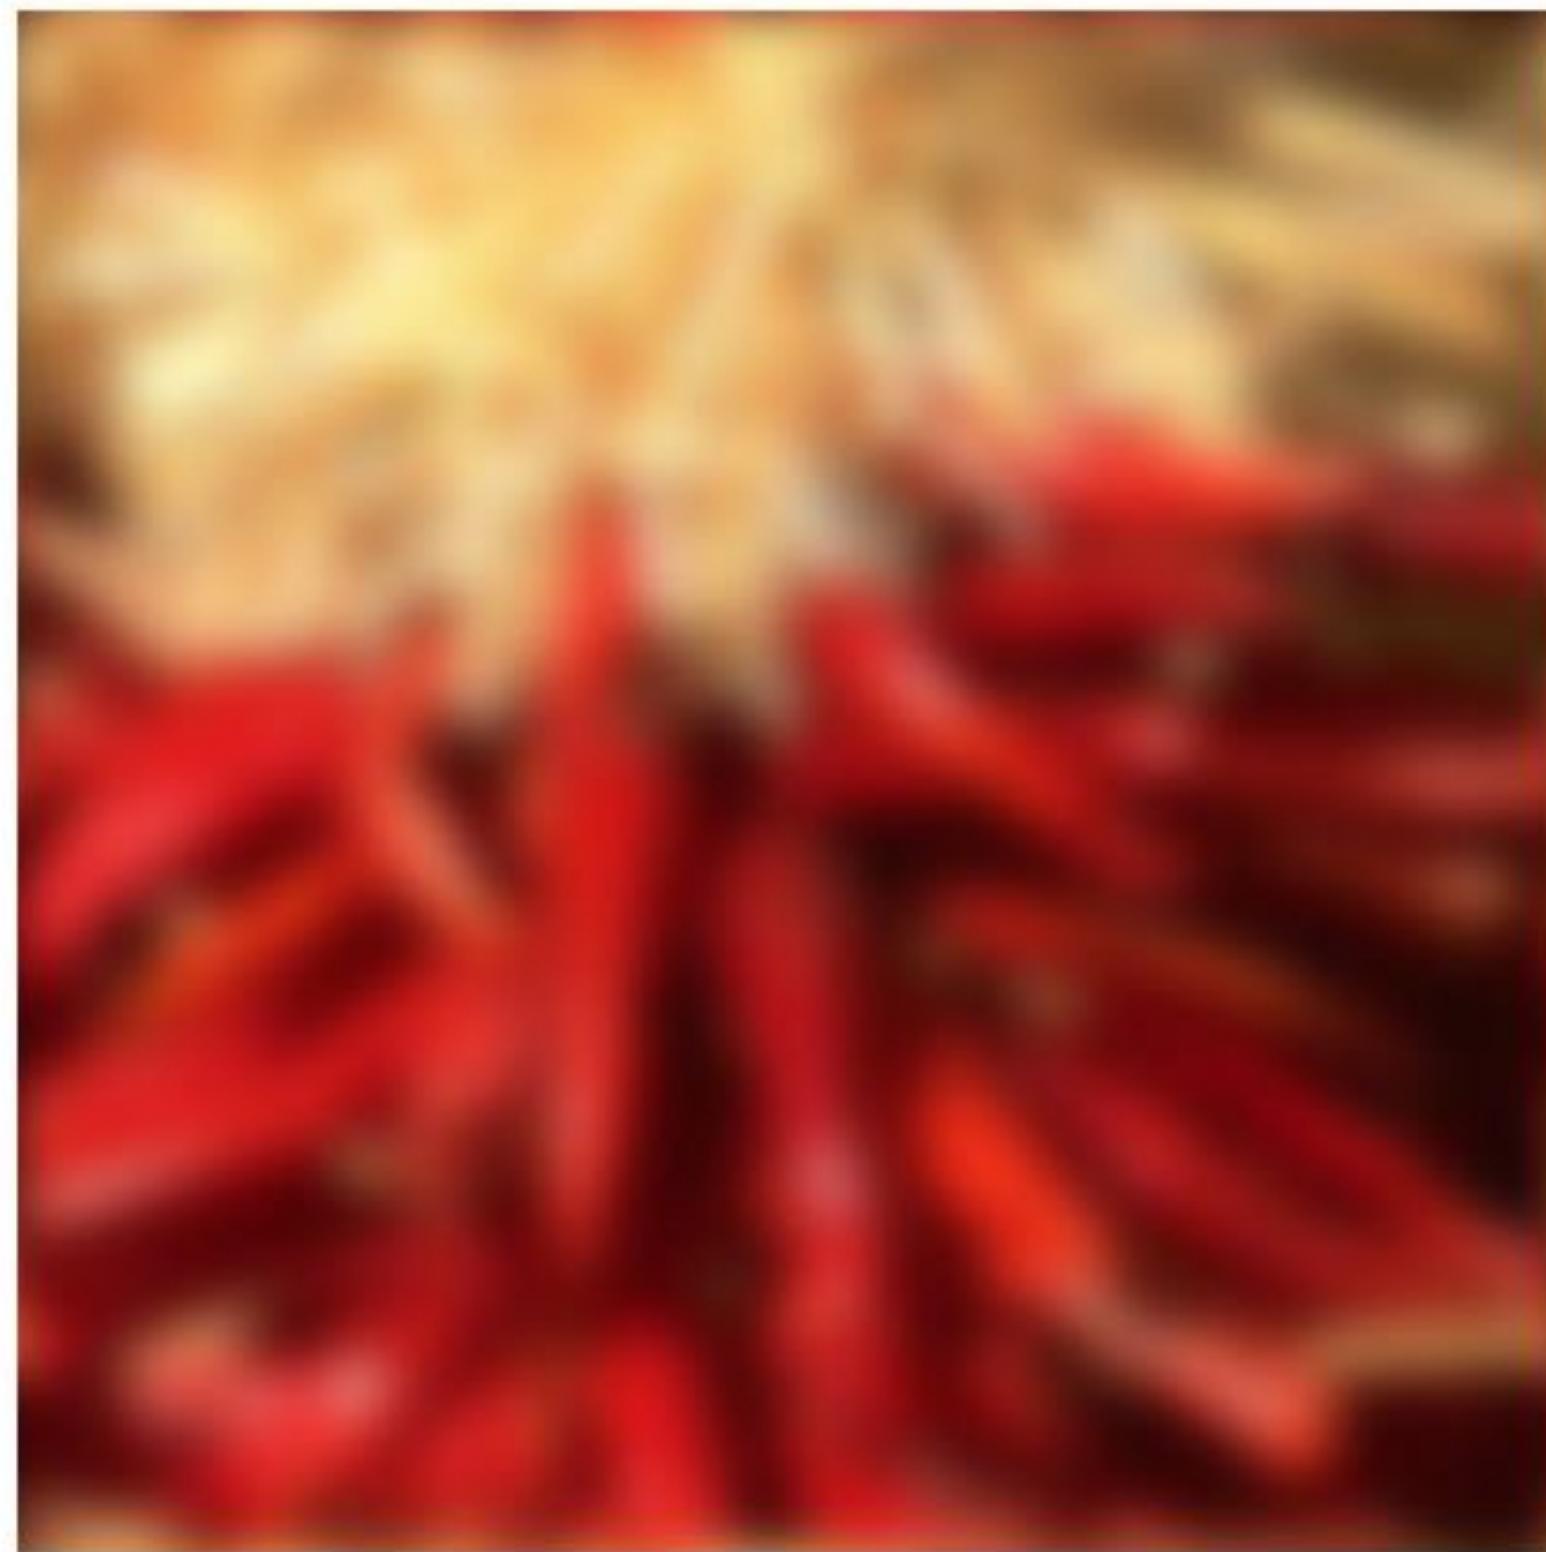
‘Boring’ Details

- When filter window falls off the edge of the image?
 - need to extrapolate
 - methods:
 - clip filter (black)
 - wrap around
 - copy edge
 - reflect across edge
 - vary filter near edge



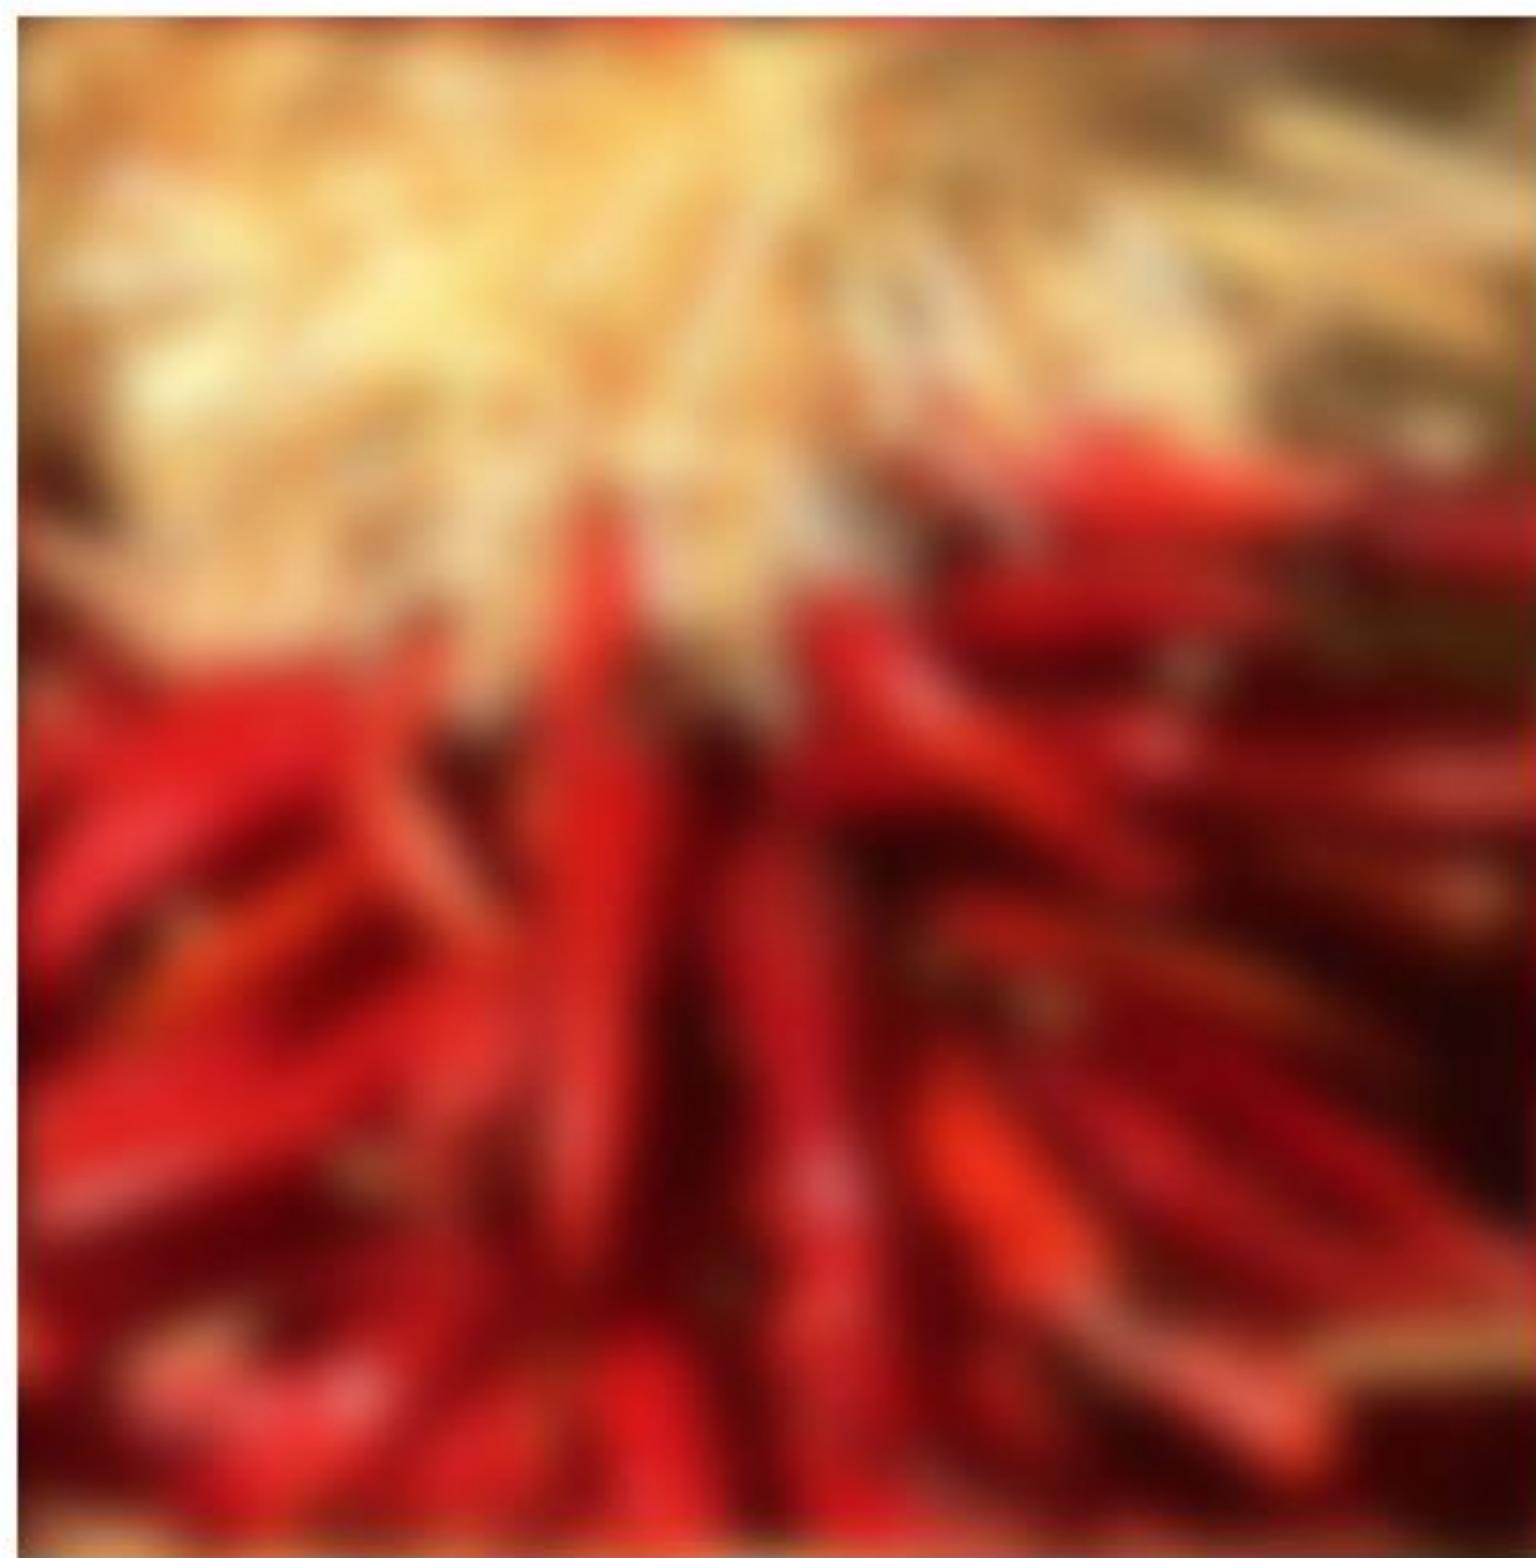
‘Boring’ Details

- When filter window falls off the edge of the image?
 - need to extrapolate
 - methods:
 - clip filter (black)
 - wrap around
 - copy edge
 - reflect across edge
 - vary filter near edge



‘Boring’ Details

- When filter window falls off the edge of the image?
 - need to extrapolate
 - methods:
 - clip filter (black)
 - wrap around
 - copy edge
 - reflect across edge
 - vary filter near edge



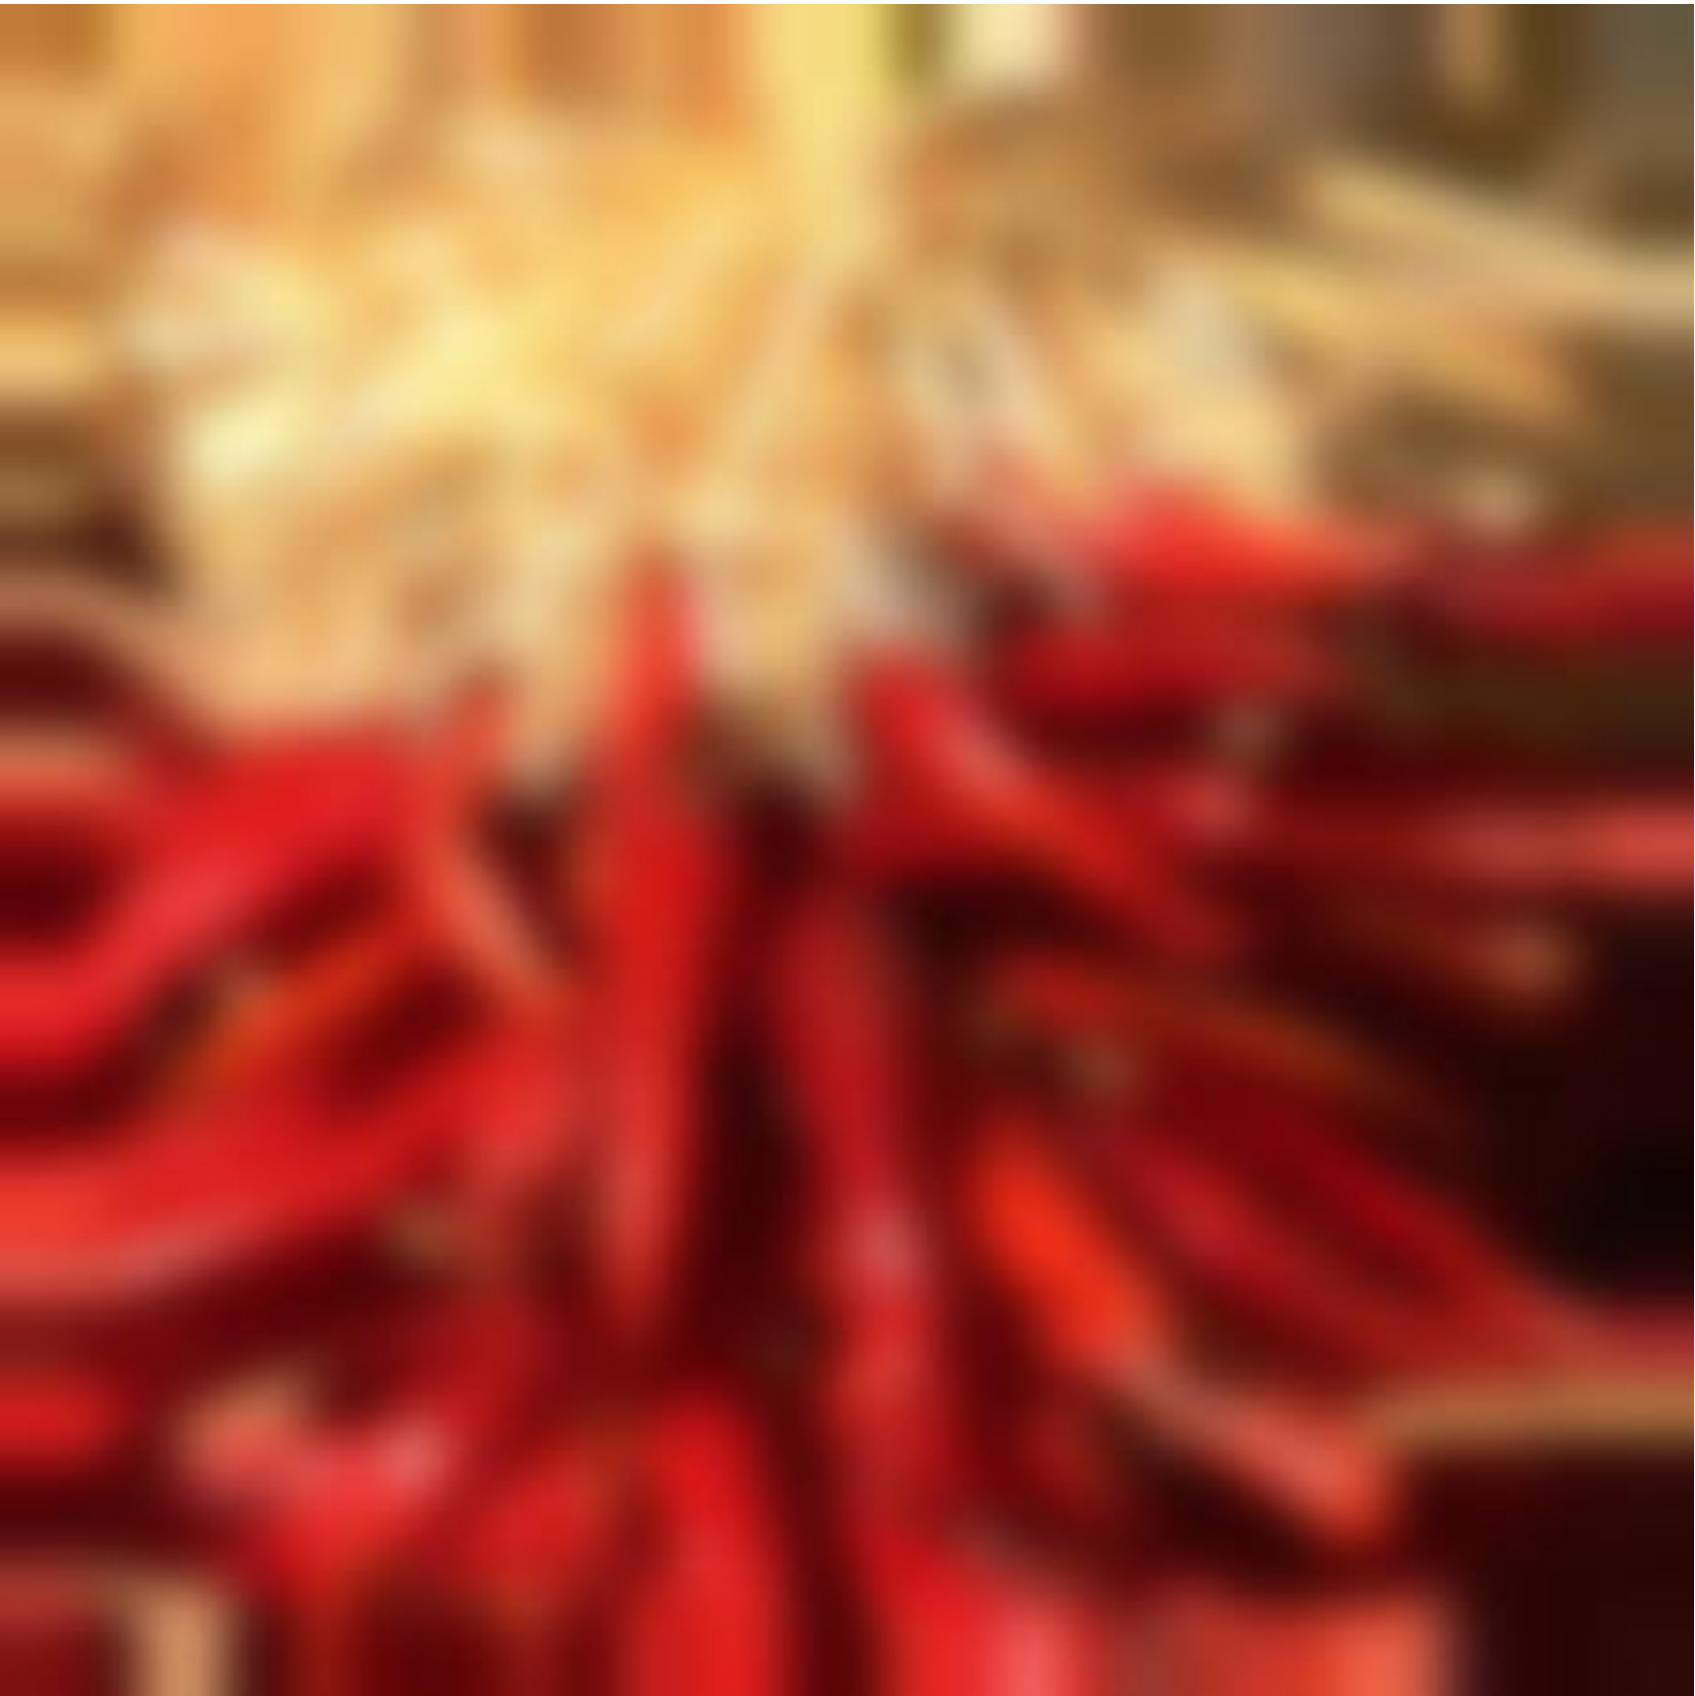
‘Boring’ Details

- When filter window falls off the edge of the image?
 - need to extrapolate
 - methods:
 - clip filter (black)
 - wrap around
 - copy edge
 - reflect across edge
 - vary filter near edge



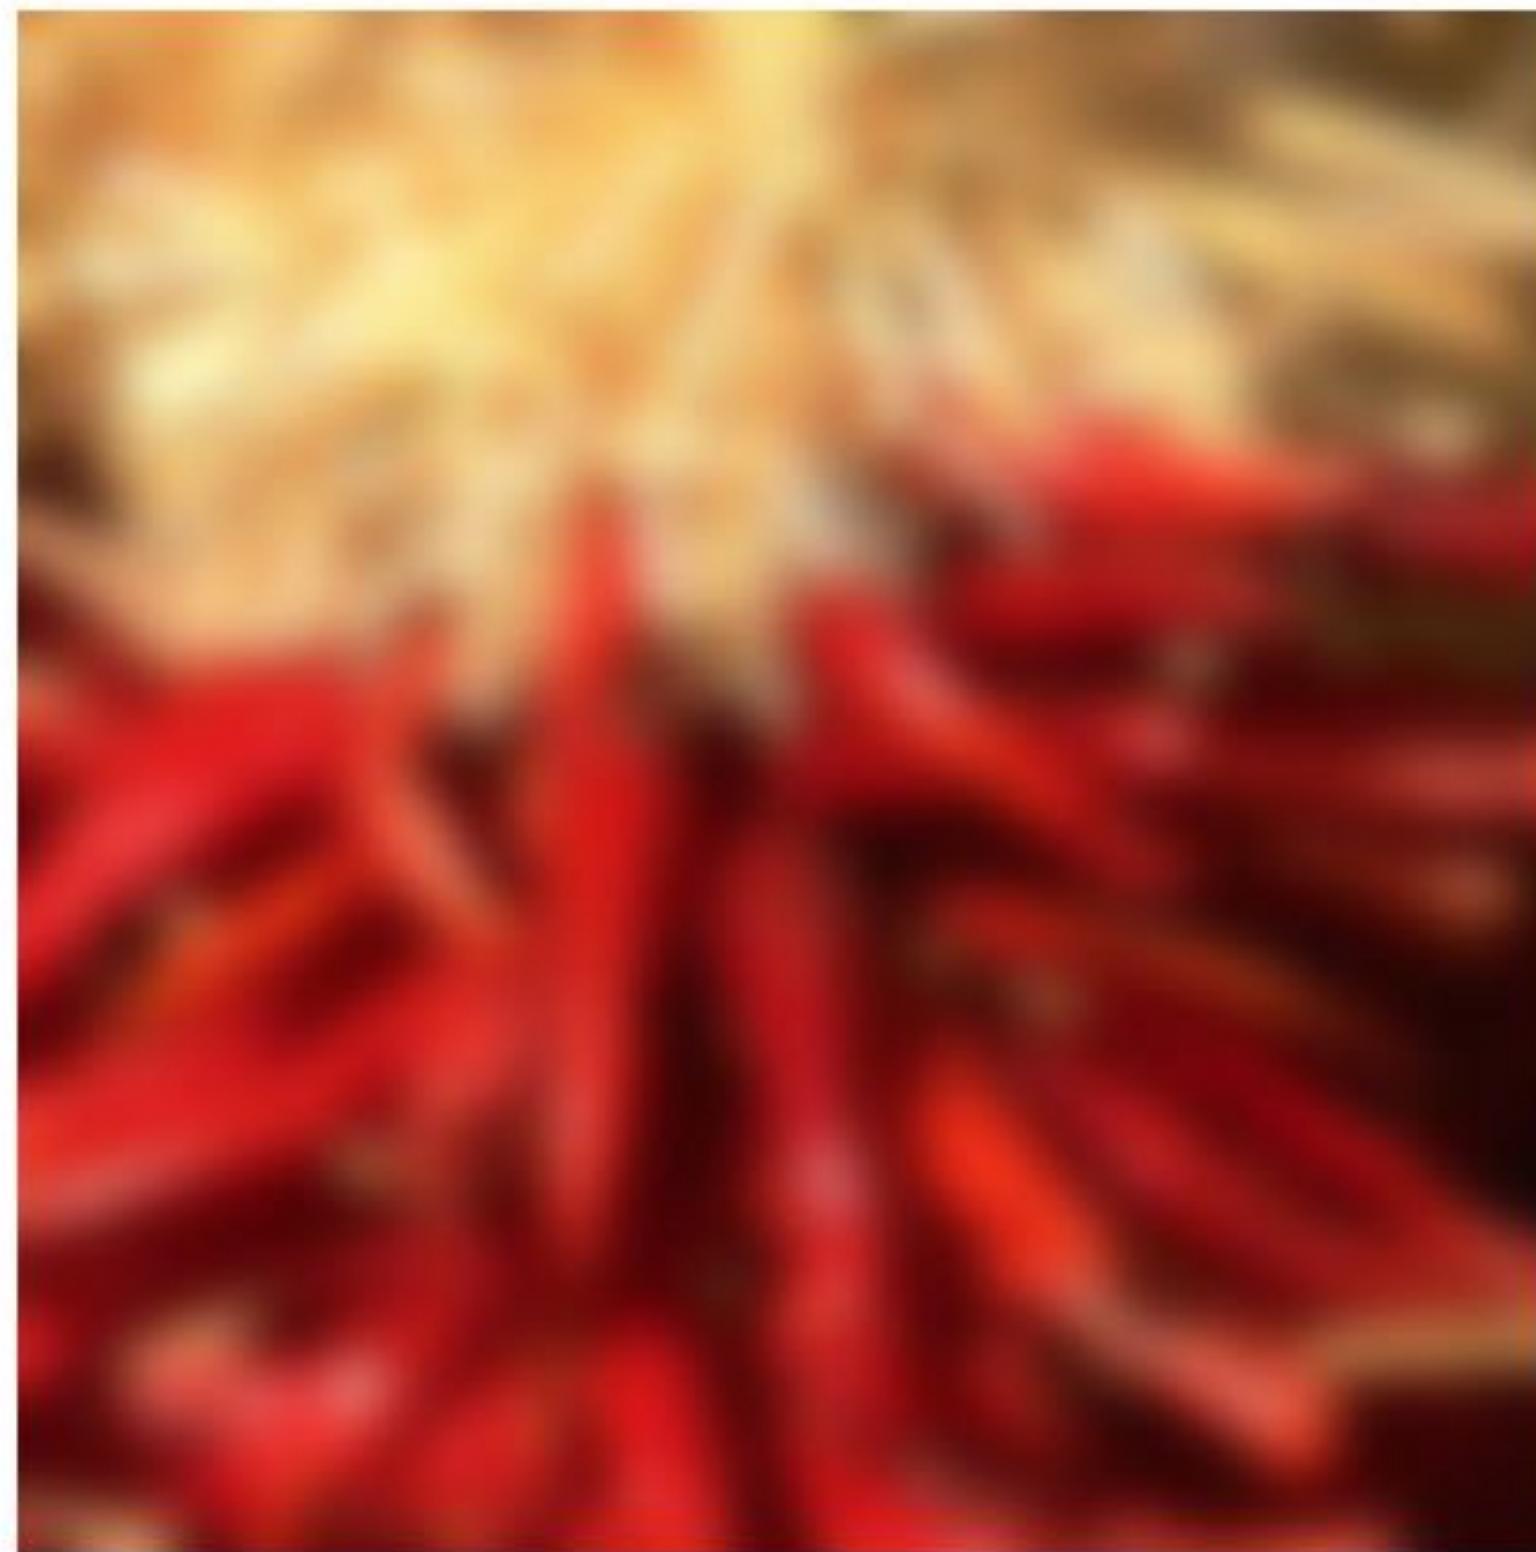
‘Boring’ Details

- When filter window falls off the edge of the image?
 - need to extrapolate
 - methods:
 - clip filter (black)
 - wrap around
 - copy edge
 - reflect across edge
 - vary filter near edge



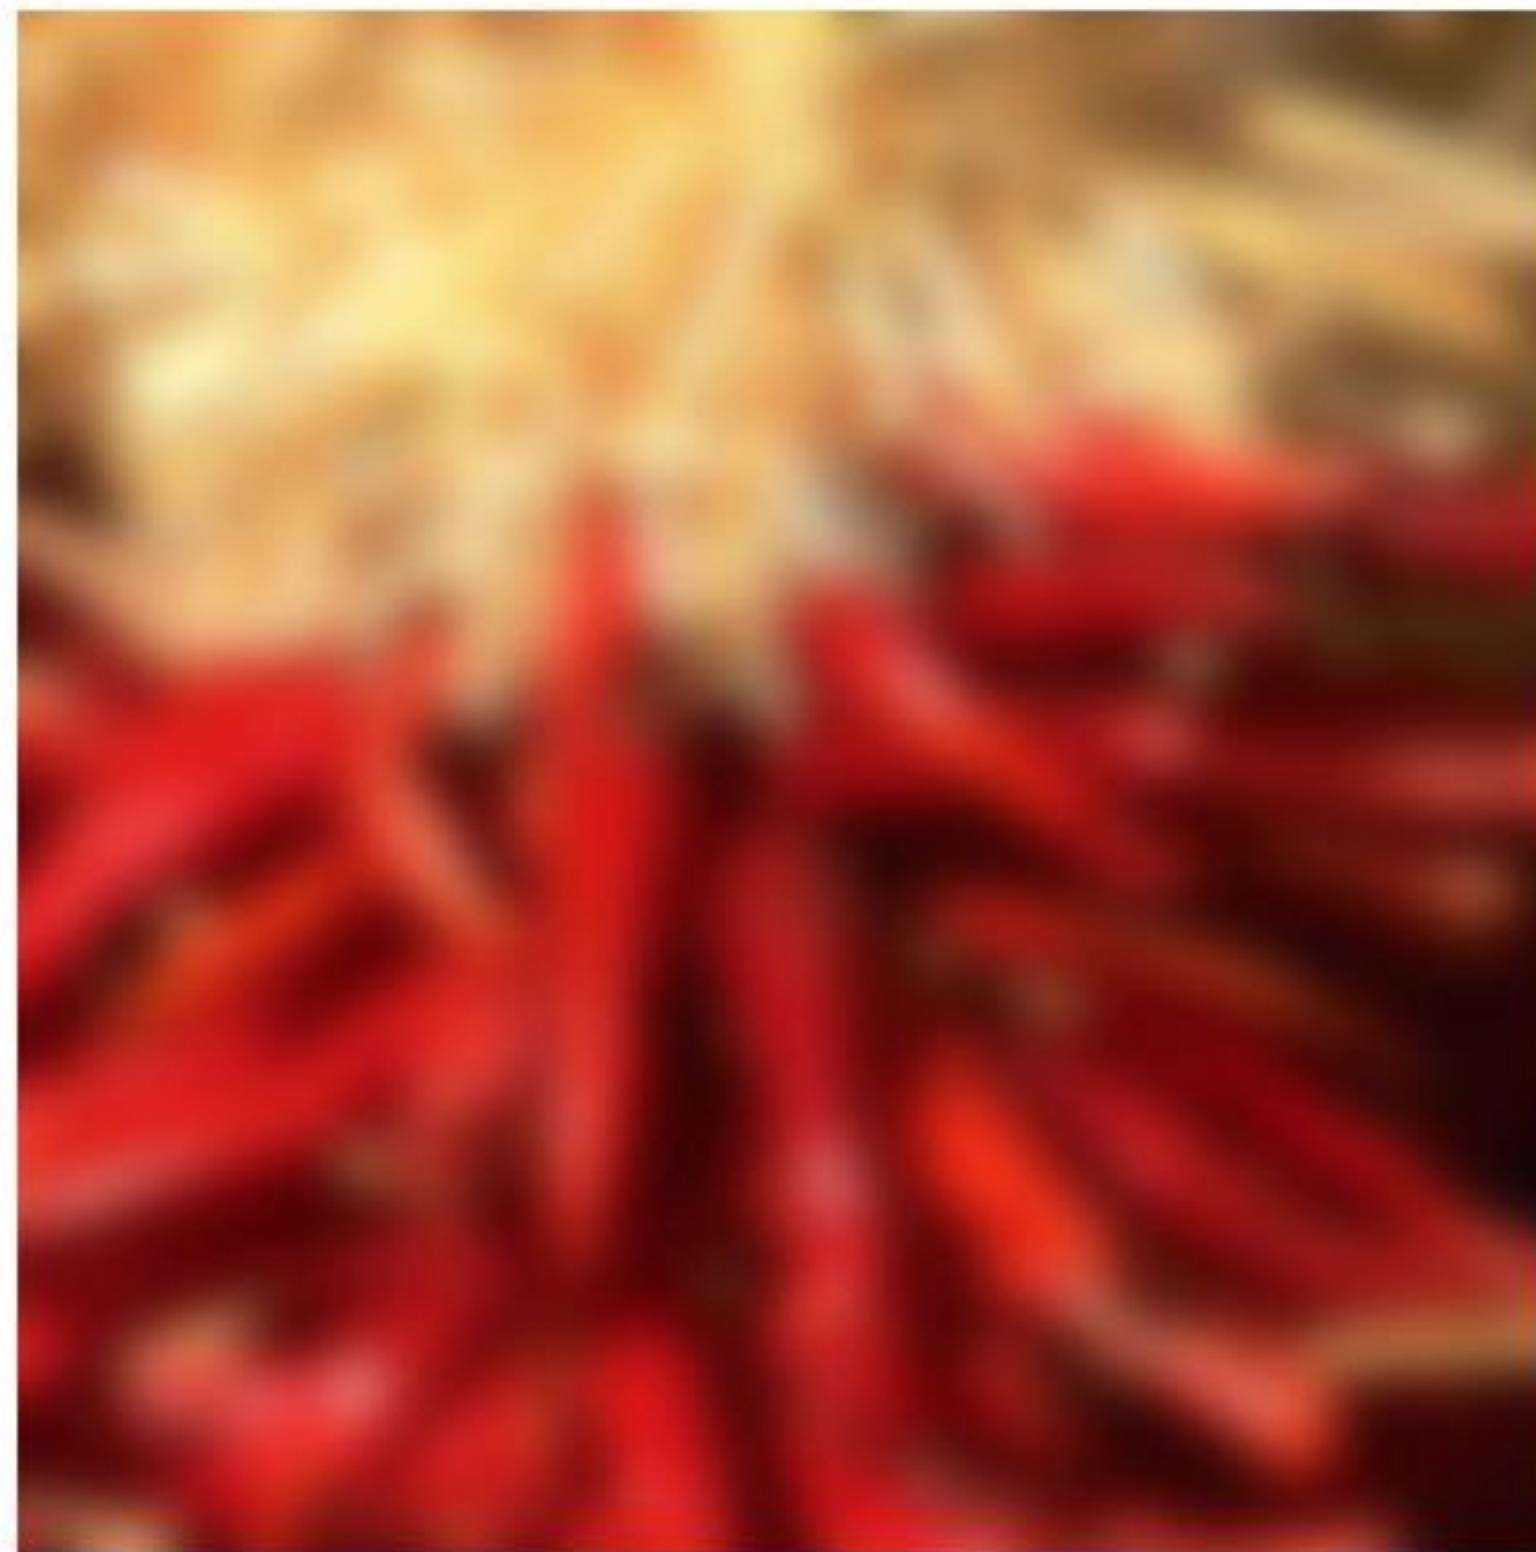
‘Boring’ Details

- When filter window falls off the edge of the image?
 - need to extrapolate
 - methods:
 - clip filter (black)
 - wrap around
 - copy edge
 - reflect across edge
 - vary filter near edge



‘Boring’ Details

- When filter window falls off the edge of the image?
 - need to extrapolate
 - methods:
 - clip filter (black)
 - wrap around
 - copy edge
 - reflect across edge
 - vary filter near edge



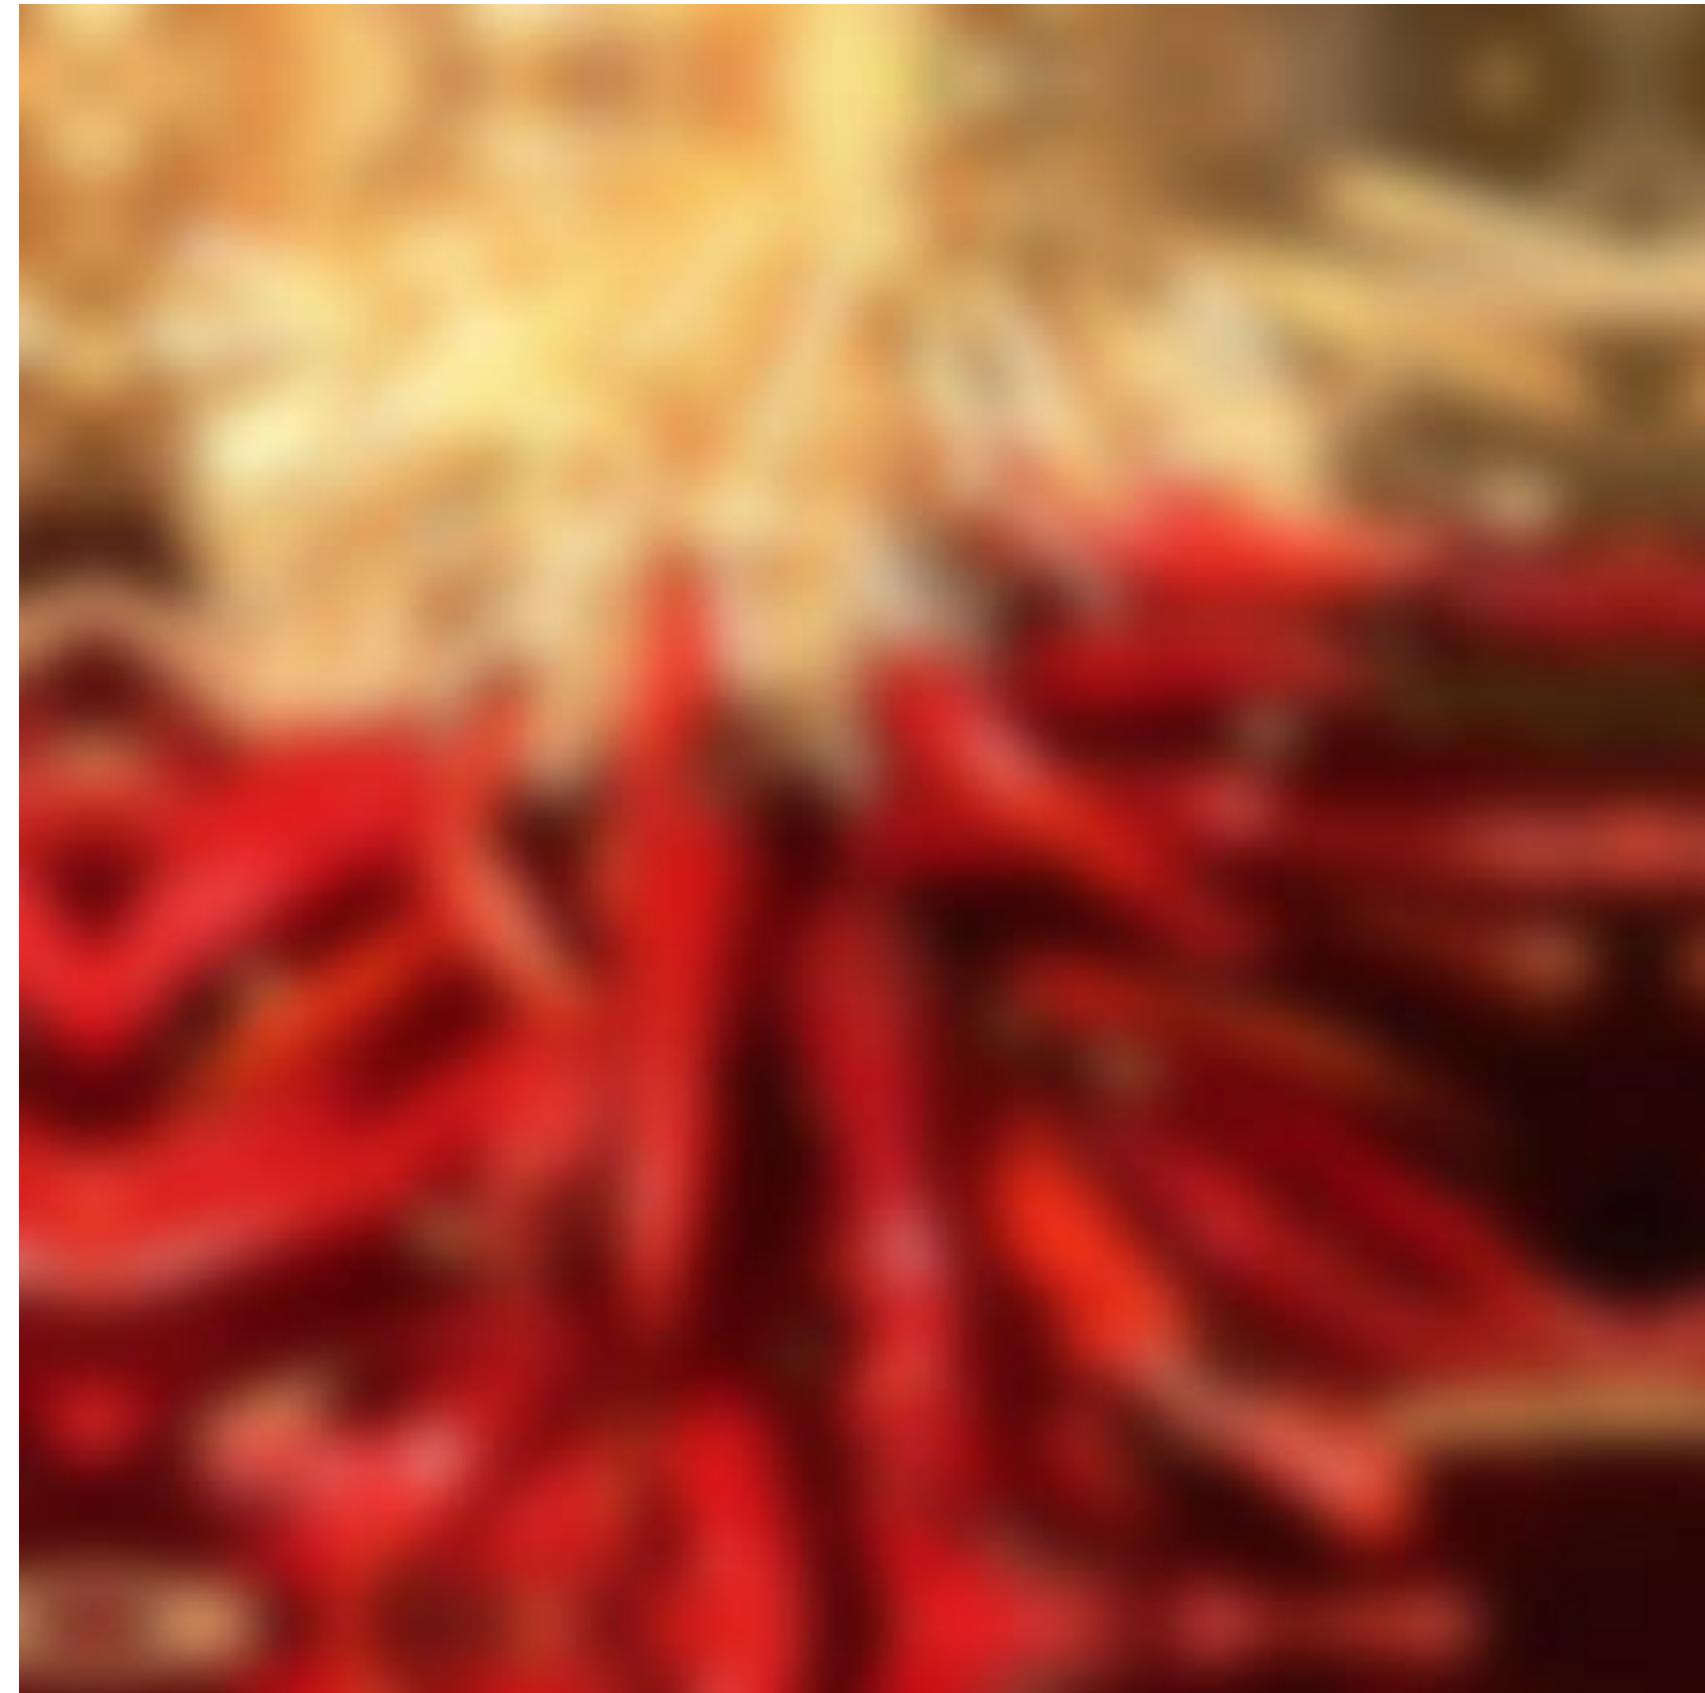
‘Boring’ Details

- When filter window falls off the edge of the image?
 - need to extrapolate
 - methods:
 - clip filter (black)
 - wrap around
 - copy edge
 - reflect across edge
 - vary filter near edge



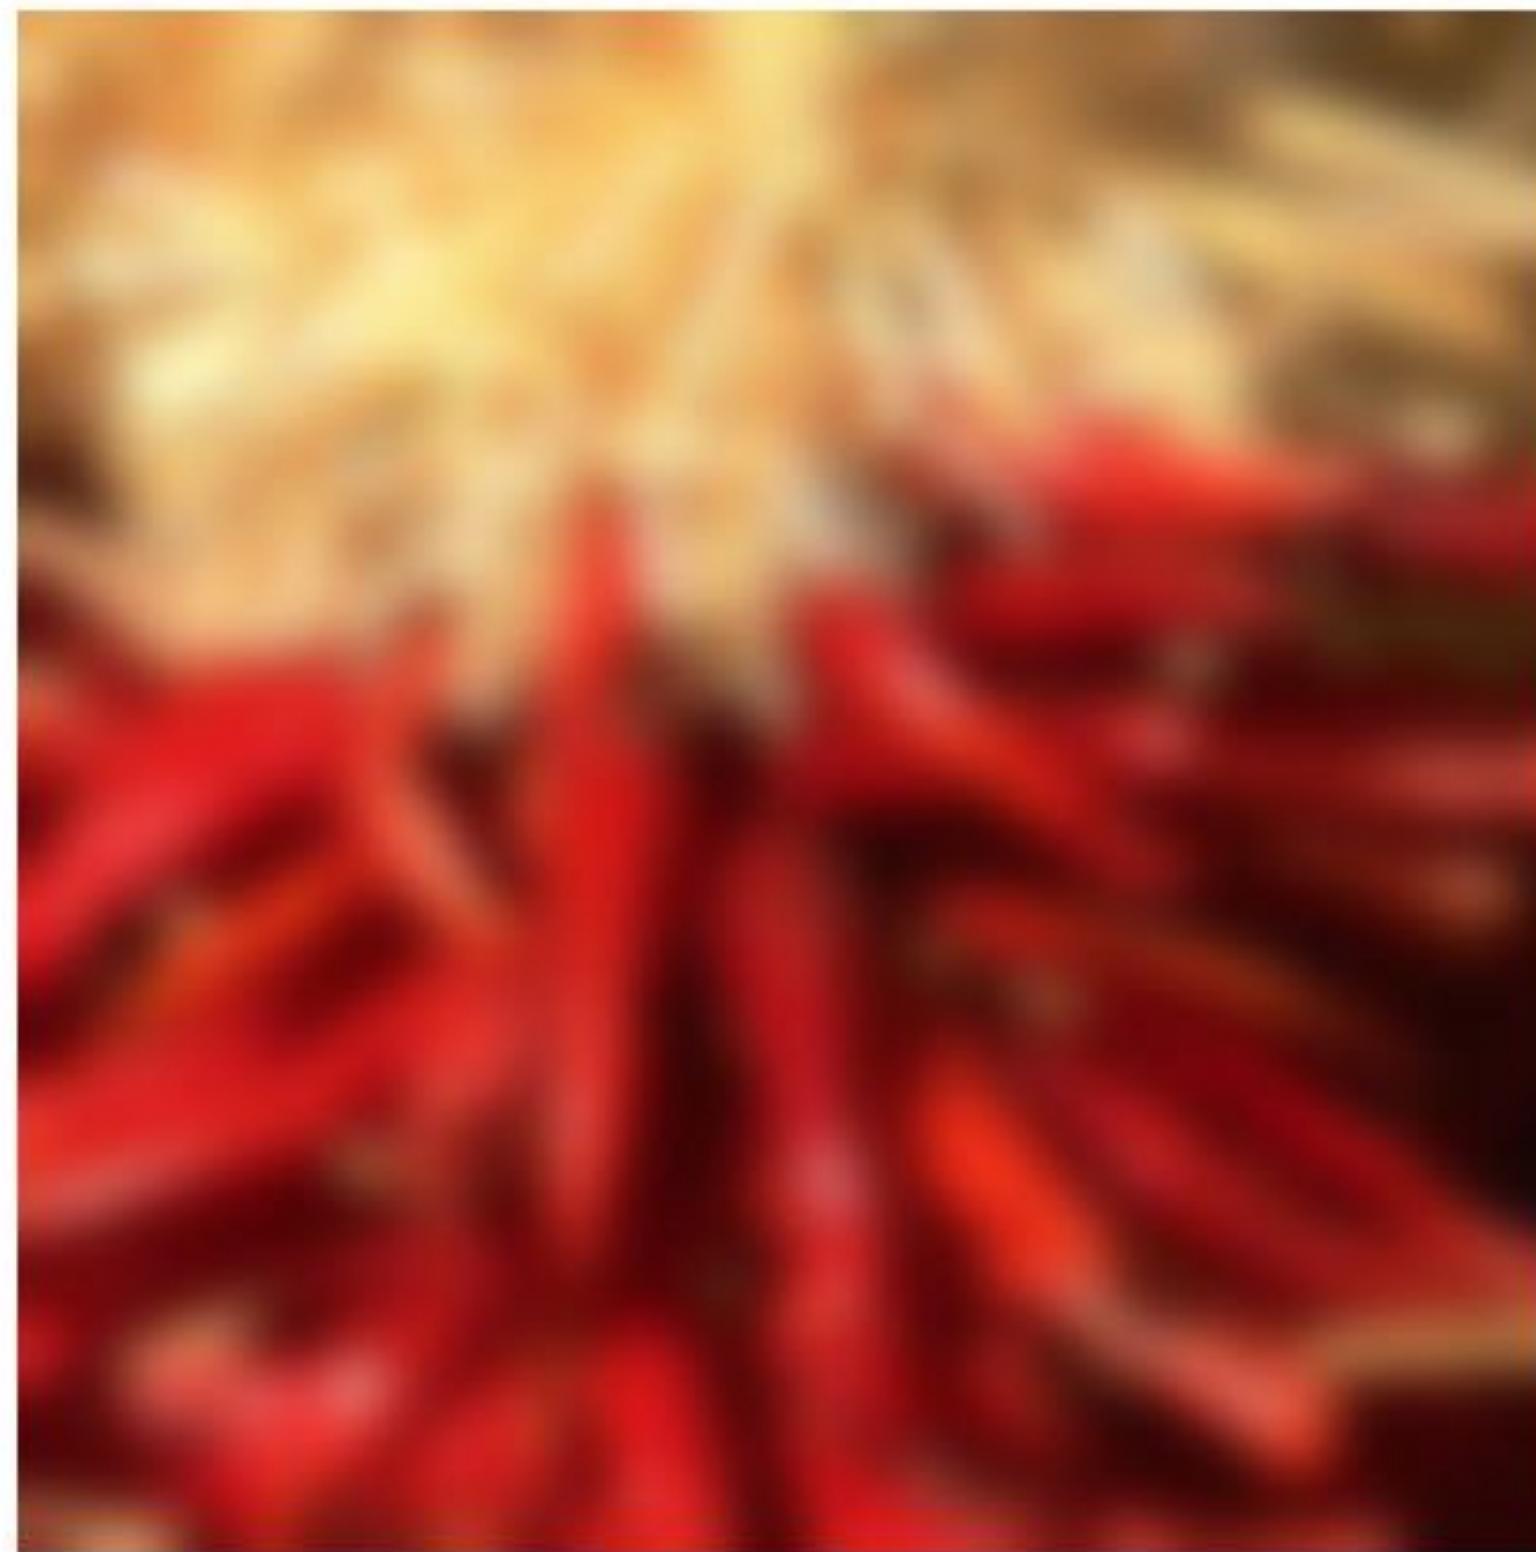
‘Boring’ Details

- When filter window falls off the edge of the image?
 - need to extrapolate
 - methods:
 - clip filter (black)
 - wrap around
 - copy edge
 - reflect across edge
 - vary filter near edge



‘Boring’ Details

- When filter window falls off the edge of the image?
 - need to extrapolate
 - methods:
 - clip filter (black)
 - wrap around
 - copy edge
 - reflect across edge
 - vary filter near edge



Smoothing Kernels (Low-pass filters)

Mean filter:

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

Smoothing Kernels (Low-pass filters)

Mean filter:

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

giving more importance
to the center &

Weighted smoothing filters:

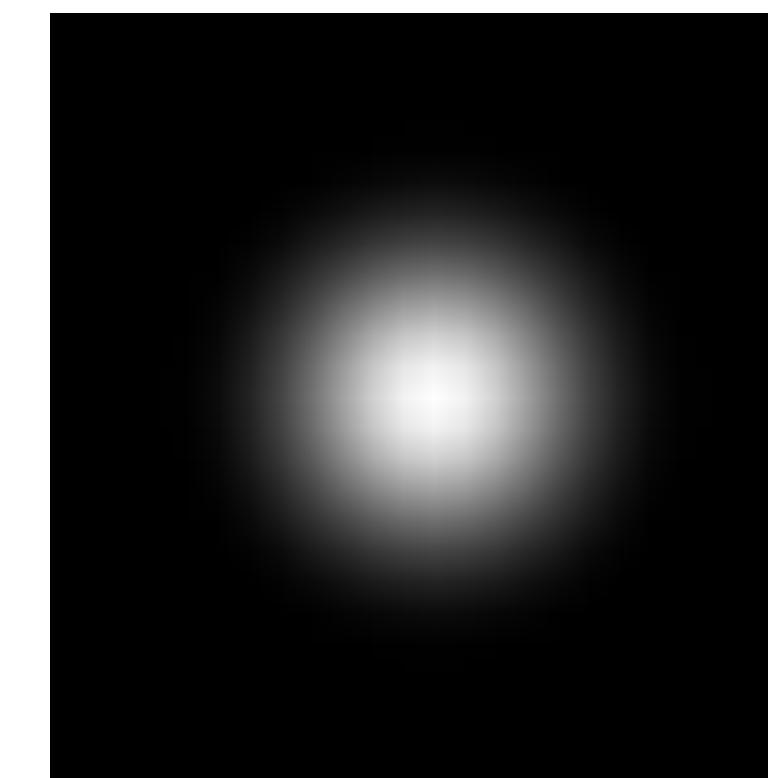
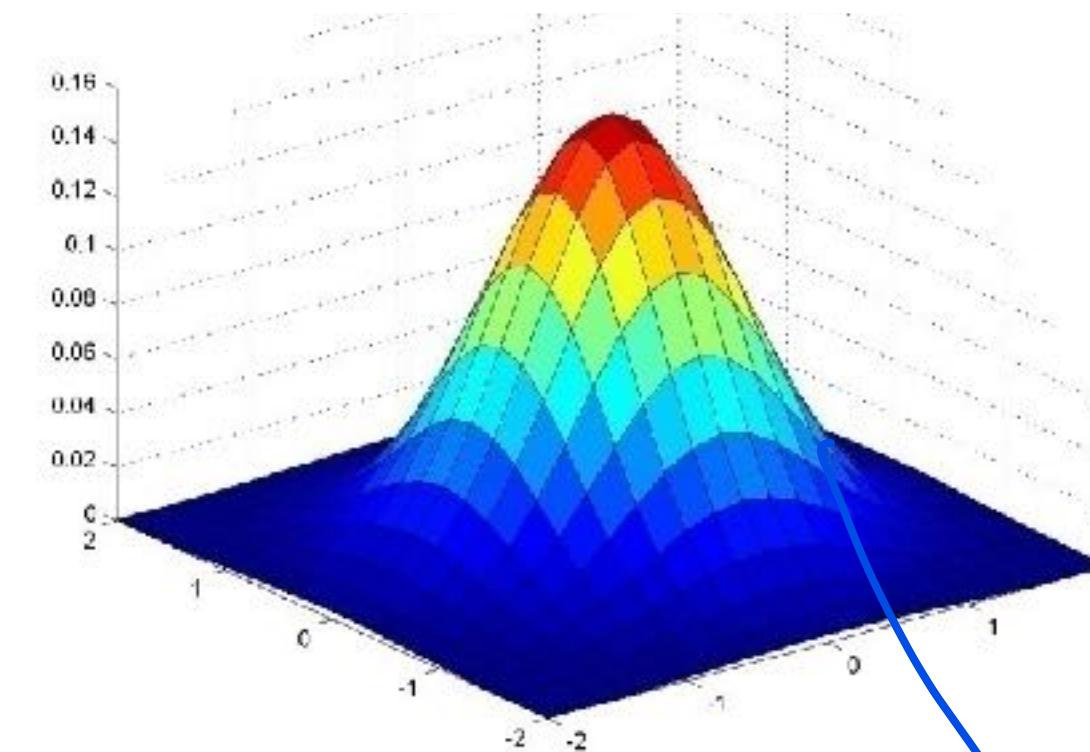
$$\frac{1}{10} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 2 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

$$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$$

(a bit like
gaussian)

Gaussian Kernel

- Weight contributions of neighbouring pixels related to nearness

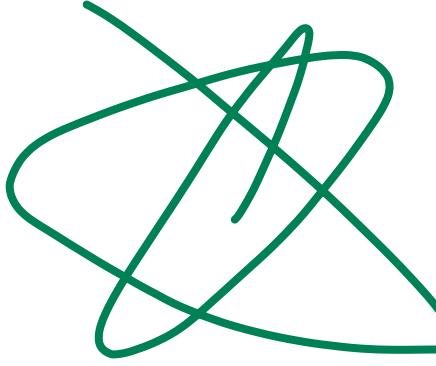


0.003	0.013	0.022	0.013	0.003
0.013	0.059	0.097	0.059	0.013
0.022	0.097	0.159	0.097	0.022
0.013	0.059	0.097	0.059	0.013
0.003	0.013	0.022	0.013	0.003

$$G_\sigma = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

- Constant factor in the front makes volume sum to 1
 - (subject to rounding errors, so better to normalize by sum of matrix)

Review of Sigma (σ)



$\sigma = \text{sigma} = \text{standard deviation} = \sqrt{\text{variance}}$

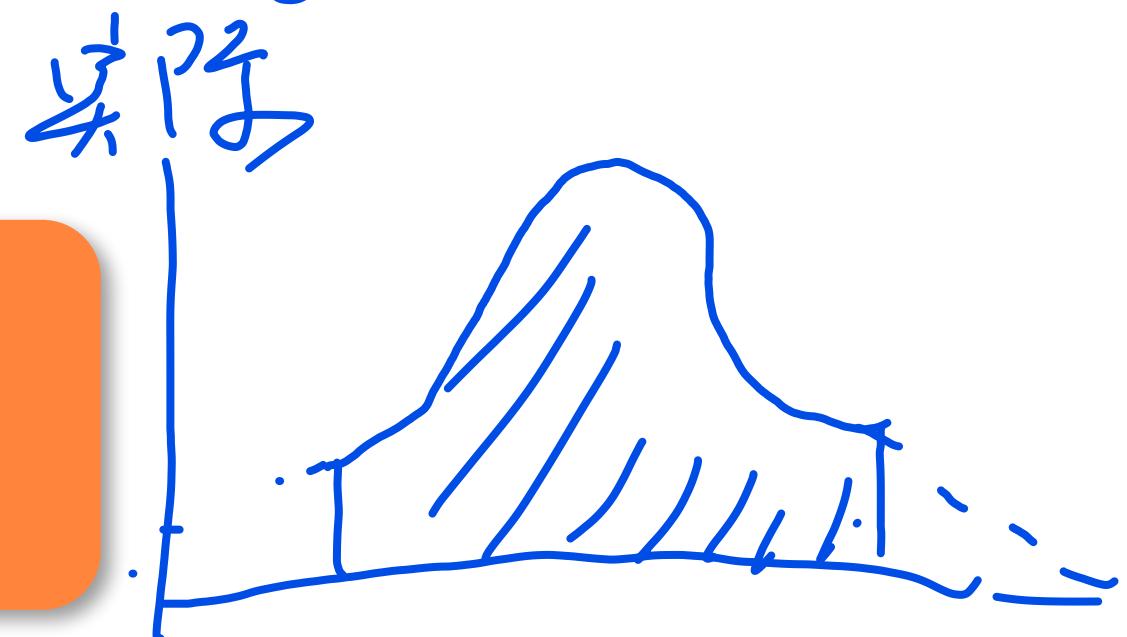
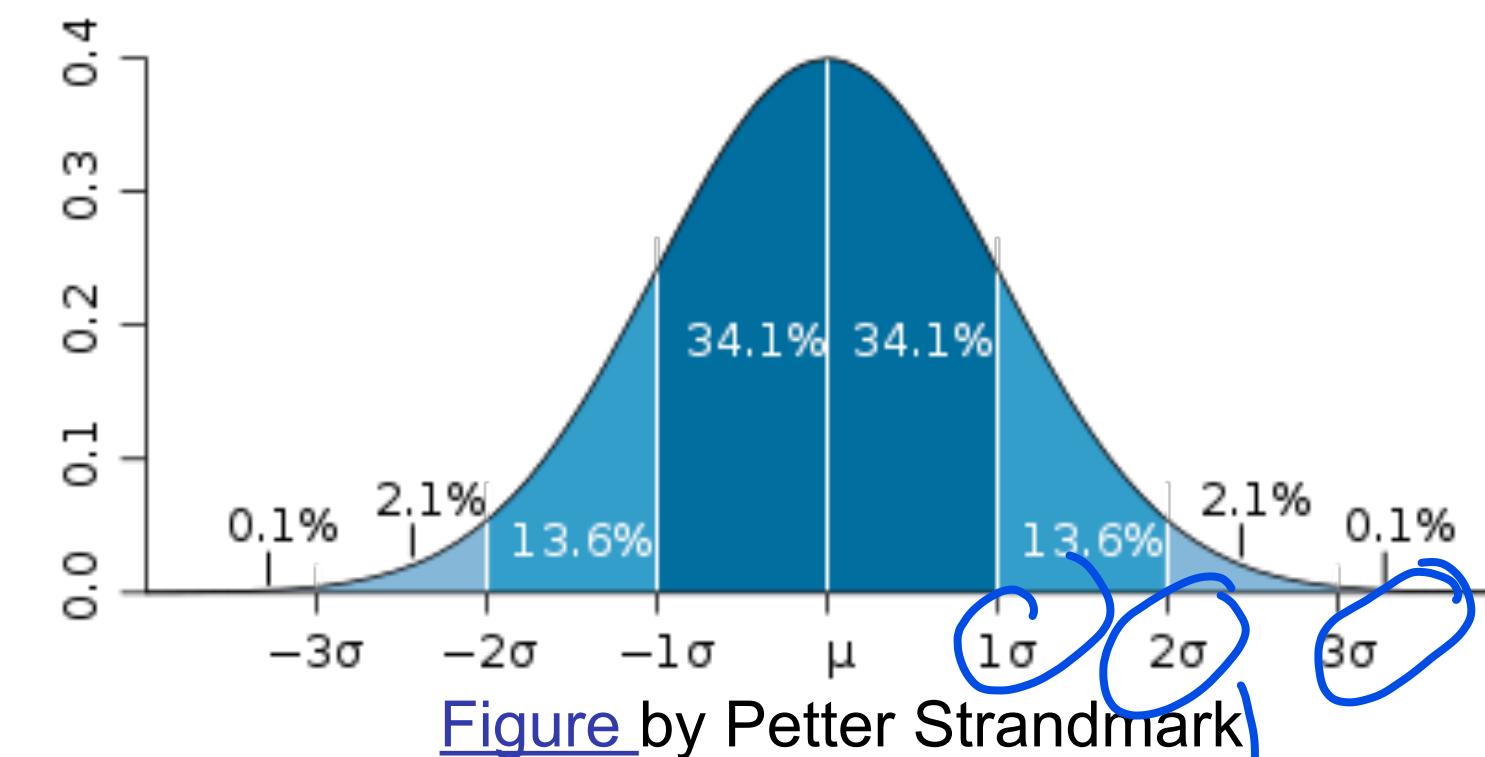
Data, $\mathbf{d} = [d_1 \ d_2 \ d_3 \ \dots \ d_n]$

Mean := $\mu = \frac{\sum_i \mathbf{d}_i}{n}$

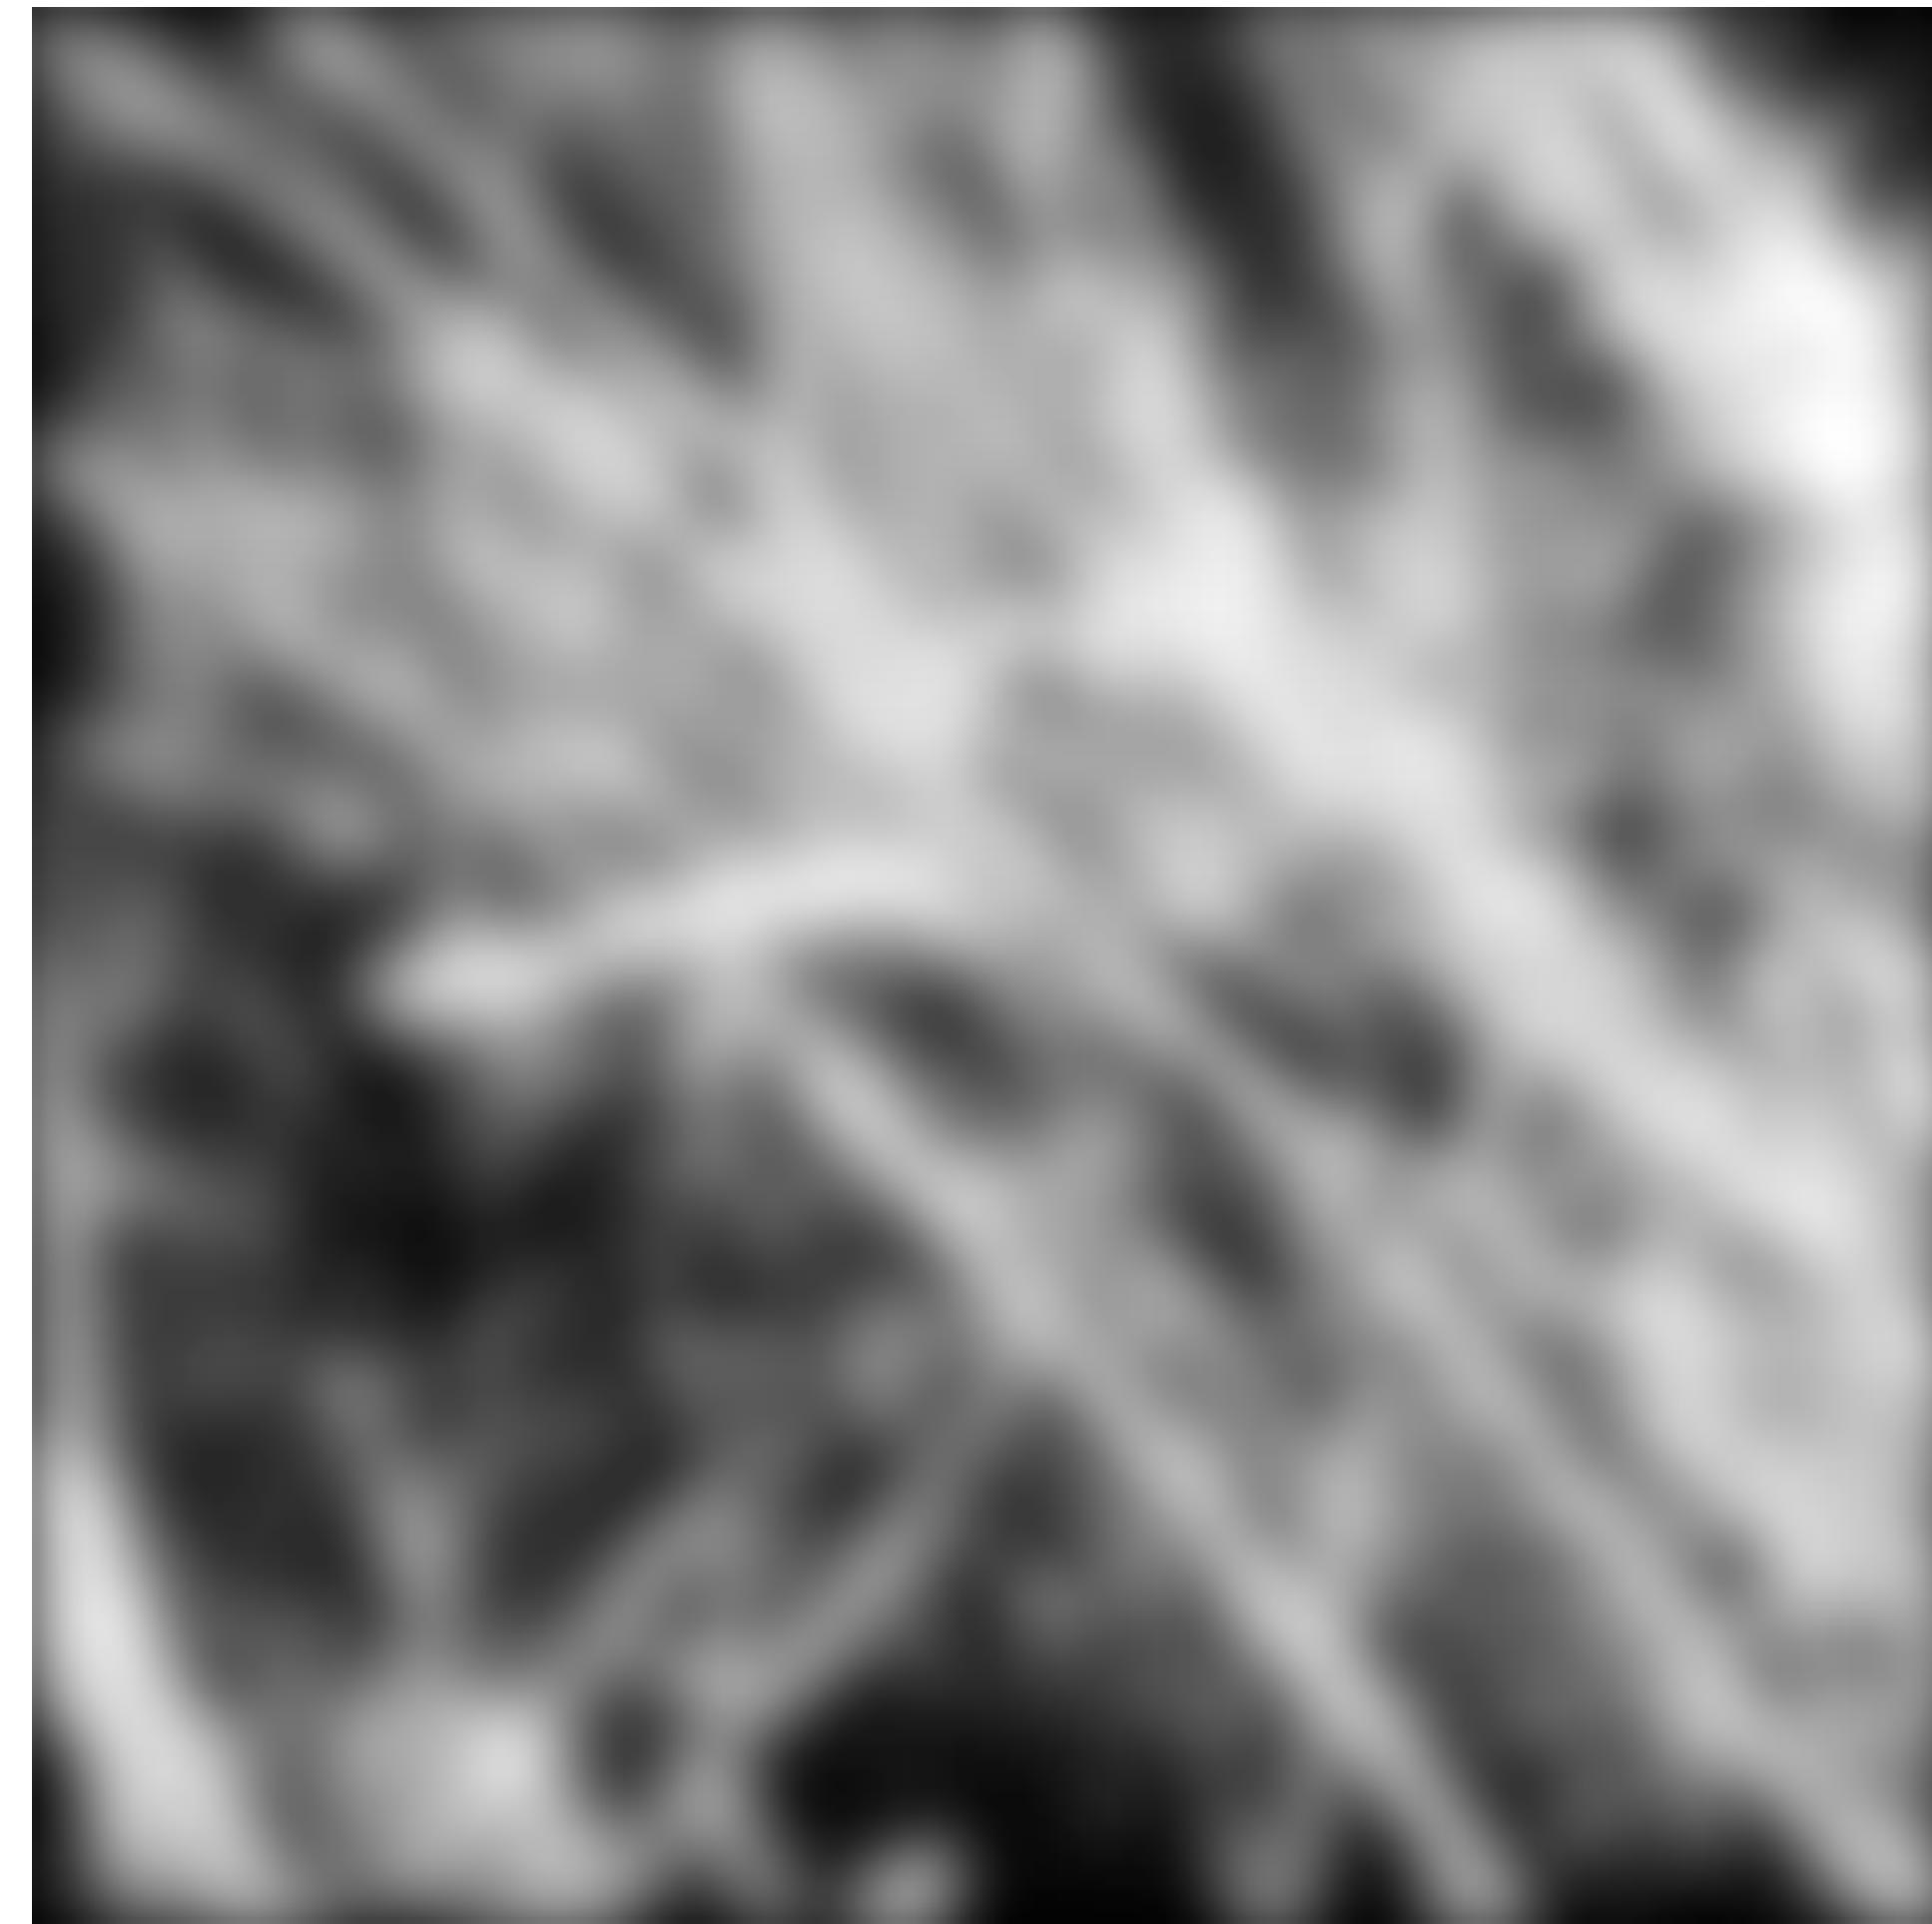
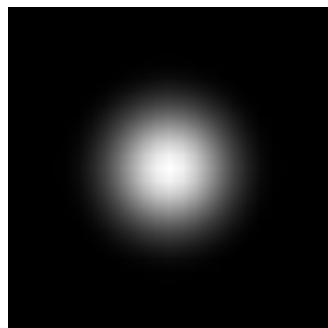
Variance := $\sum_i (\mathbf{d}_i - \mu)^2 / n$

StdDeviation := $\sigma = \sqrt{\text{Variance}}$

mean (d), var (d), std (d)

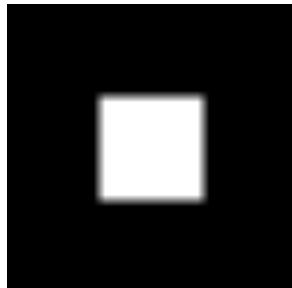


Gaussian Blur



Slide credit: D.A. Forsyth

Smoothing with Box Filter



Slide credit: D.A. Forsyth

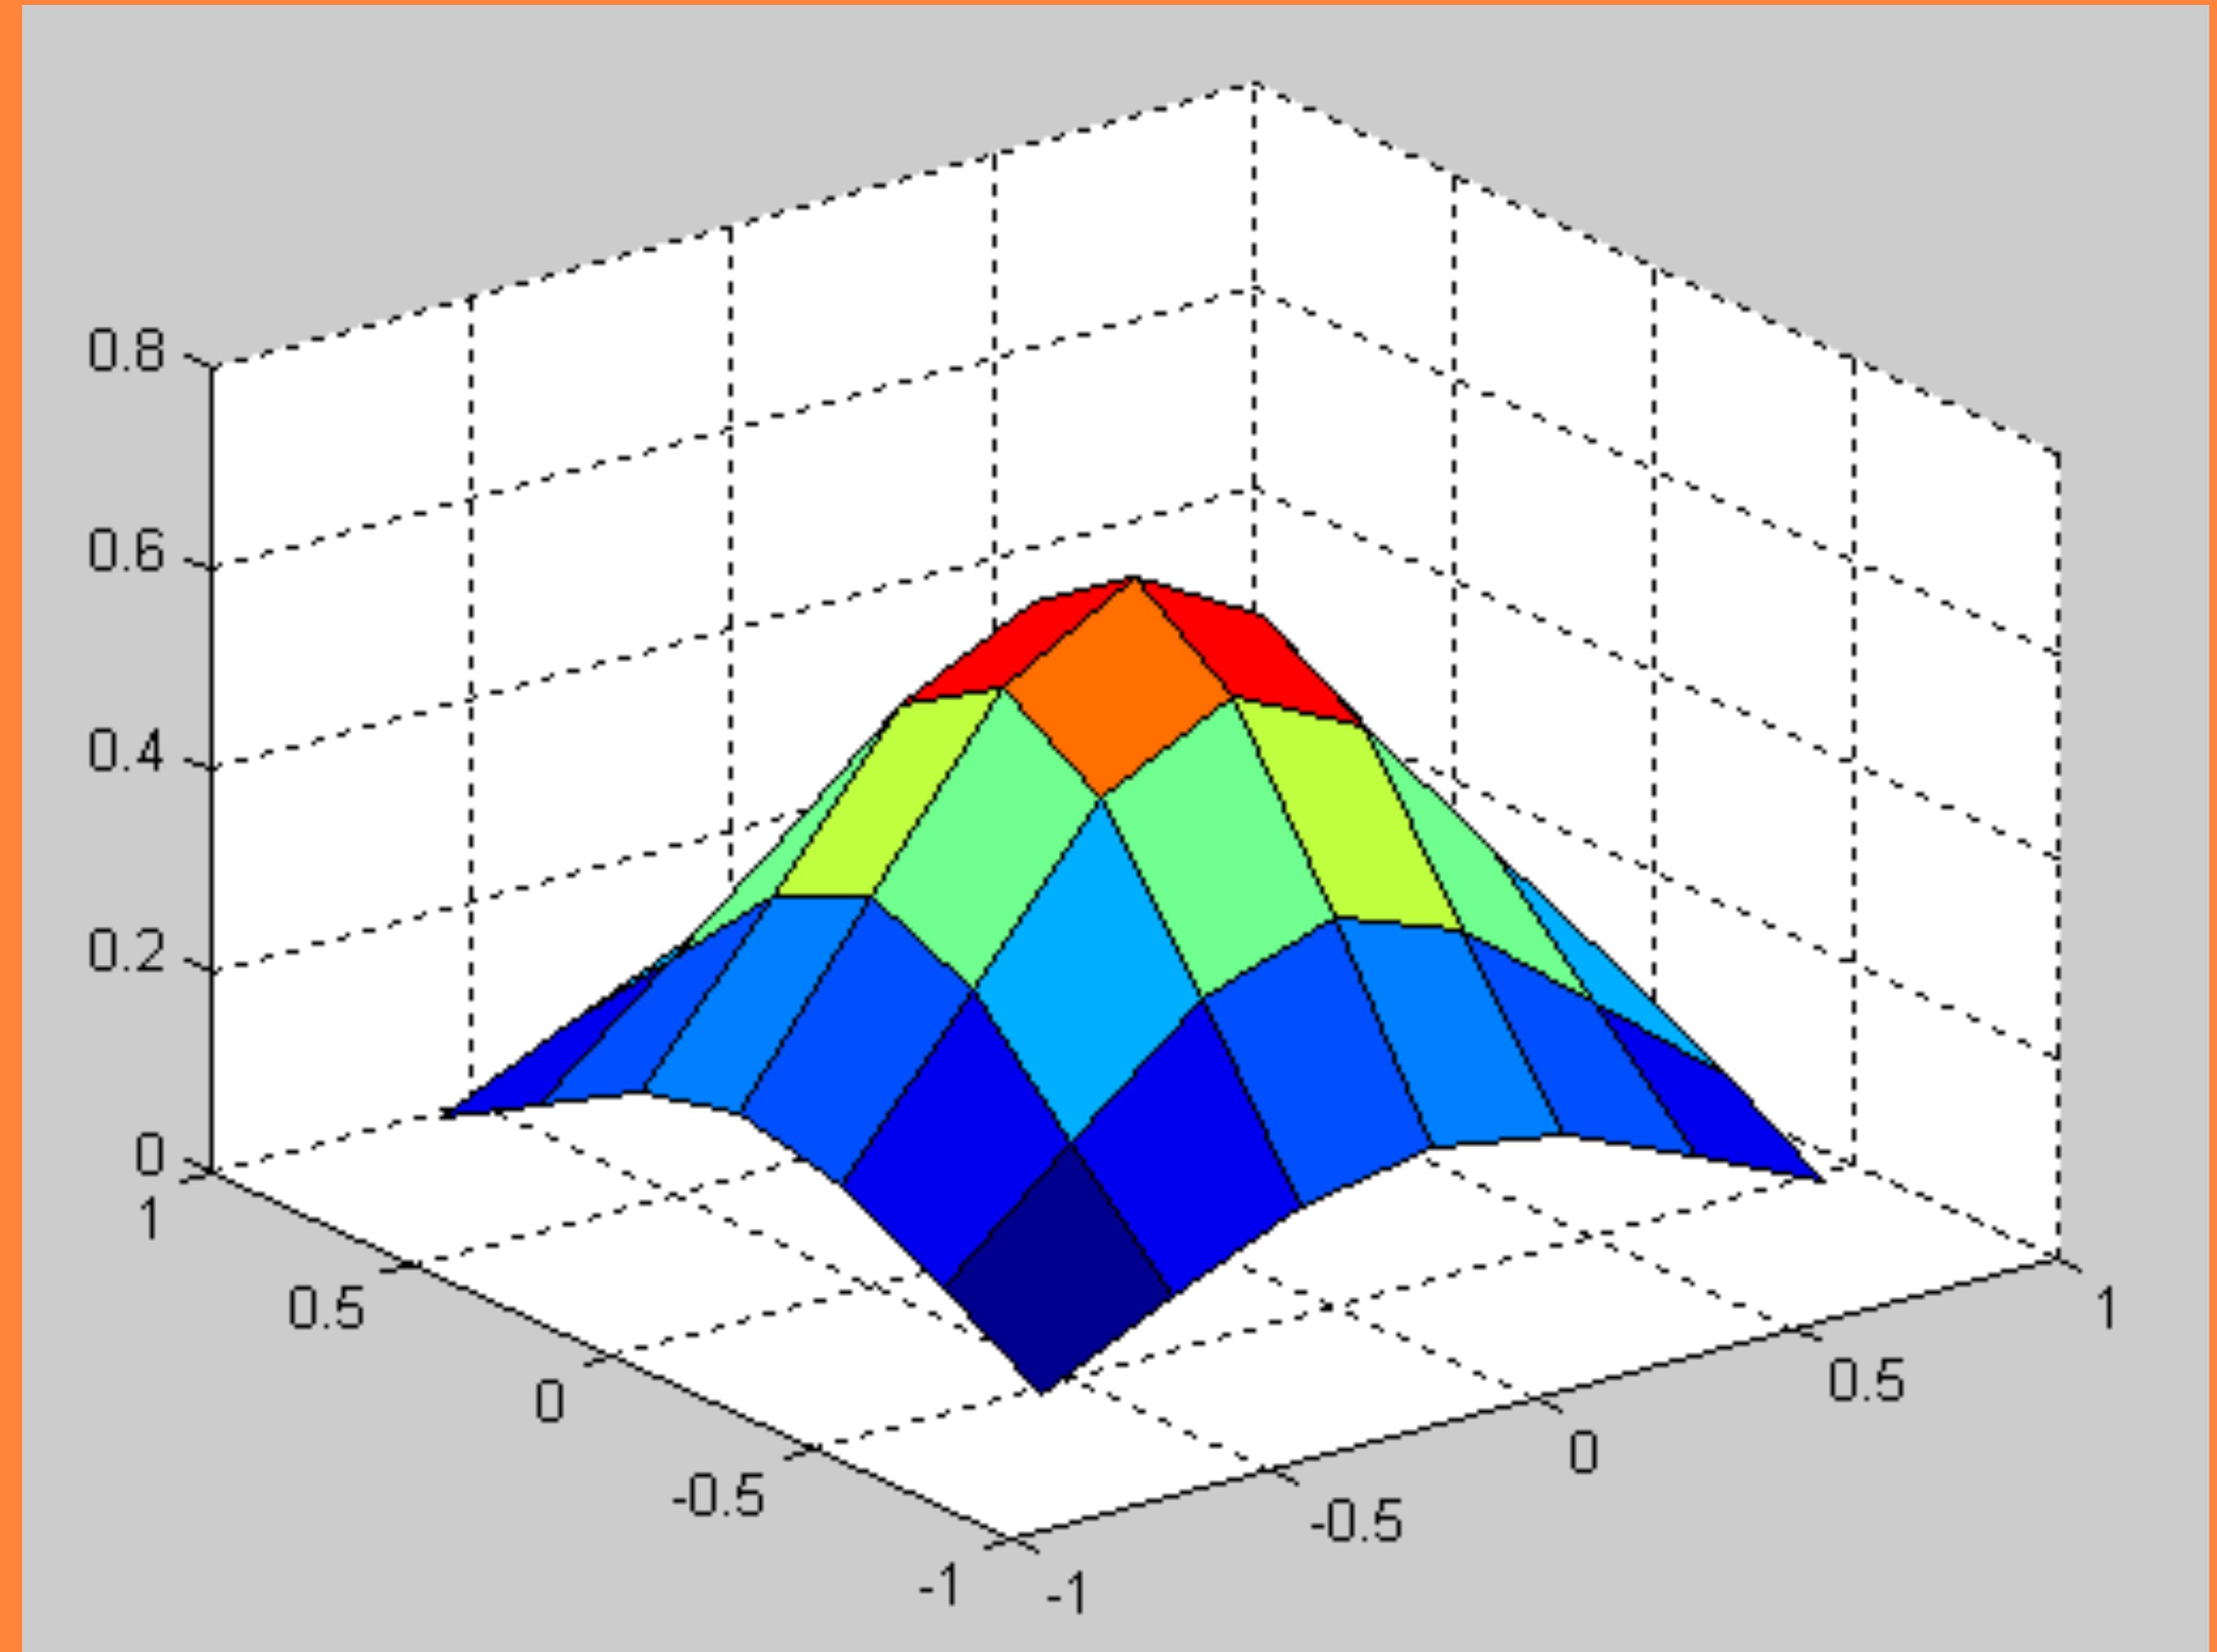
X =

-0.7500	-0.5000	-0.2500	0	0.2500	0.5000	0.7500
-0.7500	-0.5000	-0.2500	0	0.2500	0.5000	0.7500
-0.7500	-0.5000	-0.2500	0	0.2500	0.5000	0.7500
-0.7500	-0.5000	-0.2500	0	0.2500	0.5000	0.7500
-0.7500	-0.5000	-0.2500	0	0.2500	0.5000	0.7500
-0.7500	-0.5000	-0.2500	0	0.2500	0.5000	0.7500
-0.7500	-0.5000	-0.2500	0	0.2500	0.5000	0.7500

Y =

-0.7500	-0.7500	-0.7500	-0.7500	-0.7500	-0.7500	-0.7500
-0.5000	-0.5000	-0.5000	-0.5000	-0.5000	-0.5000	-0.5000
-0.2500	-0.2500	-0.2500	-0.2500	-0.2500	-0.2500	-0.2500
0	0	0	0	0	0	0
0.2500	0.2500	0.2500	0.2500	0.2500	0.2500	0.2500
0.5000	0.5000	0.5000	0.5000	0.5000	0.5000	0.5000
0.7500	0.7500	0.7500	0.7500	0.7500	0.7500	0.7500

```
>> Z = (1/(2*pi*sigma.^2)) * exp(-(X.^2 + Y.^2)/(2*sigma.^2));
>> surf(X,Y,Z)    % shown with sigma = 0.5
>>
>> % Probably want Z = Z / sum(Z(:));
```



Separable Kernels

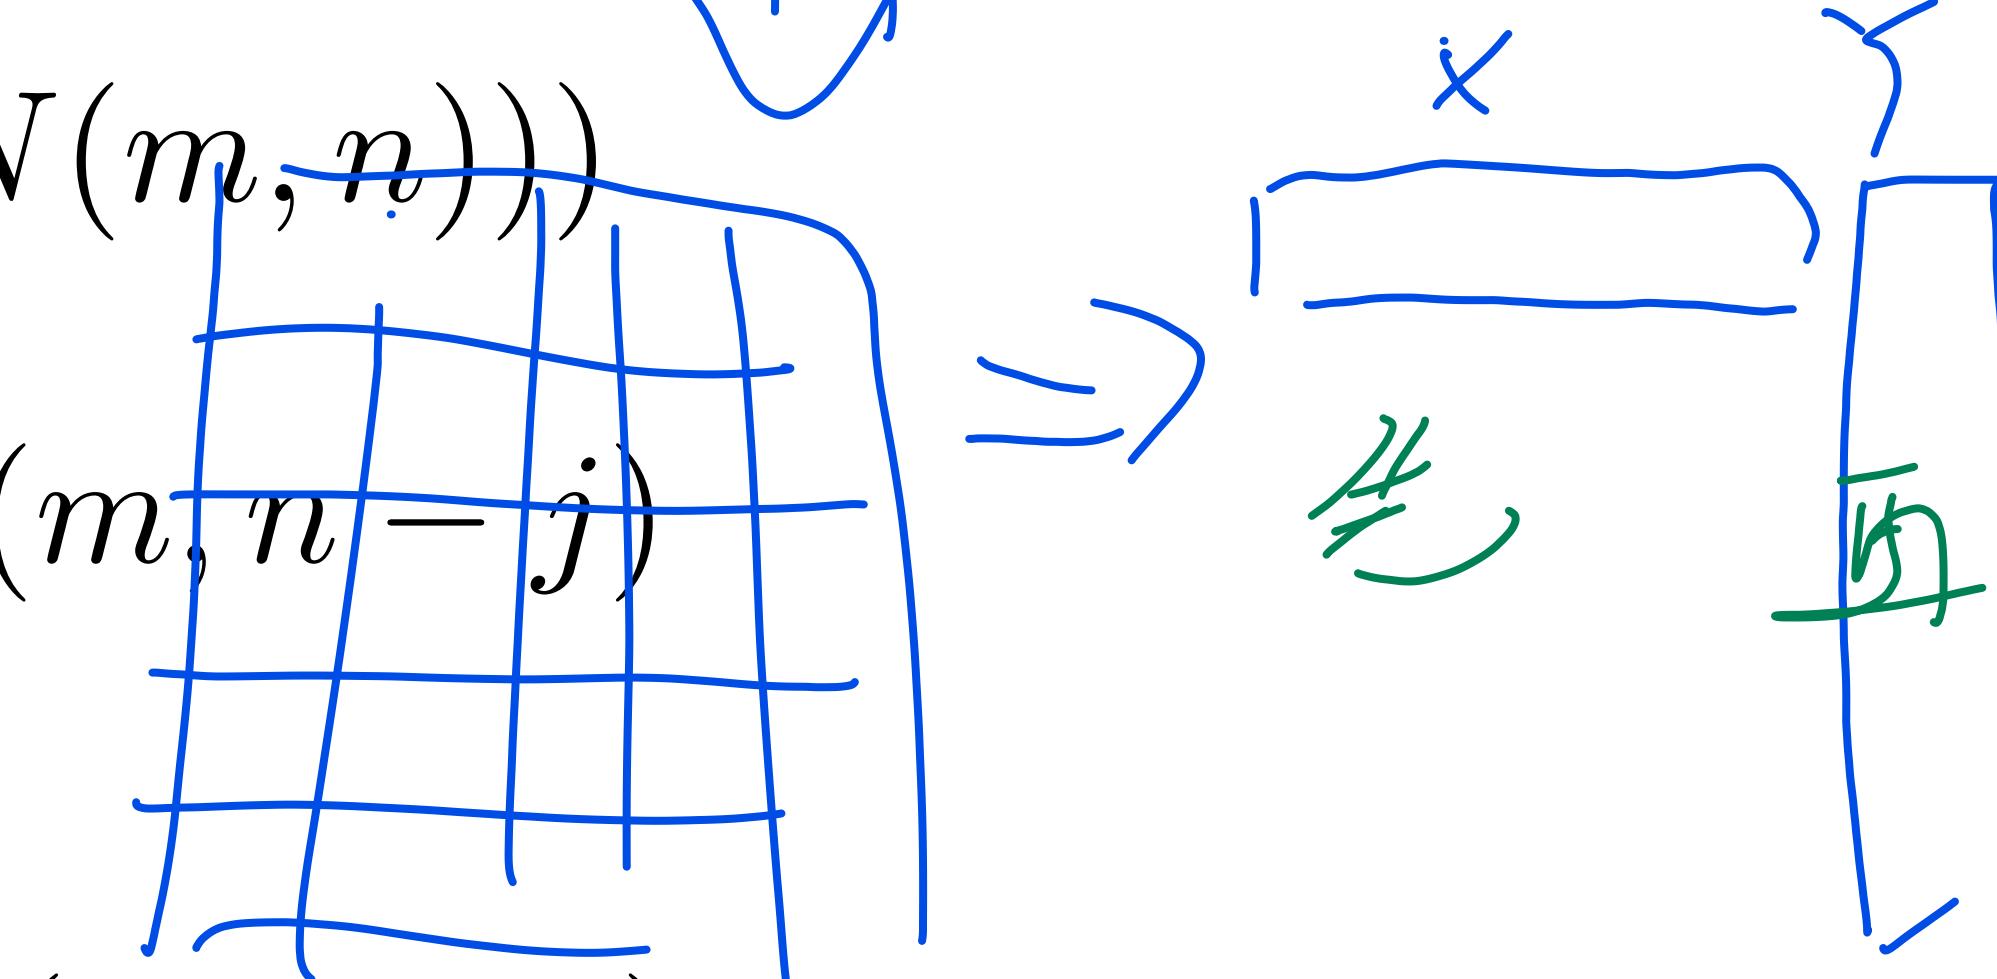
- Separable filters can be written as $K(m,n) = f(m)g(n)$
- For a rectangular neighbourhood with size $(2M+1) \times (2N+1)$,

$$I'(m, n) = f \star (g \star I(N(m, n)))$$

$$I''(m, n) = \sum_{j=-N}^N g(j)I(m, n-j)$$

$$I'(m, n) = \sum_{i=-M}^M f(i)I''(m-i, n)$$

2DF
每 pixel
(25 次变成 10 次)



Gaussian Smoothing Kernels

$$\begin{aligned} g(x, y) &= \frac{1}{2\pi\sigma^2} \exp\left(\frac{-(x^2 + y^2)}{2\sigma^2}\right) \\ &= \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(\frac{-x^2}{2\sigma^2}\right) \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(\frac{-y^2}{2\sigma^2}\right) \\ &= g(x)g(y) \end{aligned}$$

Gaussian Smoothing Kernels

$$\begin{aligned}g(x, y) &= \frac{1}{2\pi\sigma^2} \exp\left(\frac{-(x^2 + y^2)}{2\sigma^2}\right) \\&= \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(\frac{-x^2}{2\sigma^2}\right) \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(\frac{-y^2}{2\sigma^2}\right) \\&= g(x)g(y)\end{aligned}$$

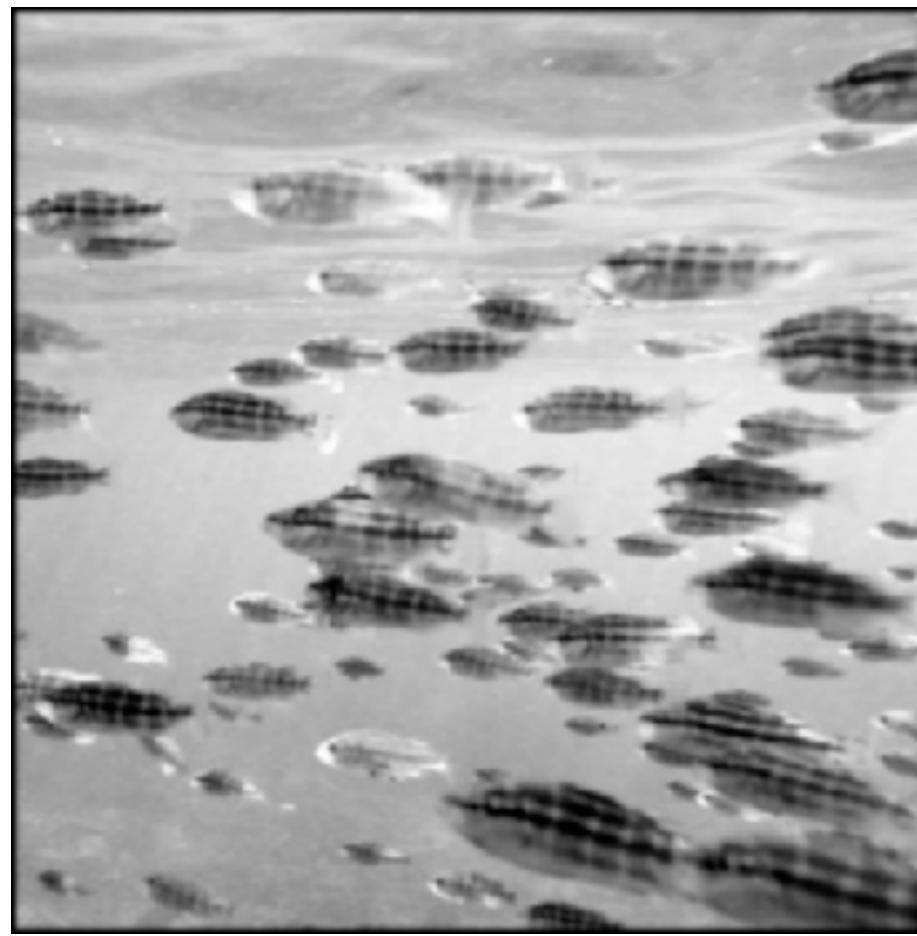
Separable!

To see how svd() can be used to obtain the separated kernels of a rank 1 kernel.

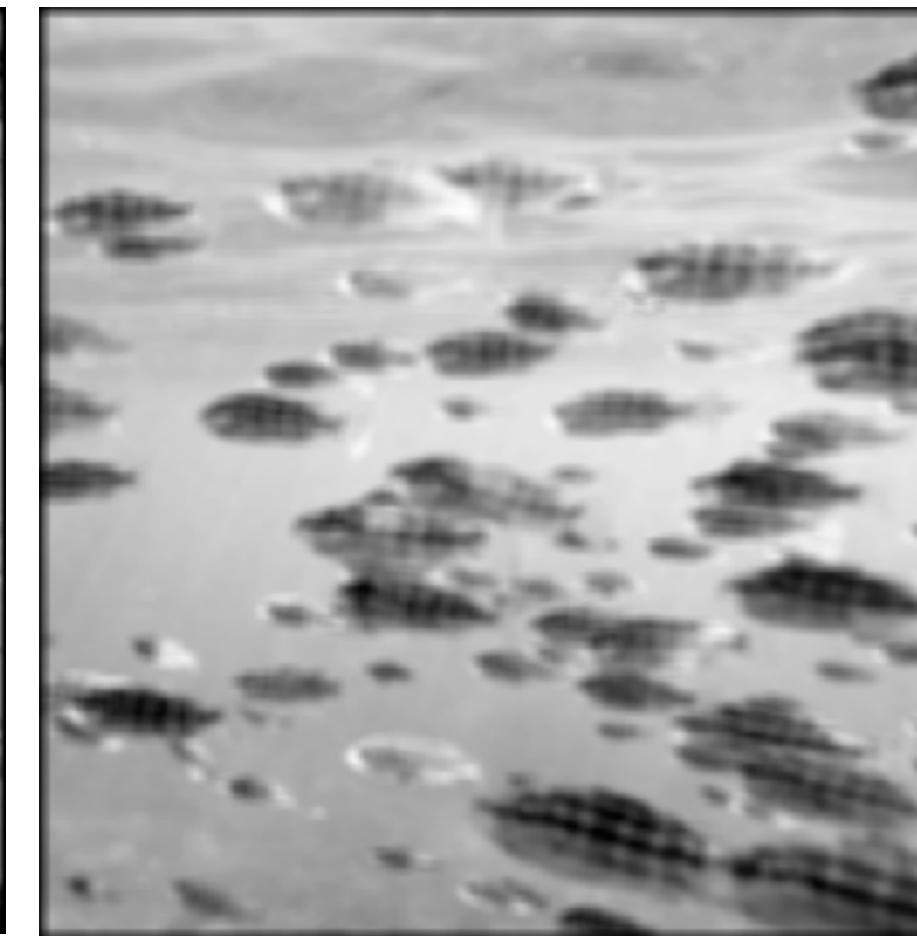
Gaussian Smoothing Kernels

- Amount of smoothing depends on σ and window size.
- Width $> 3\sigma$

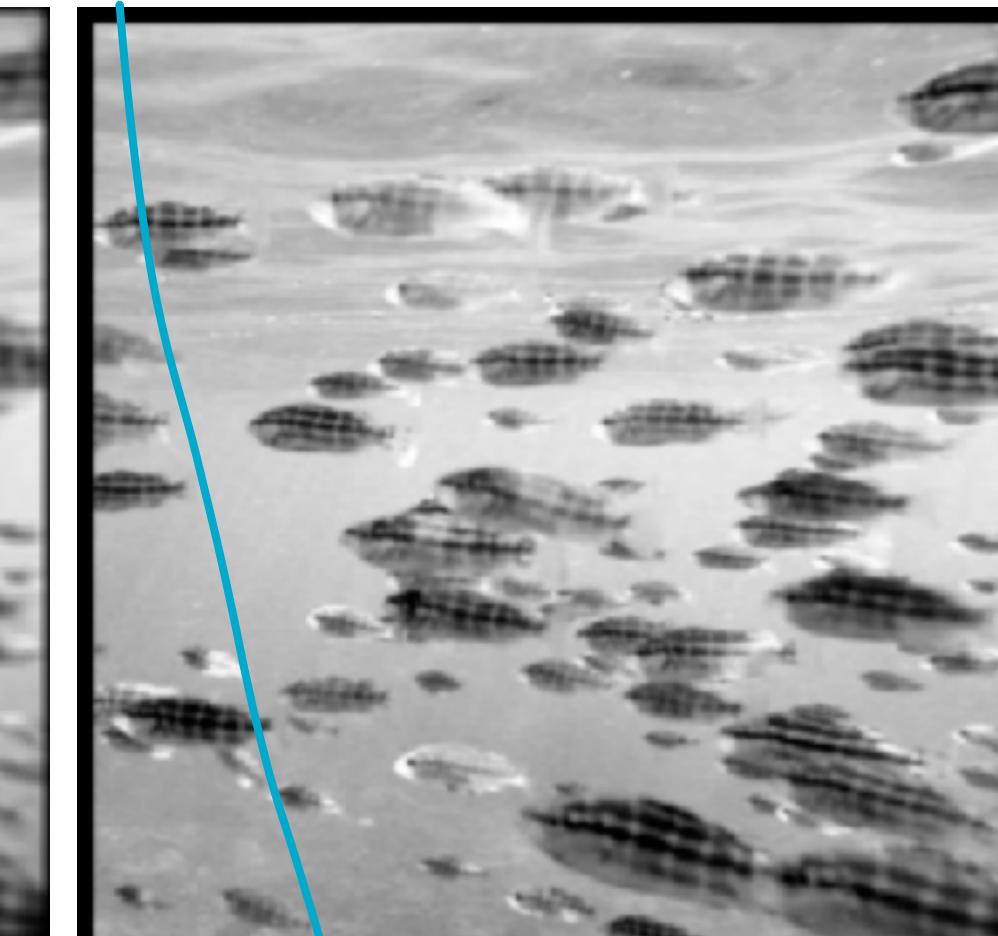
7x7; $\sigma = 1.$



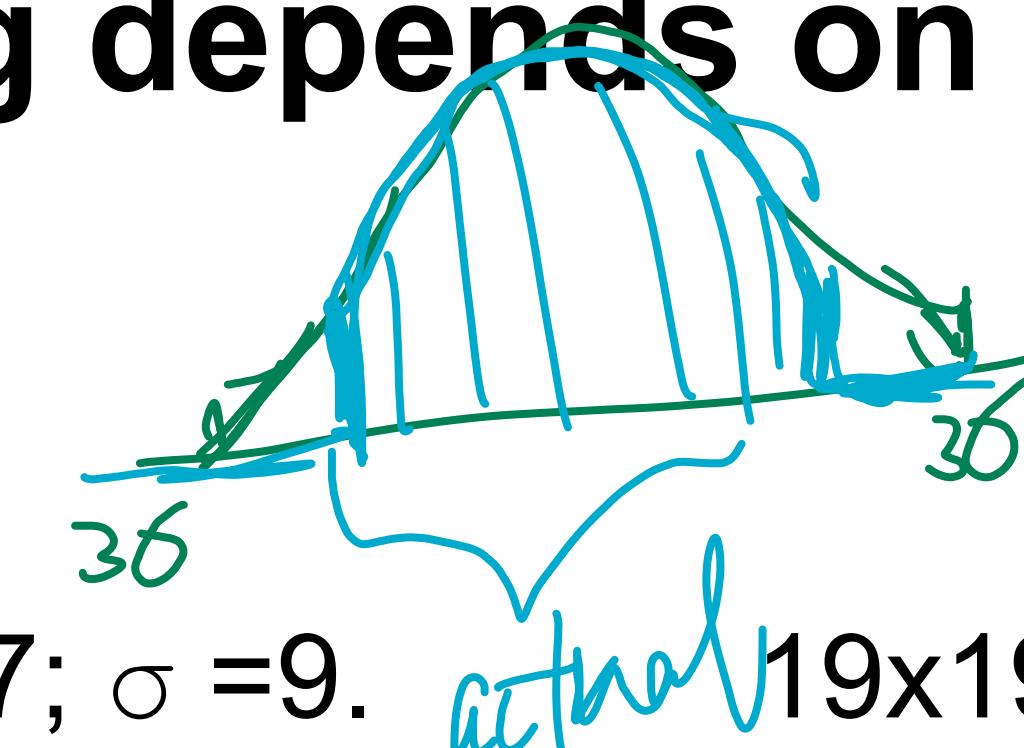
7x7; $\sigma = 9.$



actual 19x19; $\sigma = 1.$

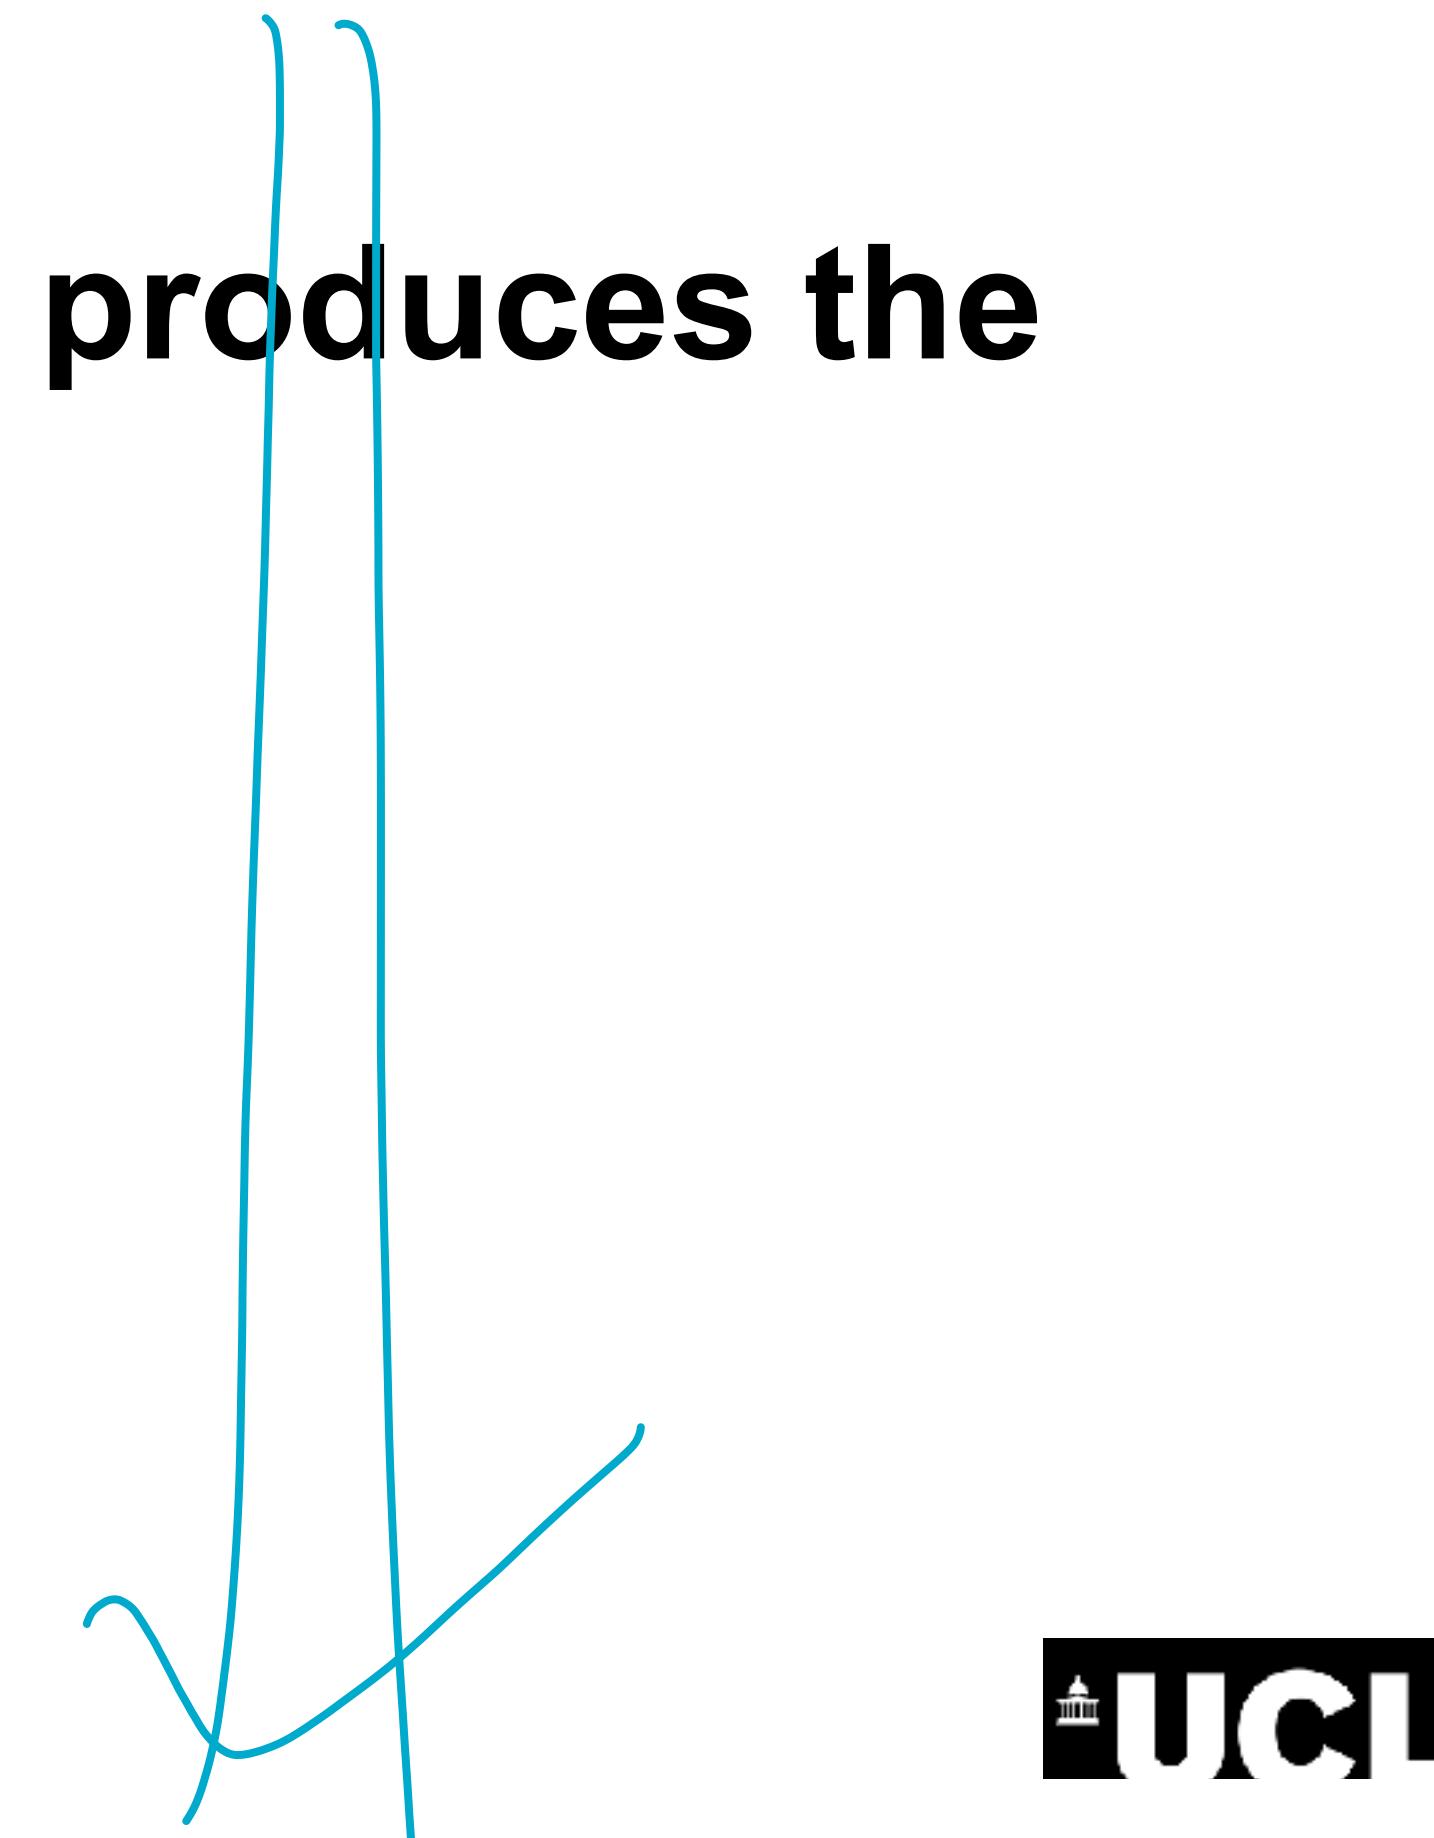


19x19; $\sigma = 9.$



Scale Space

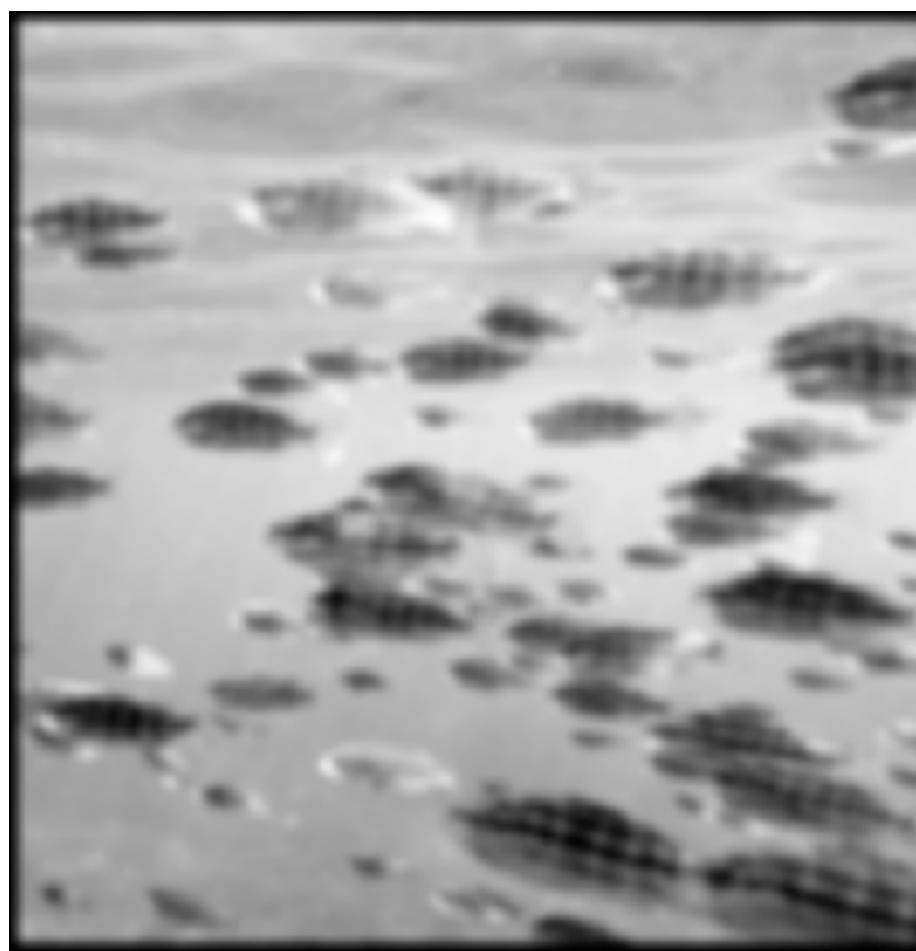
- Convolution of a Gaussian with standard deviation σ with itself, is a Gaussian with standard deviation $\sigma\sqrt{2}$.
- Repeated convolution by a Gaussian filter produces the *scale space* of an image.



Scale Space Example

11x11; $\sigma = 3$.

1



repeated
Grass
on the
left result 4

2



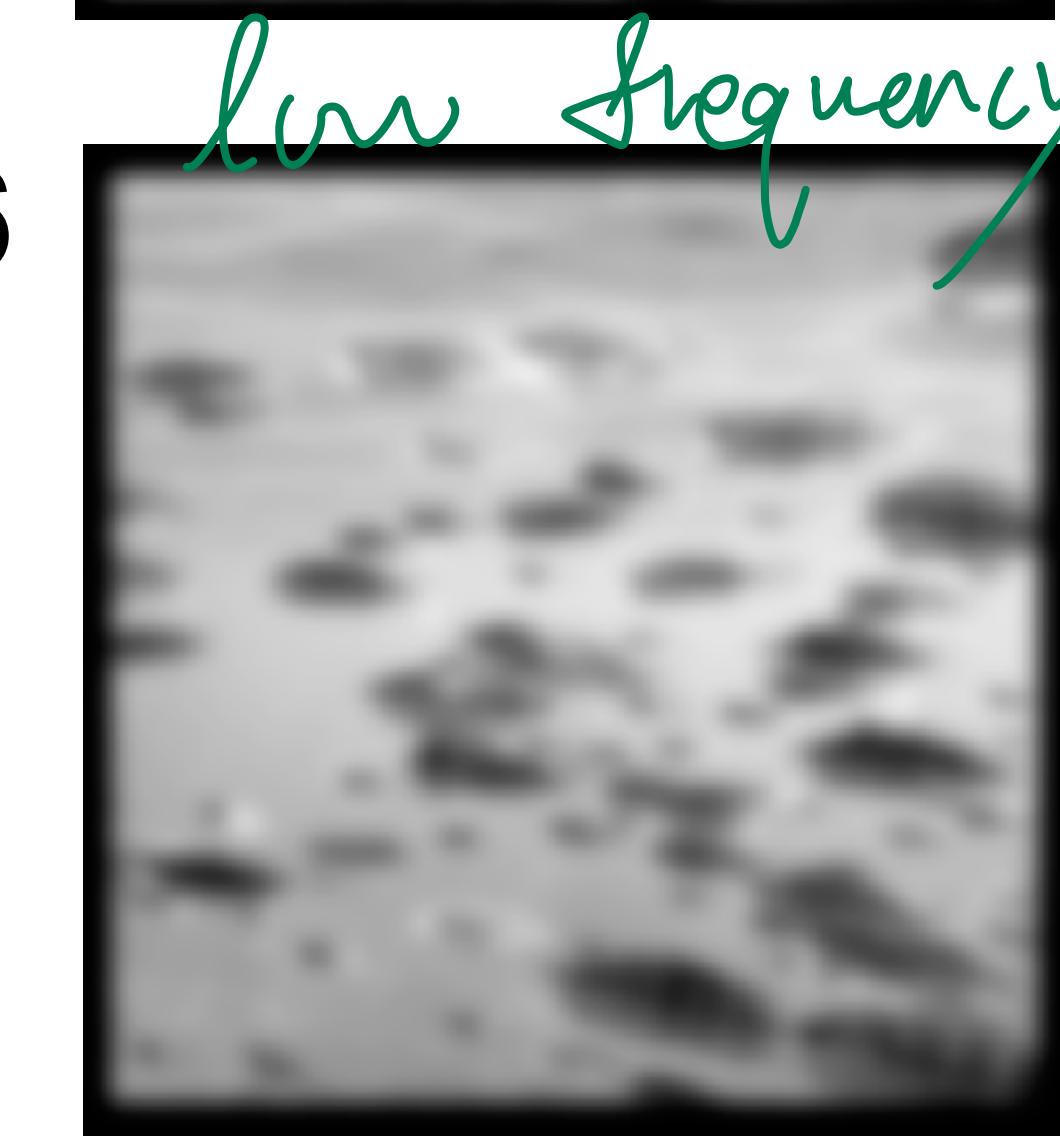
3



5



6

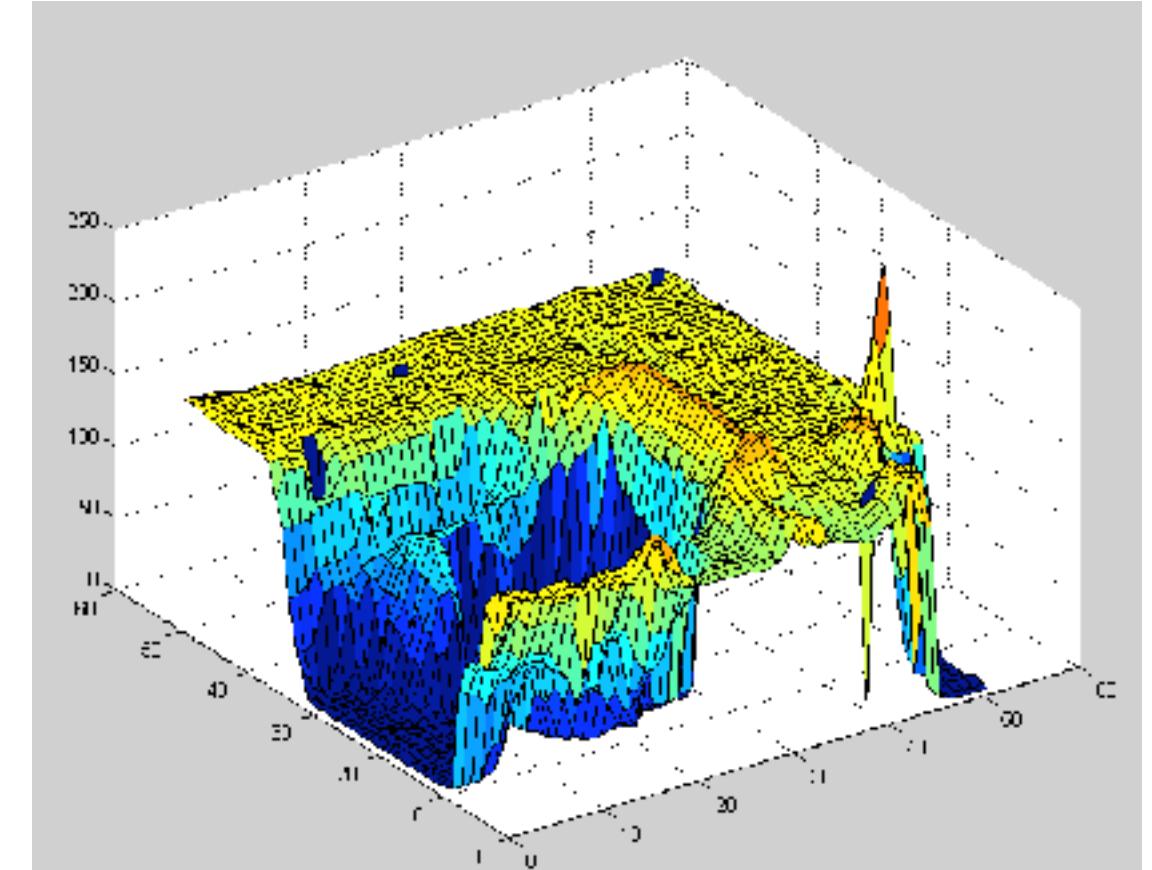


Gaussian Smoothing Kernel Top-5

1. Rotationally symmetric
2. Has a single lobe/mode
 - Neighbour's influence decreases monotonically
3. Still one lobe in frequency domain
 - No corruption from high frequencies
4. Simple relationship to σ
5. Easy to implement efficiently

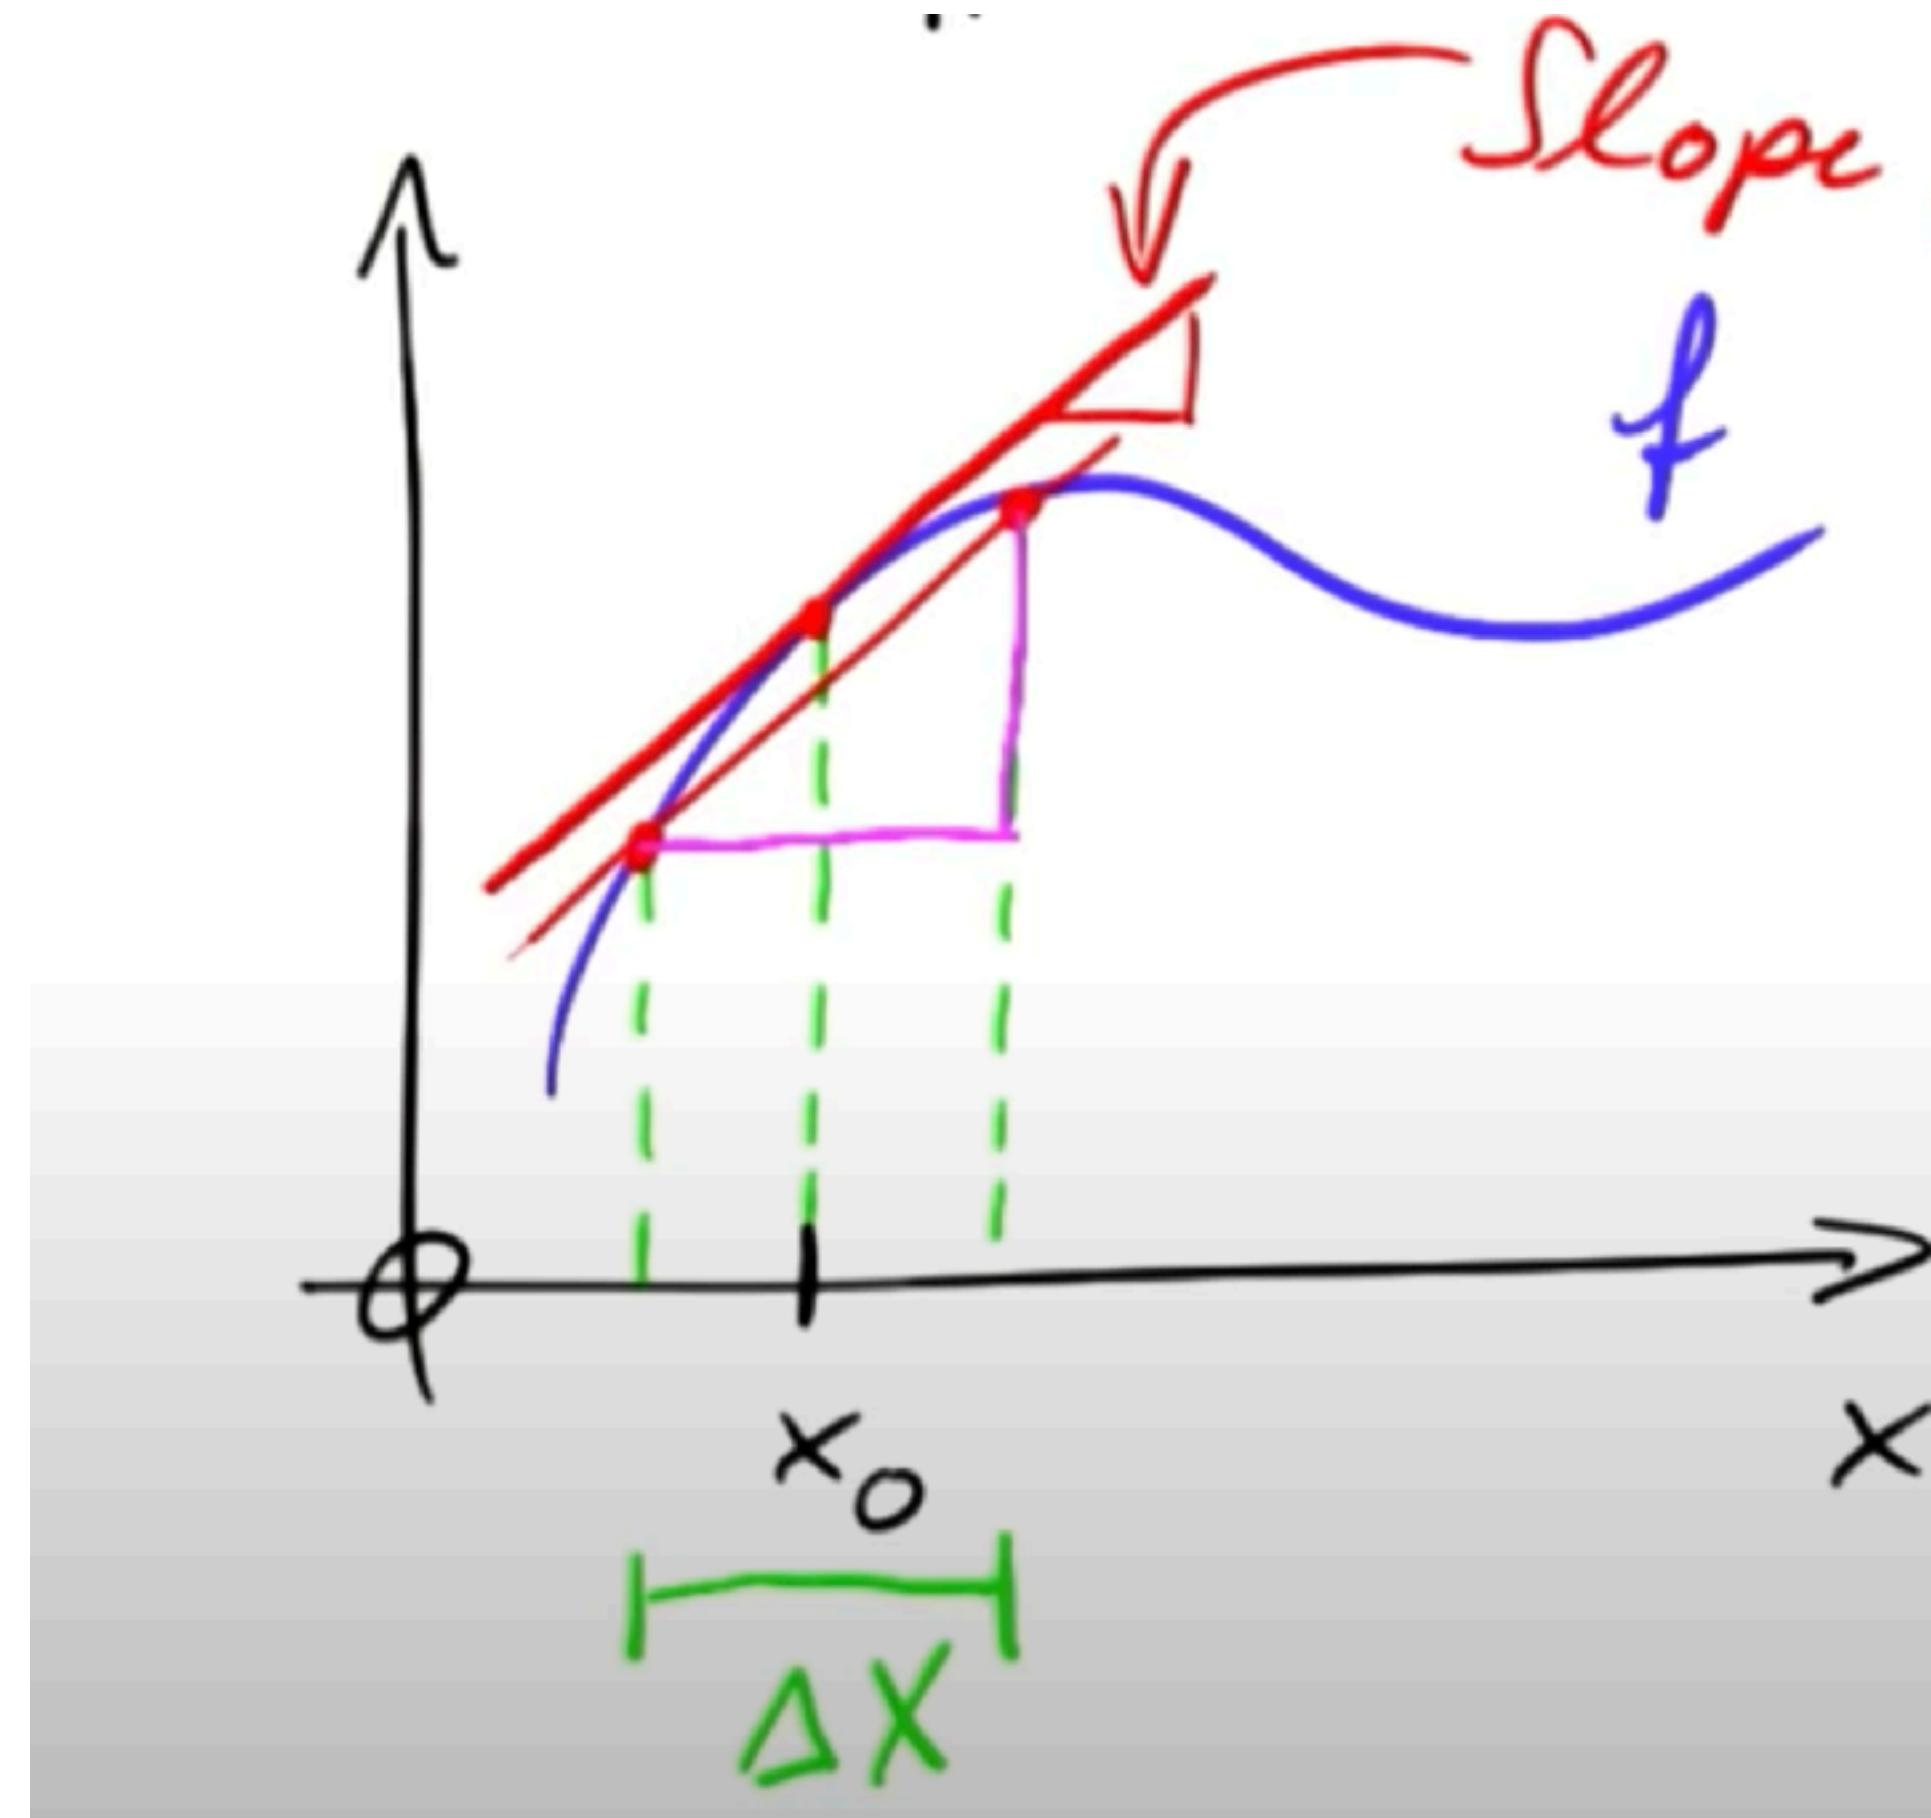
Smoothing vs. Sharpening

- Smoothing \leftrightarrow Integration, i.e.,
 - Summing
 - Sand-papering the height field's slope
- Sharpening \leftrightarrow Differentiation, i.e.,
 - Subtracting
 - Accentuating the height field's slope



$z = I'$'s
intensity

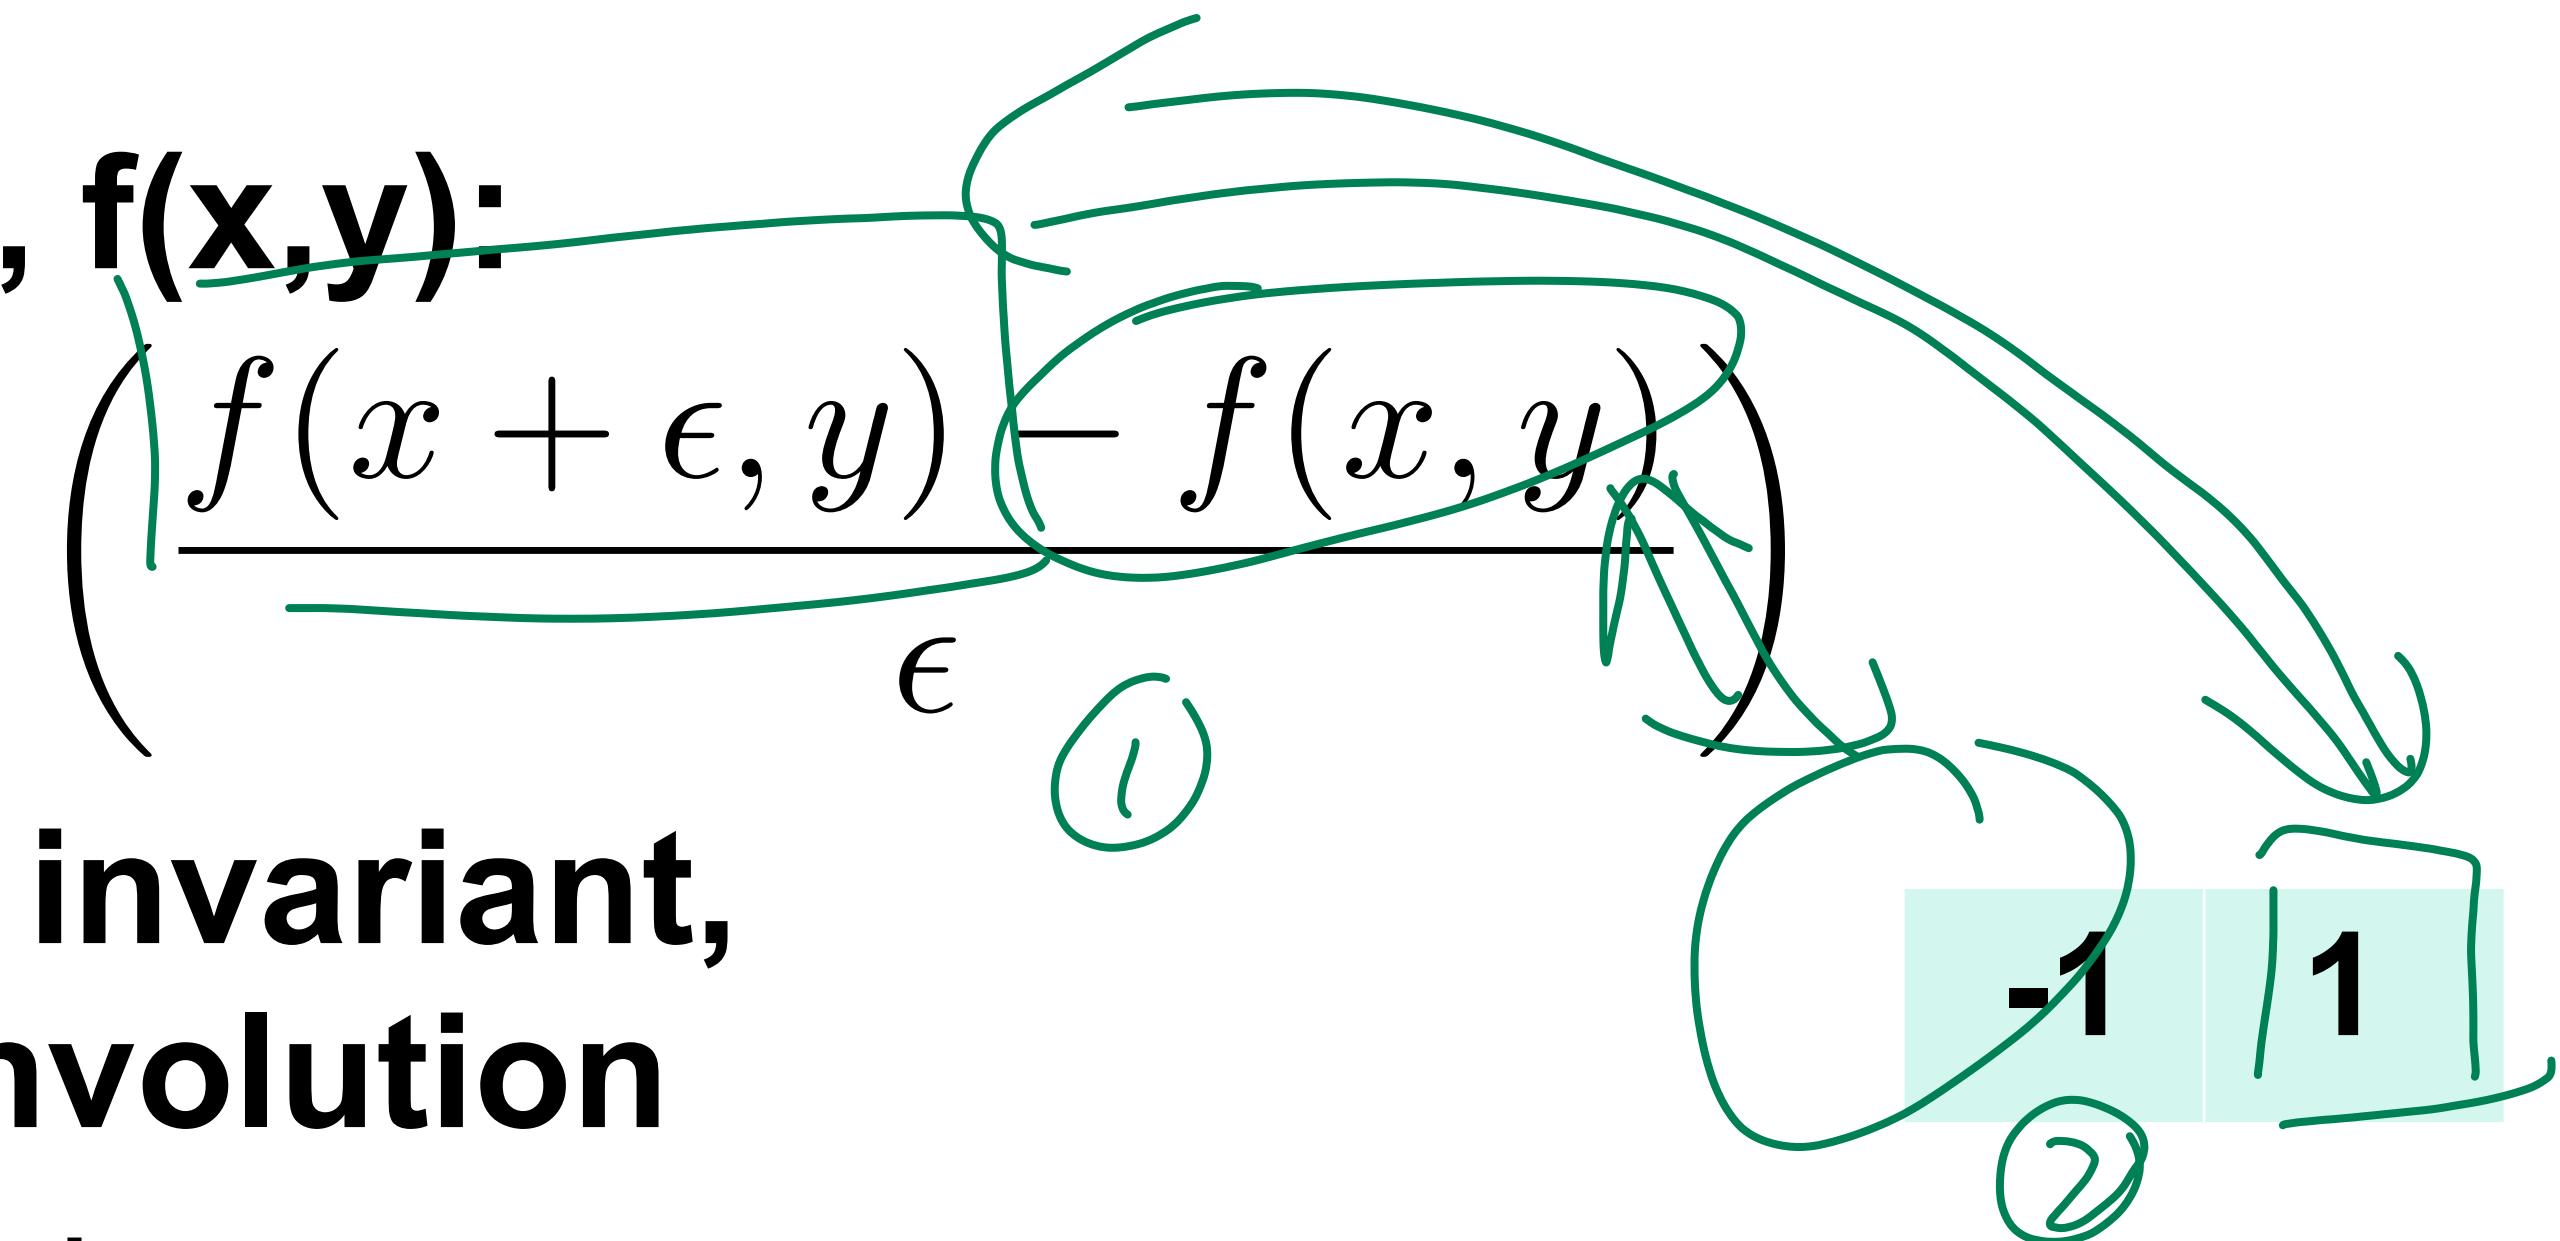
1st and 2nd Derivative



Differentiation and Convolution

- Recall, for 2D function, $f(x,y)$:

$$\frac{\partial f}{\partial x} = \lim_{\epsilon \rightarrow 0}$$



- This is linear and shift invariant, so the result of a convolution

- We could approximate this as (which is obviously a convolution)

$$\frac{\partial f}{\partial x} \approx \frac{f(x_{n+1}, y) - f(x_n, y)}{\Delta x}$$

∴ ① 可以完全用
x 像
可从 edge detected
Image Filtering
COMP0026: Image Processing

1st and 2nd Derivative

The derivative of a function f at a point x is defined by the limit

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x + h) - f(x)}{h}$$

Approximation of the derivative when h is small

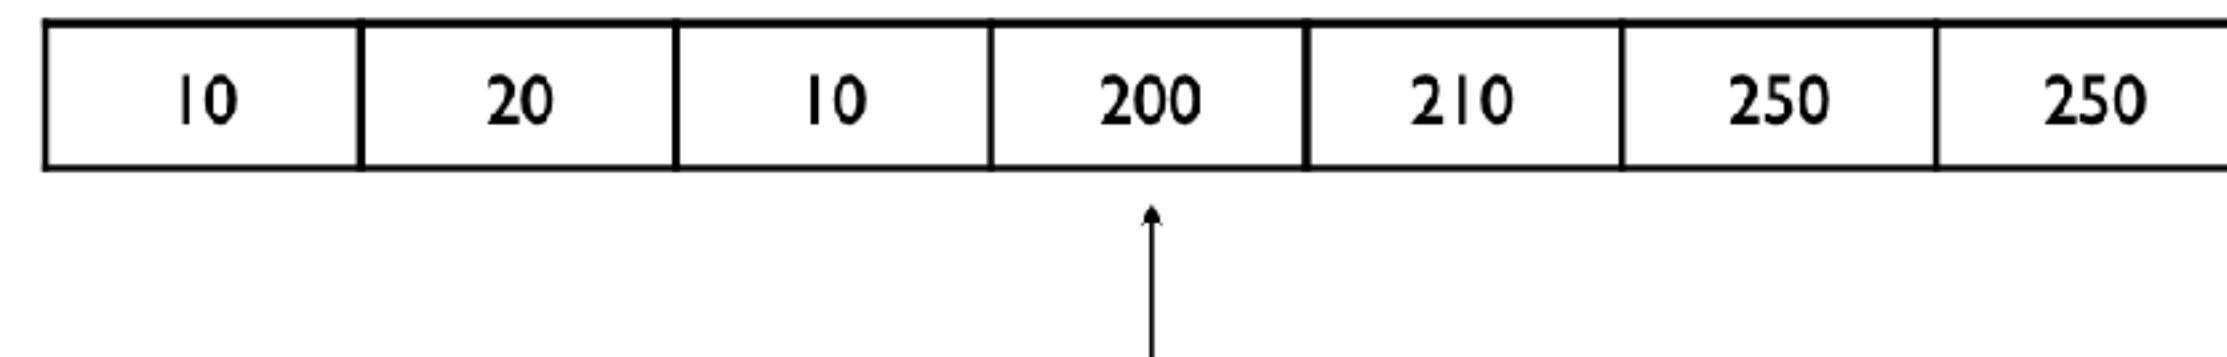
This definition is based on the ‘forward difference’ but ...

1st and 2nd Derivative

it turns out that using the ‘central difference’ is more accurate

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x + 0.5h) - f(x - 0.5h)}{h}$$

How do we compute the derivative of a discrete signal?



1st and 2nd Derivative

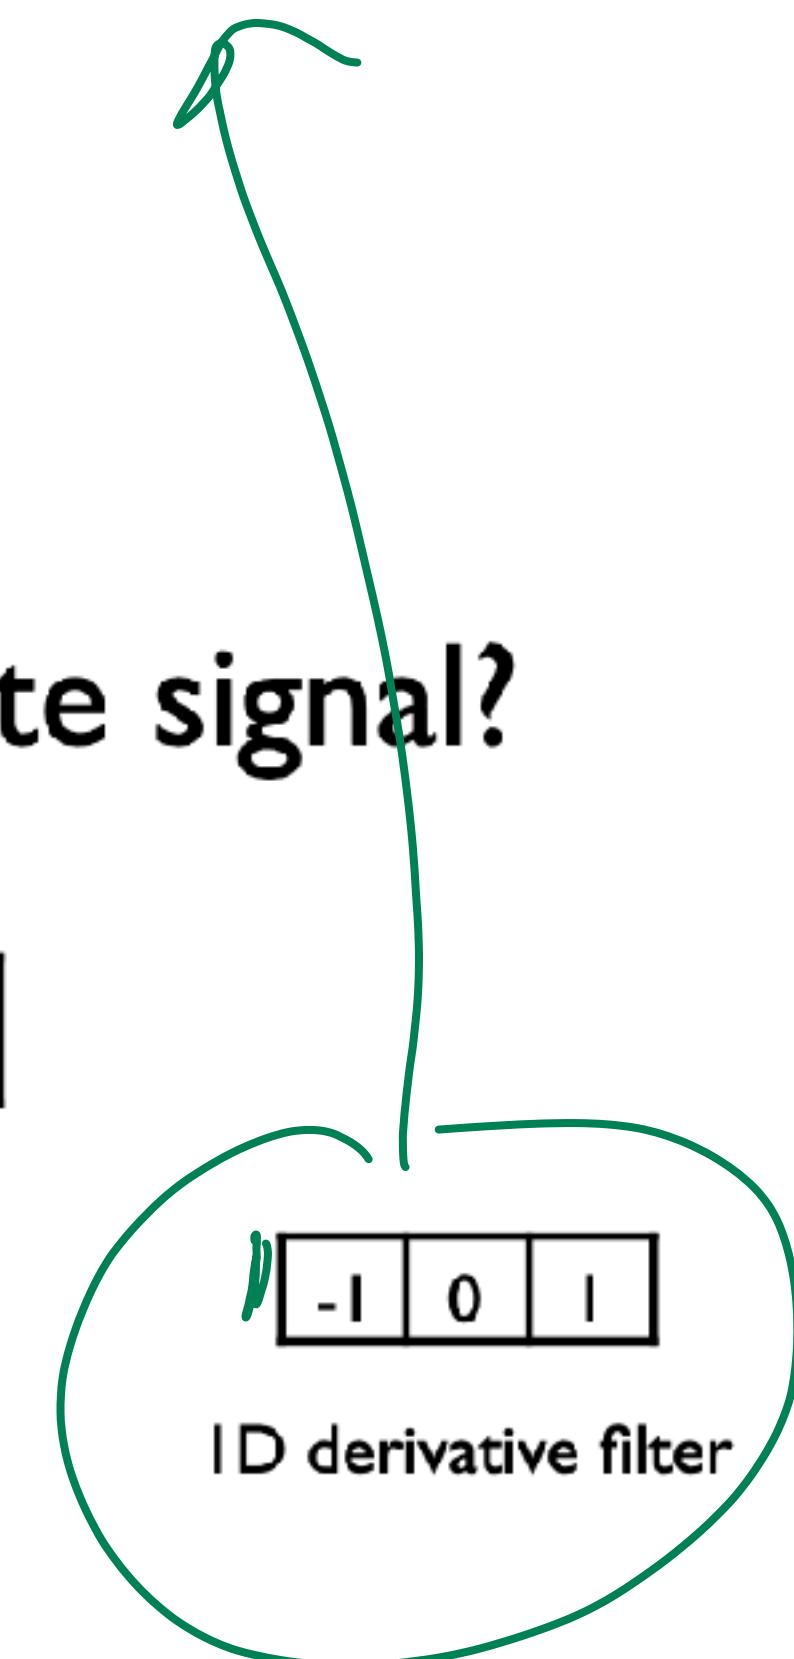
it turns out that using the 'central difference' is more accurate

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x + 0.5h) - f(x - 0.5h)}{h}$$

How do we compute the derivative of a discrete signal?

10	20	10	200	210	250	250
----	----	----	-----	-----	-----	-----

$$f'(x) = \frac{f(x+1) - f(x-1)}{2} = \frac{210 - 10}{2} = 100$$



2nd Derivative

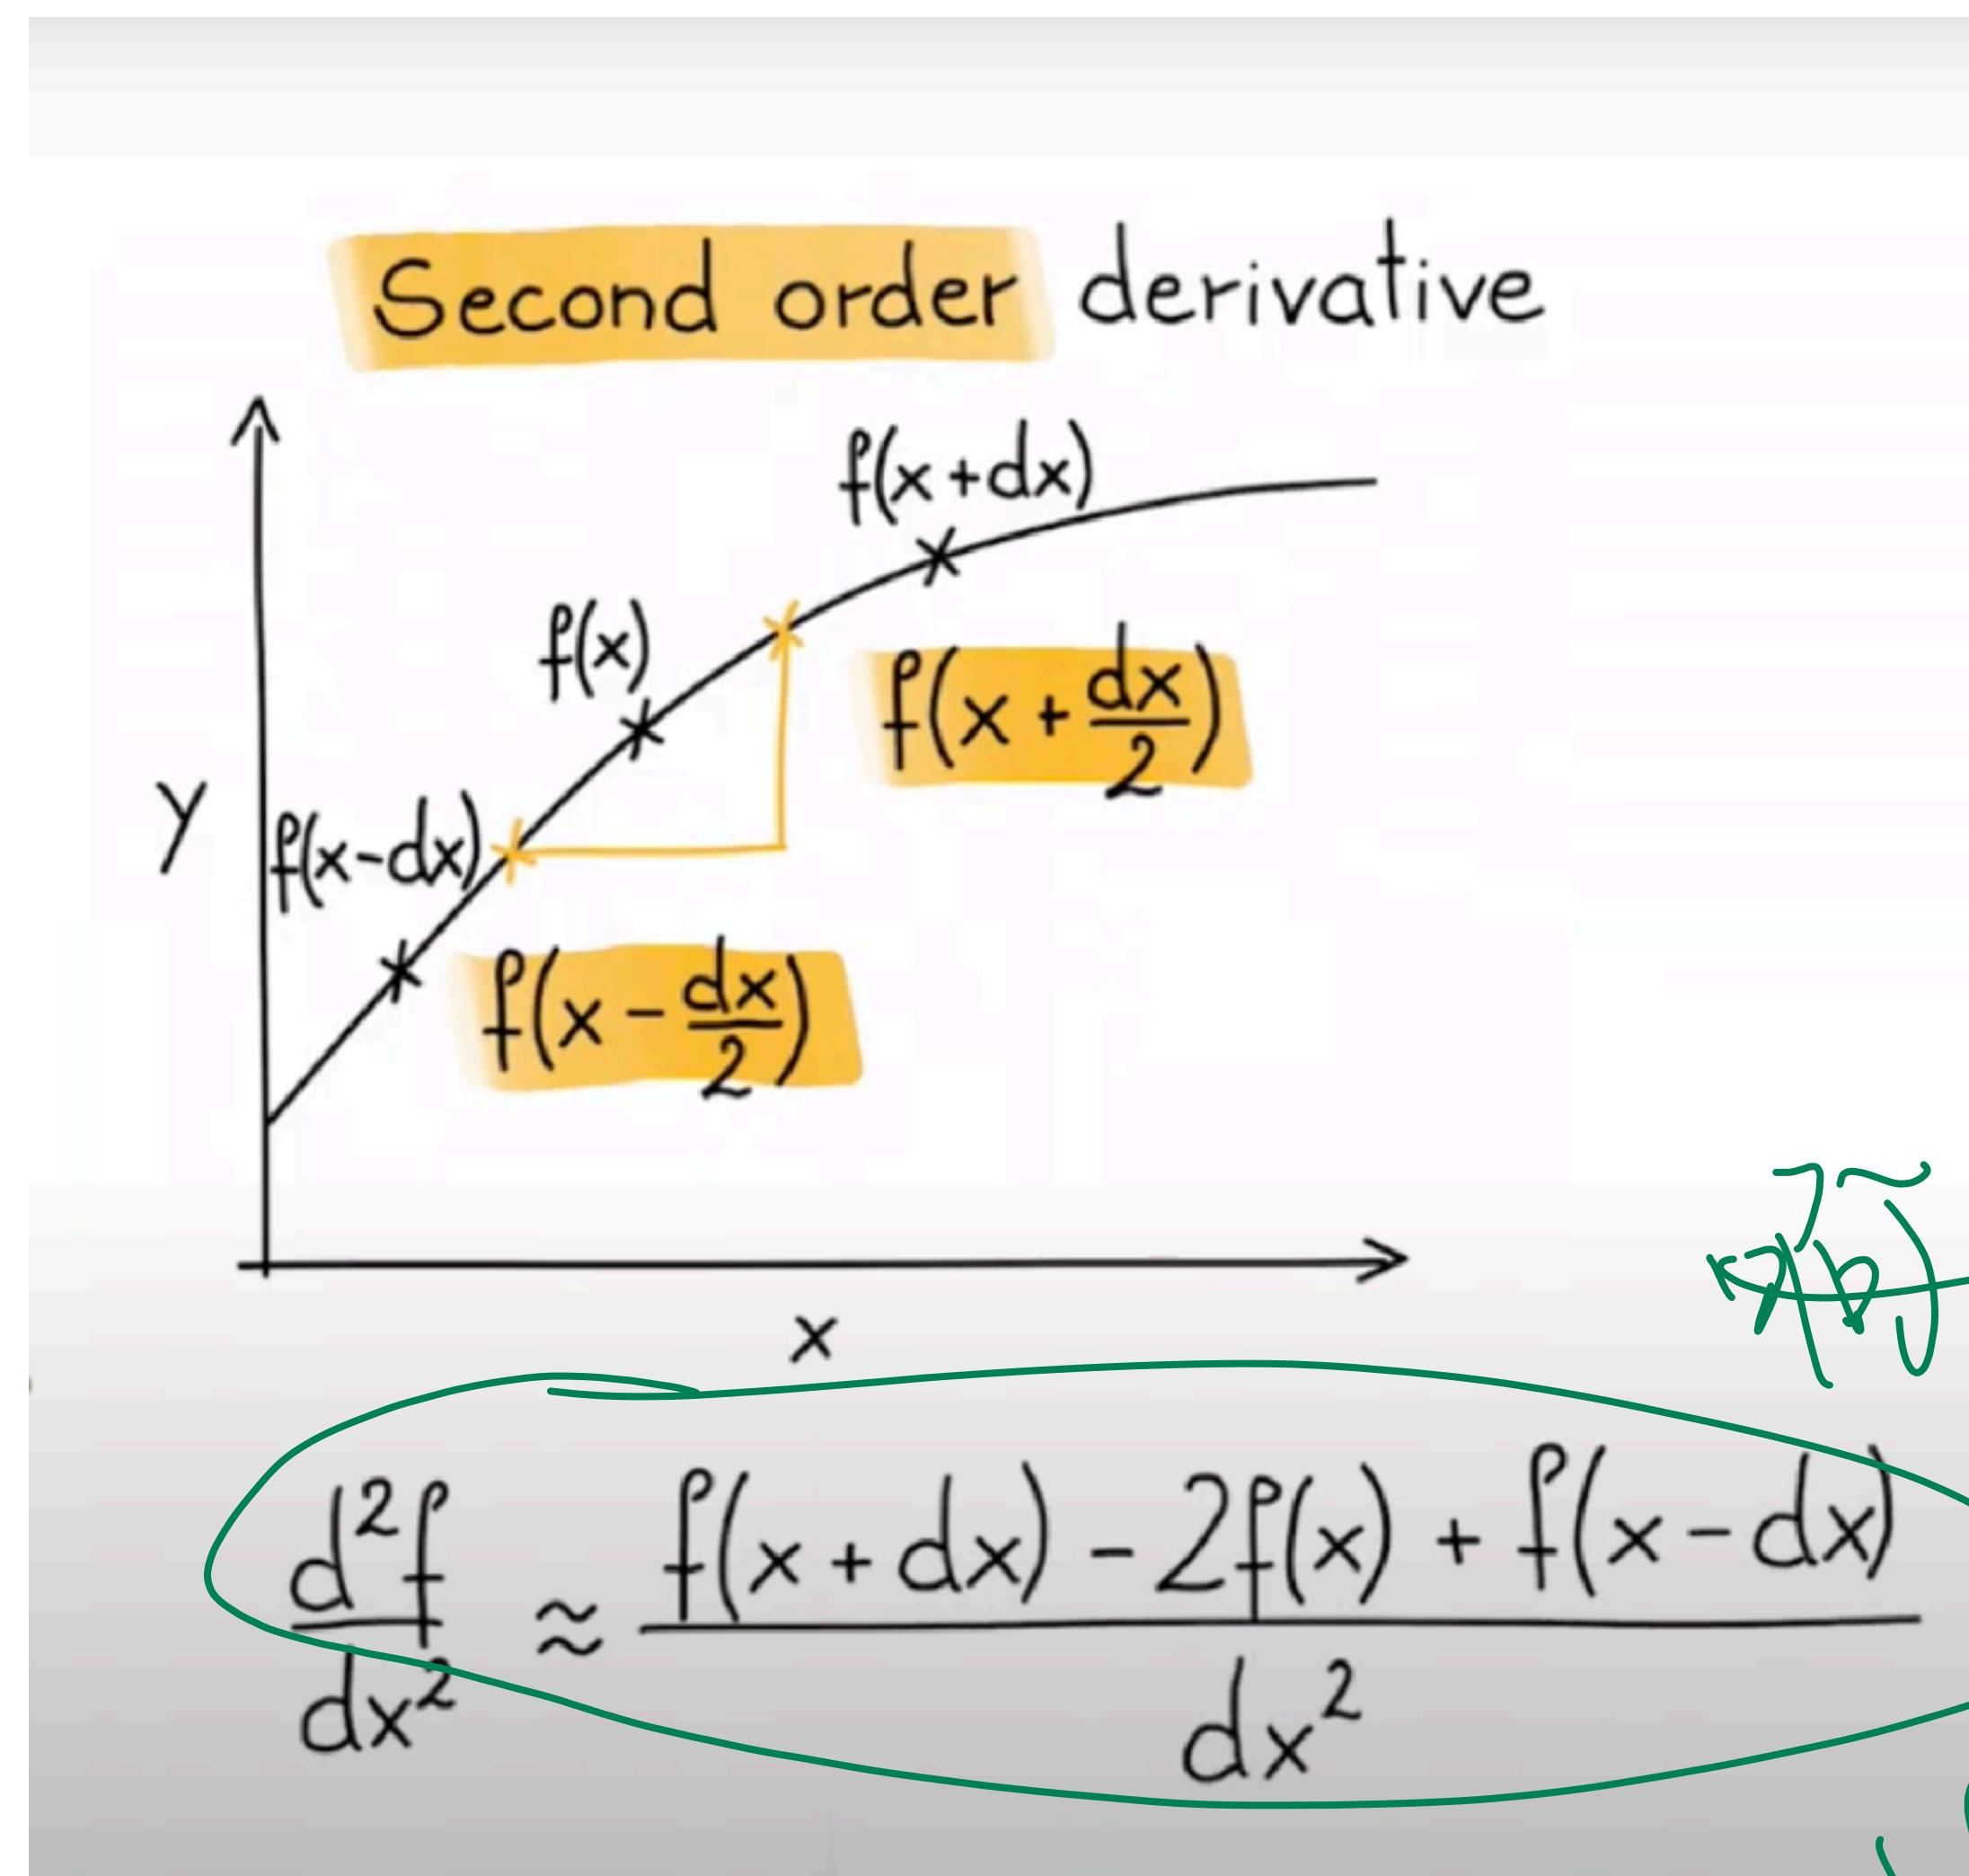


Image Filtering

1st and 2nd Derivative

$$\frac{\partial f}{\partial x} = f(x+1) - f(x)$$

(Look ahead)

$$\frac{\partial^2 f}{\partial x^2} = f(x+1) + f(x-1) - 2f(x)$$

[1 - 2 1]

(Look ahead and look back)

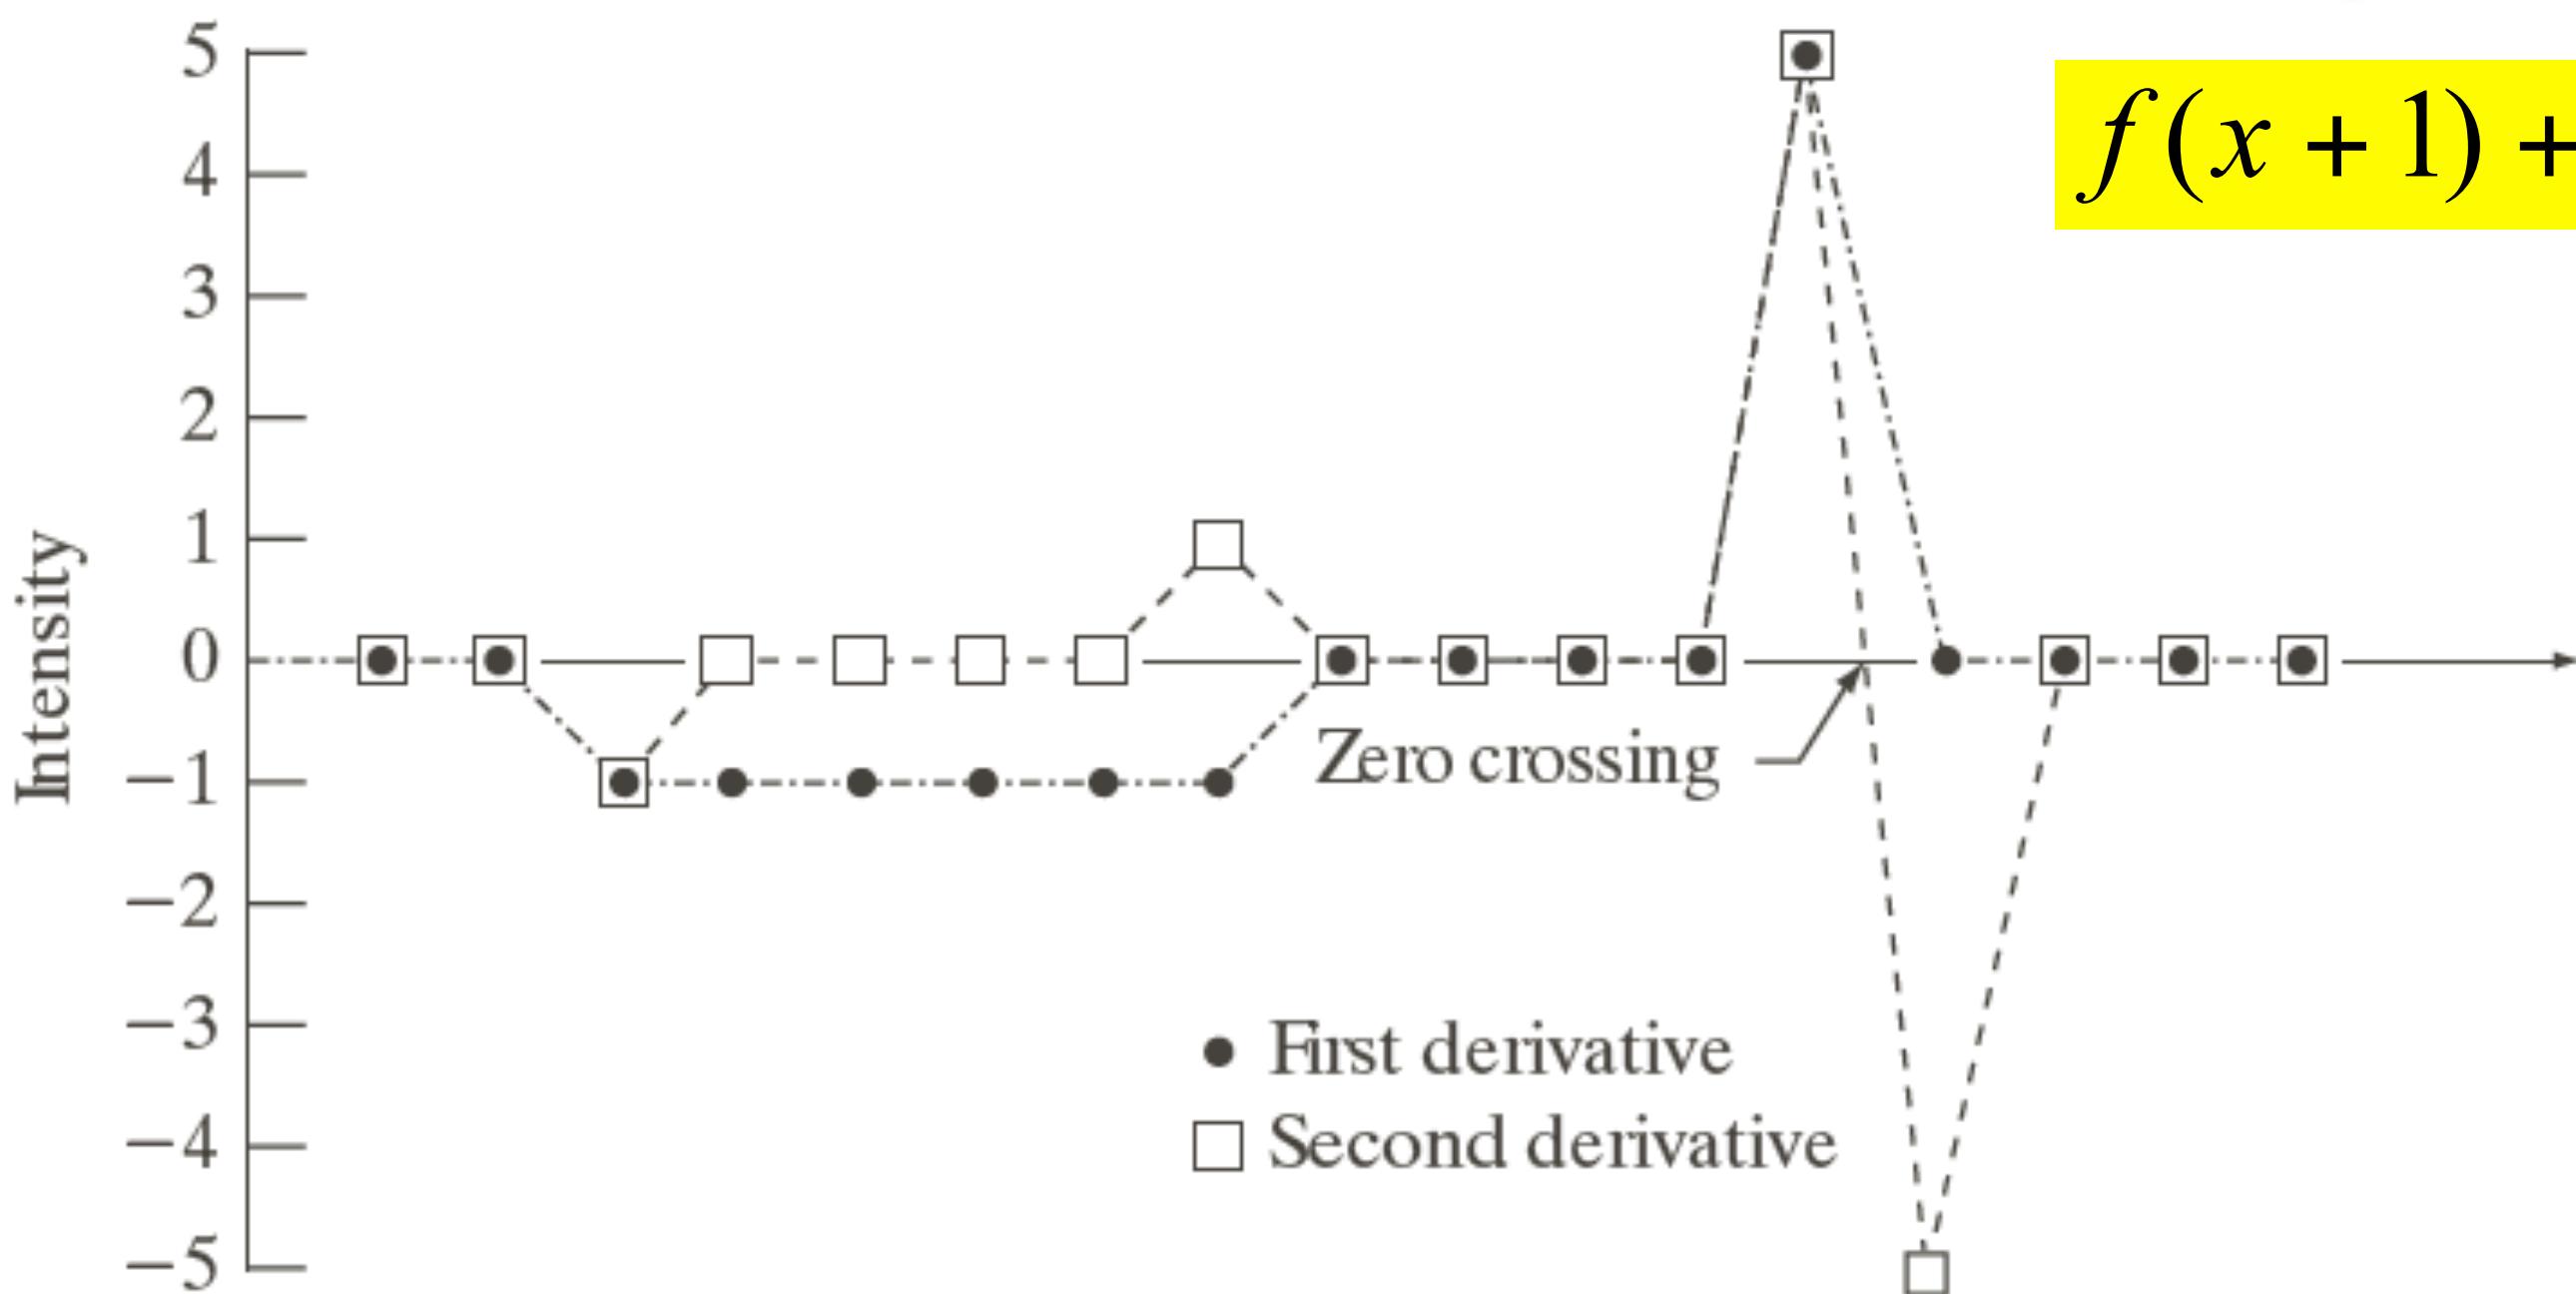
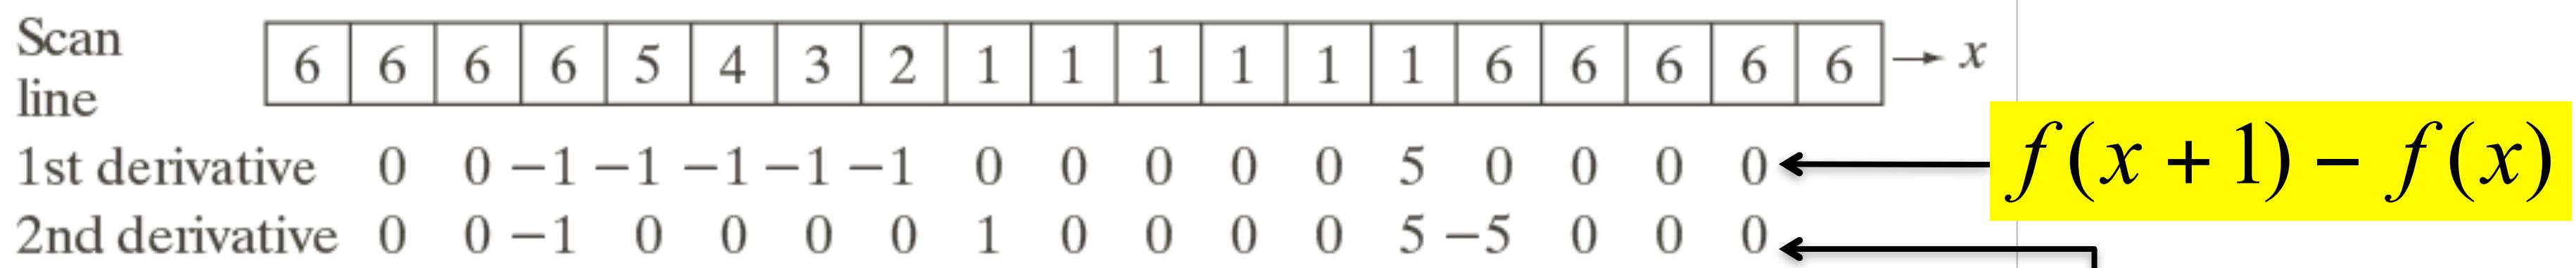
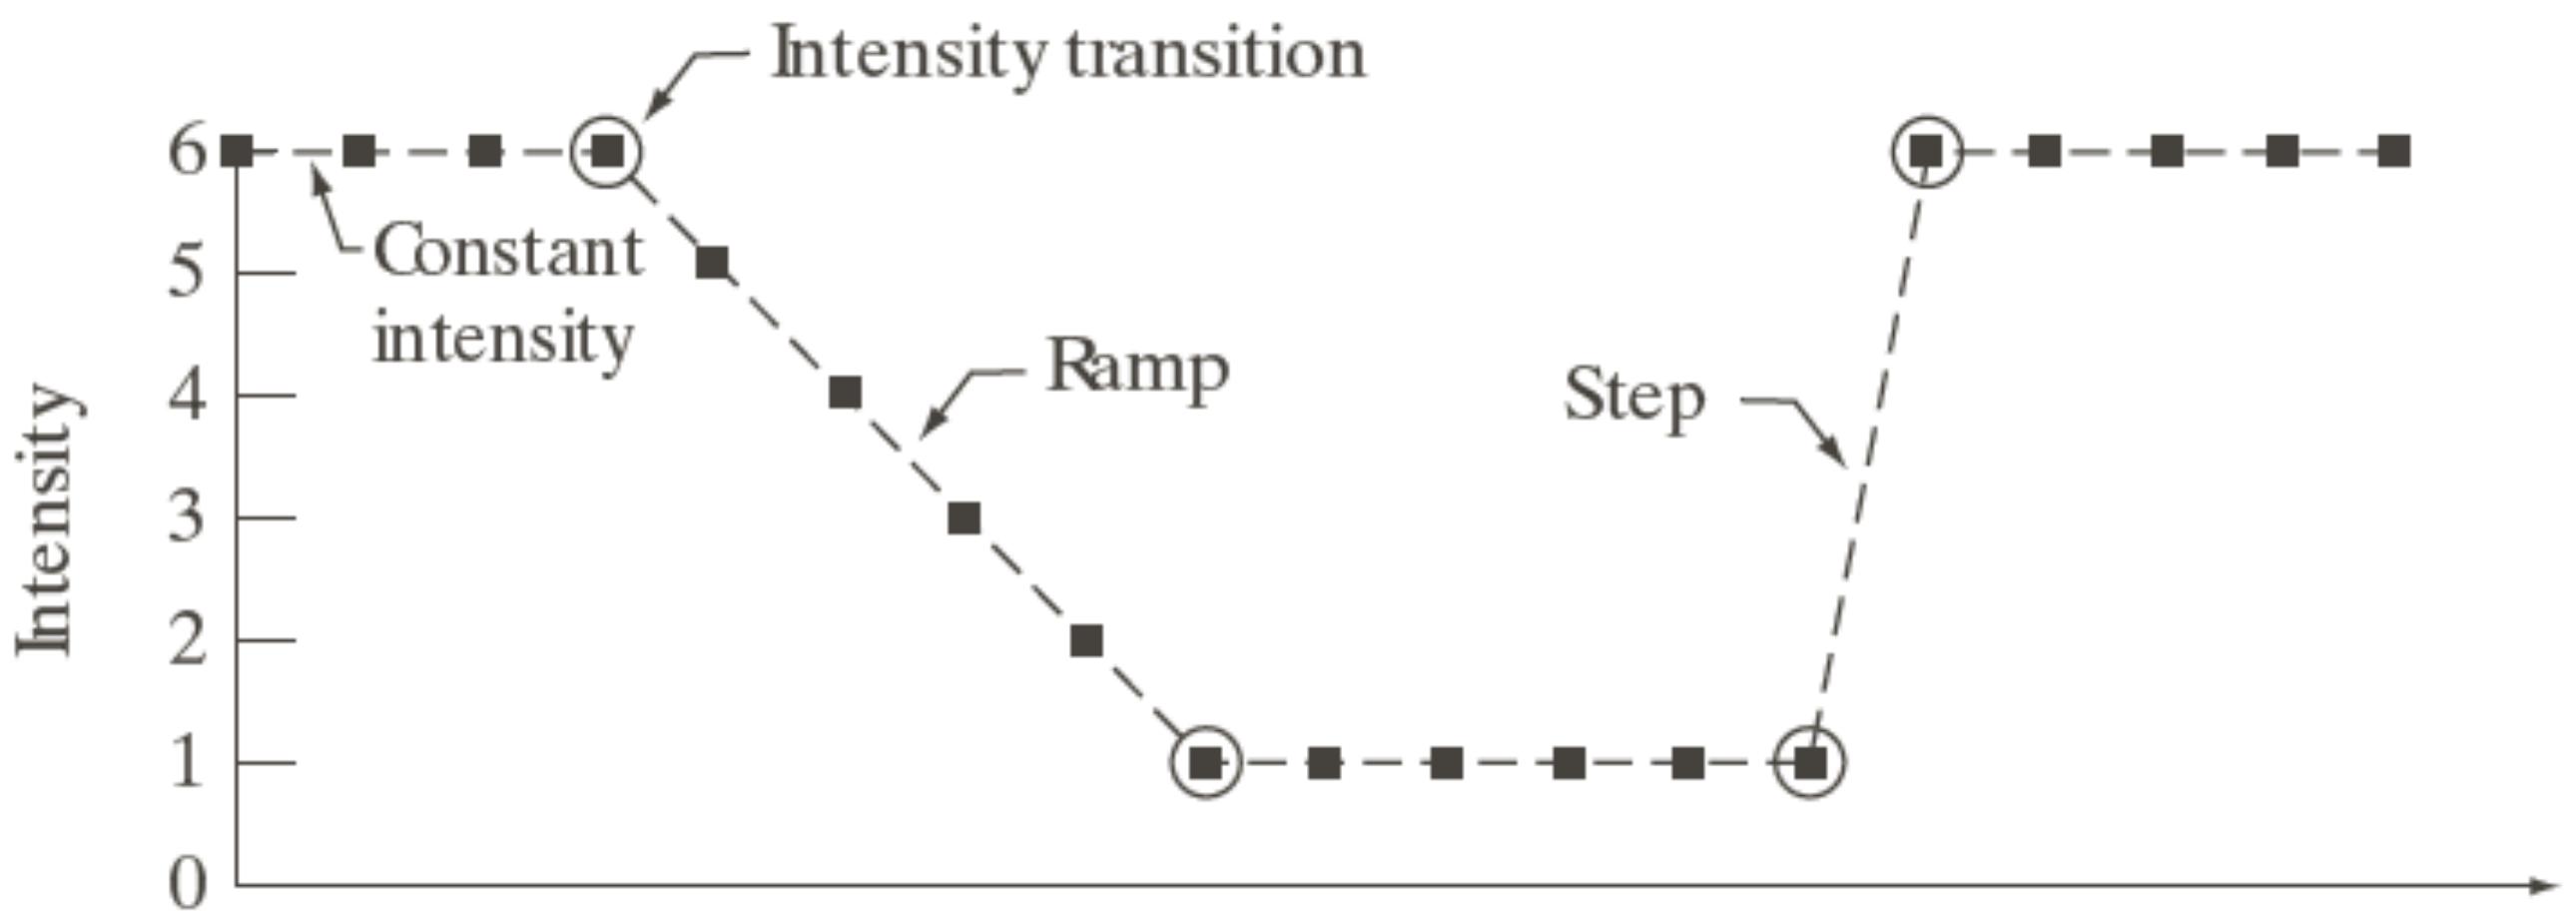


Figure from [DIP](#) by Gonzalez+Woods, 2008

Vertical Gradients from Finite Differences



Differential Filters

Prewitt operator:

$$\begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$$

give more weight to pixels in the middle

derivative respect to \downarrow



Sobel operator:

$$\begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

Sobel ∇

Differential Filters

Decomposing the Sobel filter

$$\begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline 2 & 0 & -2 \\ \hline 1 & 0 & -1 \\ \hline \end{array} = \begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 1 \\ \hline \end{array} \quad \begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline \end{array}$$

1st Derivative: From 1D to 2D

$$\frac{\partial f}{\partial x} = f(x+1) - f(x)$$

in 1D 2D:

$$\nabla f = \begin{bmatrix} \frac{\partial f}{\partial x} \\ \frac{\partial f}{\partial y} \end{bmatrix} = \begin{bmatrix} g_x \\ g_y \end{bmatrix}$$

- Gradient is anisotropic
- $mag(\nabla f) = \sqrt{g_x^2 + g_y^2}$ is isotropic

- Orientation of gradient: $\alpha(x, y) = \tan^{-1} \left[\frac{g_y}{g_x} \right]$

Imagine g_x and g_y as vectors that form two sides of a right triangle.
The result of vector-addition points in the direction of the gradient!

2nd Derivative: From 1D to 2D

$$\frac{\partial^2 f}{\partial x^2} = f(x+1) + f(x-1) - 2f(x) \quad \text{in 1D} \quad \boxed{?} \quad \text{2D: } \nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

$$\left. \begin{aligned} \frac{\partial^2 f}{\partial x^2} &= f(x+1, y) + f(x-1, y) - 2f(x, y) \\ \frac{\partial^2 f}{\partial y^2} &= f(x, y+1) + f(x, y-1) - 2f(x, y) \end{aligned} \right\}$$

0	1	0
1	-4	1
0	1	0

1	1	1
1	-8	1
1	1	1

-1	-1	-1
-1	8	-1
-1	-1	-1

Incorporating diagonals,
isotropic in increments of 45°

Laplacian: use this (negative) version
of kernel for convenience

Image Sharpening

- Also known as *enhancement*
- Increases the high frequency components to enhance edges
- $I' = I + \alpha(k^*I)$, where k is a high-pass filter kernel and α is a scalar in $[0,1]$

Image Sharpening: Example 1



$$I' = I + \alpha \nabla^2 I = I + \alpha(K * I)$$

A diagram illustrating a 3x3 kernel for image sharpening. The kernel is represented by a 3x3 grid of numbers:

-1	-1	-1
-1	8	-1
-1	-1	-1

An upward-pointing arrow is positioned above the grid, indicating its application to the image.

Figure from NASA, obtained on [DIP](#)

Image Sharpening: Example 1

$$I' = I + \alpha \nabla^2 I = I + \alpha(K * I)$$

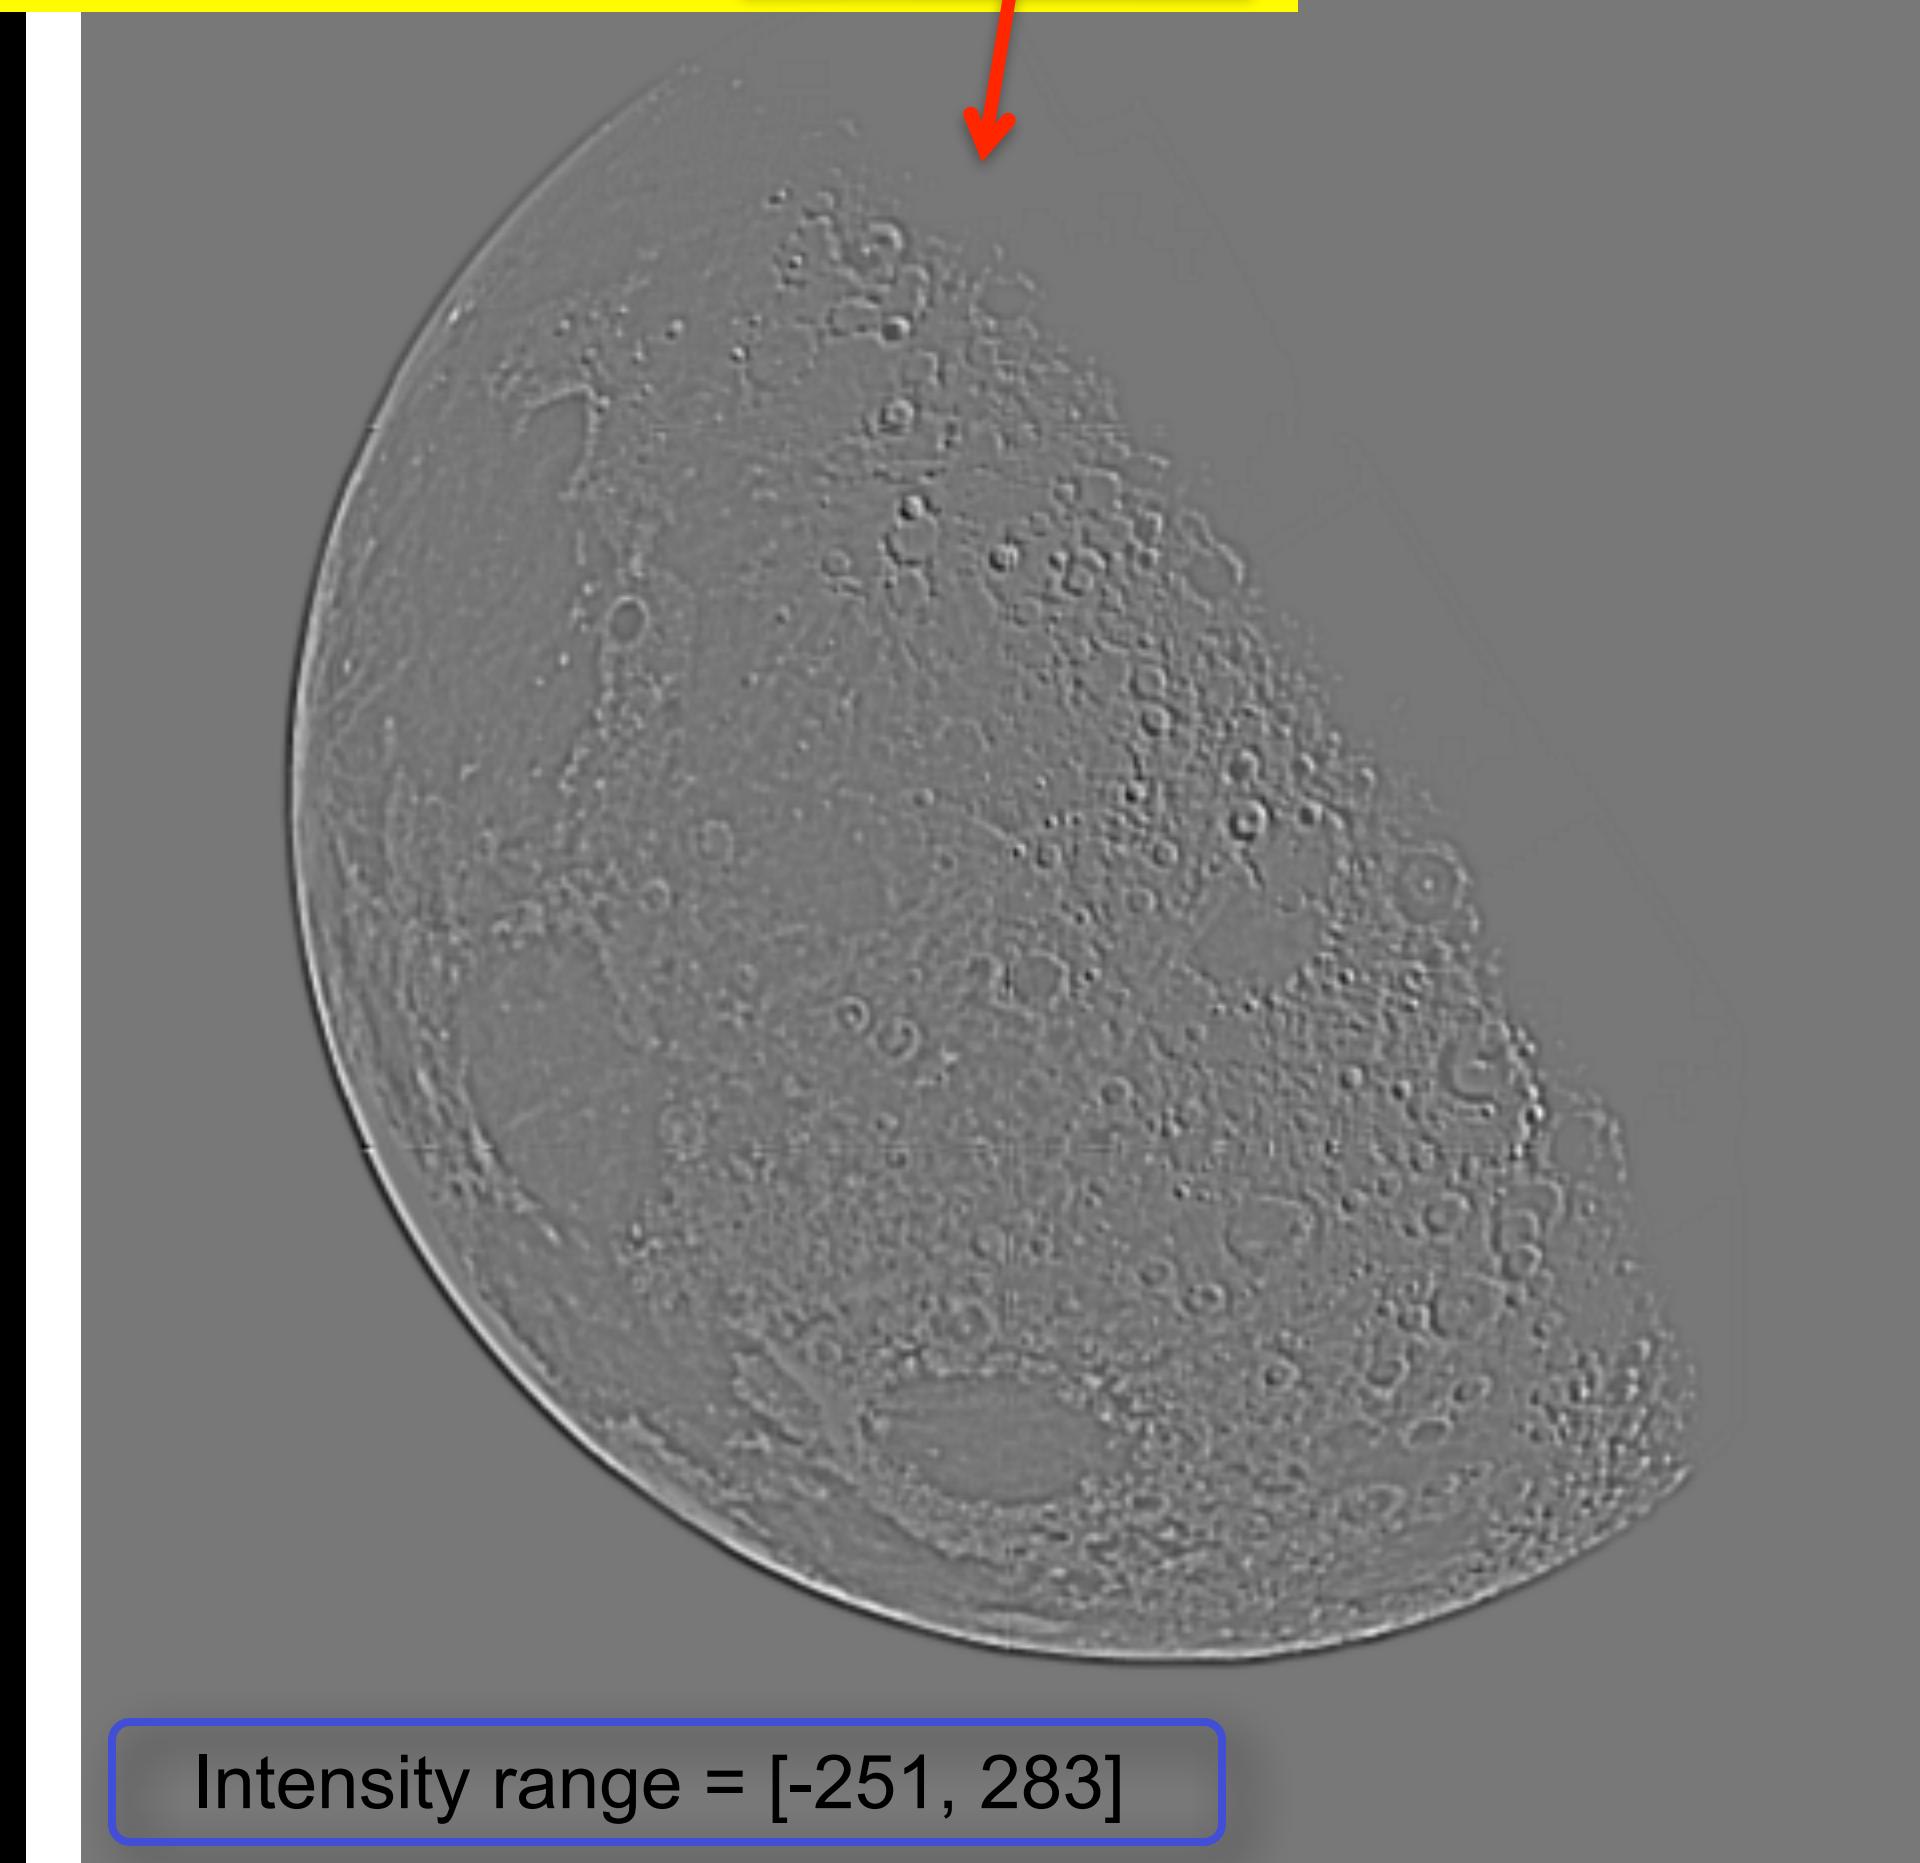
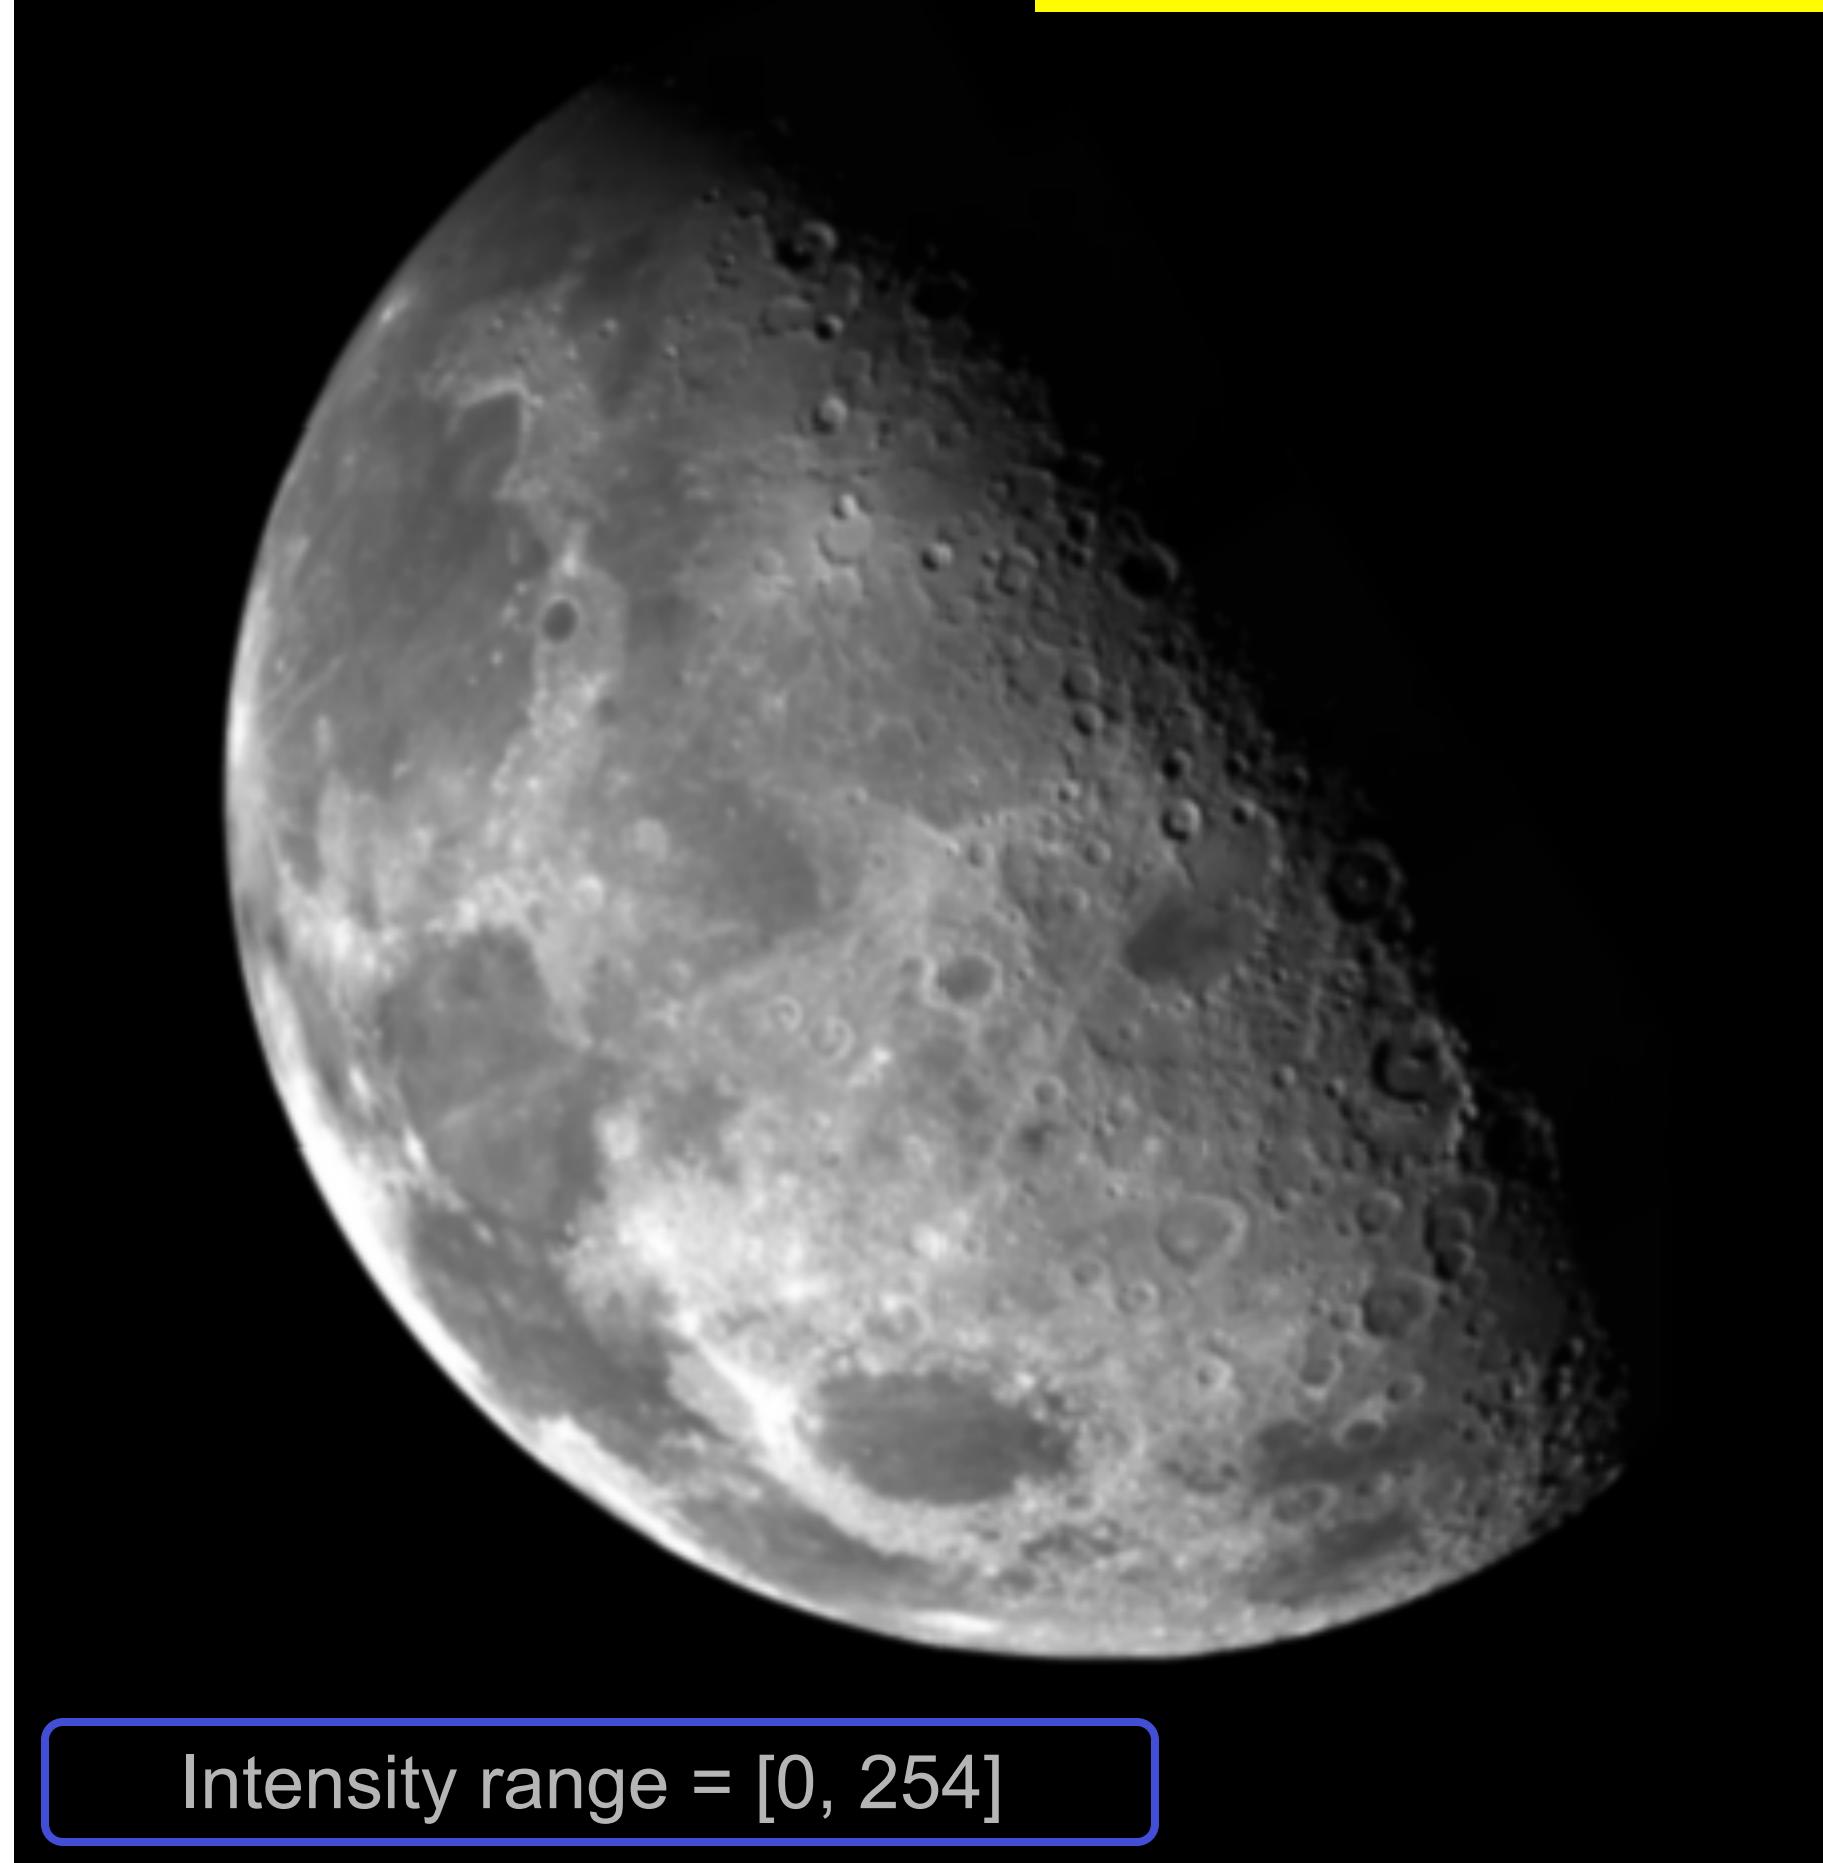


Image Sharpening: Example 1



/

COMP0026: Image Processing

Image Filtering

/'



Intensity range = [-225, 473]

Image Sharpening: Example 1



/

COMP0026: Image Processing

Image Filtering



I'

Intensity range = [-225, 473]

Image Sharpening: Example 2

$$I' = I + \alpha(I - K * I) \Leftarrow \text{"Unsharp Mask"}$$

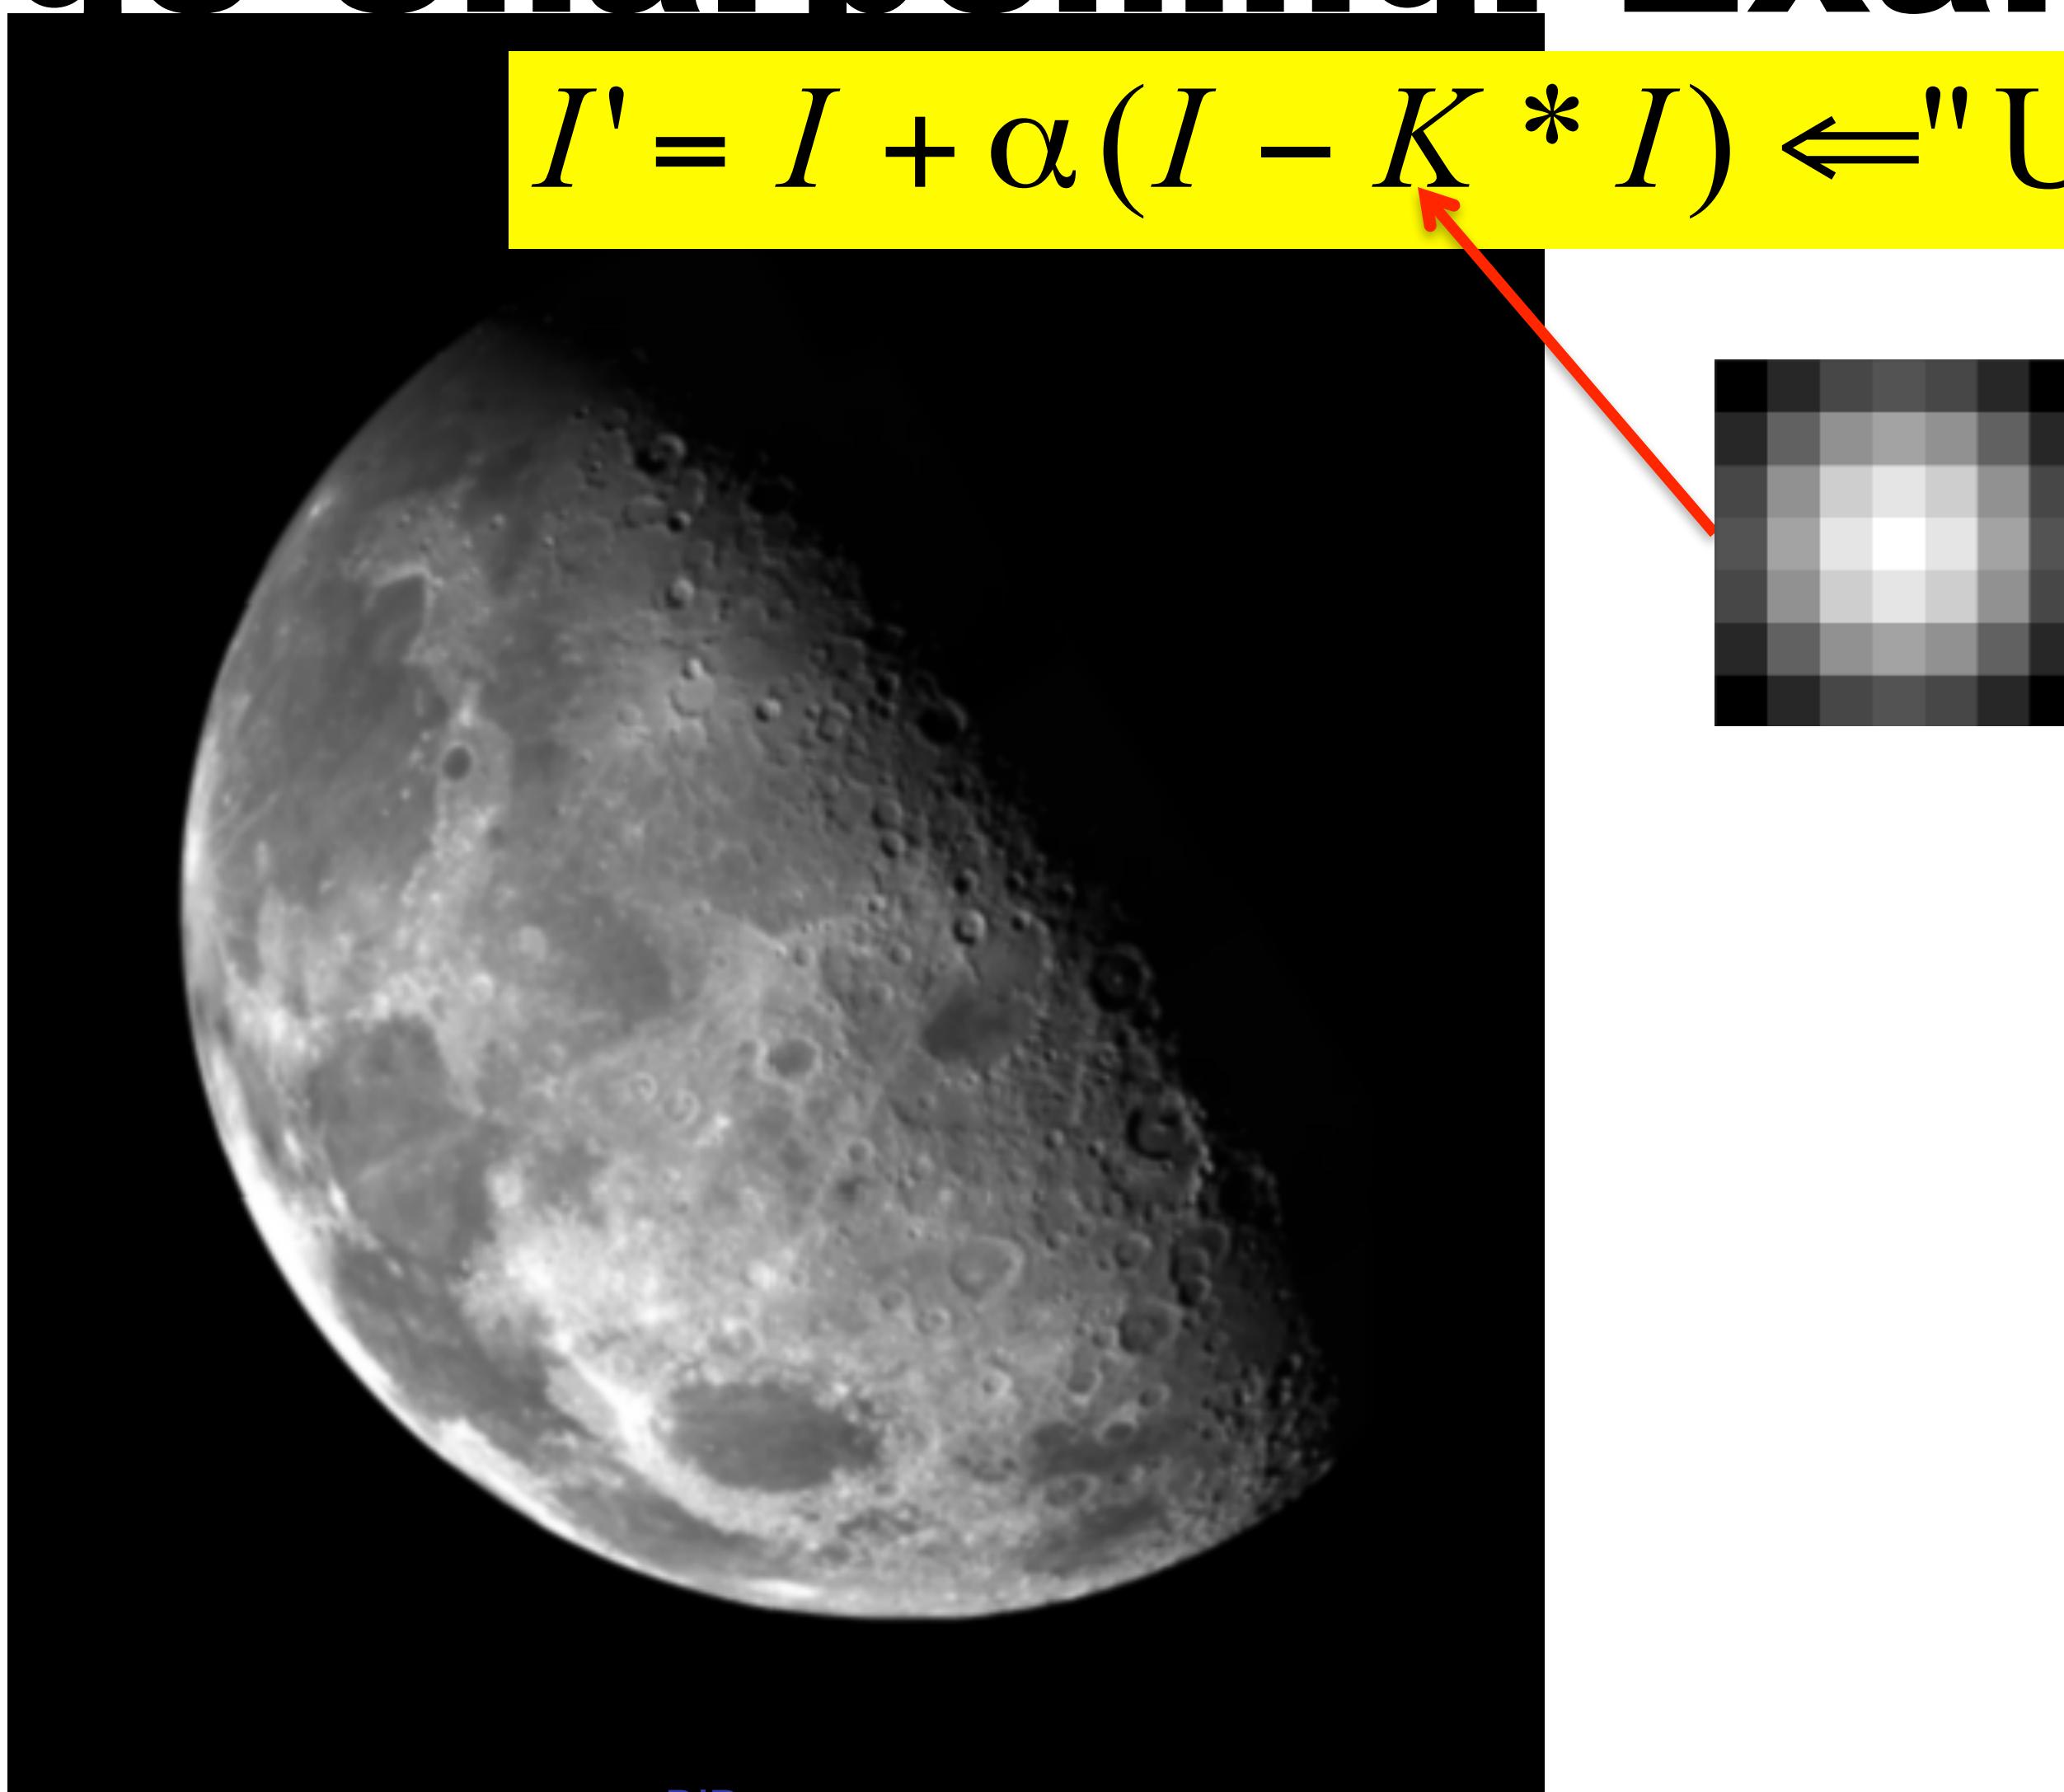


Figure from NASA, obtained on [DIP](#)

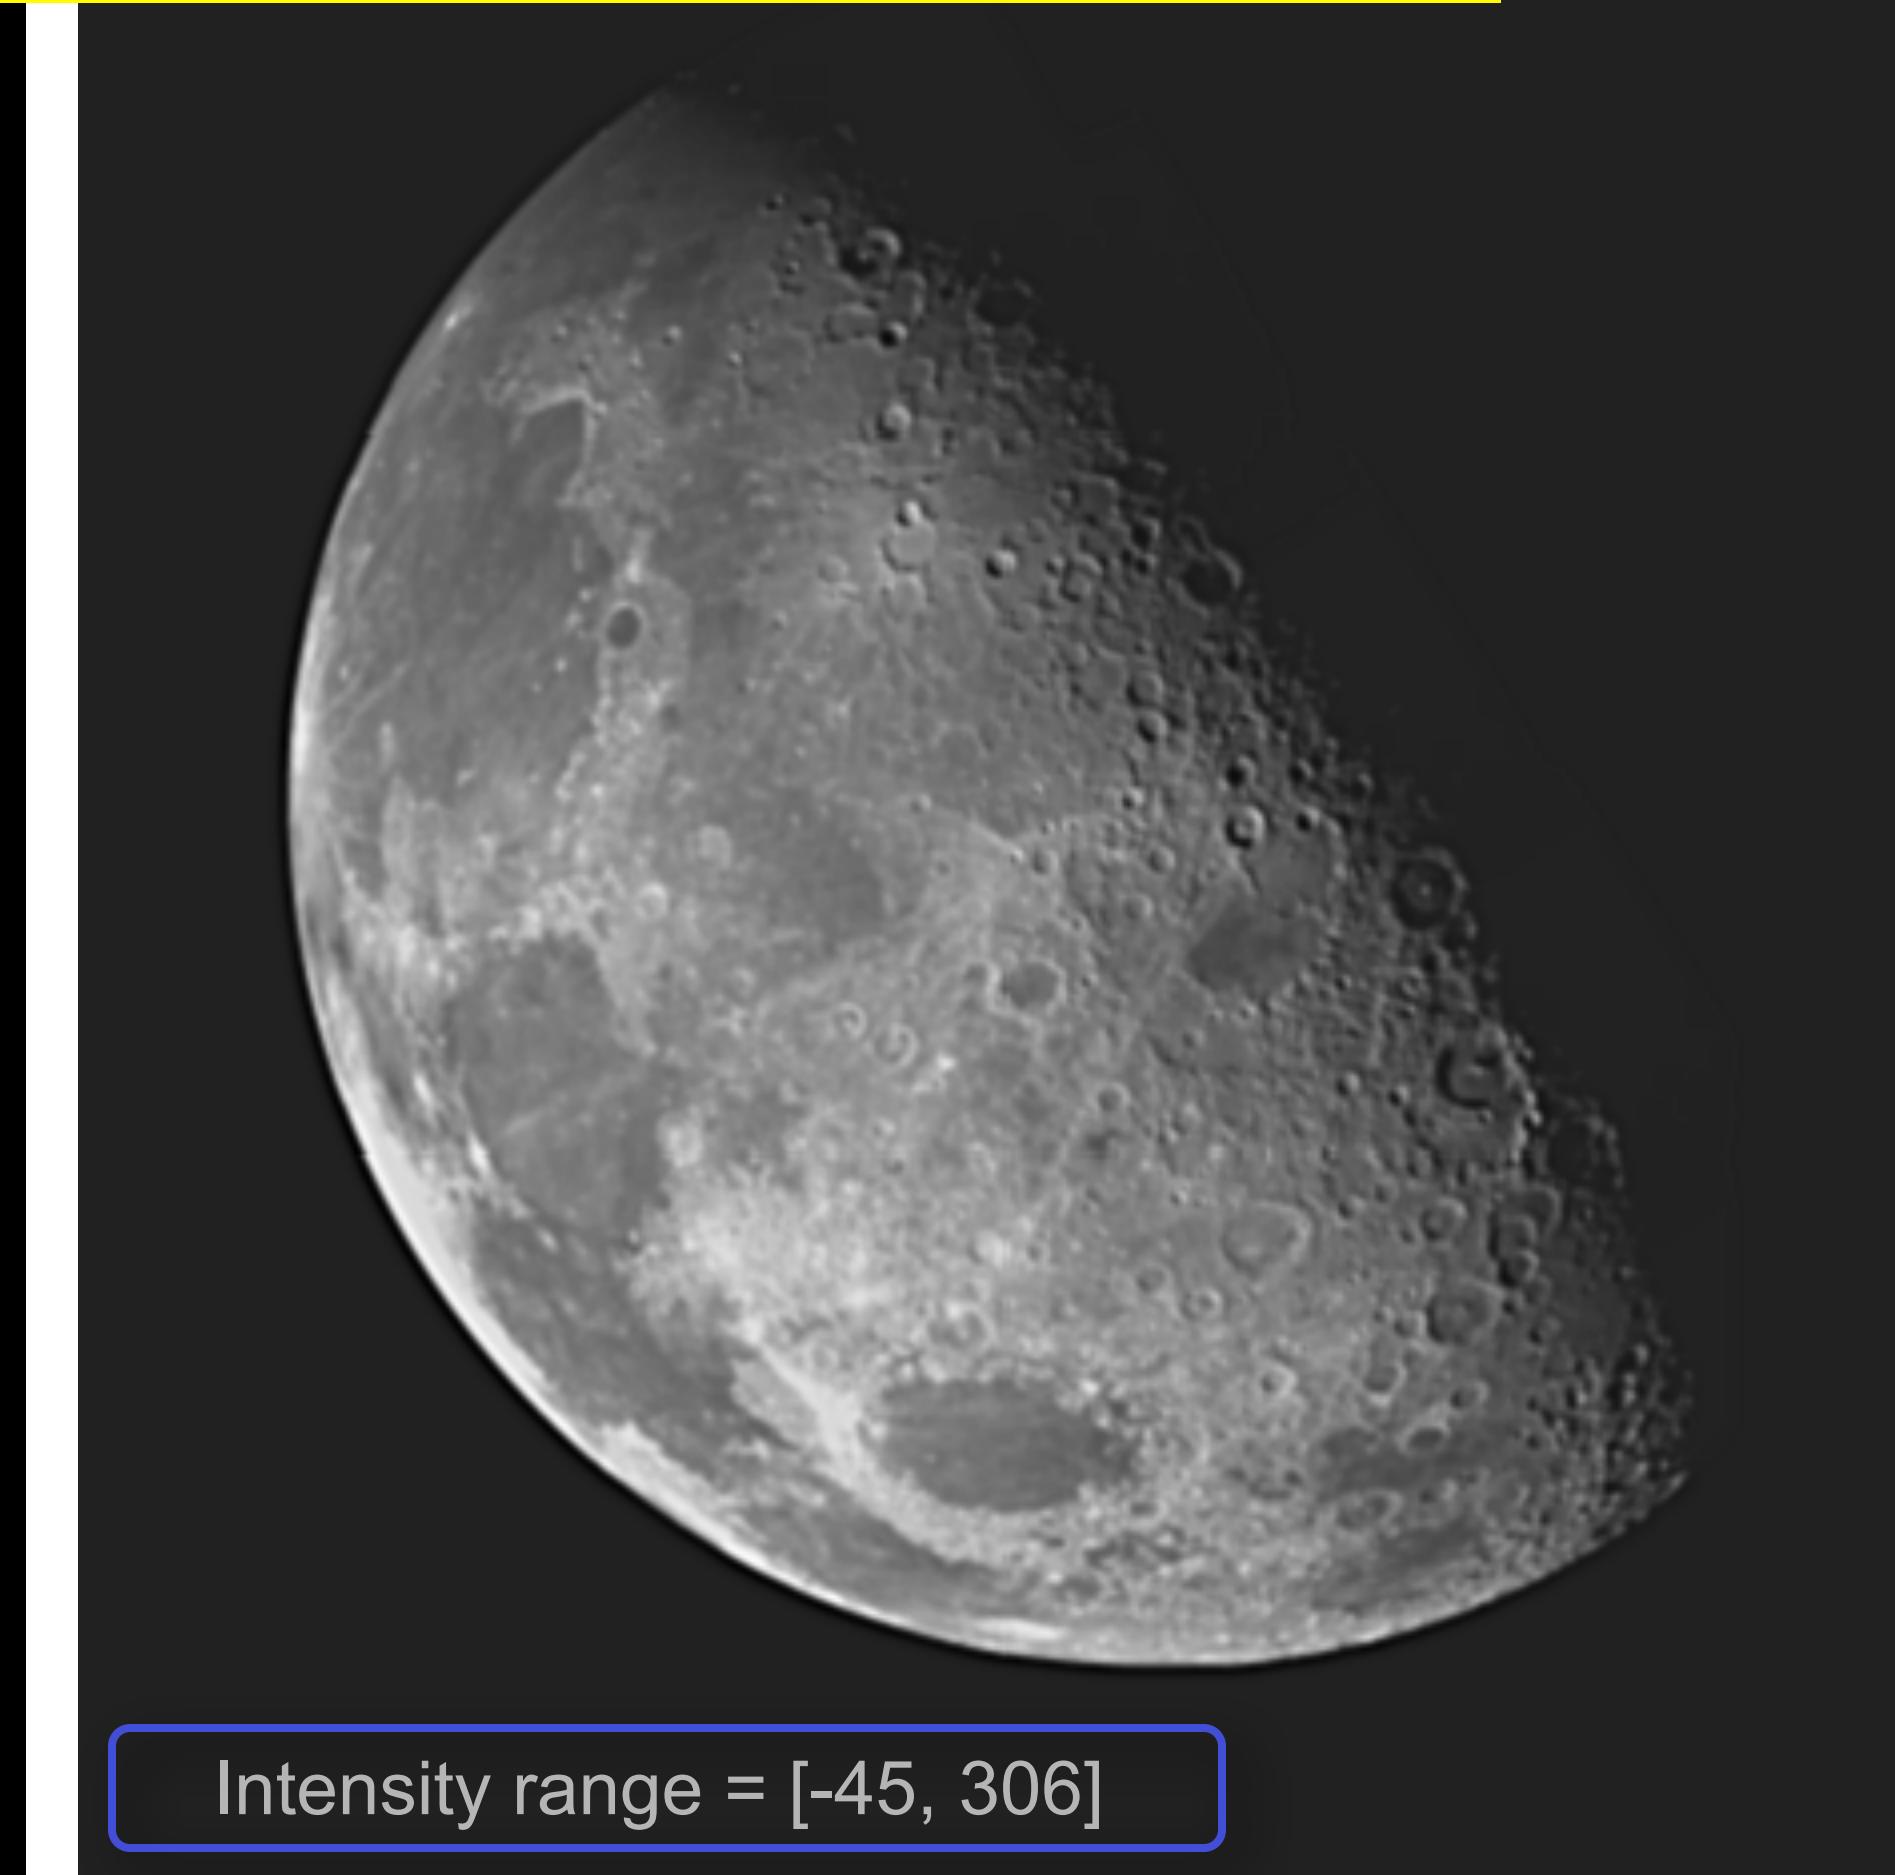
Image Sharpening: Example 2

$$I' = 2I - K * I \Leftarrow \text{"Unsharp Mask"}$$



/

/'



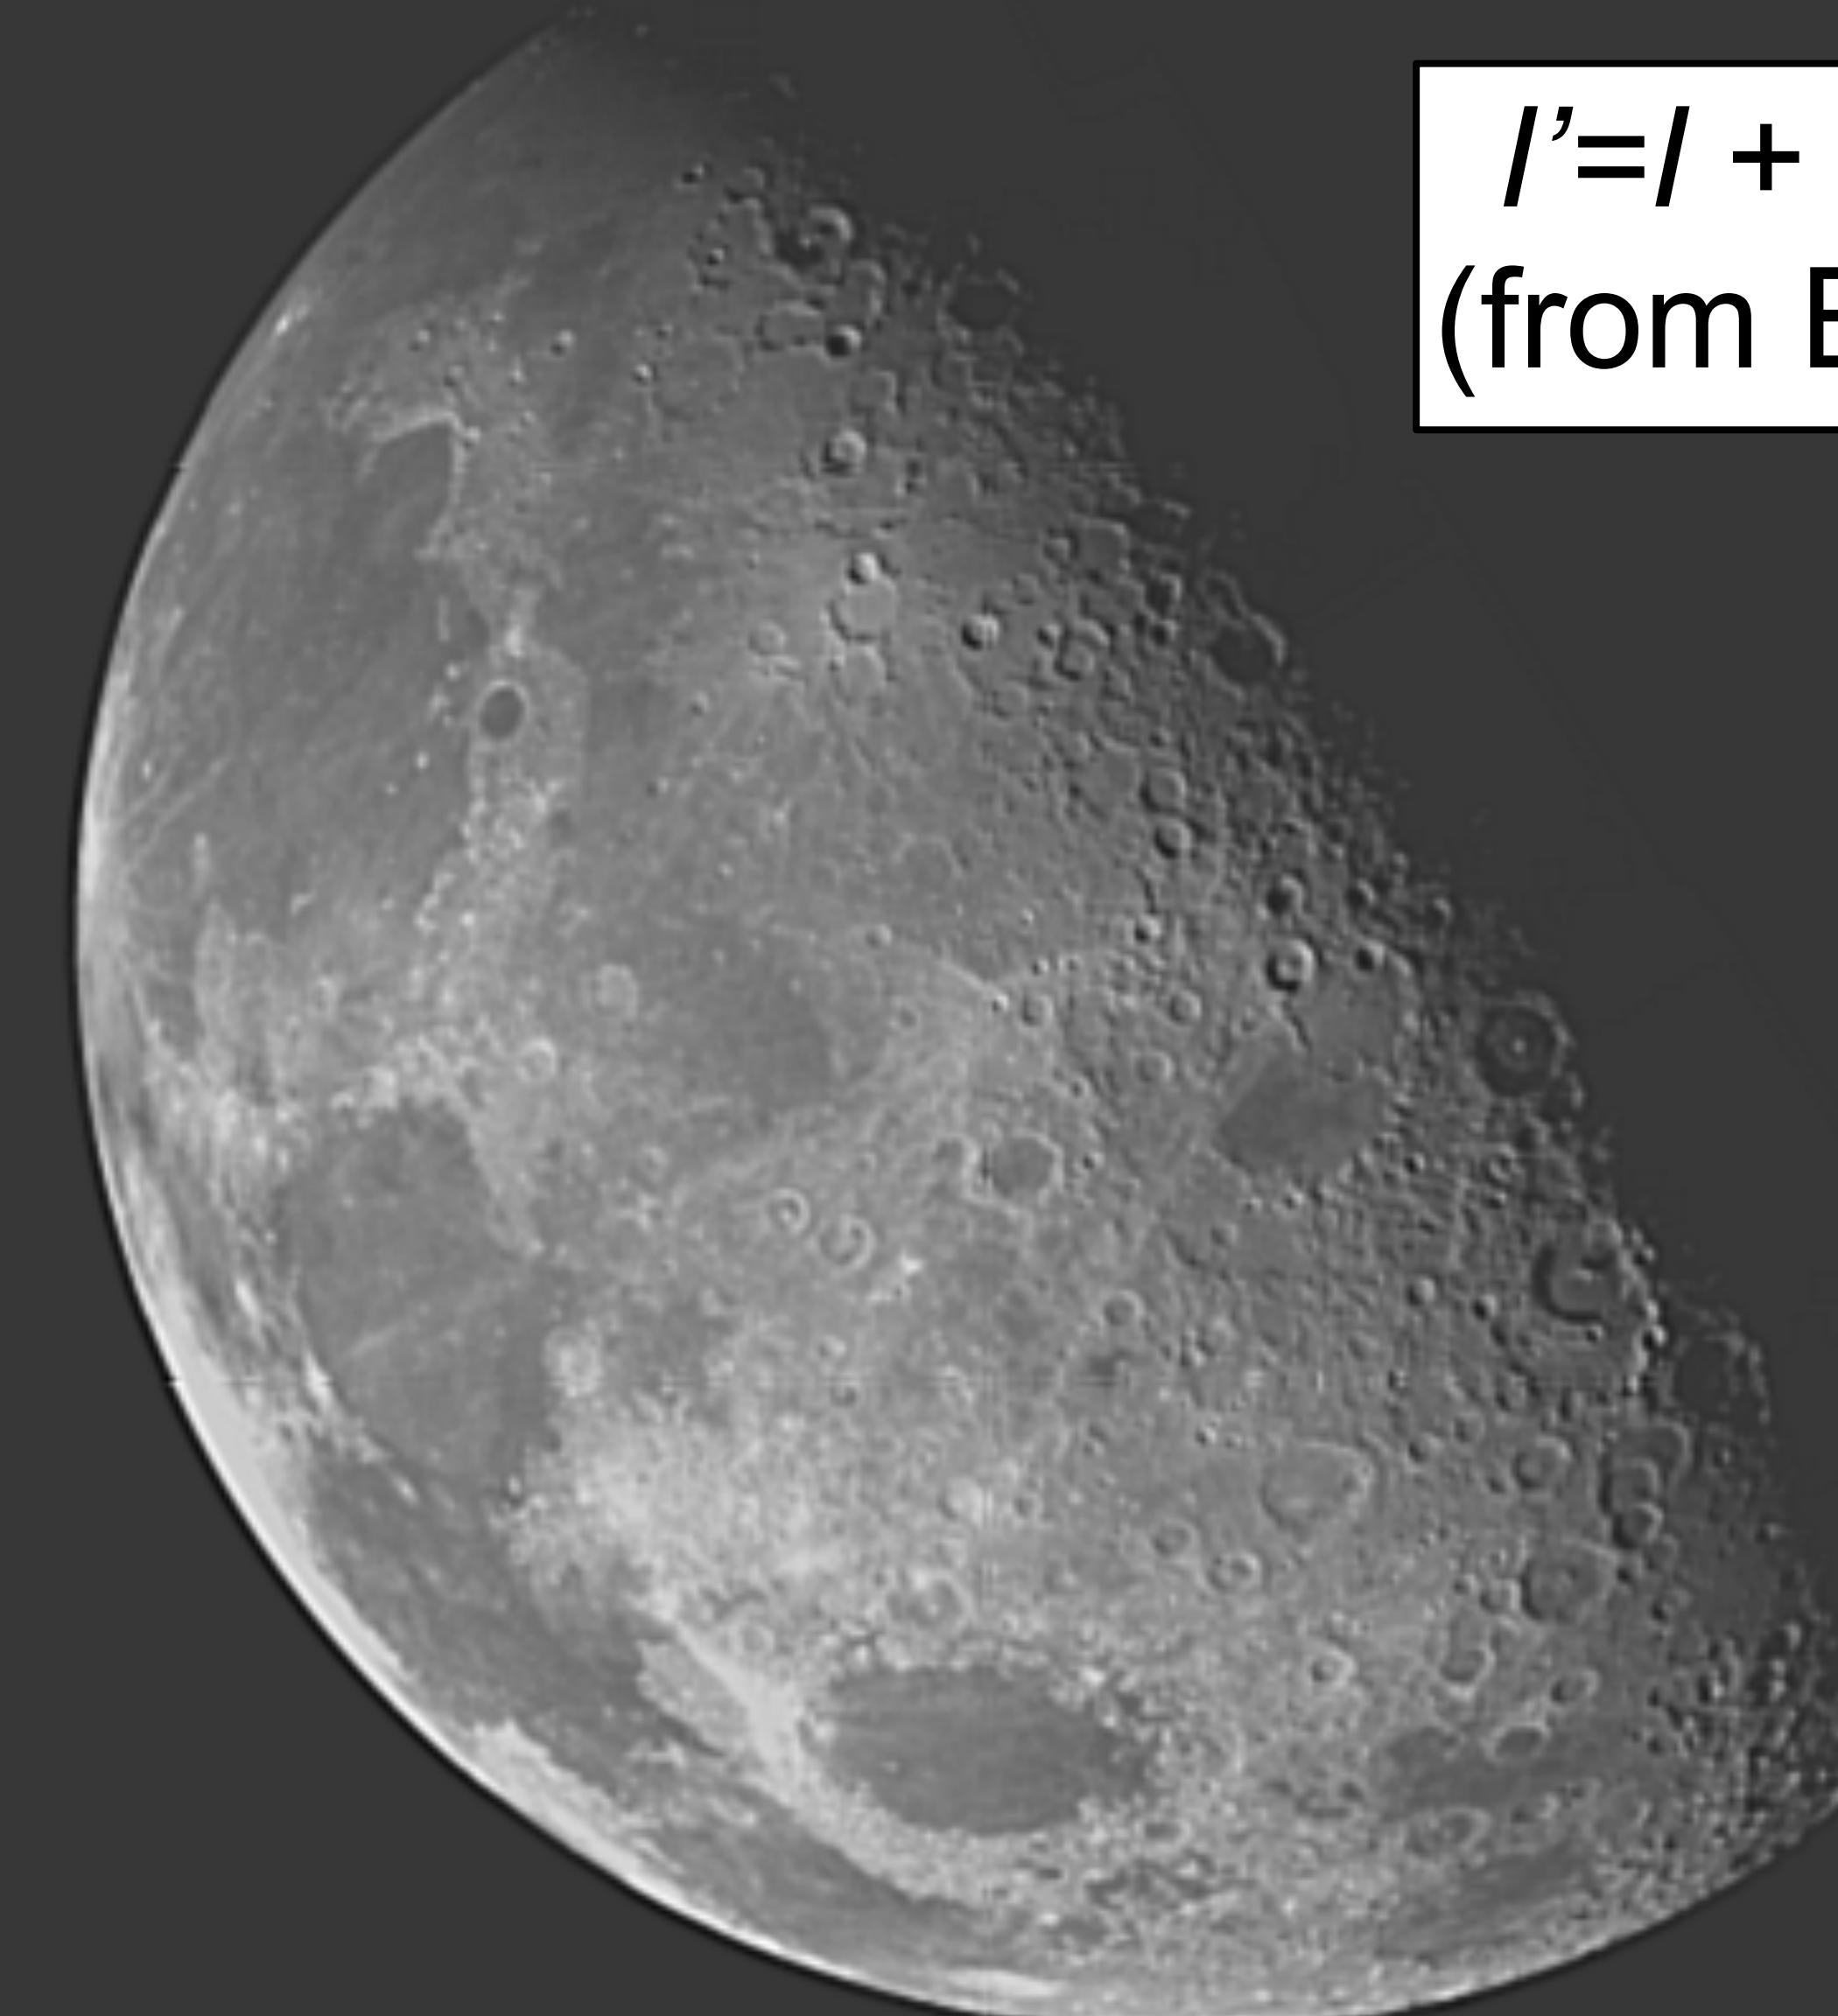
Intensity range = [-45, 306]





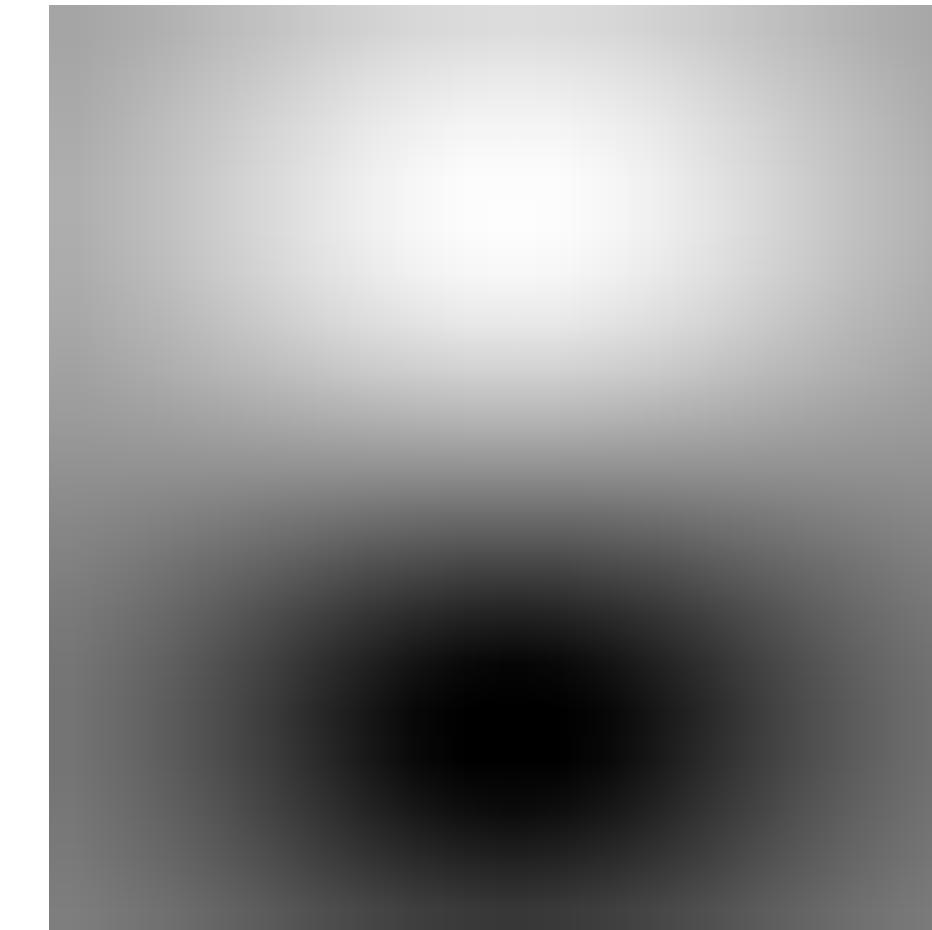
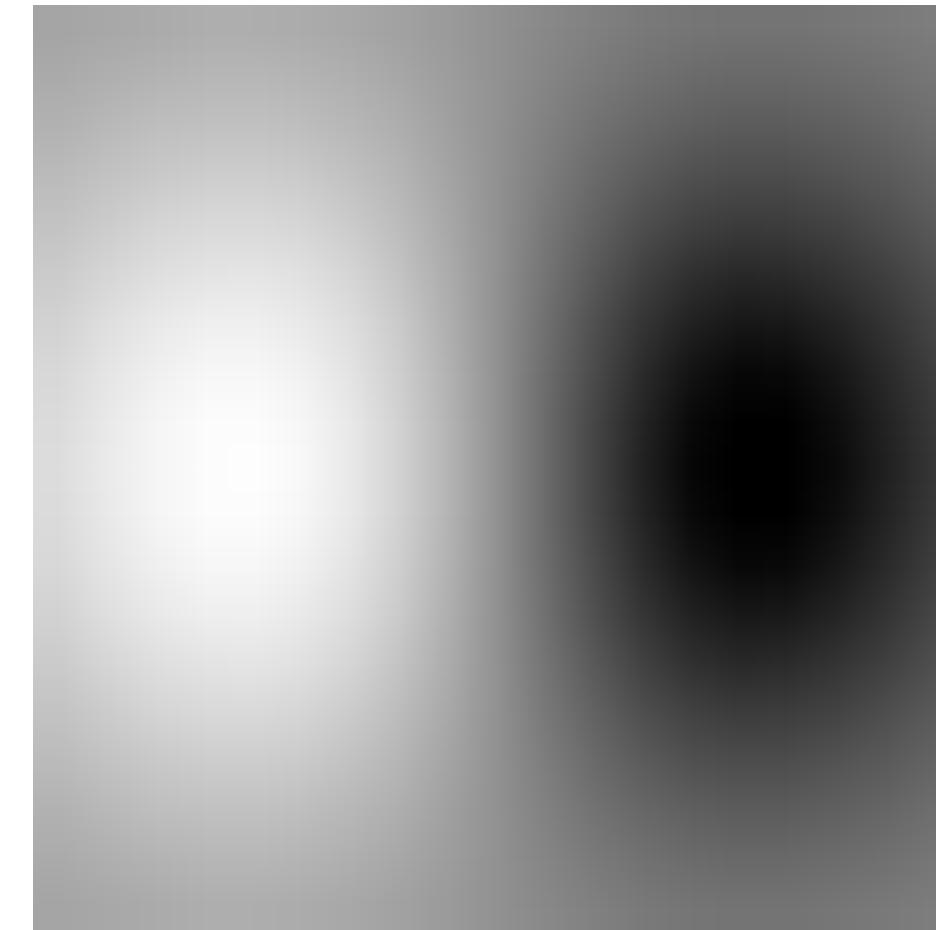
I' = Unsharp Mask
(from Example 2)

$I' = I + \text{Laplacian}$
(from Example 1)



Filters are Templates

- Filter at some point can be seen as taking a dot-product between the image and some vector
- Filtering the image is a set of dot products
 - filters look like the effects they are intended to find
 - filters find effects they look like



Slide credit: D.A. Forsyth

Scaled Representations

- Find *correlations/convolutions* at all scales
 - e.g., when finding hands or faces, we don't know what size they will be in a particular image
 - Template size is constant, but image size changes
- Efficient search for *correspondence*
 - look at coarse scales, then refine with finer scales
 - much less cost, but may miss best match
- Examining all *levels of detail*
 - Find edges with different amounts of blur
 - Find textures with different spatial frequencies (*levels of detail*)

Slide credit: D.A. Forsyth