Deep Learning for Visual Computing (COMP0169)

Convolutional Neural Networks

Niloy Mitra **Tobias Ritschel**

UCL

Introduction

- Idea
- Evolution

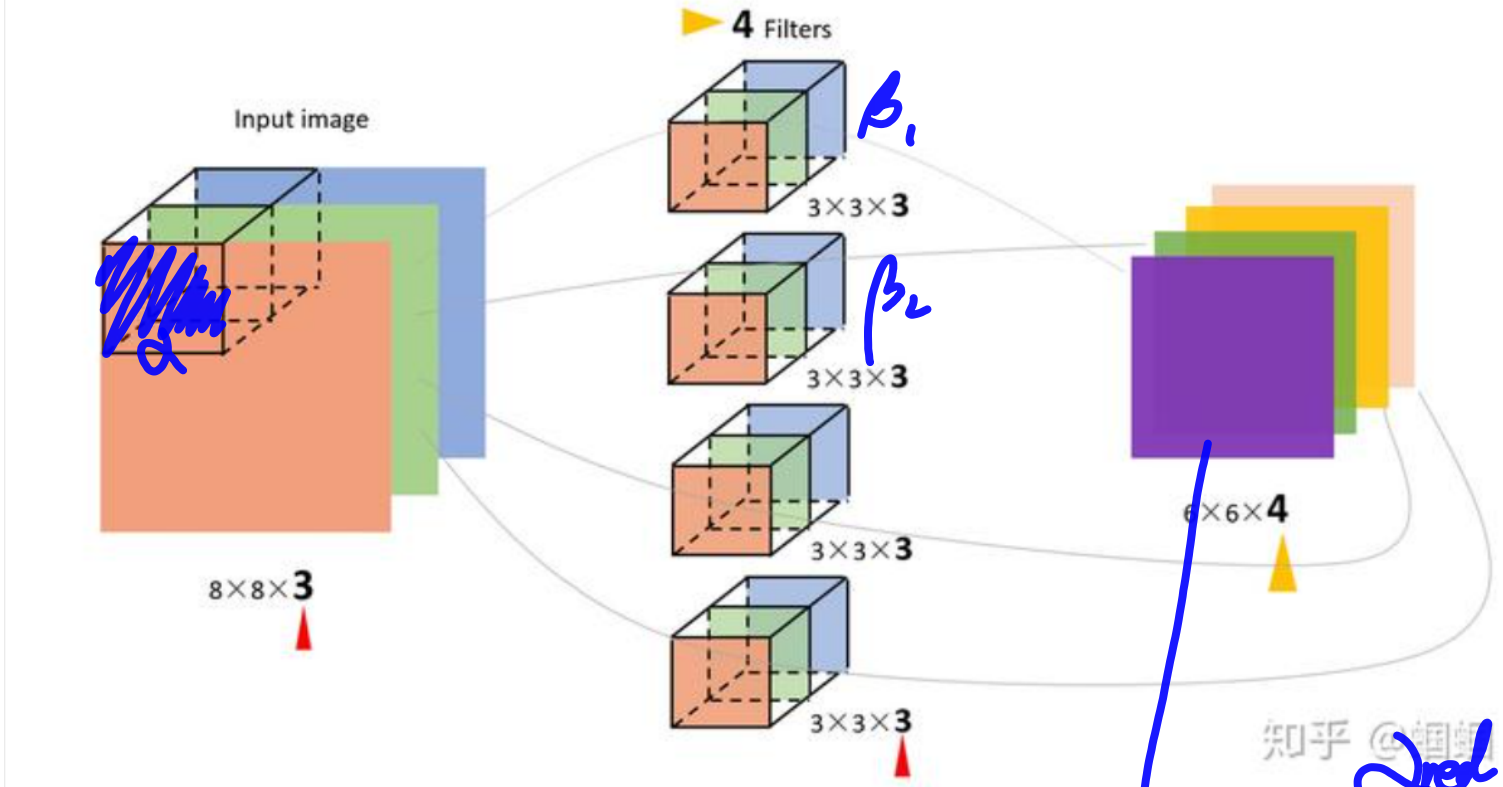

Where do filters come from?

- So far we said we will code those filters ourselves
- Which is the right filter for a task?

Quiz

I say “  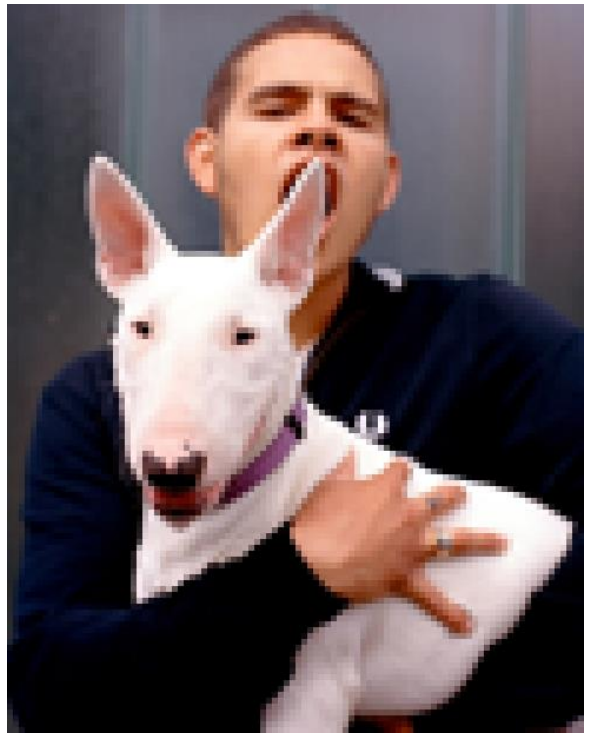 “? You say (1, 1, 1)/3

Input Output

More quiz

I say “  “?  You say (-1,0,1)/2

Input

Output

More quiz, computer helping

I say “   “?

Input

Output

Computer says ... a lot of numbers

Finding a filter by optimization

- Optimize the filter kernel to produce a desired result

$$\arg\min_\theta ||f \odot g_\theta - h||_2$$

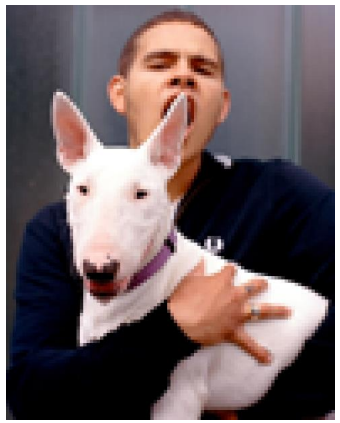

kernel

$$\begin{matrix} g_{11} & g_{12} & g_{13} \\ g_{21} & g_{22} & g_{23} \\ g_{31} & g_{32} & g_{33} \end{matrix}$$

$$\theta \in \mathbb{R}^{3\times 3} = \{g_{11}, \ldots, g_{33}\}$$

Linear filtering

- Convolution is a linear operation, so this can be solved quite easily

$$\arg\min_\theta ||f \odot g_\theta - h||_2$$

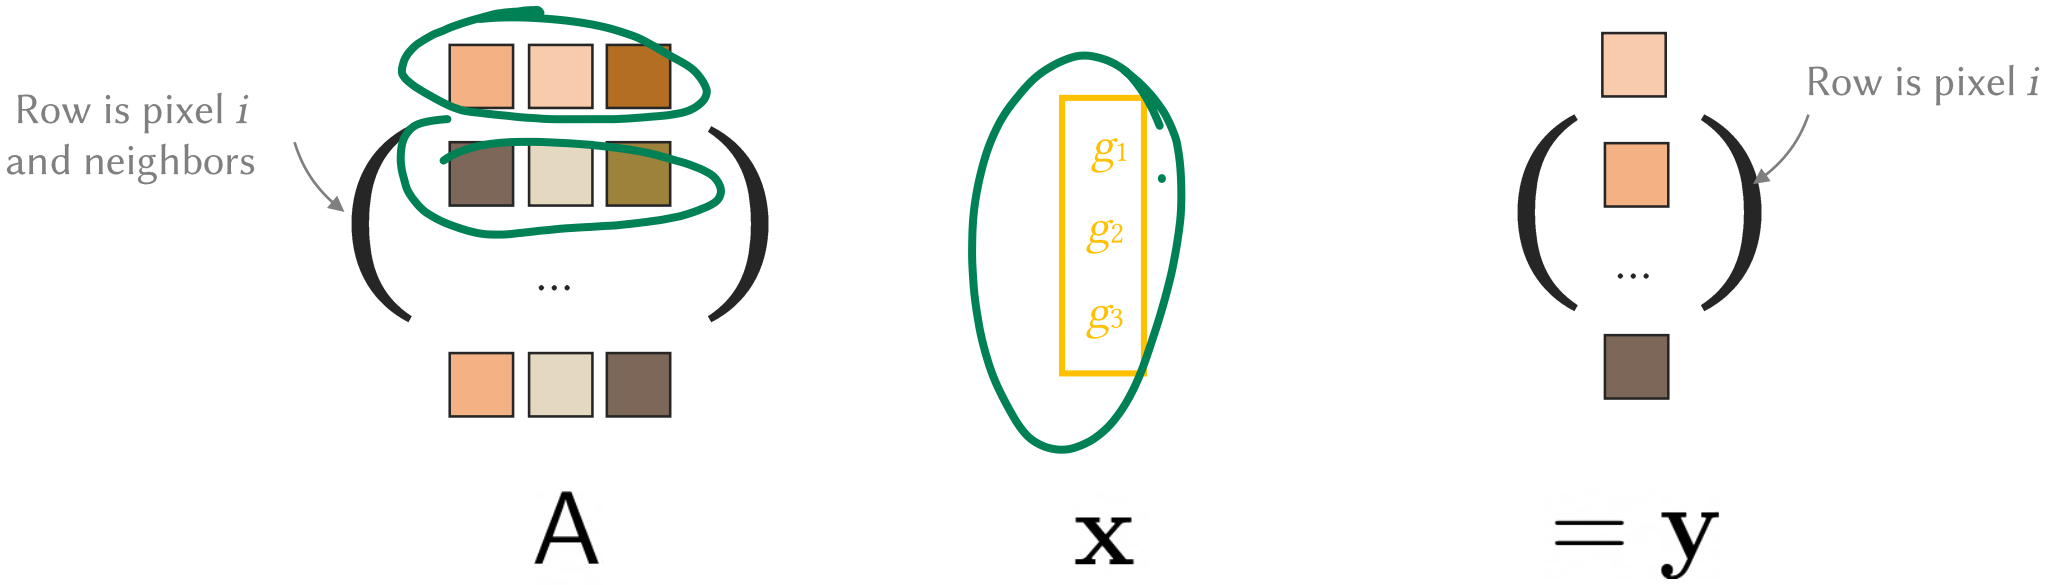

The cheetah and the zebra

- Could we use this to classify images?

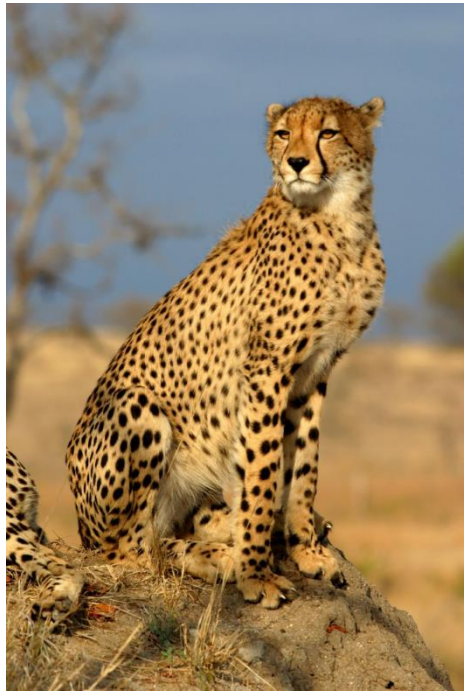

cheetah

zebra

Probably relevant

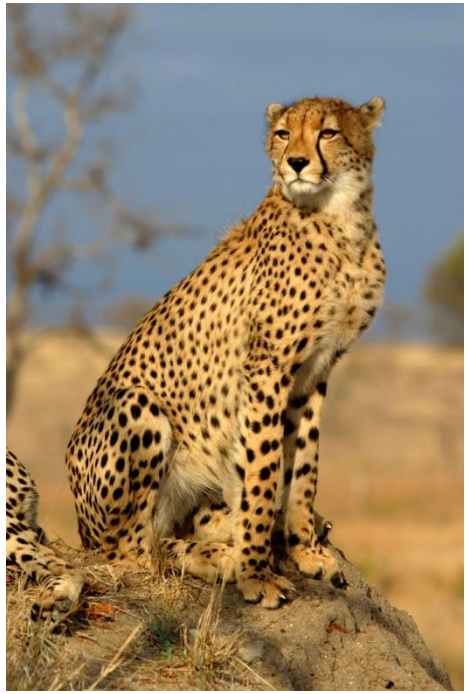

cheetah

zebra

The cheetah and the zebra

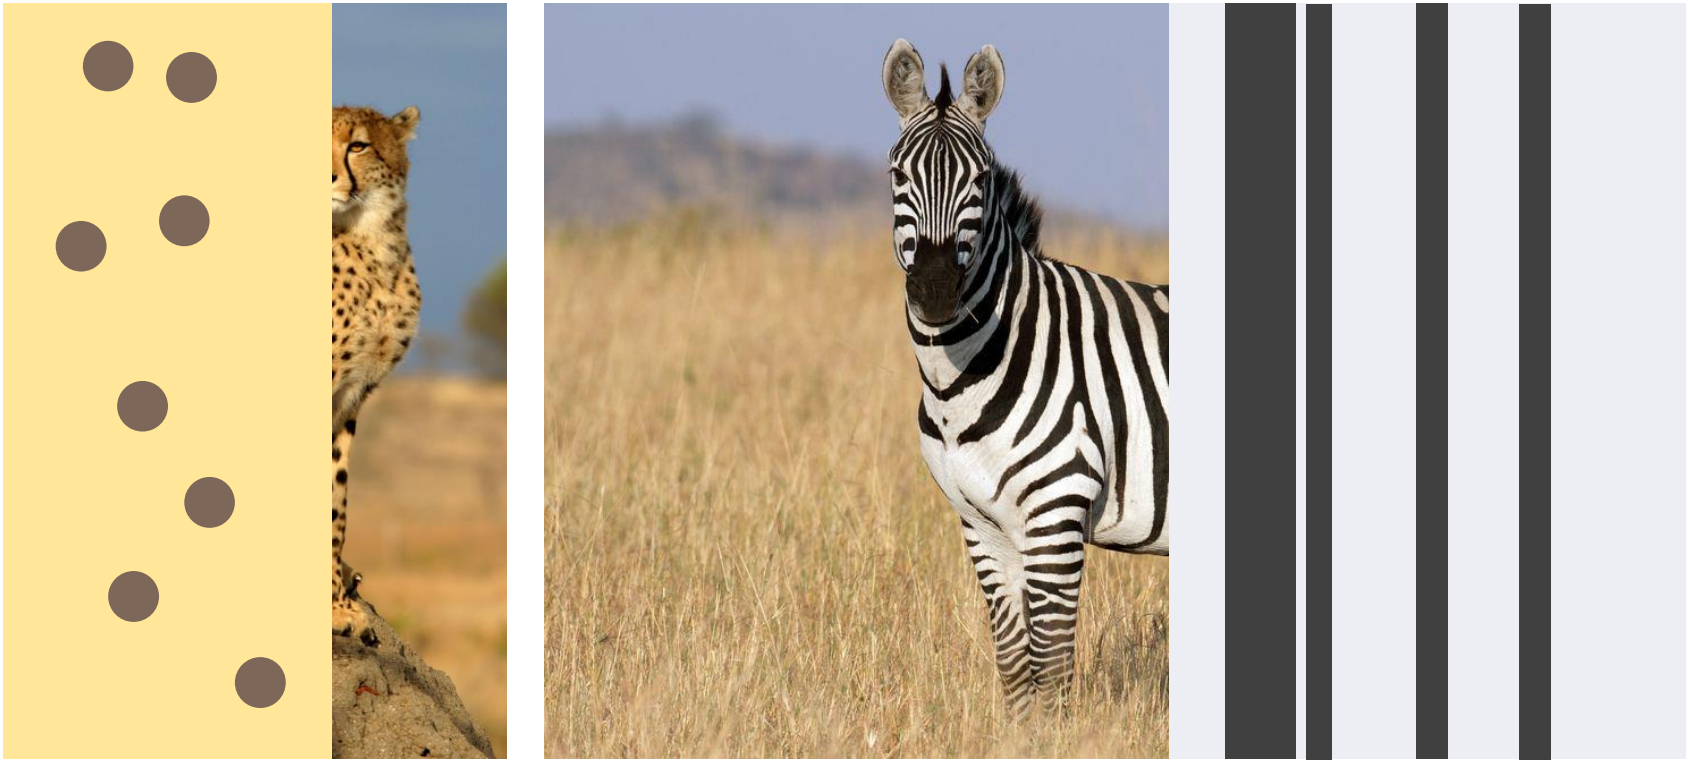

Points, edges, non-points, non-edges

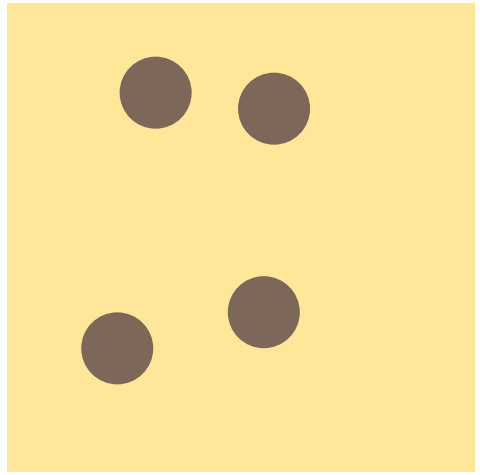

$$\odot f_{\theta_1} = 1$$

$$\odot f_{\theta_2} = 1$$

$$\odot f_{\theta_1} = -1$$

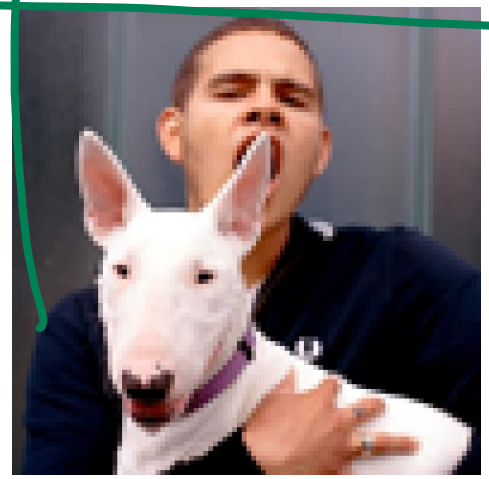

$$\odot f_{\theta_2} = -1$$

Example solution

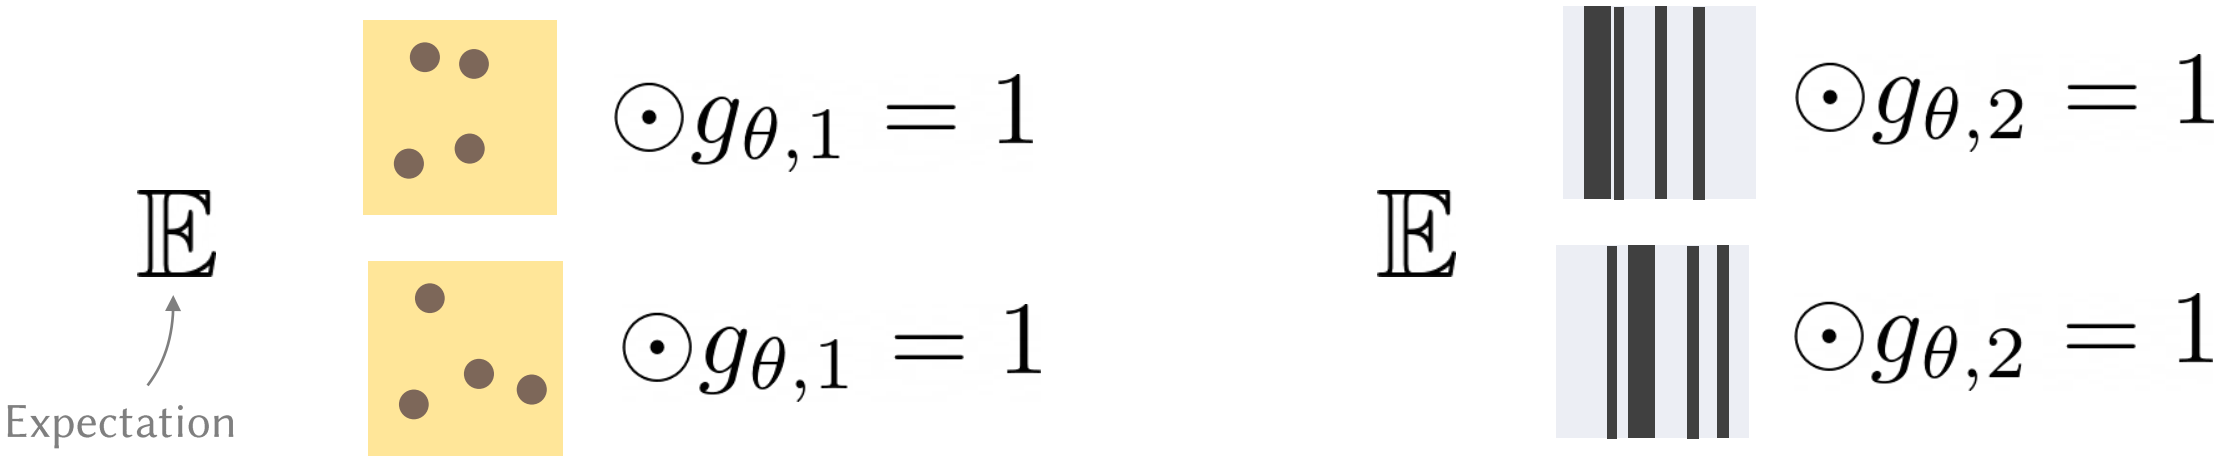

-1	-1	-1
-1	8	-1
-1	-1	-1

-1	2	-1
-1	2	-1
-1	2	-1

Non-linearity

- As we saw in the classification and NN lecture, **non-linearity** is important

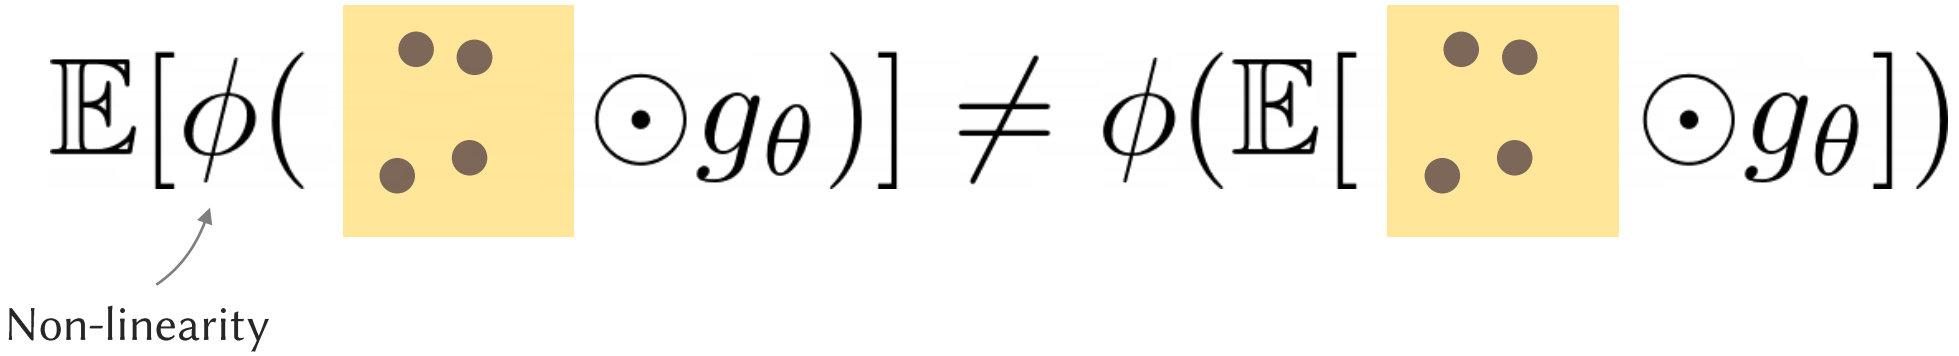

Non-linearity

- And after that no-linearity, we could again do a convolution
- We call this **layers**

$$\mathbb{E}[\phi_1(\ \odot g_{\theta_1})\odot g_{\theta_2}] \neq$$
$$\mathbb{E}[\phi_1(\ \odot(g_{\theta_1}\odot g_{\theta_2}))]$$

Depth + non-linearity

- What this looks like

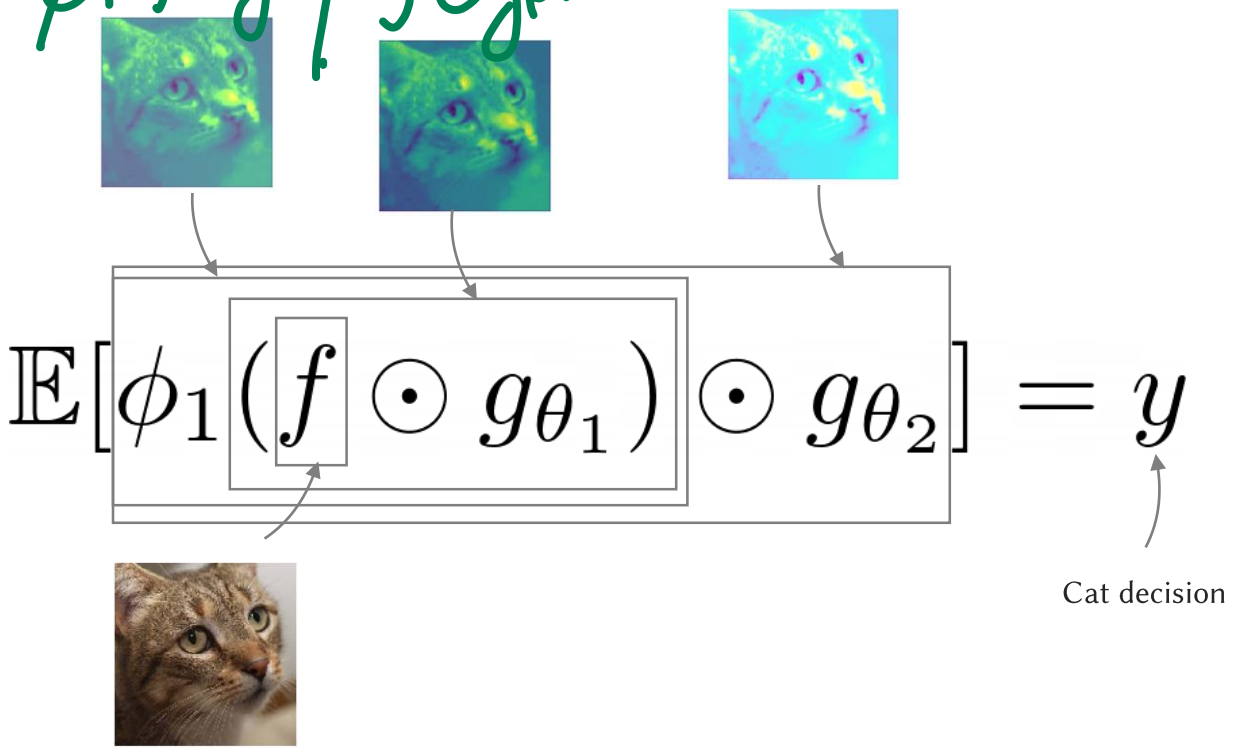

Multi-resolution

- We recall that finding edges and any feature is best done on a pyramid
- CNN so far one resolution
- Solution: insert **pooling** reducing resolution after each convolution

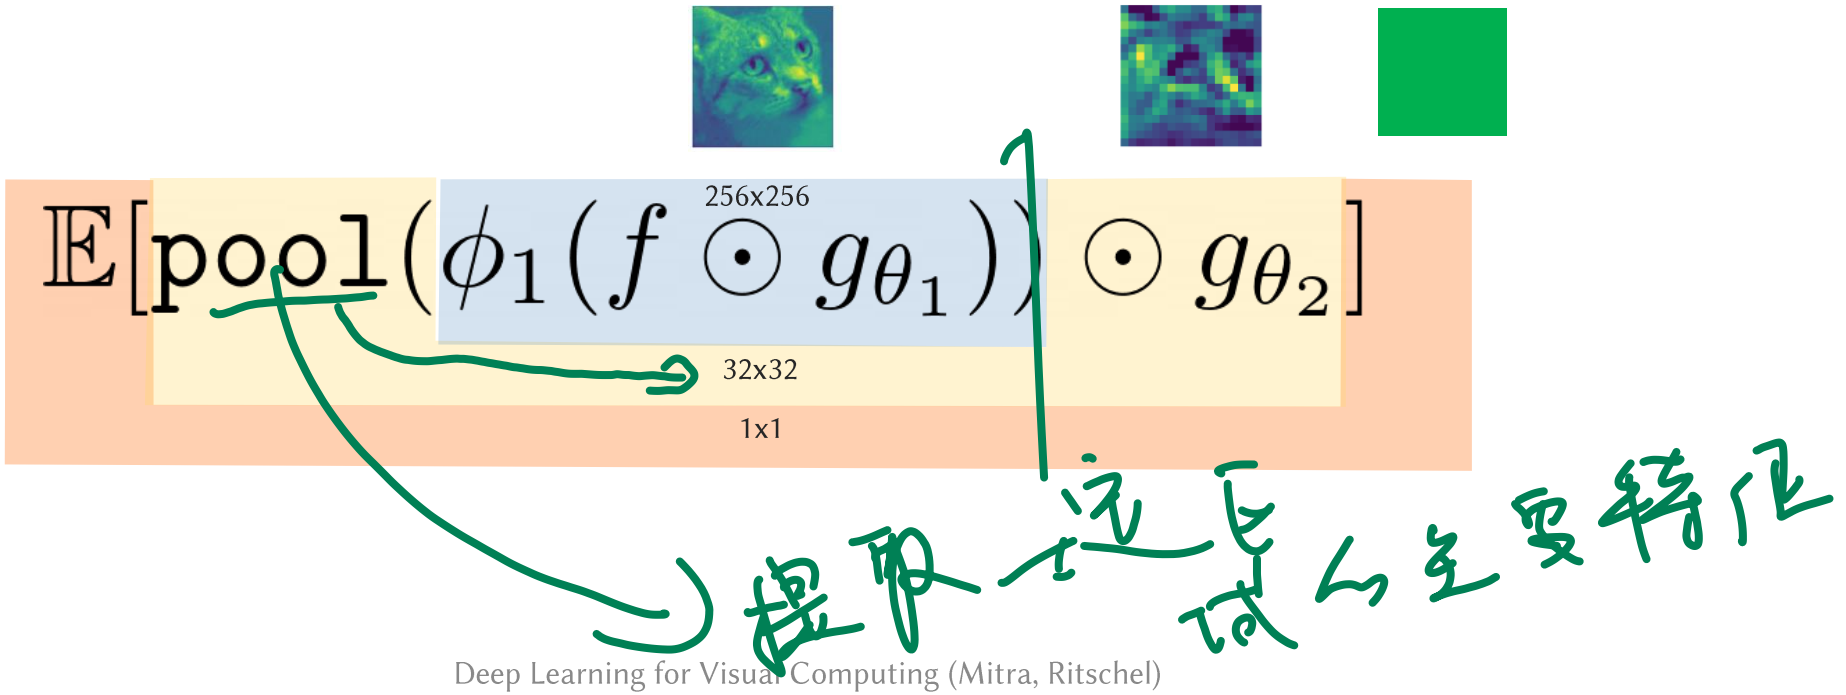

Fully-connected layer

- Instead of expectation, take all pixels and feed them to a classic NN

Map c-times-n vector to number

32×32

$$\psi(\texttt{stack}(\texttt{pool}(\phi_1(f \odot g_{\theta_1})) \odot g_{\theta_2})) = y$$

Make c-times-n-vector from all n pixels and all c chans

Handwriting recognition: LeNet

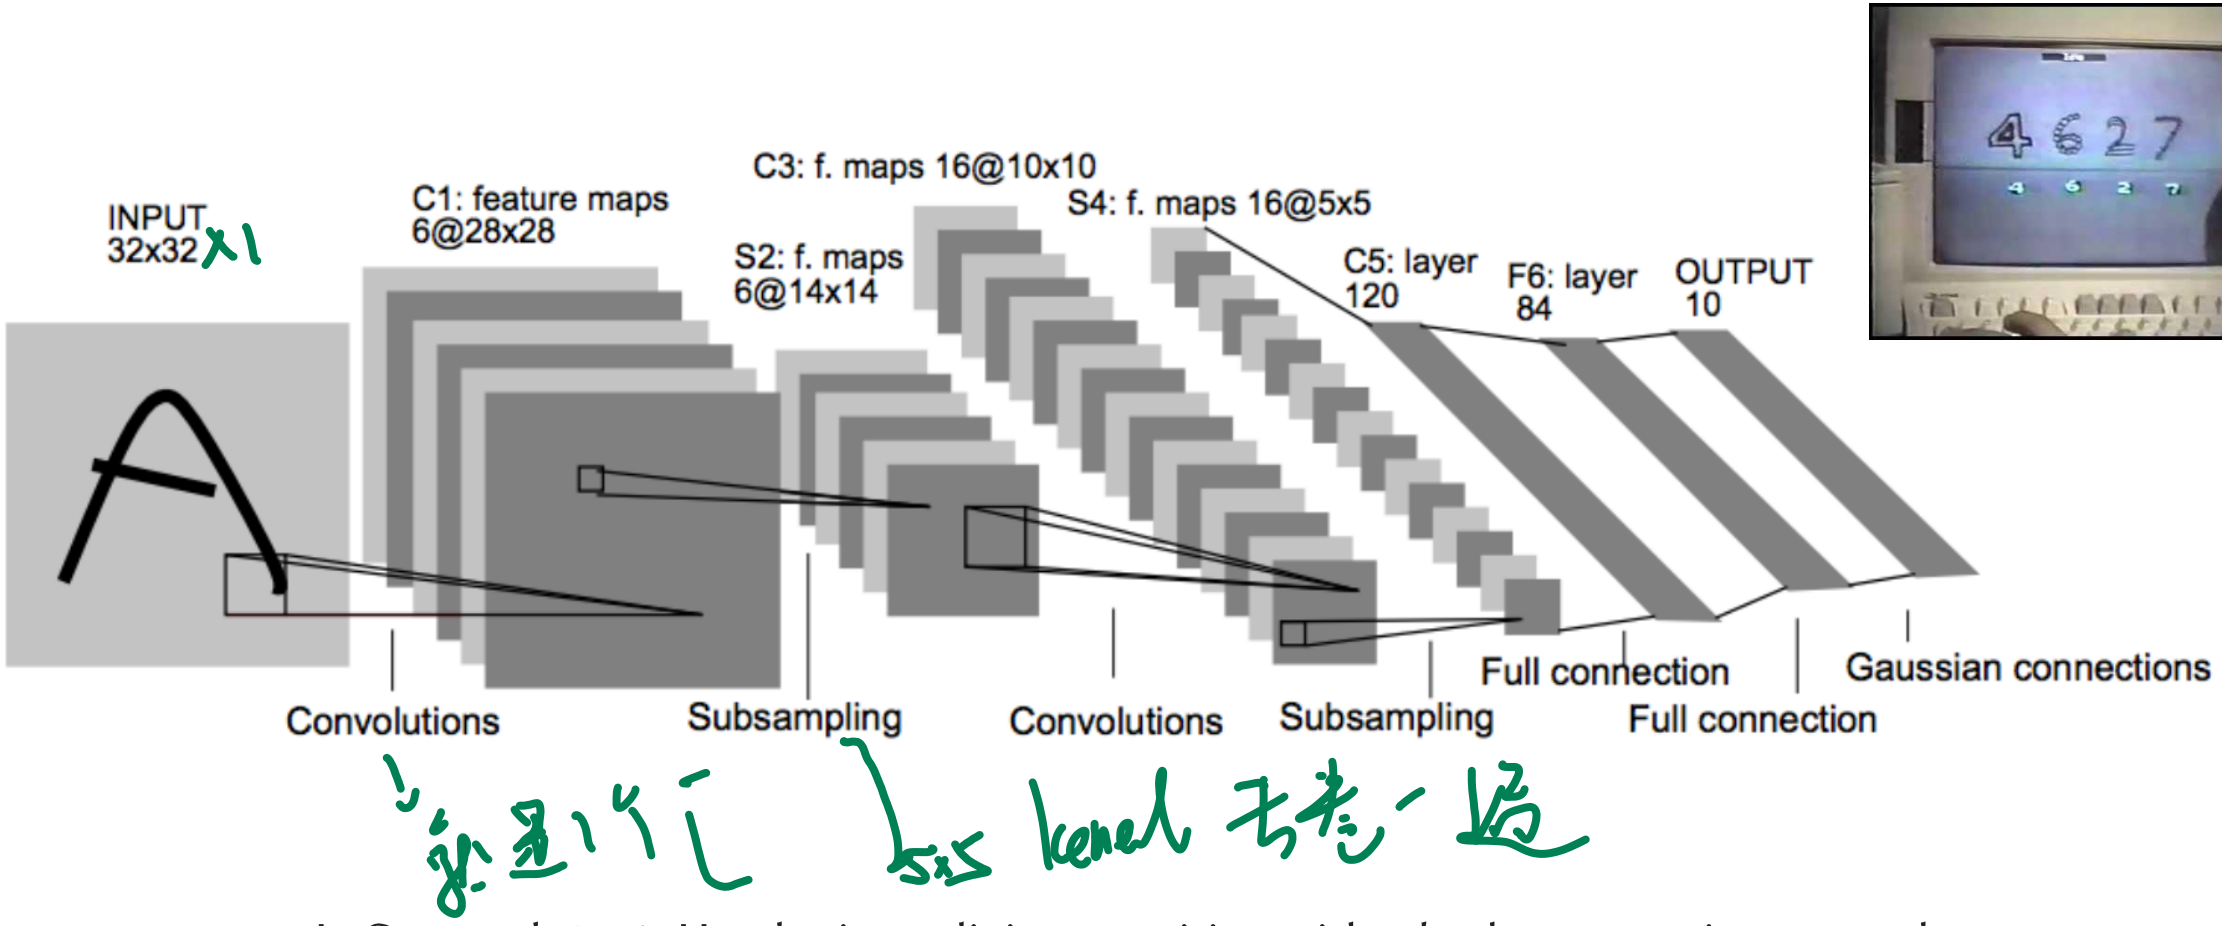

LeCun et al, 1990: Handwritten digit recognition with a back-propagation network

Image classification: AlexNet

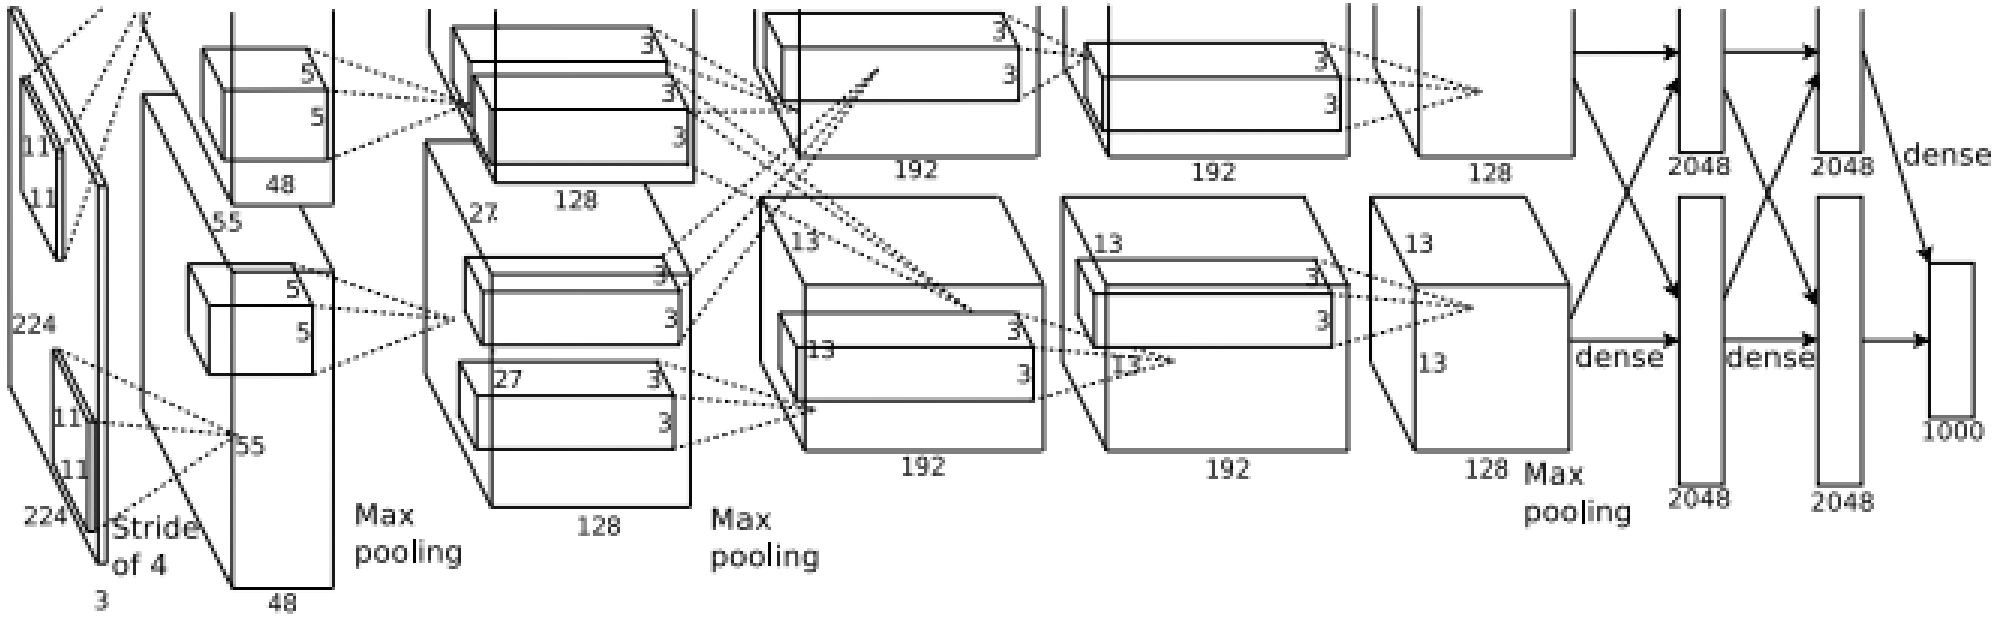

Krizhevsky et al, 2012: *ImageNet Classification with Deep Convolutional Neural Networks*

What happened between 1992 and 2012?

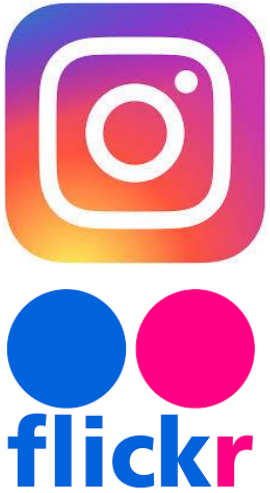

Image data

60k vs
60M

GPUs

Pentium II 0.2 GFLOPS
Nvidia GTX 680 2600 GFLOPS

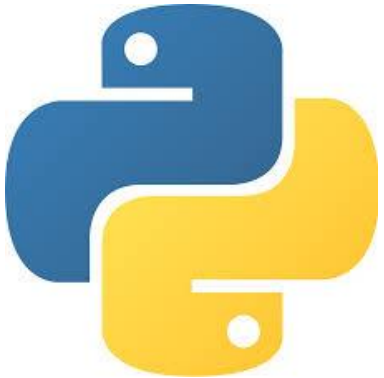

Scripting

C++ vs
CUDA/Python

More architectures

- Not always relevant or clear what happened

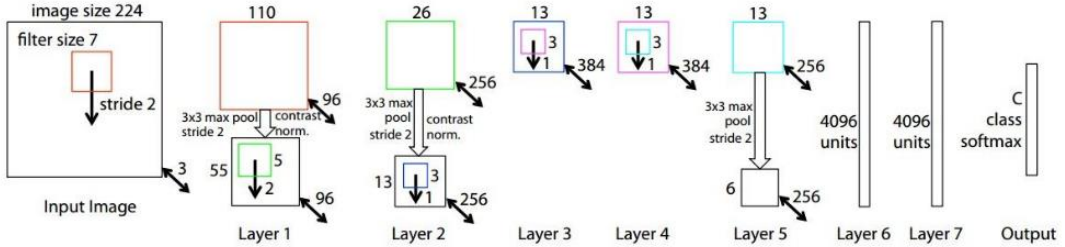

GoogLeNet/Inception (2014)

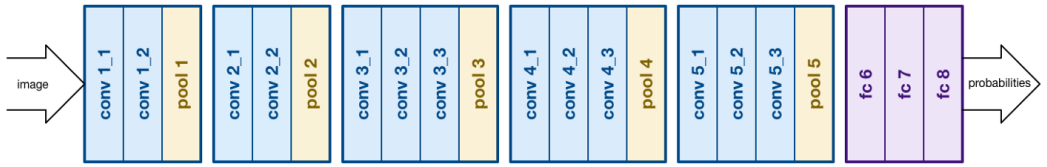

VGG Net (2014)

ResNet, 2015

- Usual sequence makes propagating gradients hard

$$\psi(\phi_1(f \odot g_{\theta_1}) \odot g_{\theta_2})$$

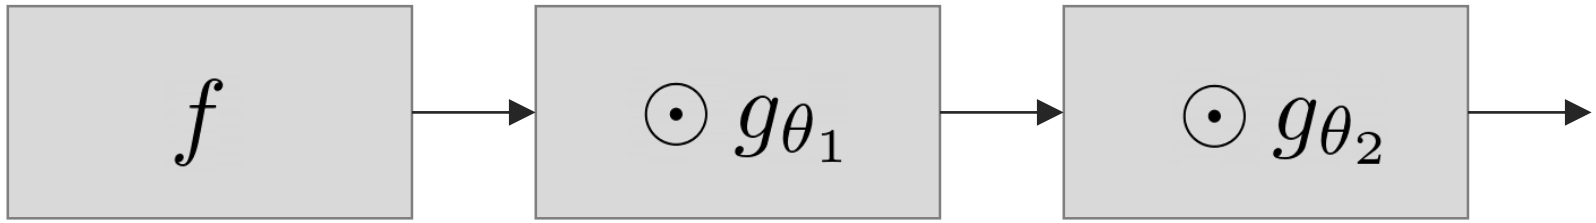

ResNet, 2015

- Usual sequence makes propagating gradients hard

$$\psi((\phi_1(f \odot g_{\theta_1}) + f) \odot g_{\theta_2})$$

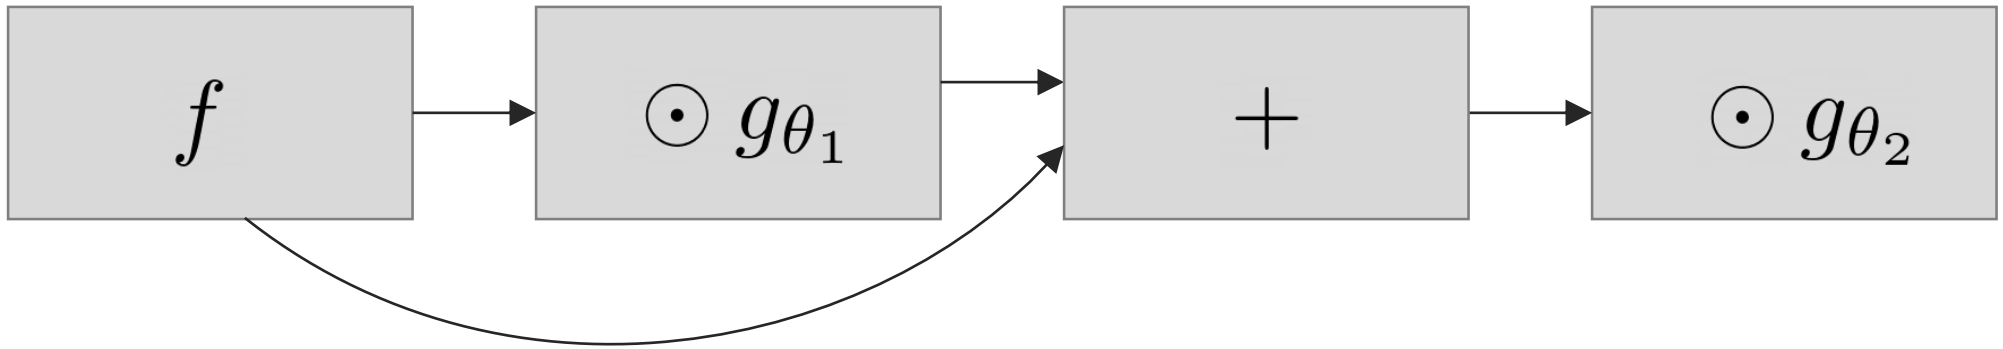

Object detection task

- Map image to bounding boxes
- Challenge: Number **unknown**

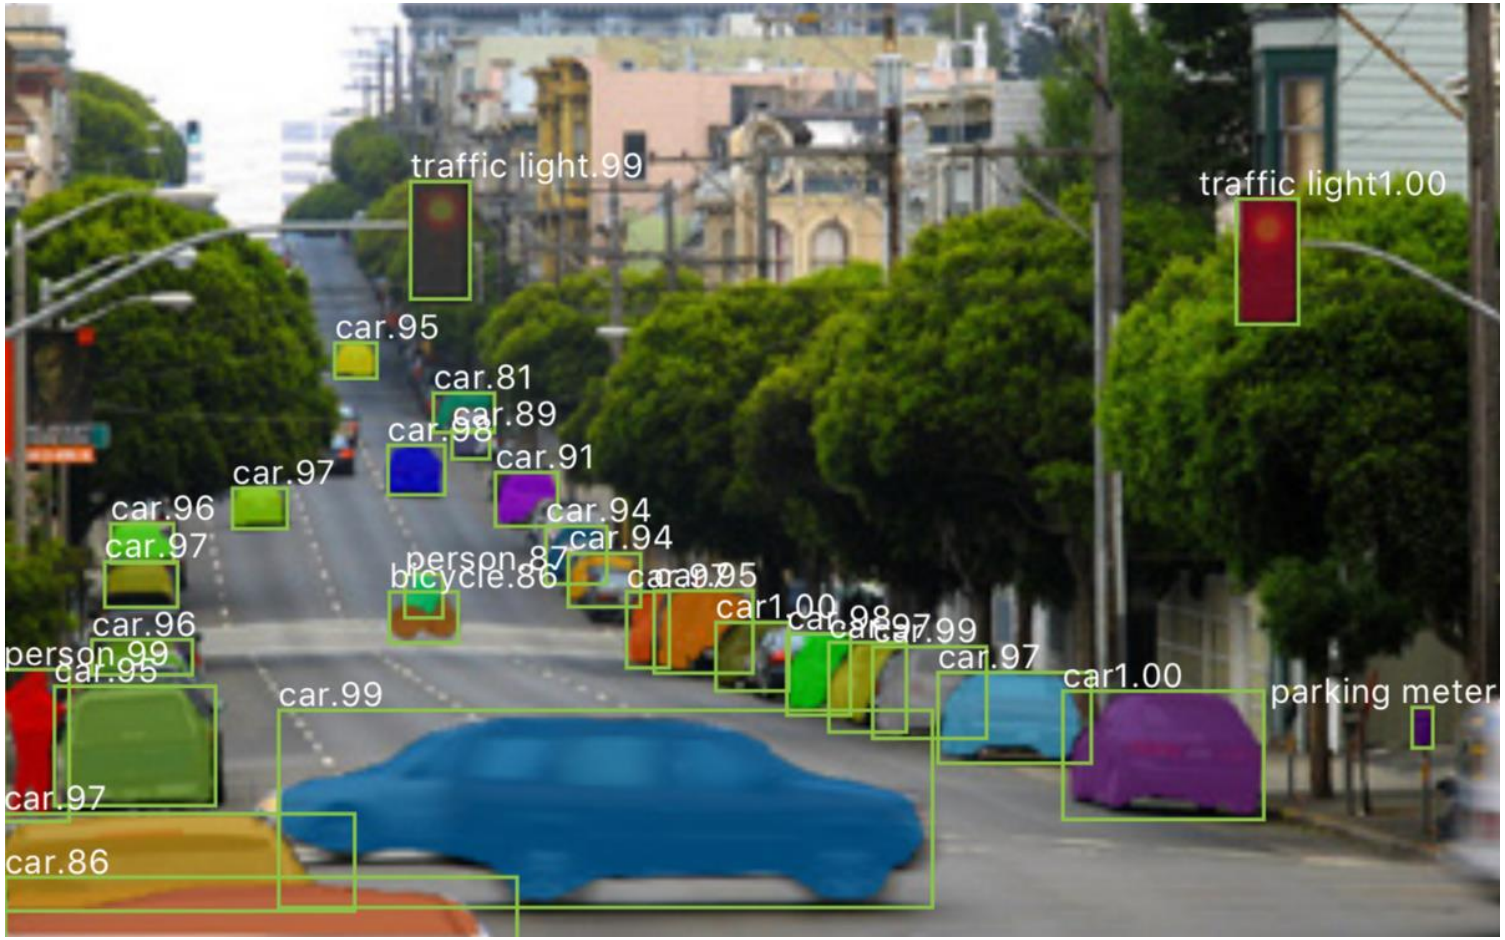

3D Object detection

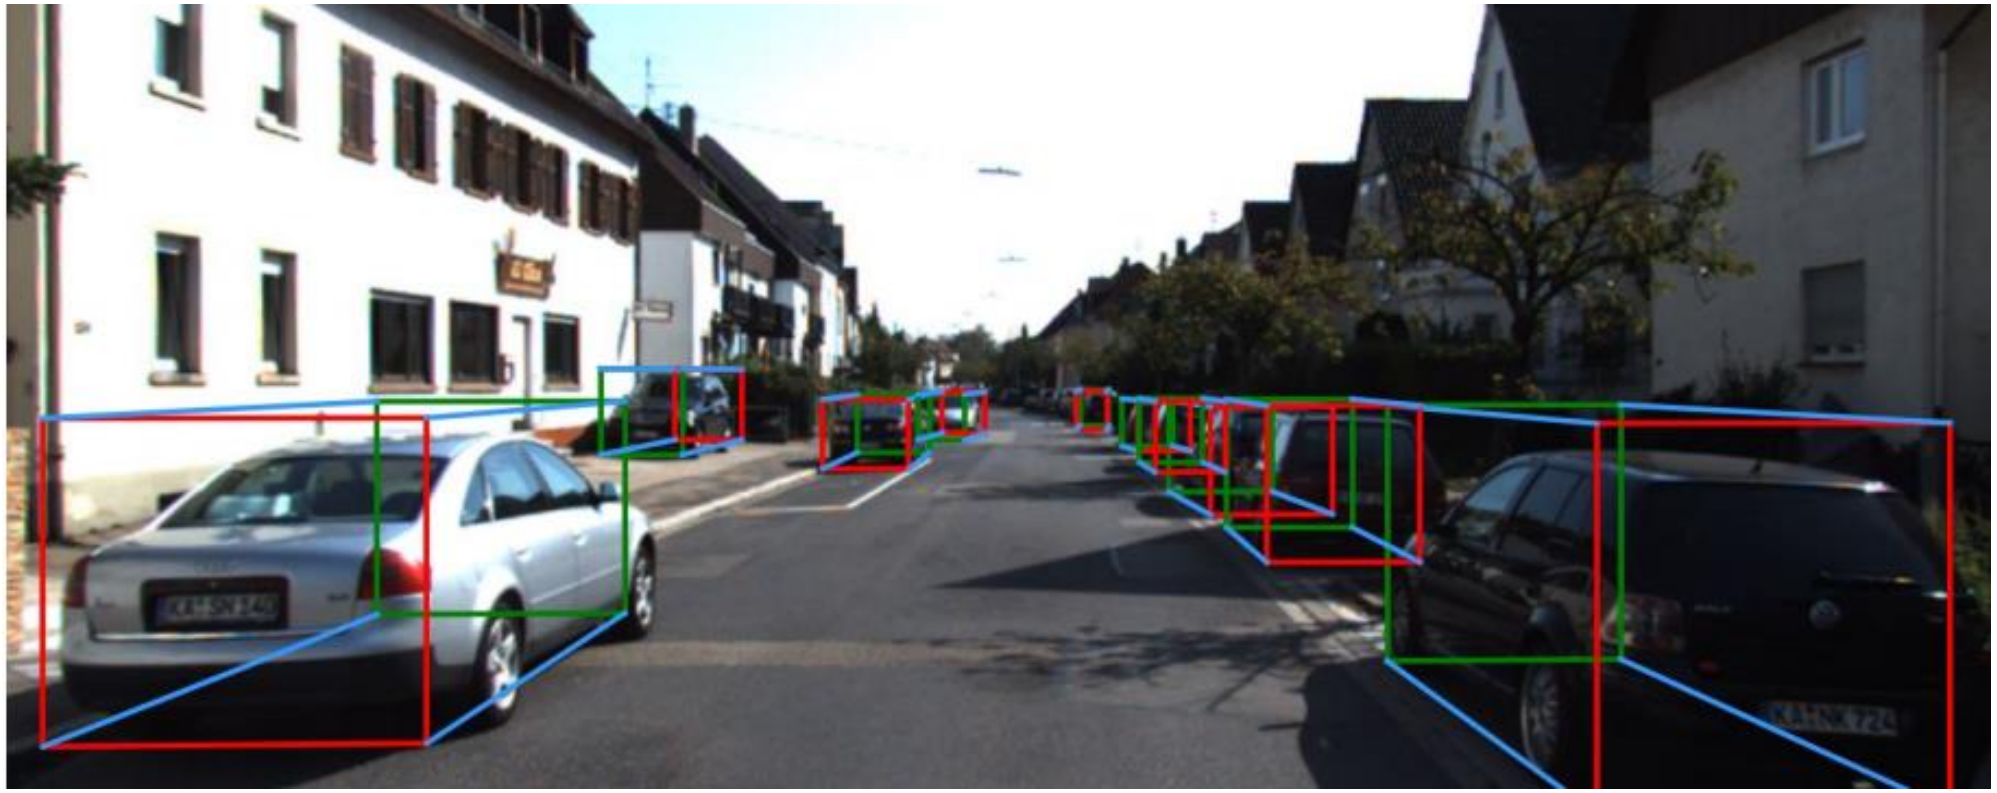

Image-to-image: Segmentation

RGB

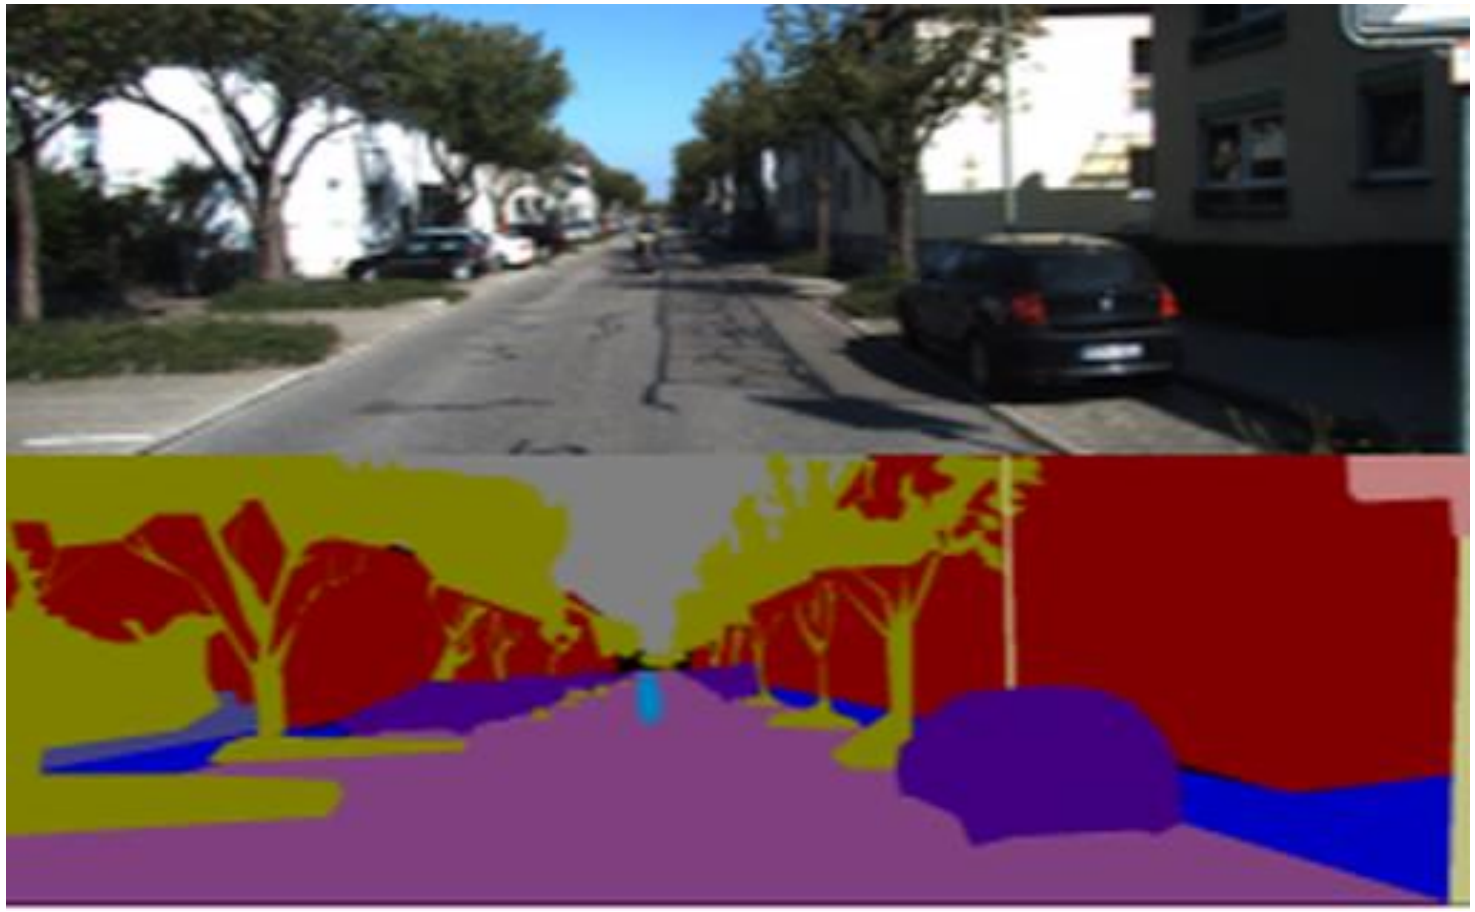

Semantic segments

Image-to-image

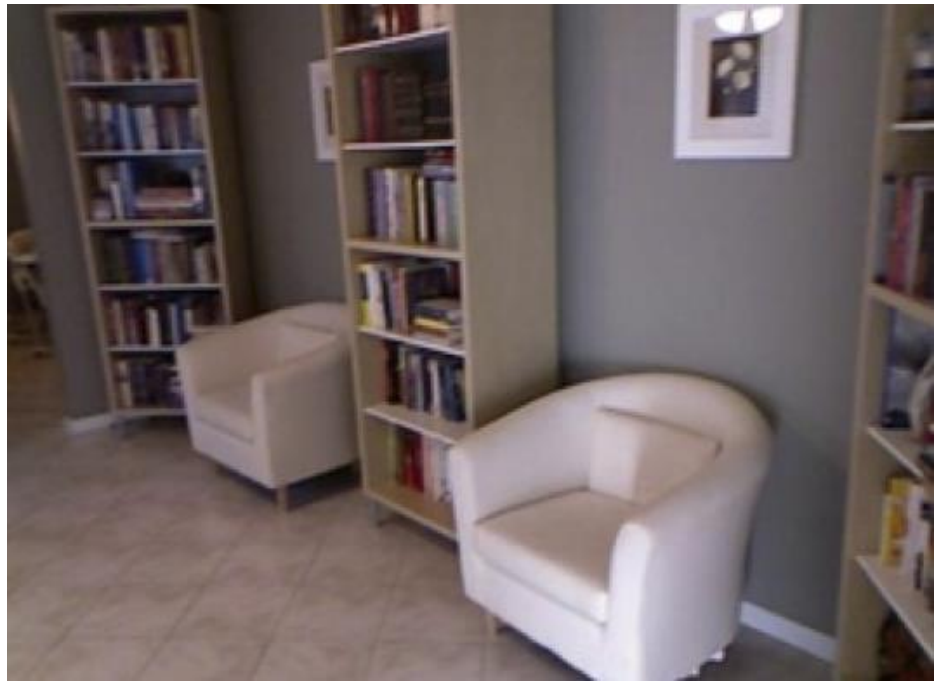

RGB

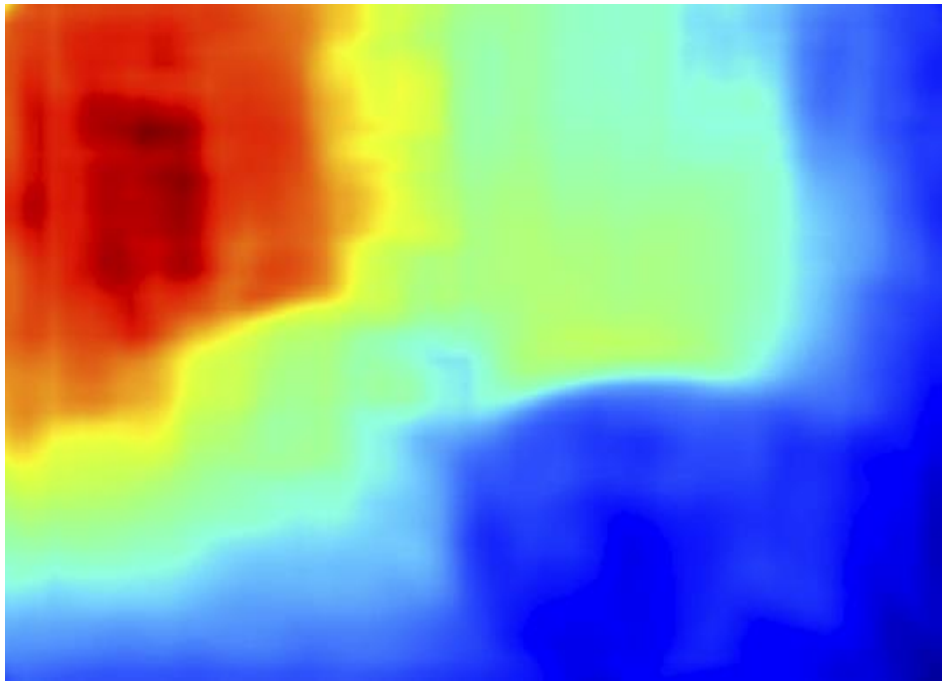

Depth

Image-to-image

- Image in, number out: Encoder-decoder
- How about image out?

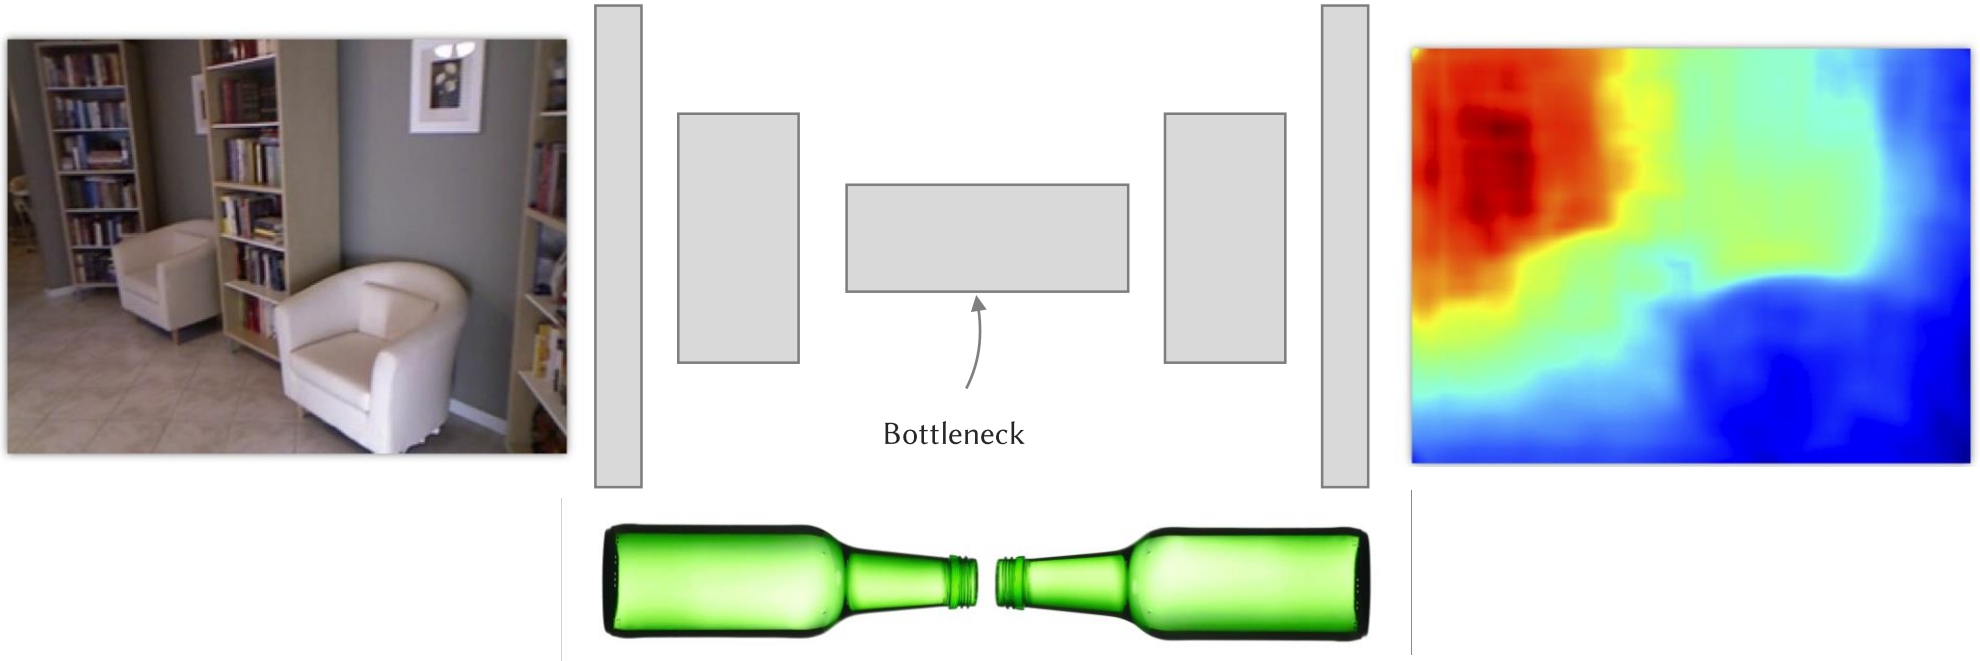

Image-to-image

- No chance to get back details
- Idea: **skip connections,** allow decoder to access encoder state (Unet)

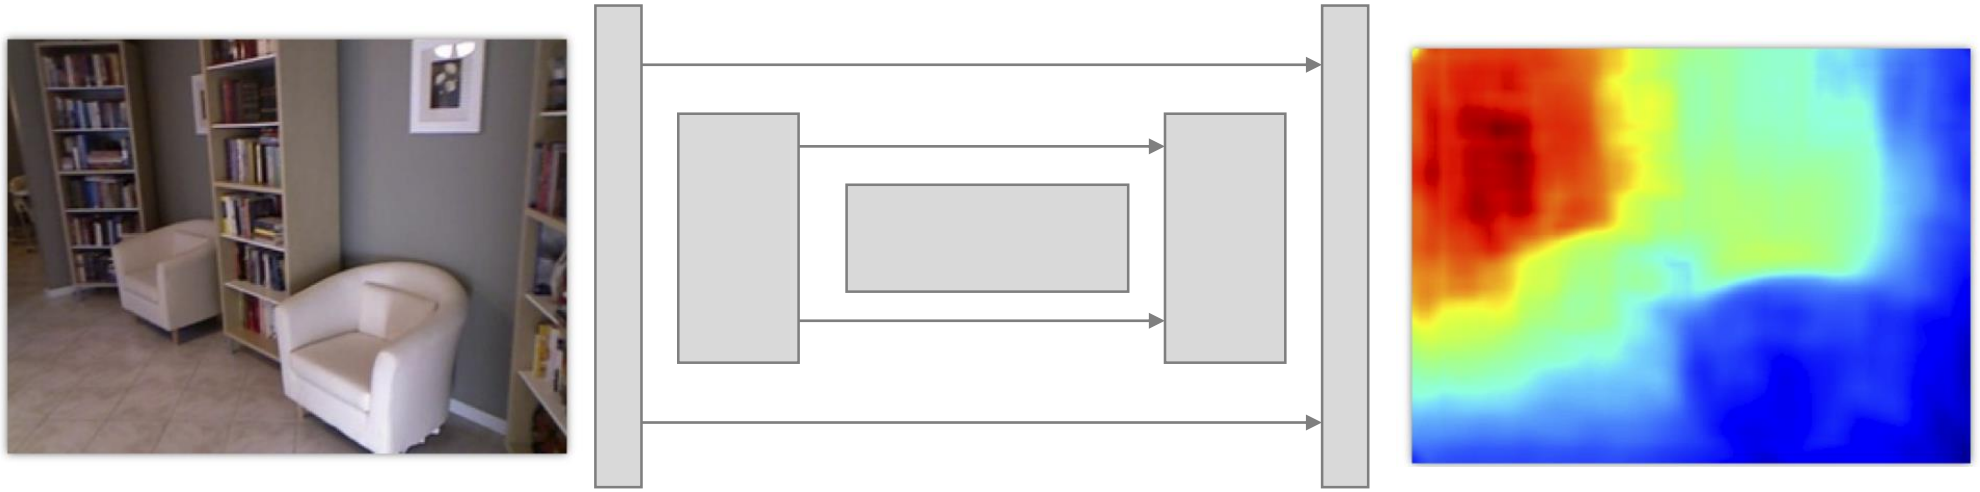

Semantic image synthesis

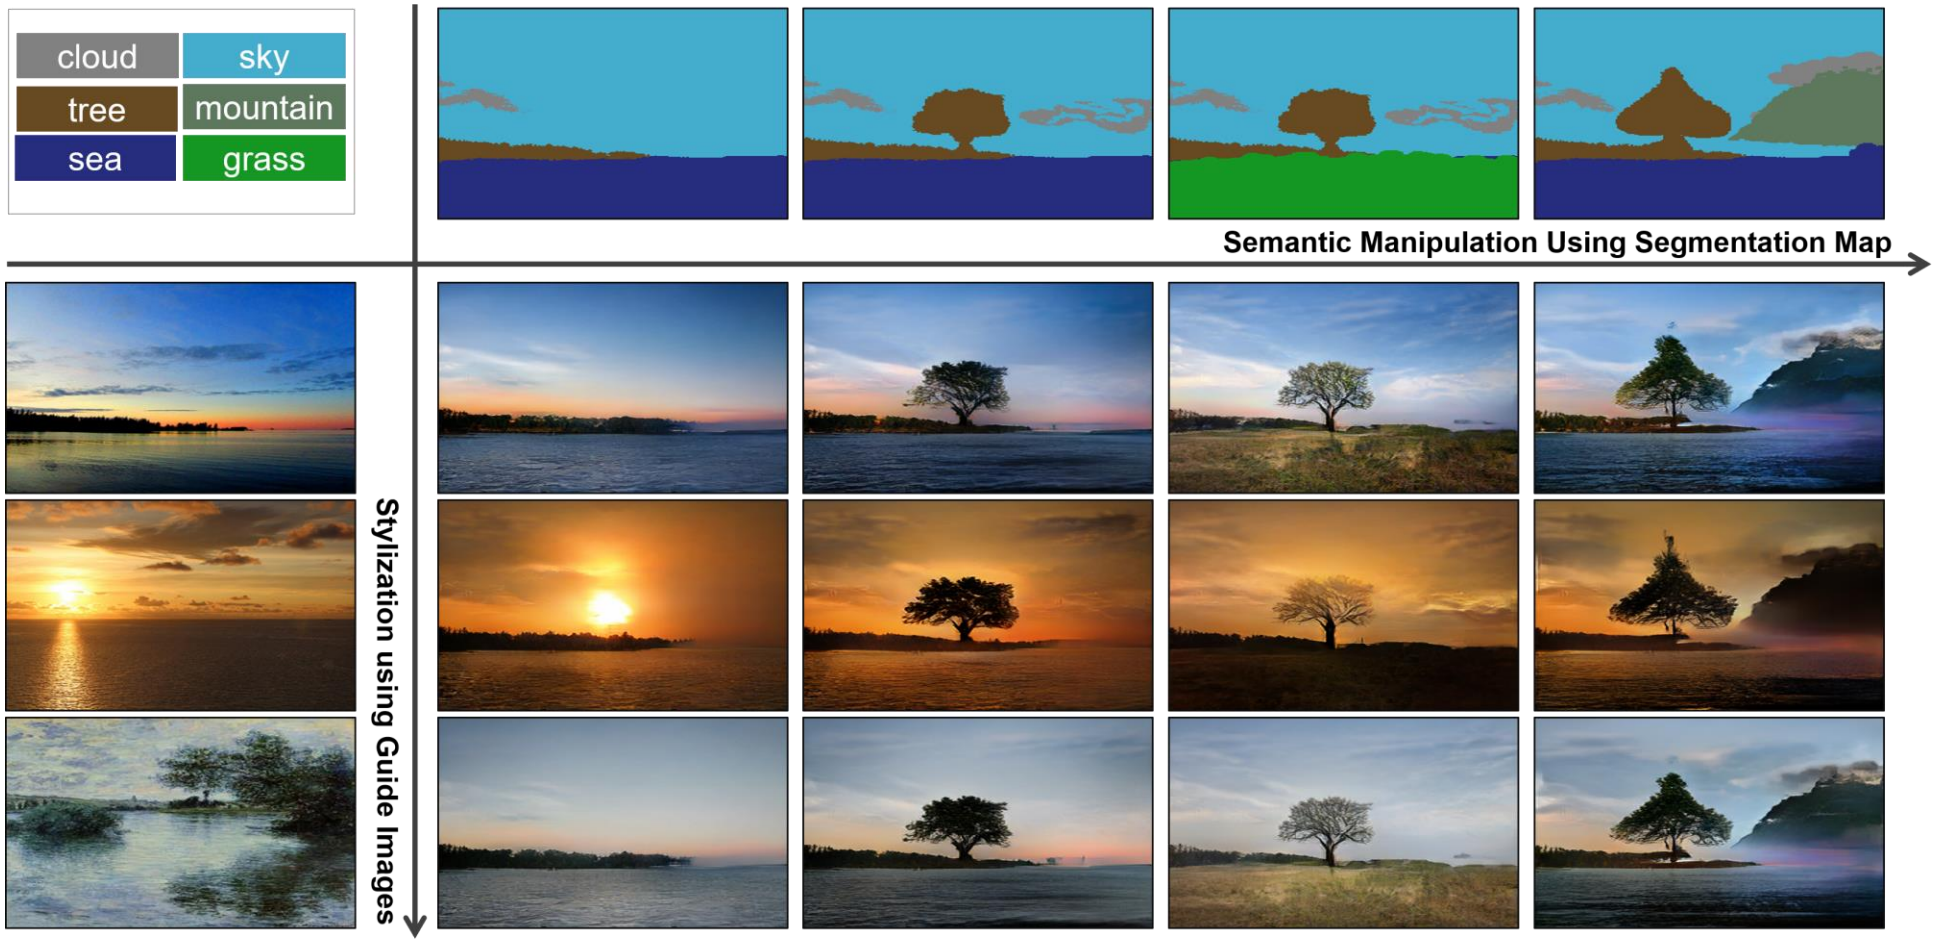

Style transfer

Encoder-decoder for video

- Insert flow-compensated difference into loss

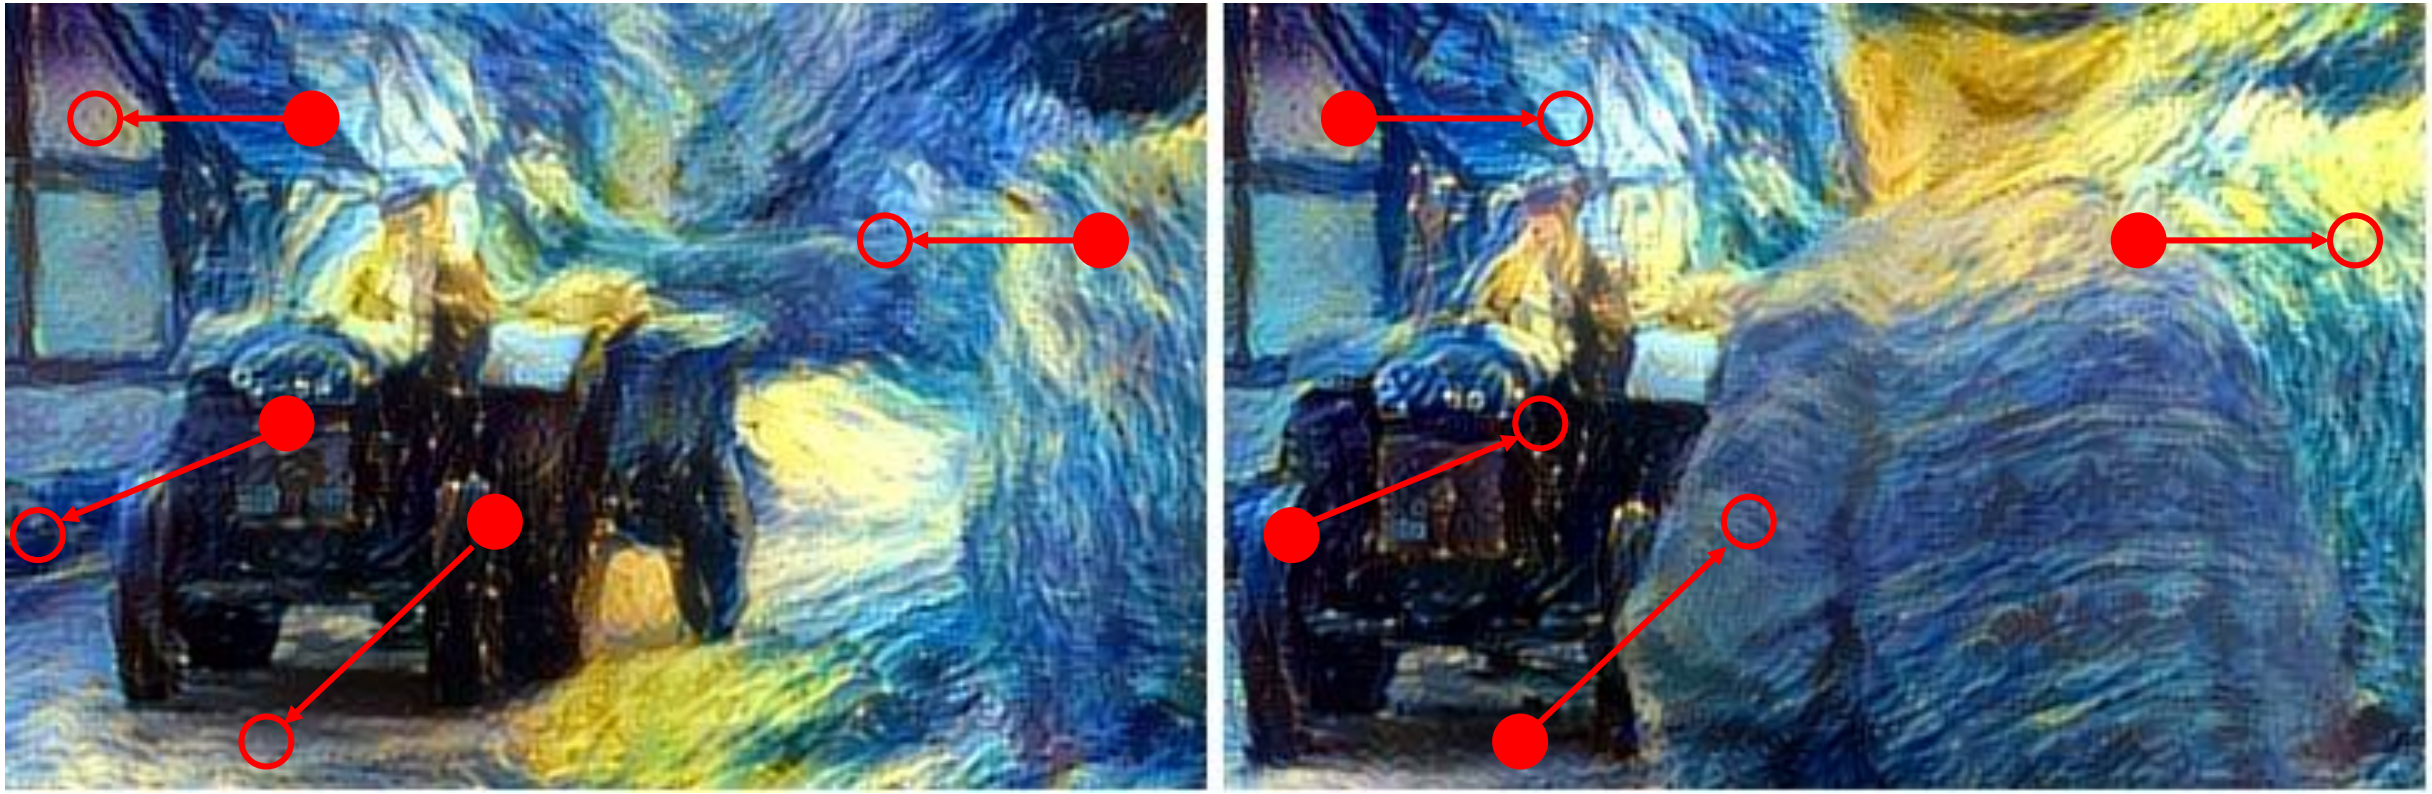

Encoder-Decoder for music

- Wavenet
- Like U-Net, just 1D

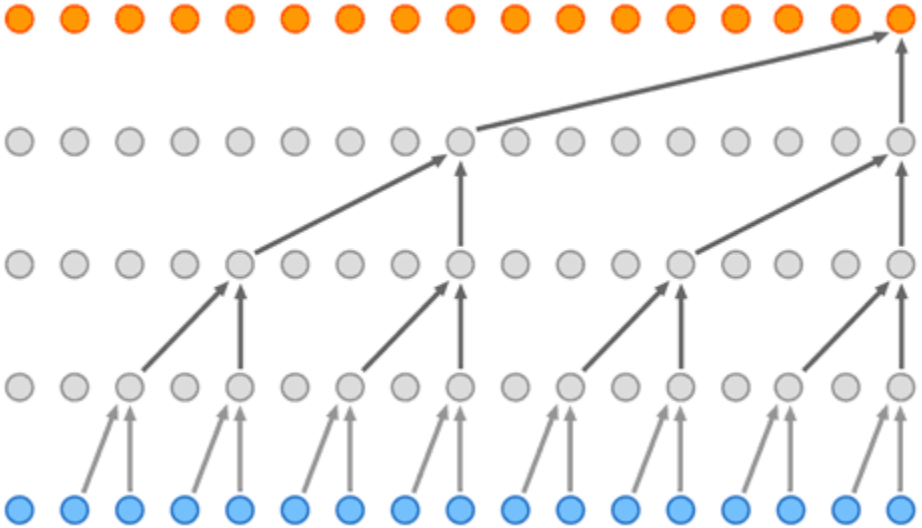

Failure

- Adversarial examples
- Small perturbation put detection off
- Found by optimizing for it

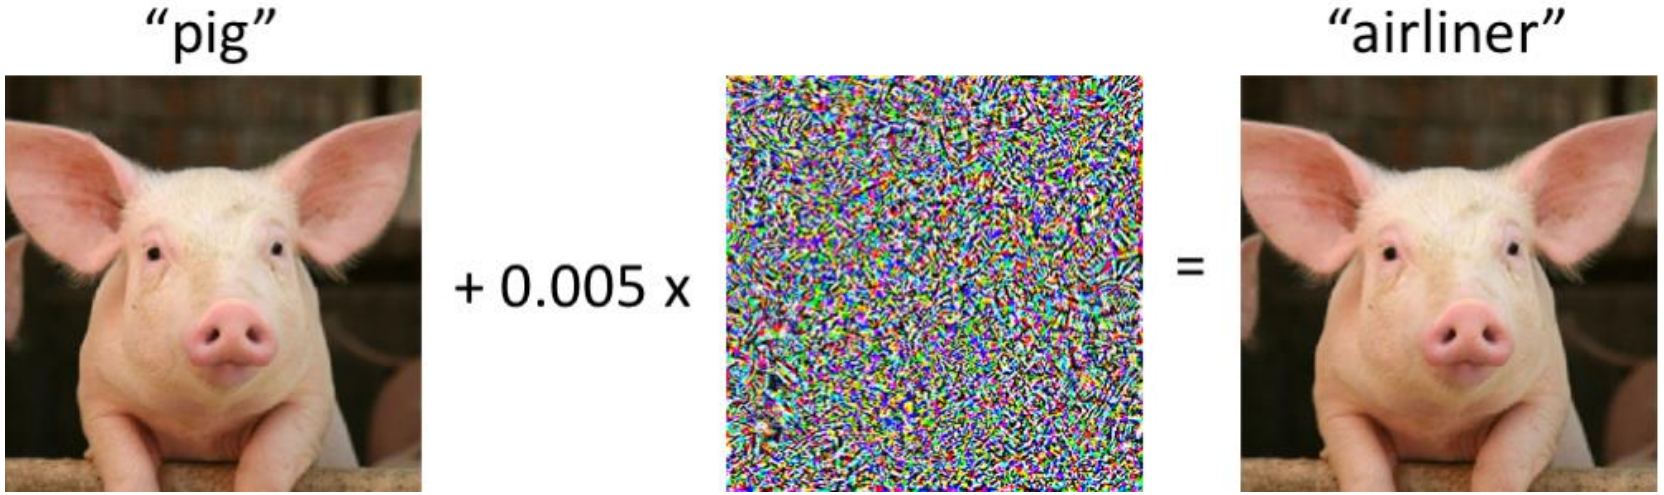

Dreaming

Conclusion

- Don’t write filters yourself
- Let the computer find them
- Stack them with non-linearities
- Add a few tricks (skip, res)
- And it does many things