Computer Graphics (COMP0027) 2022/23

Shadows

Tobias Ritschel

Challenges

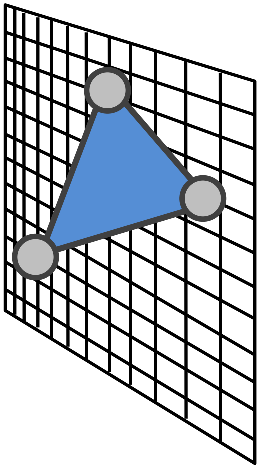

Projection

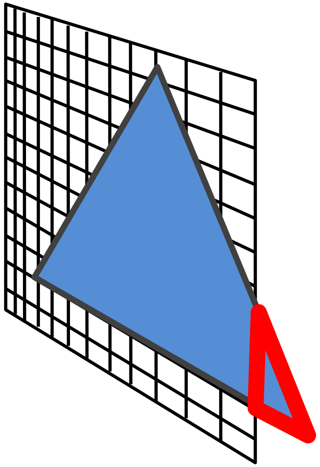

Clipping

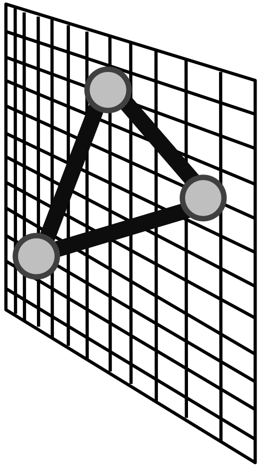

Rasterization

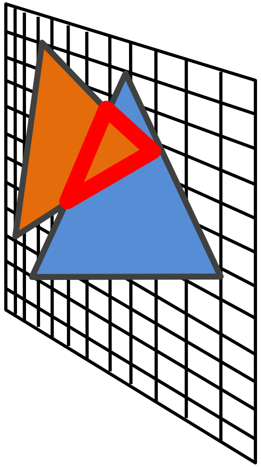

Visibility

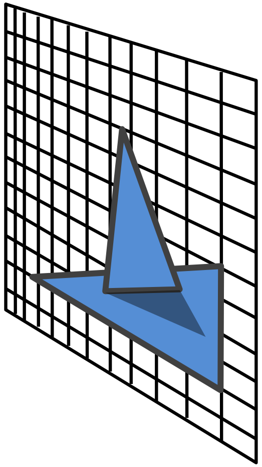

Shading

Pipeline

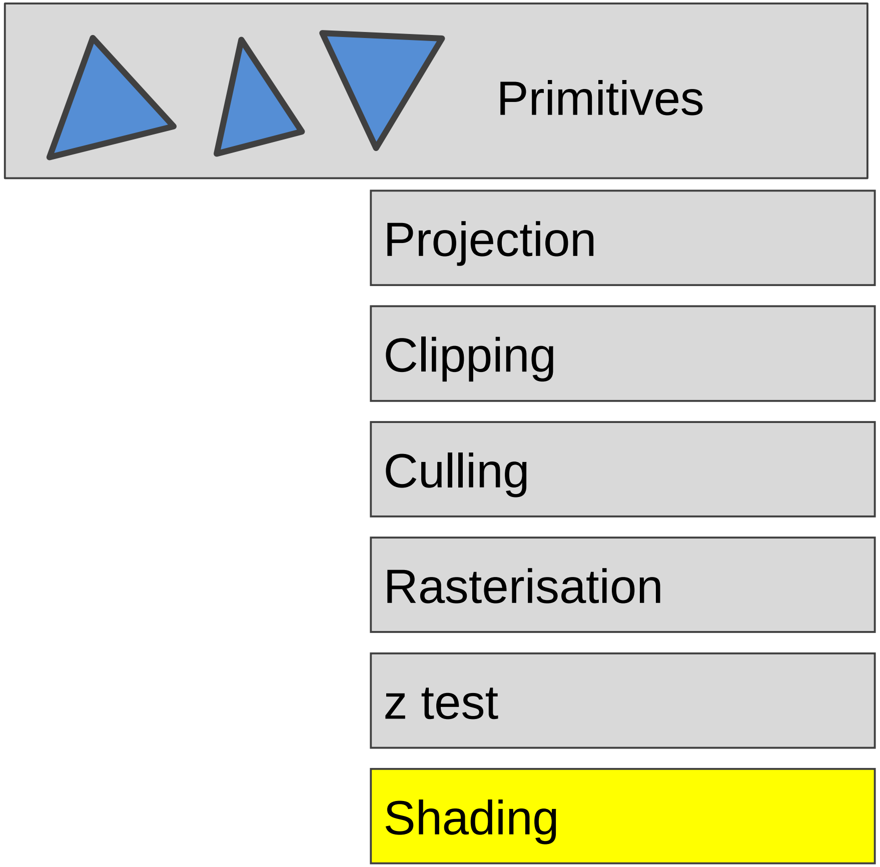

Primitives

Projection

Clipping

Culling

Rasterisation

z test

Shading

Outline

- Introduction
- Sharp shadows
- Soft shadows
- Conclusion

Why are Shadows Important?

- Depth cue
- Scene Lighting
- Realism
- Contact points

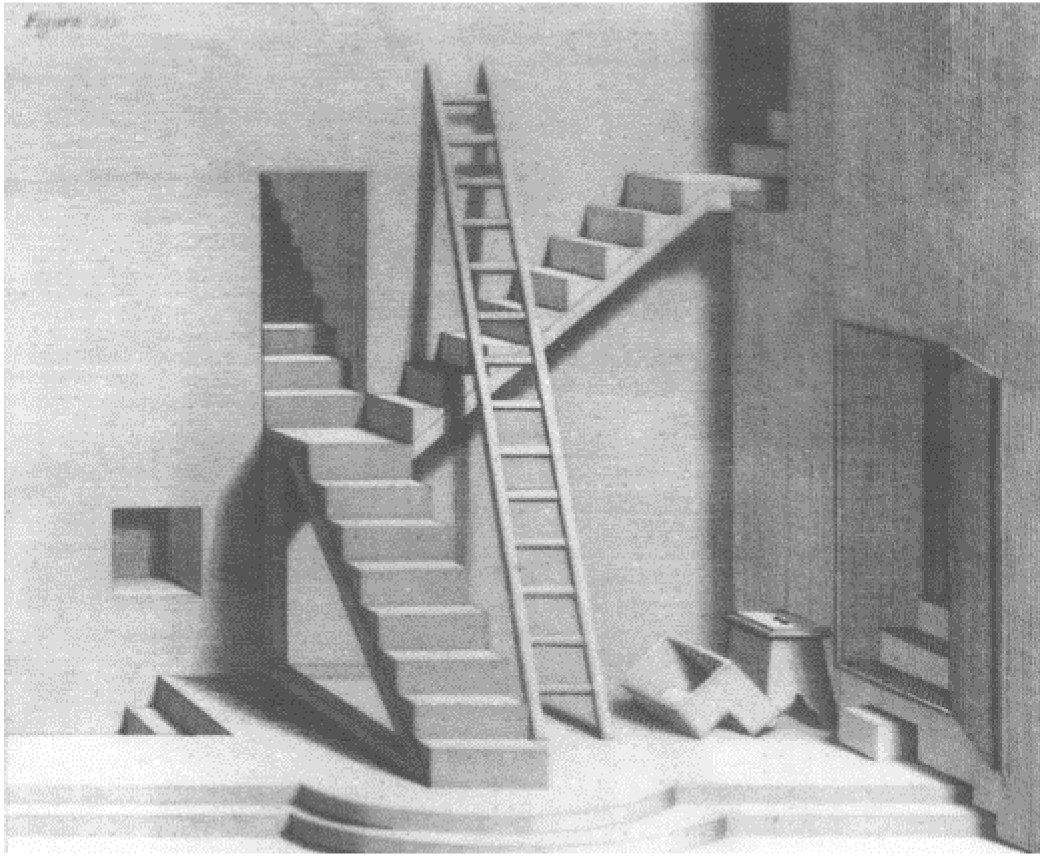

Spatial cue

空间线索

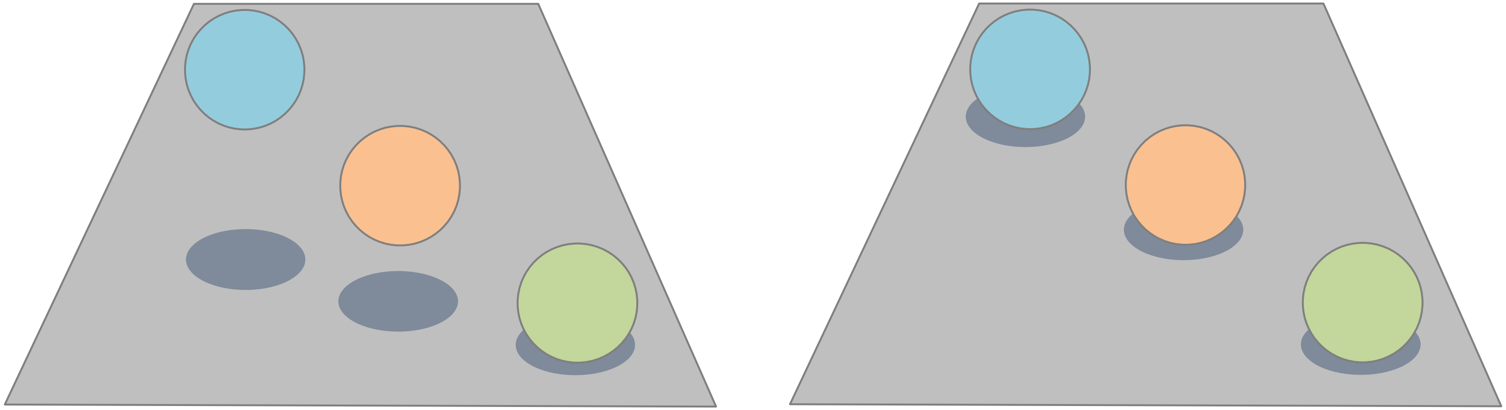

Light position cue (sundial)

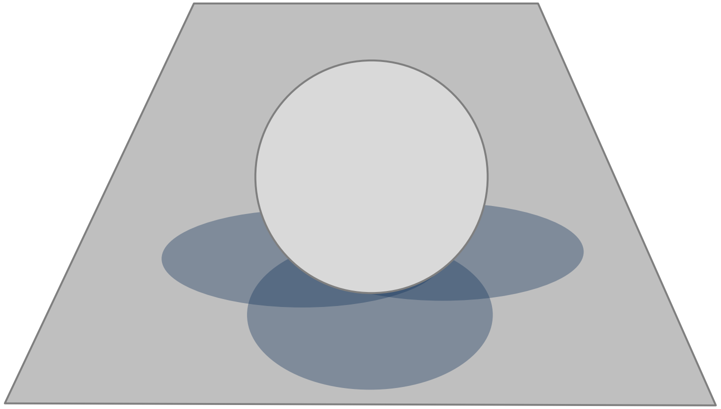

Light color

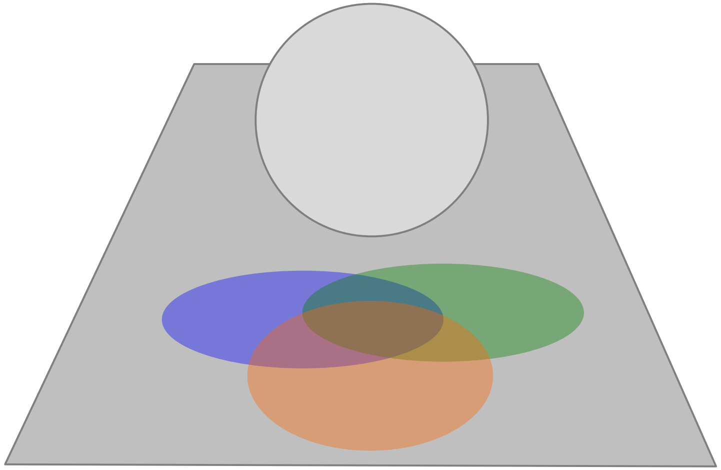

← shadow blue

是因为

天空是blue的

$(K_{ambient} = blue)$

↑

Light distance

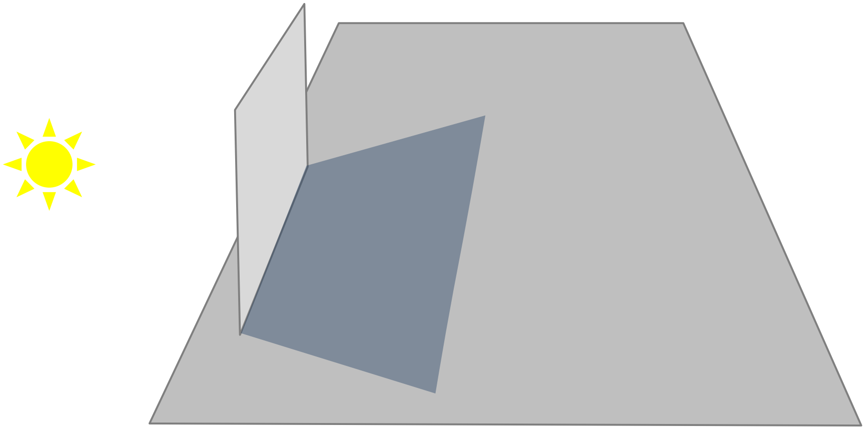

Light distance (far)

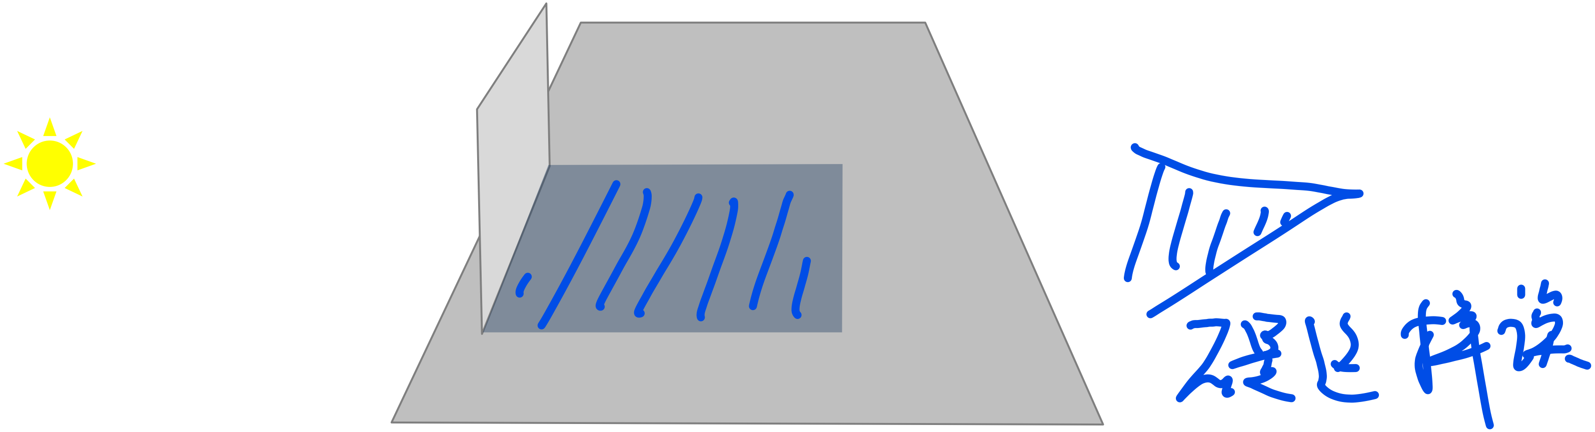

Light distance (near)

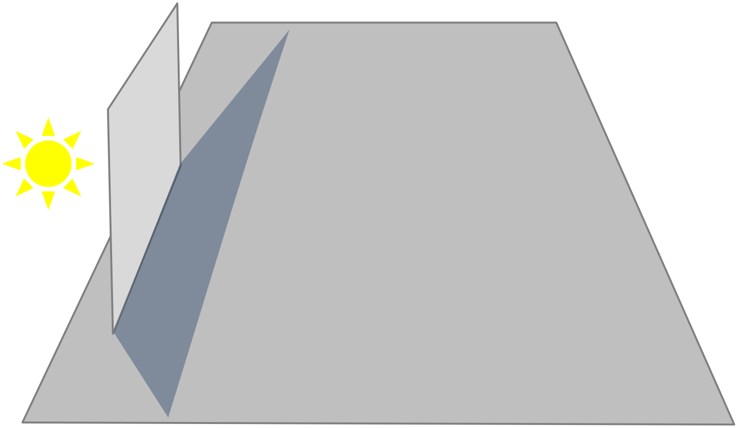

Light size (small)

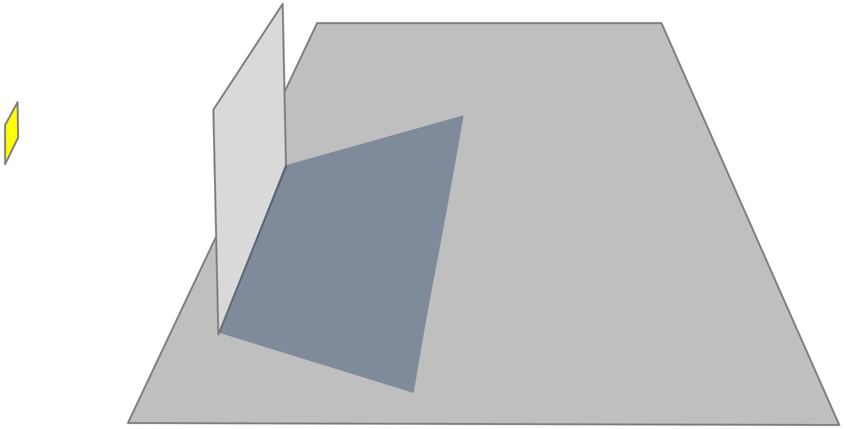

Light size: large

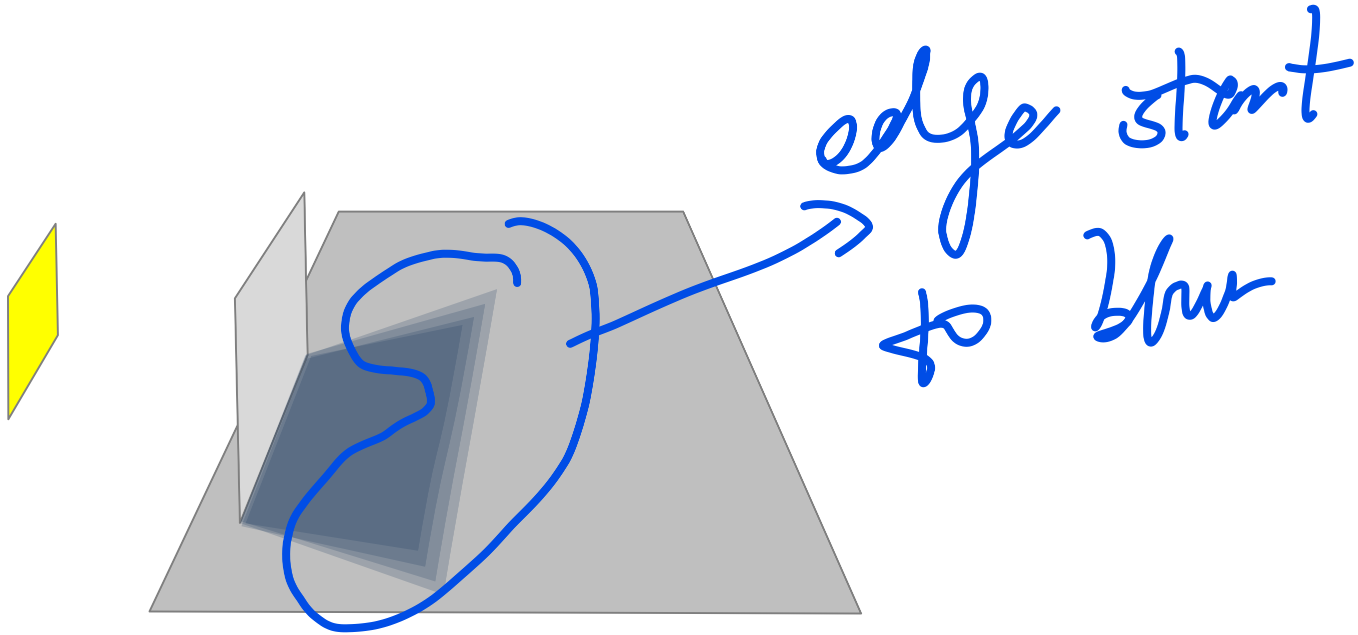

edge start to blur

Good book

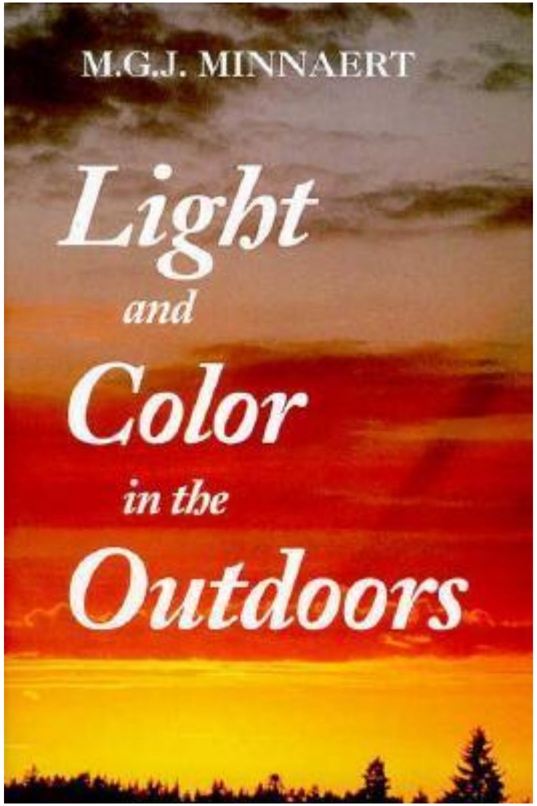

(Will not help you with passing this course)

Shadows are Complex

- In the real world sources of light are not points
- The intensity within a shadow is not constant
  - **Umbra**
    part that sees nothing of the source
  - **Penumbra**
    part that receives some light

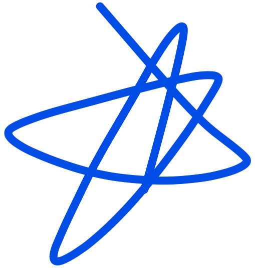

terminology

soft shadows

Definition

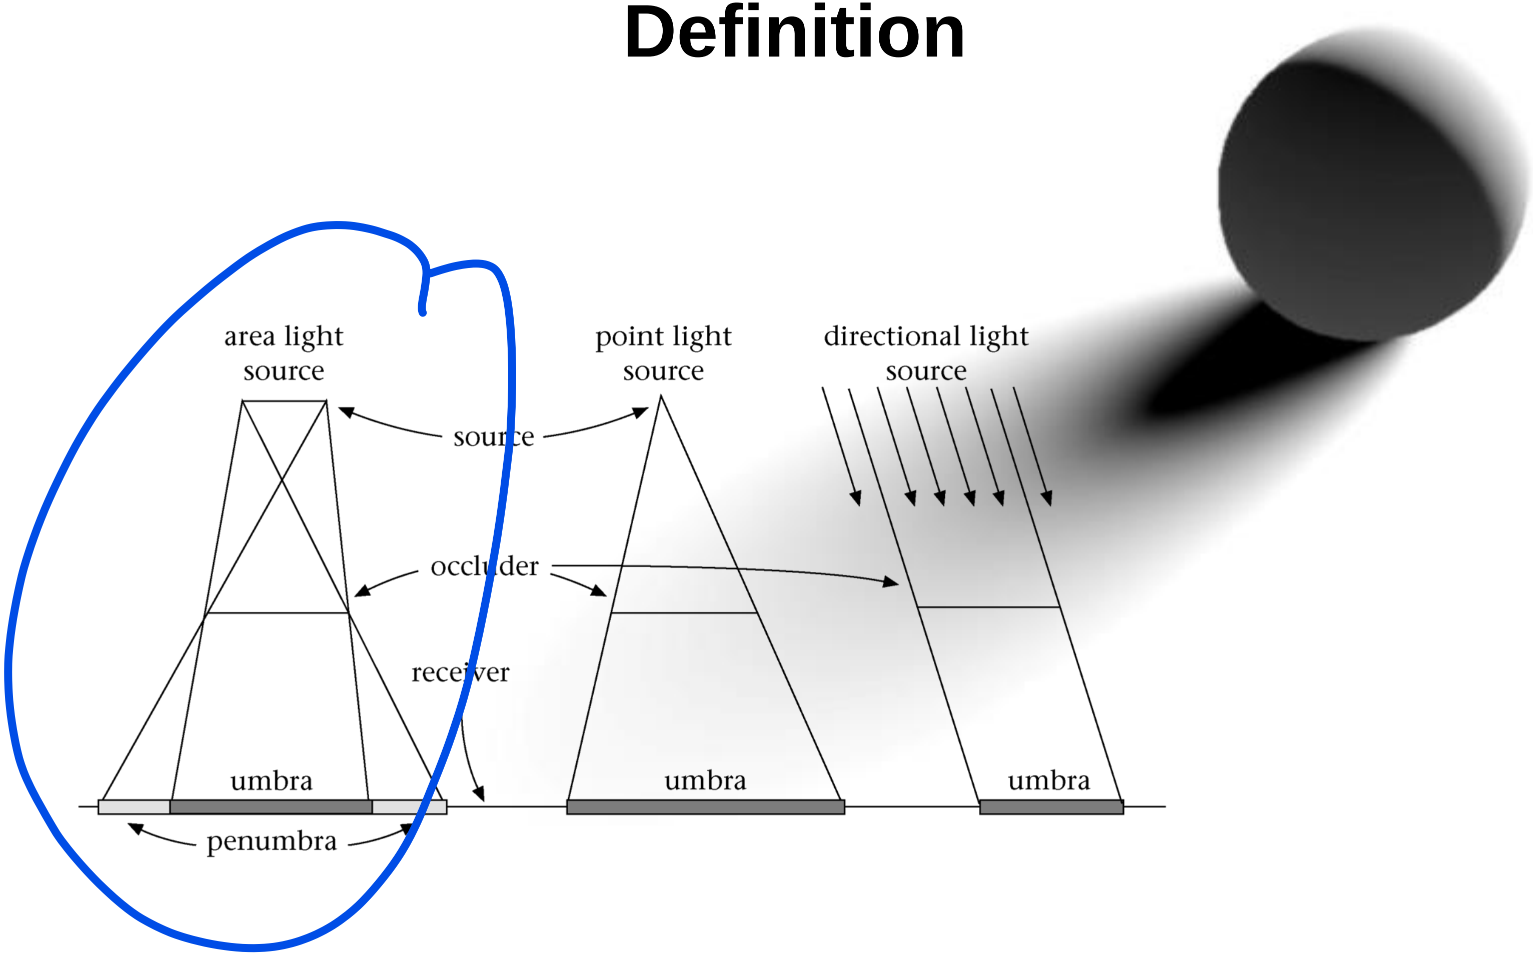
area light
source
point light
source
directional light
source
source
occluder
receiver
umbra
umbra
umbra
penumbra

Shadow in ray-tracing

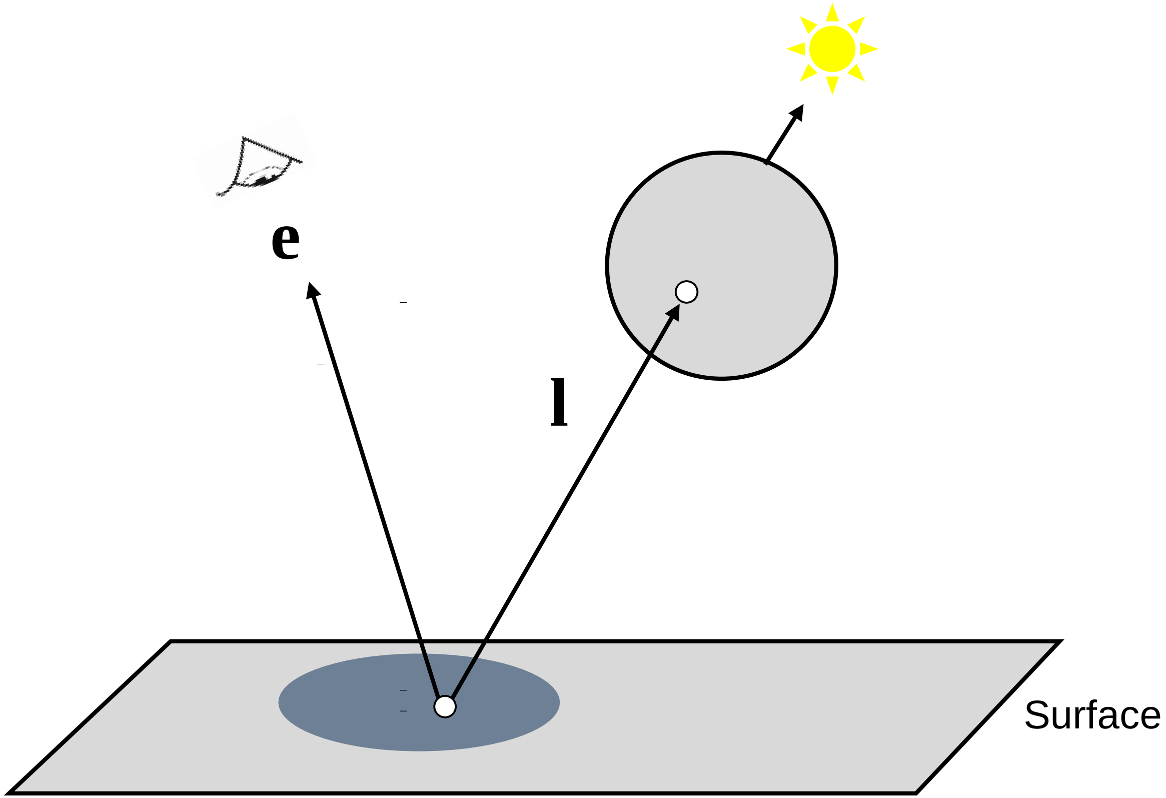

Soft shadows in ray-tracing

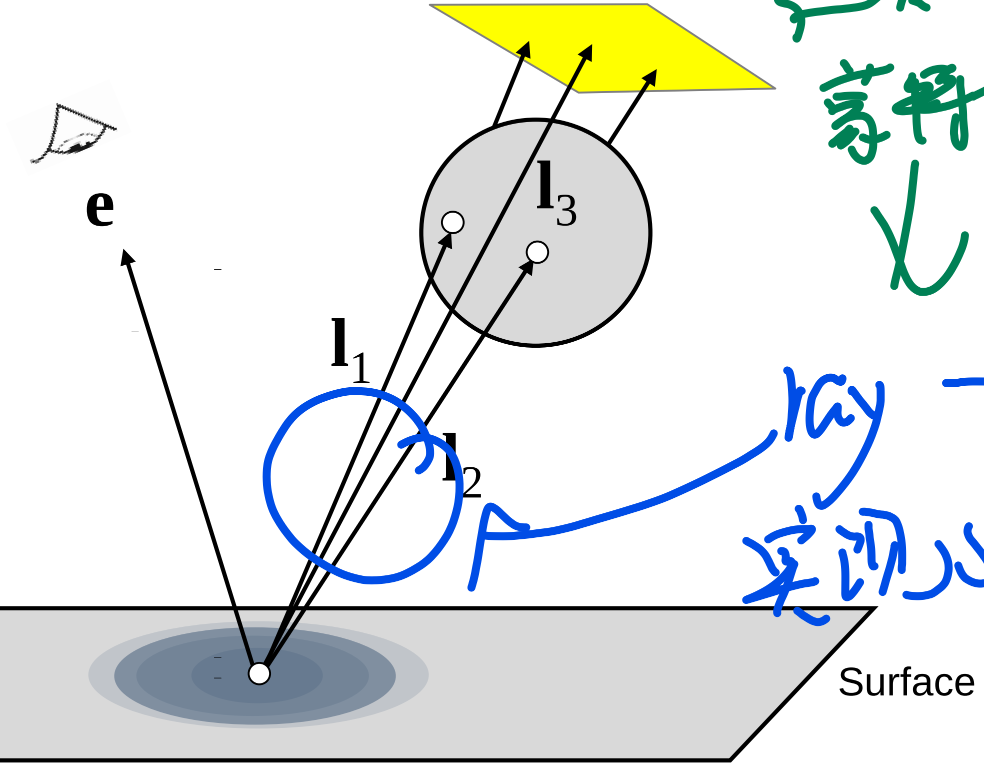

但实际用
蒙特卡洛做
✓

ray trace
实现 soft shadow
的方式

Current Shadowing Methods

- There exist a large number of methods
- We are interested in methods suitable for interactive walkthroughs, speed is crucial
- We will classify them by complexity:

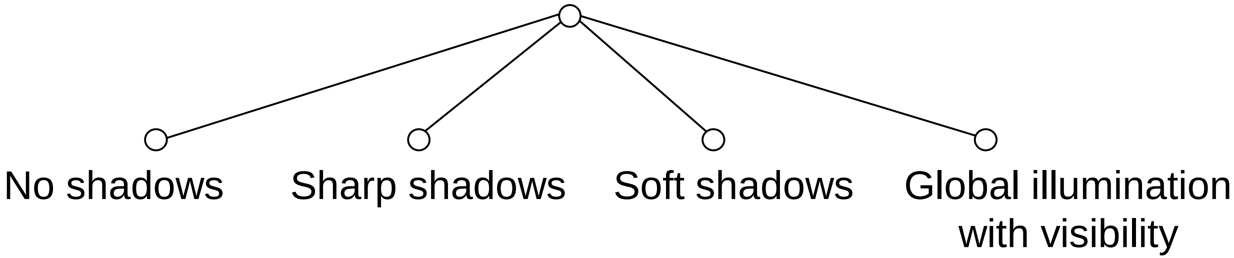

No shadows

Sharp shadows

Soft shadows

Global illumination with visibility

Sharp shadows

Sharp Shadows

- Source is assumed to be a point or direction

light & models

can move

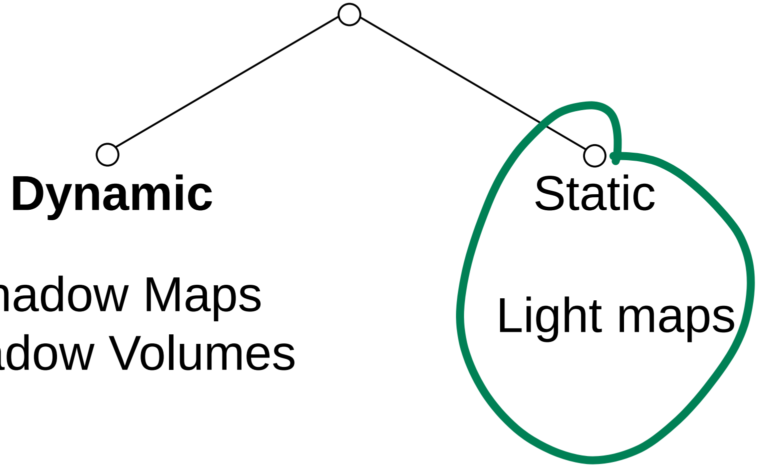

Dynamic

Shadow Maps
Shadow Volumes

Static

Light maps

之前学的

Shadow Maps

发现map B 中这里是别人的

Light / shadow duality:
A point is in **not** visible from the light source

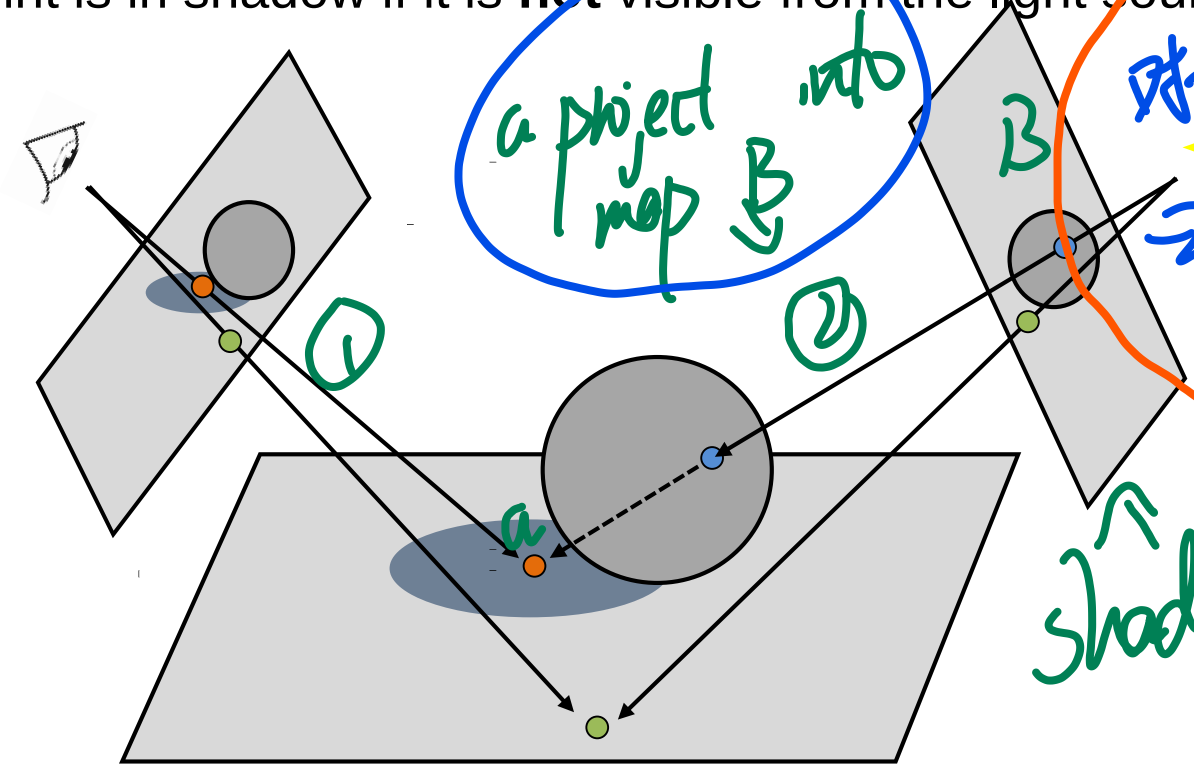

Two passes

- Shadow map pass
  - Use the light source as a view point (light space)
  - Render scene
  - Store depth information in a shadow z-buffer = **shadow map**
- Shading pass
  - Render scene as usual form the cameras view point
  - Each pixel's position (x_v, y_v, z_v) is transformed to light space (x_s, y_s, z_s) ,
  - If
    - the z_s value is less or equal to the shadow map at x_s,y_s it is **lit**.
    - else it is **shadowed**

Shadow Maps

Shadow Maps

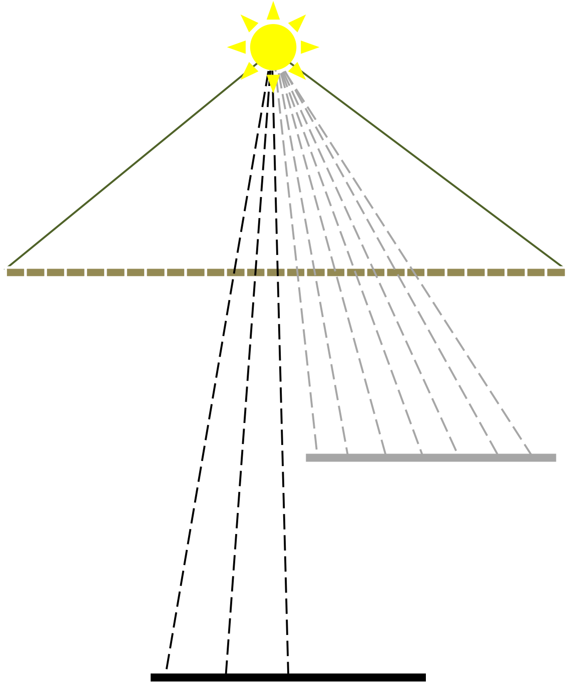

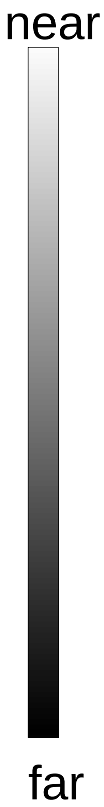

Shadow Maps

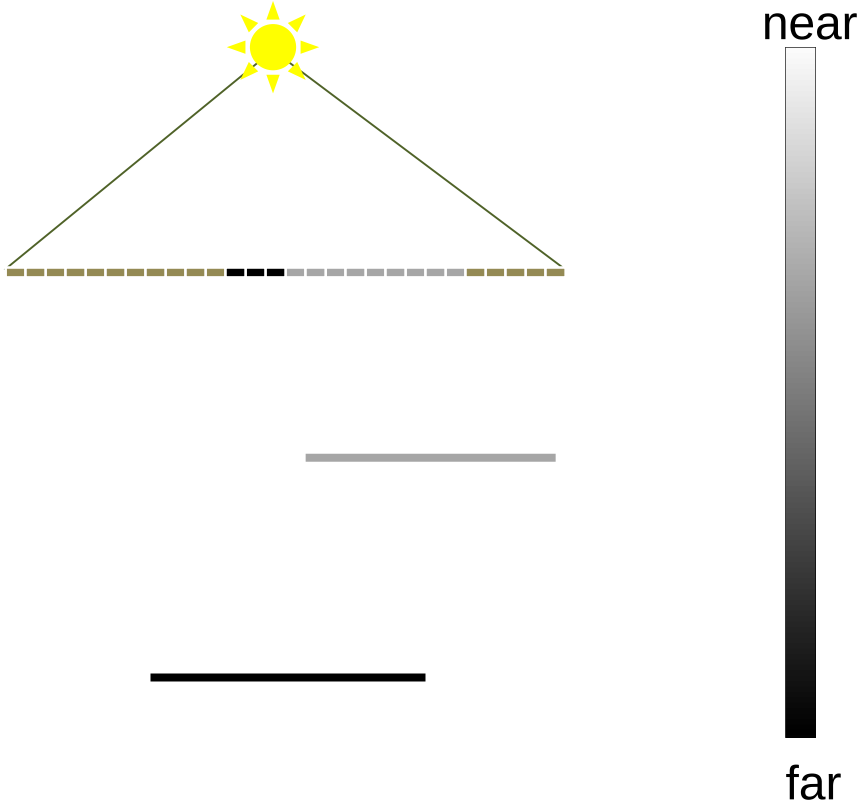

Shadow Maps

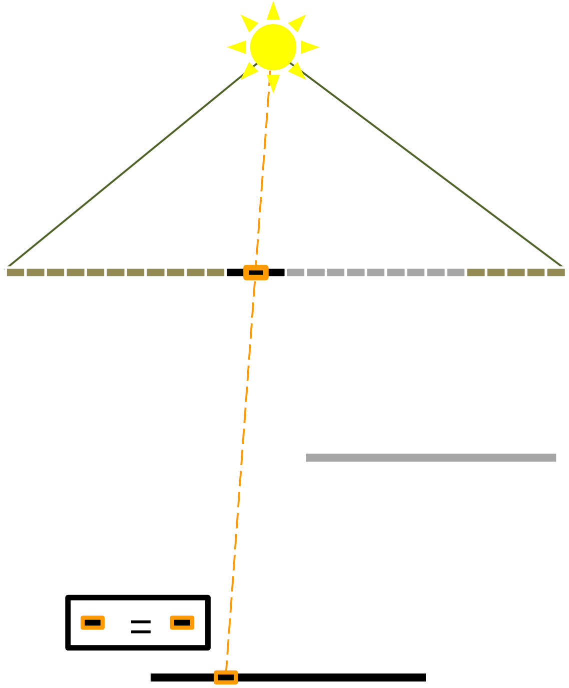

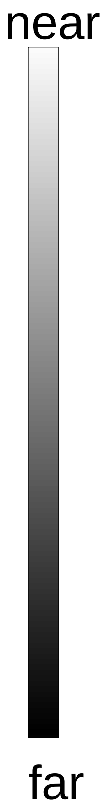

Shadow Maps

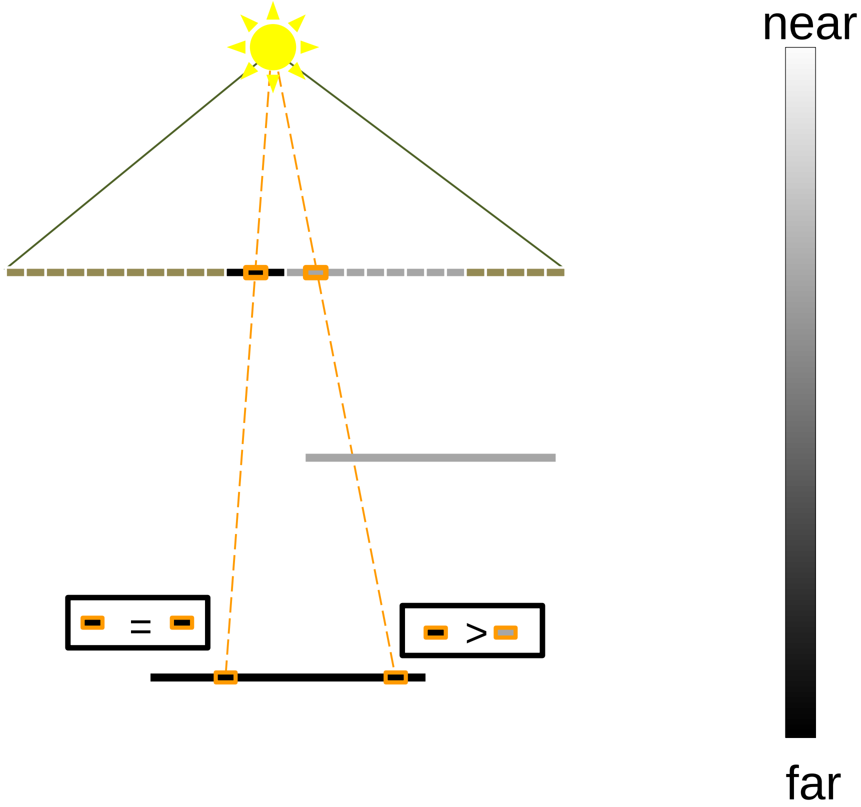

Shadow Map Filtering

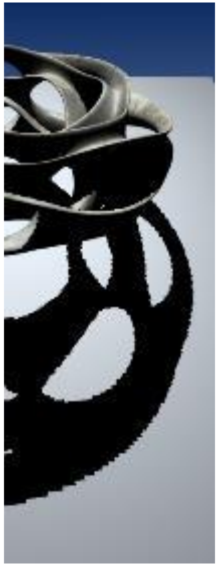

unfiltered

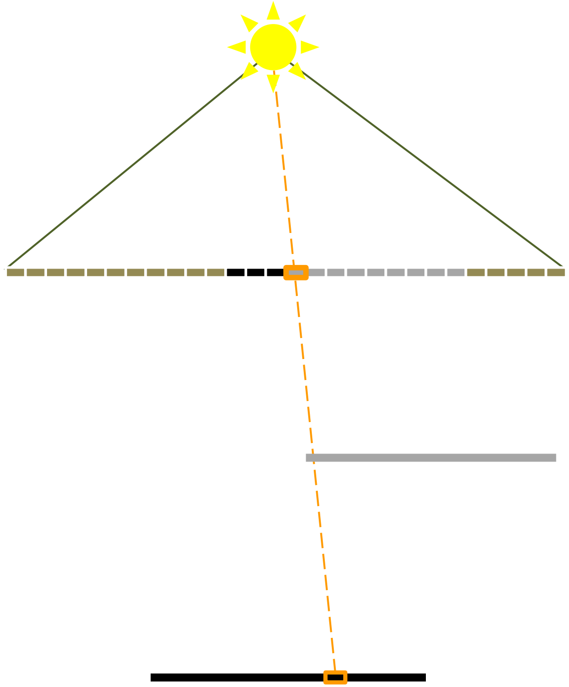

Shadow Map Filtering

多次采样 ShadowMap

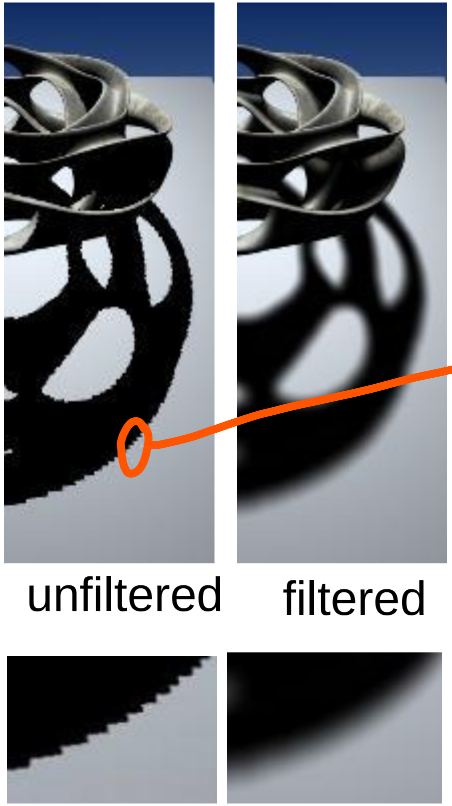
unfiltered

filtered

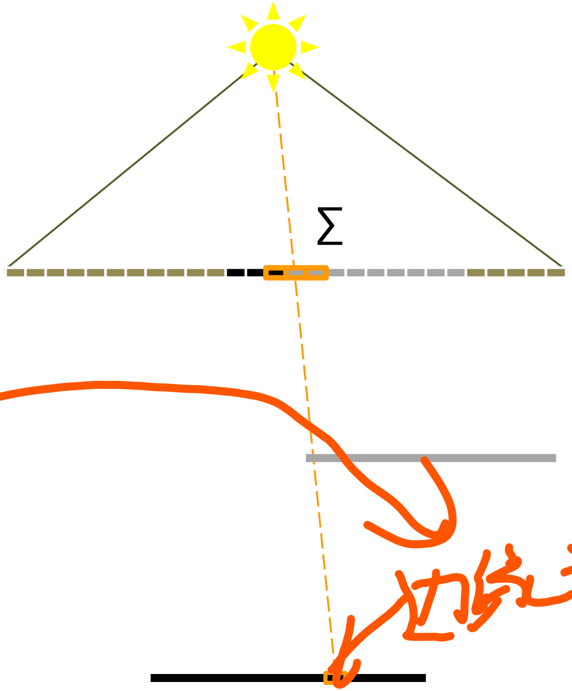

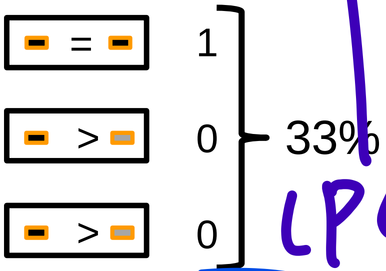

(PCF)

边缘 ⇒ also take into its neighbours

考虑周边遮挡情况

Percentage Closer Filtering
[Reeves et al. '87]

Shadow Map Filtering

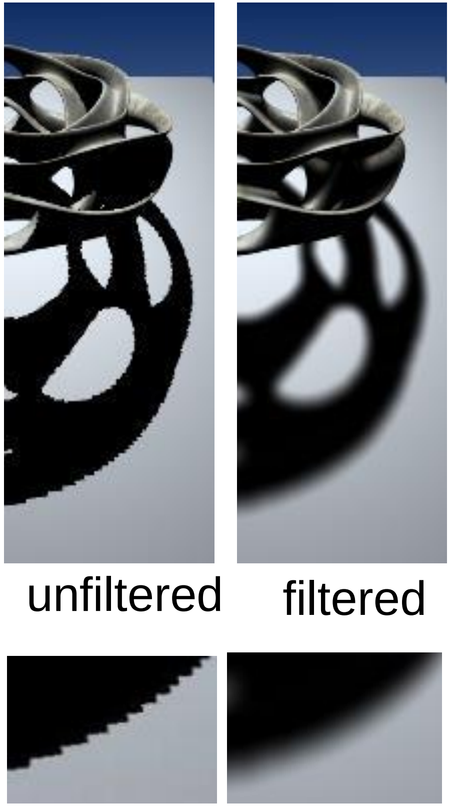

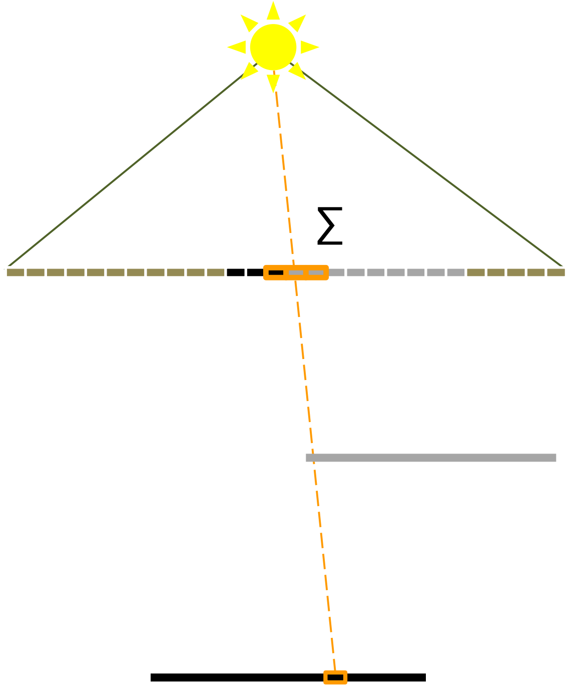

Percentage Closer Filtering [Reeves et al. ‘87]

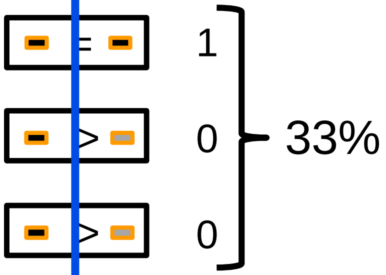

compare first, then filter!

↓

Expensive

问题是expensive 还进行 soft 的

Self shadowing: Disaster

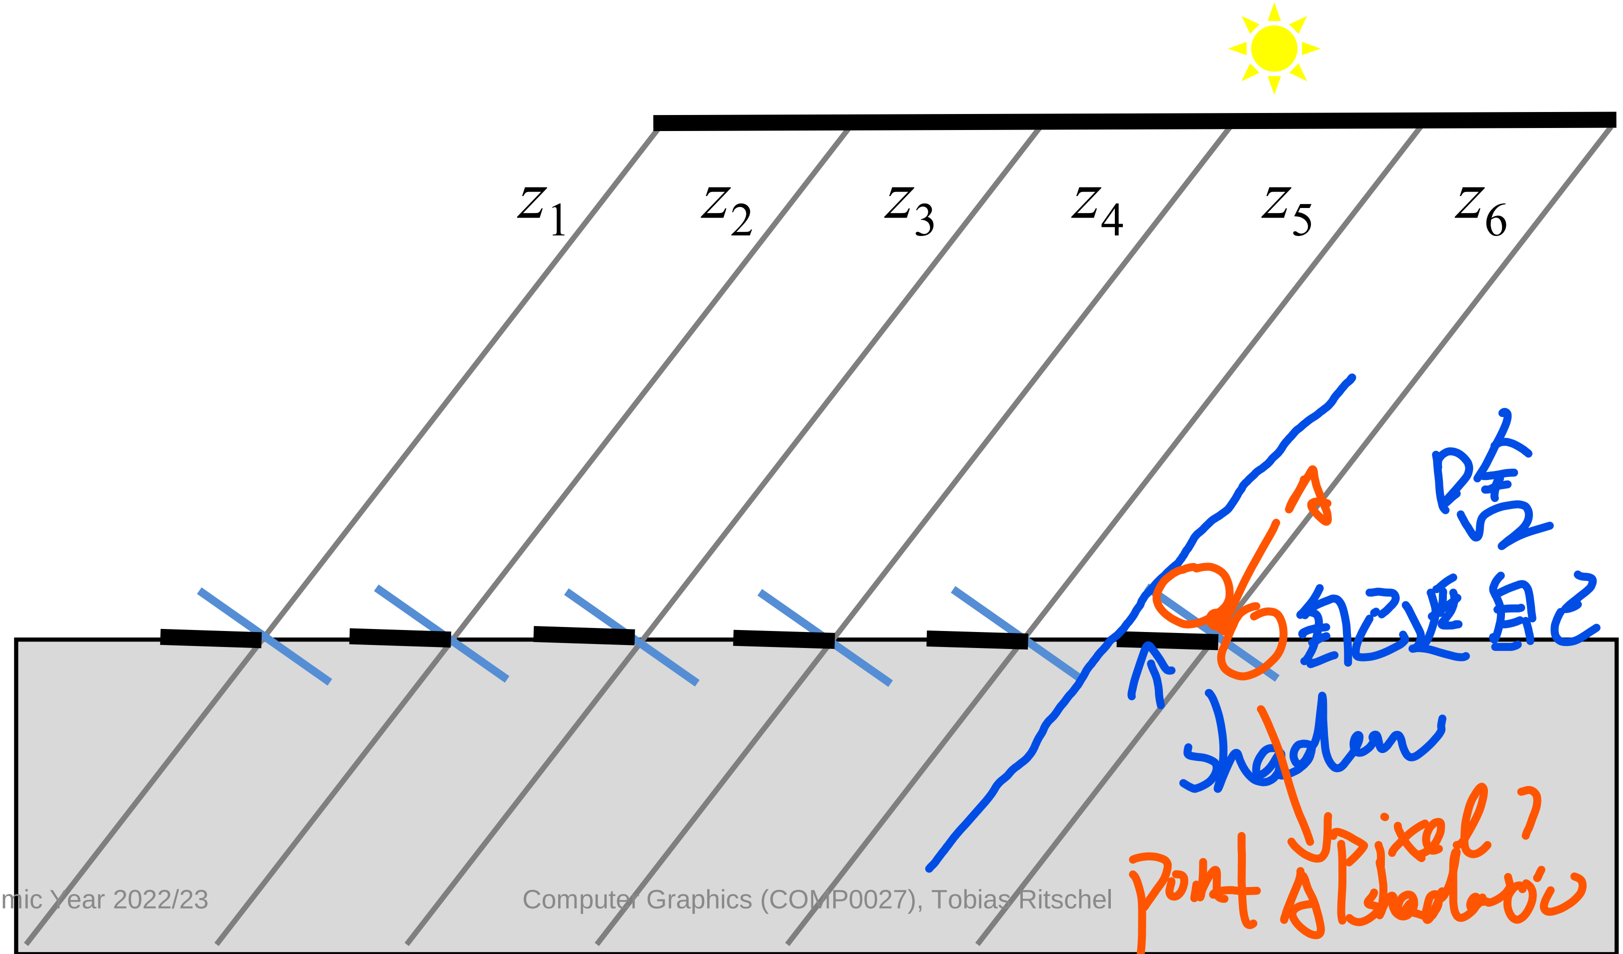

Self shadowing fix

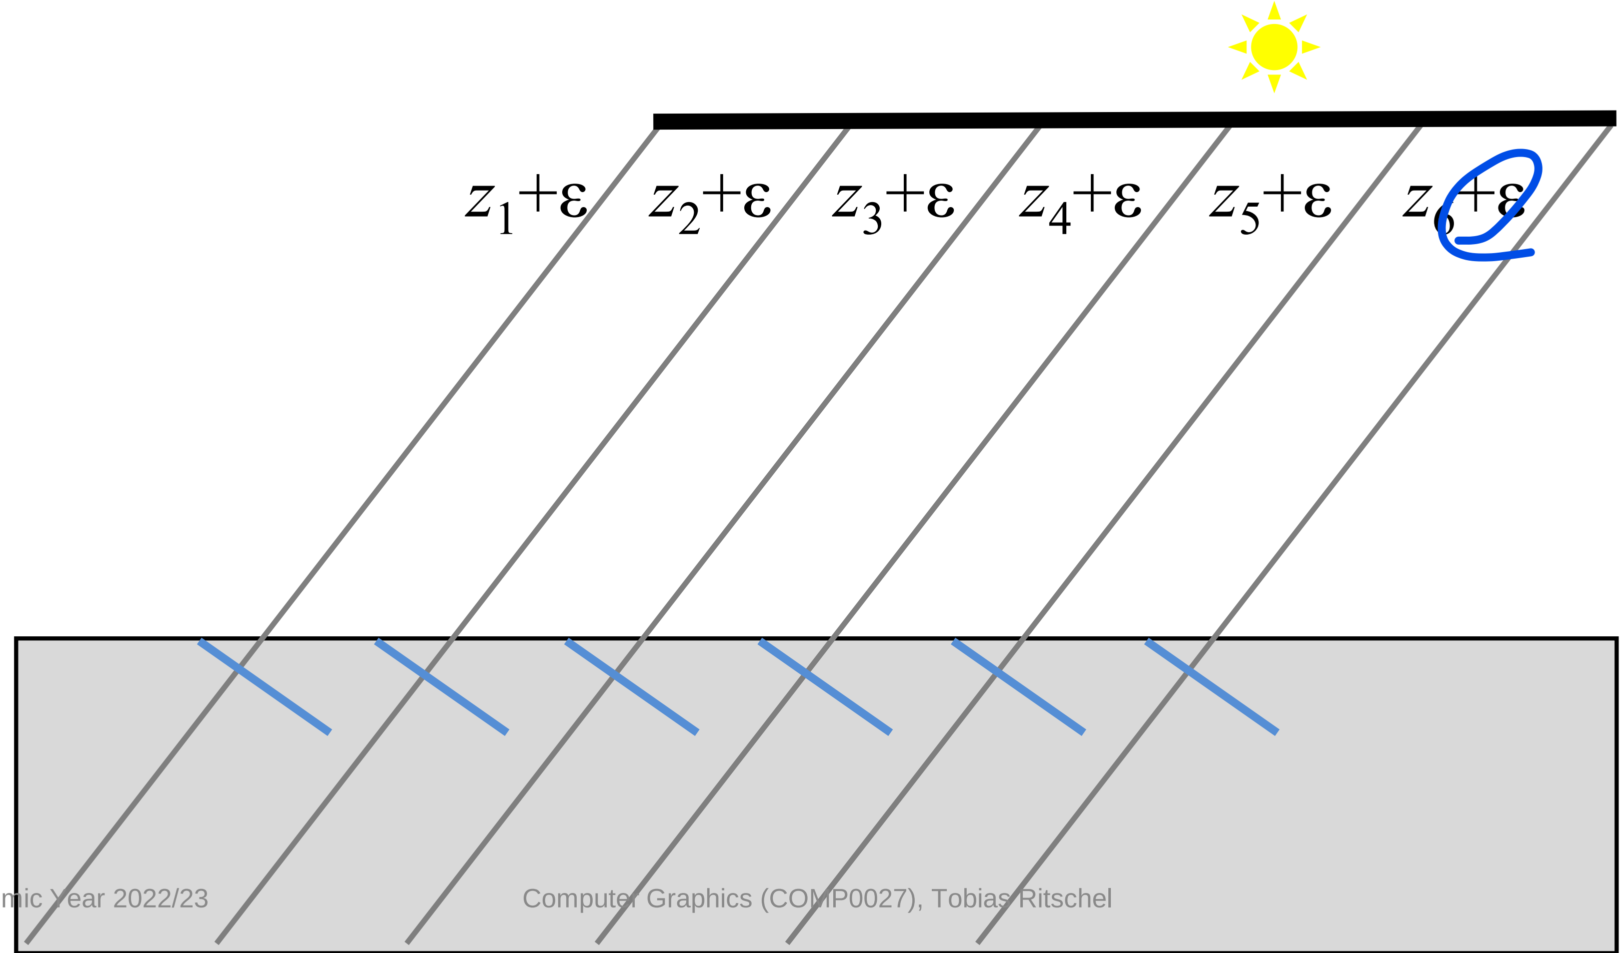

Bias (Epsilon) for Shadow Maps

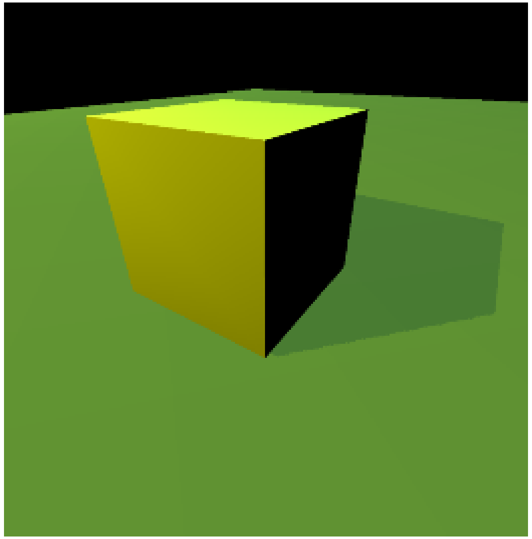

Correct result

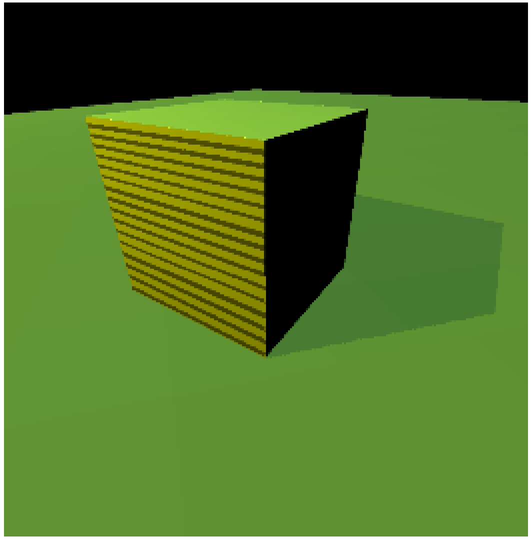

Not enough bias

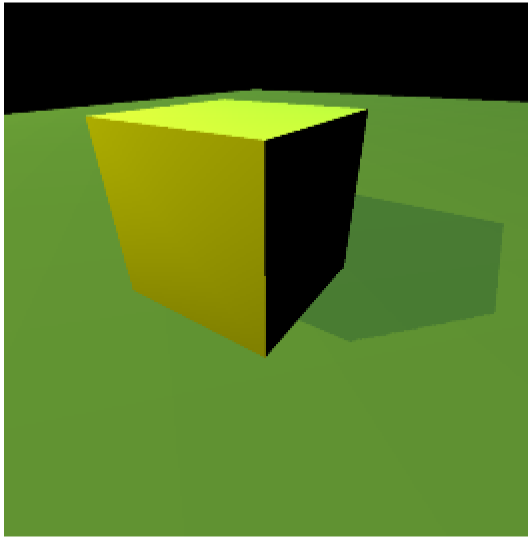

Too much bias

Shadow Maps

- “Less than or equal” test is imprecise
  - Gives rise to “shadow acne”
- Often found in hardware now
  - Otherwise high-cost operation
- Imprecise since it is only accurate in the image space of the light
  - Imagine a shadow throw over complex objects or long distances
- Quality depends on resolution (jagged edges)
  - Percentage-closer filtering helps
- FOV of shadow map?

Shadow Volume Method

妹[illegible]过吕

- Shadow volume (SV) is the volume of space below a polygon that cannot see the source (a culled pyramid)
- During rendering of image, the line from a point visible through a pixel to the eye is intersected with all object SVs
- The number of intersections indicates if the point is in shadow or not

check point whether in shadow volumn

Shadow Volumes

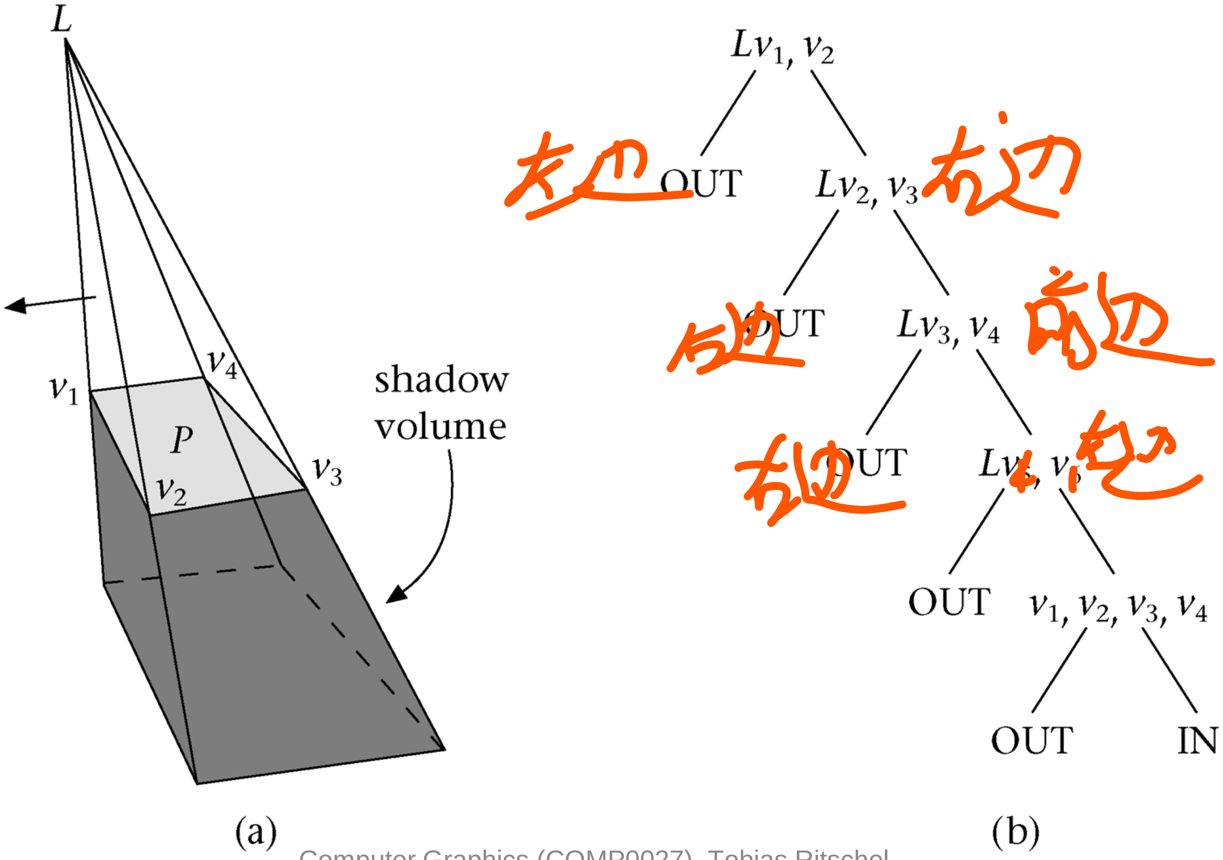

Shadow Volumes

- Just like a polygon - you are inside a volume if you need to cross a surface to exit it
- General idea of shadow volumes is count the number of shadow planes you cross
  - +1 for front facing
  - -1 for back facing
- If total is >0 you are in shadow ← 奇数 (进入次数)
- Special case if the eye itself is in shadow

↑ 不能这么算

↓

Shadow Volumes

Two stages:

1) Volume **construction**
   - Find all planes of the shadow volume and their plane equations

2) Volume **test**
   - Determine shadow plane count per pixel
   - Use a scan-line method OR stencil test

Shadow Volume Example

volumn
进圈一蓝一橙

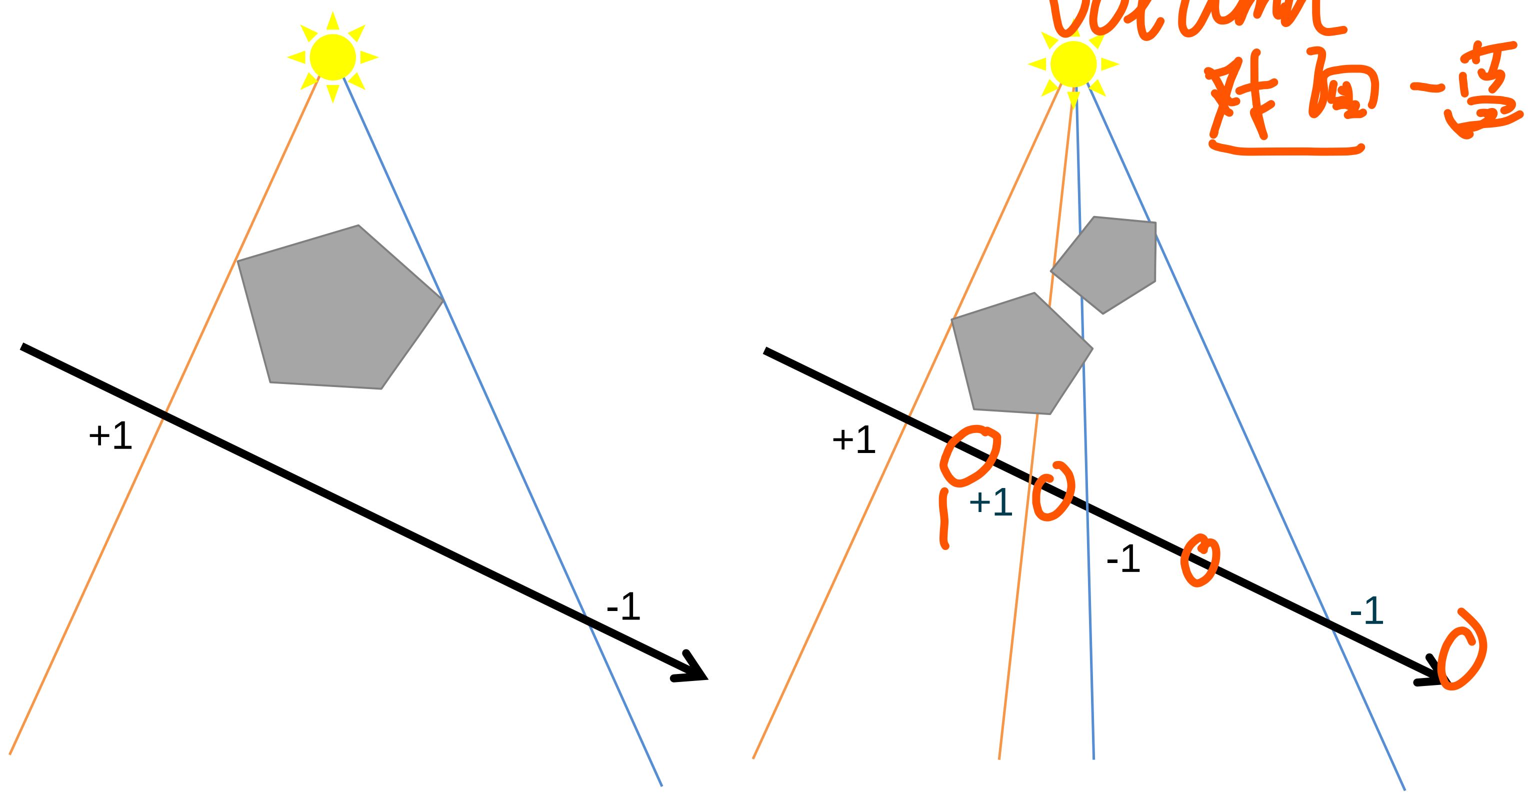

Shadow

Shadow

Shadow Volumes with OpenGL

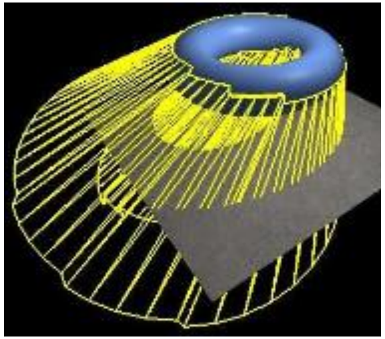

- Shadow volumes are rendered at each frame
- The stencil buffer is used for counting how many SV are crossed
- Sometimes not all objects are used for casting shadows

Shadow Volumes & Stencil Test

模板测试

- A stencil buffer is screen sized buffer (1-8 bit) that stores a flag about a rendering operation
  - E.g., `stencil[x, y]` is negated if `zbuffer[x, y]` is less than current z value (i.e. stencil is set if and only if *z* buffer-test passes)
- Many uses in graphics

所以 stencil Test 是用 shadow map 做的!?

Shadow Volumes & Stencil Test

- Render the scene into the RGB and z-buffer
- Turn z-buffer **writing** off
- Render all shadow volume polygons with the stencil buffer
  - Increment stencil count for front-facing
  - Decrement for back facing
- Re-render scene with lighting **off** and only render pixels where stencil is non-zero

No-shadow example

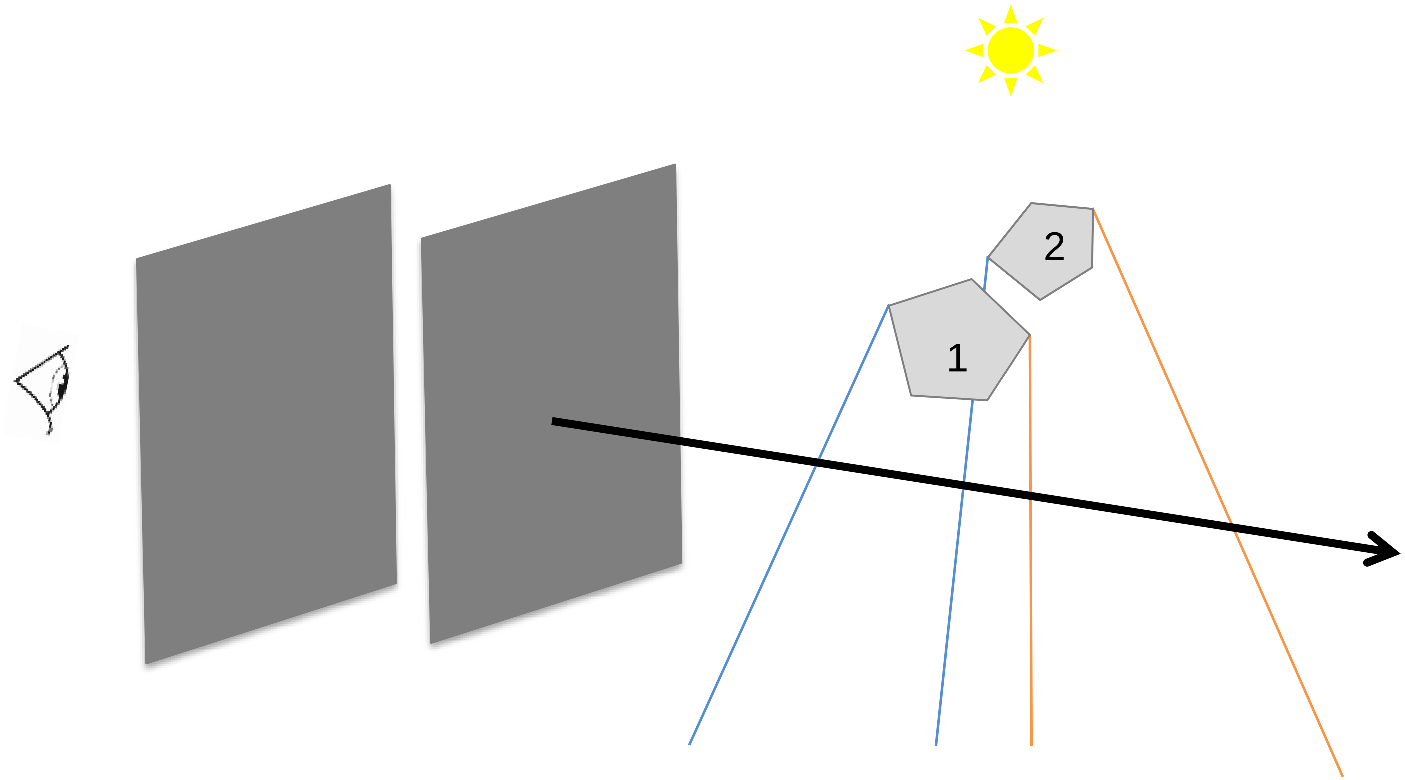

No-shadow example

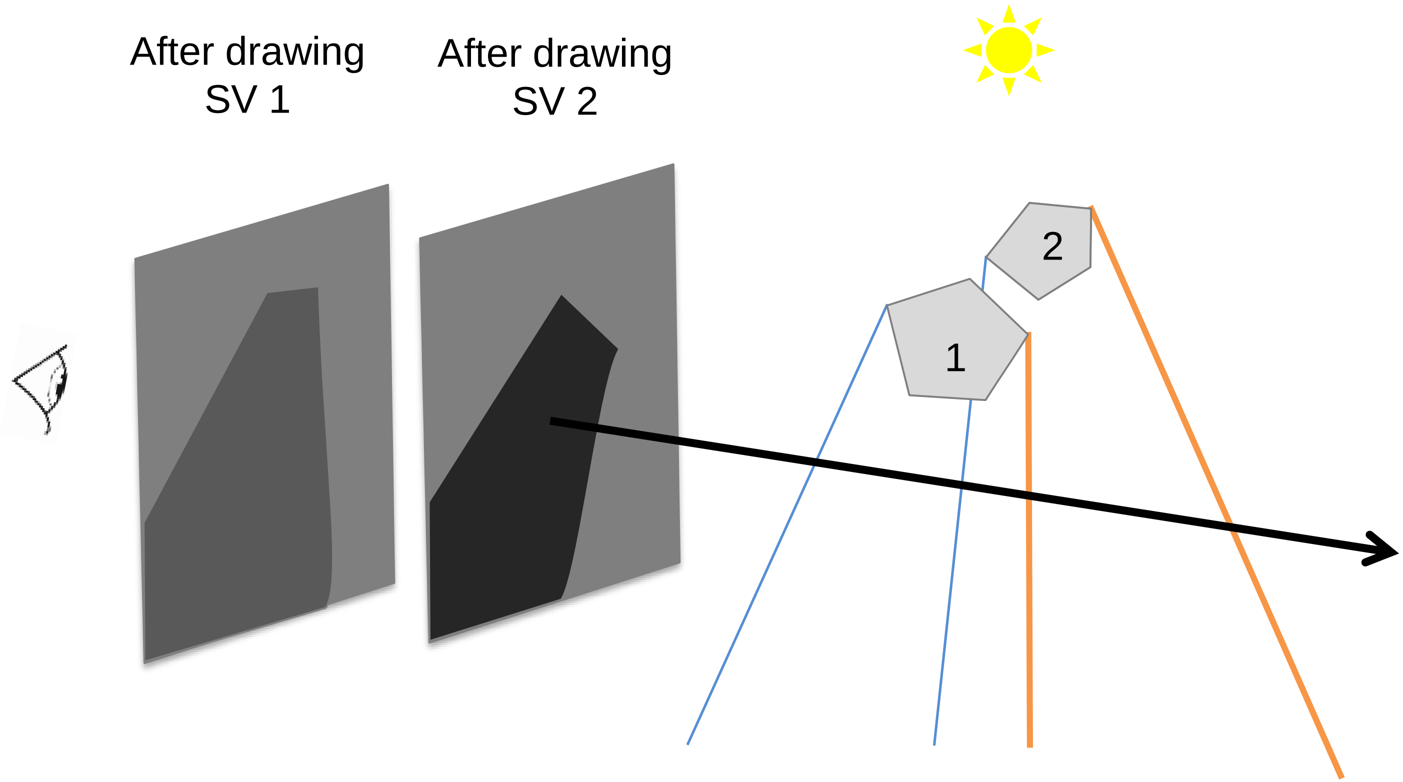

After drawing SV 1

After drawing SV 2

No-shadow example

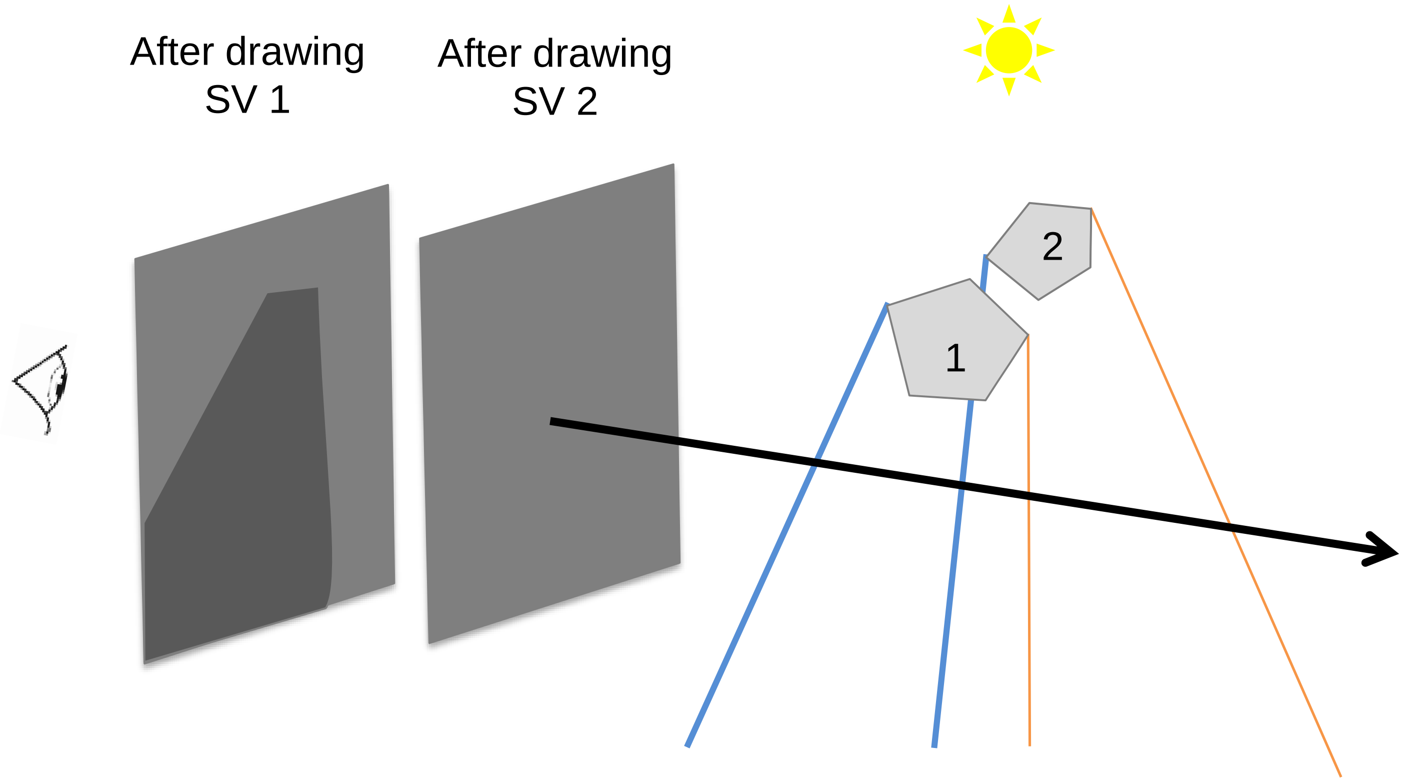

Has-shadow example

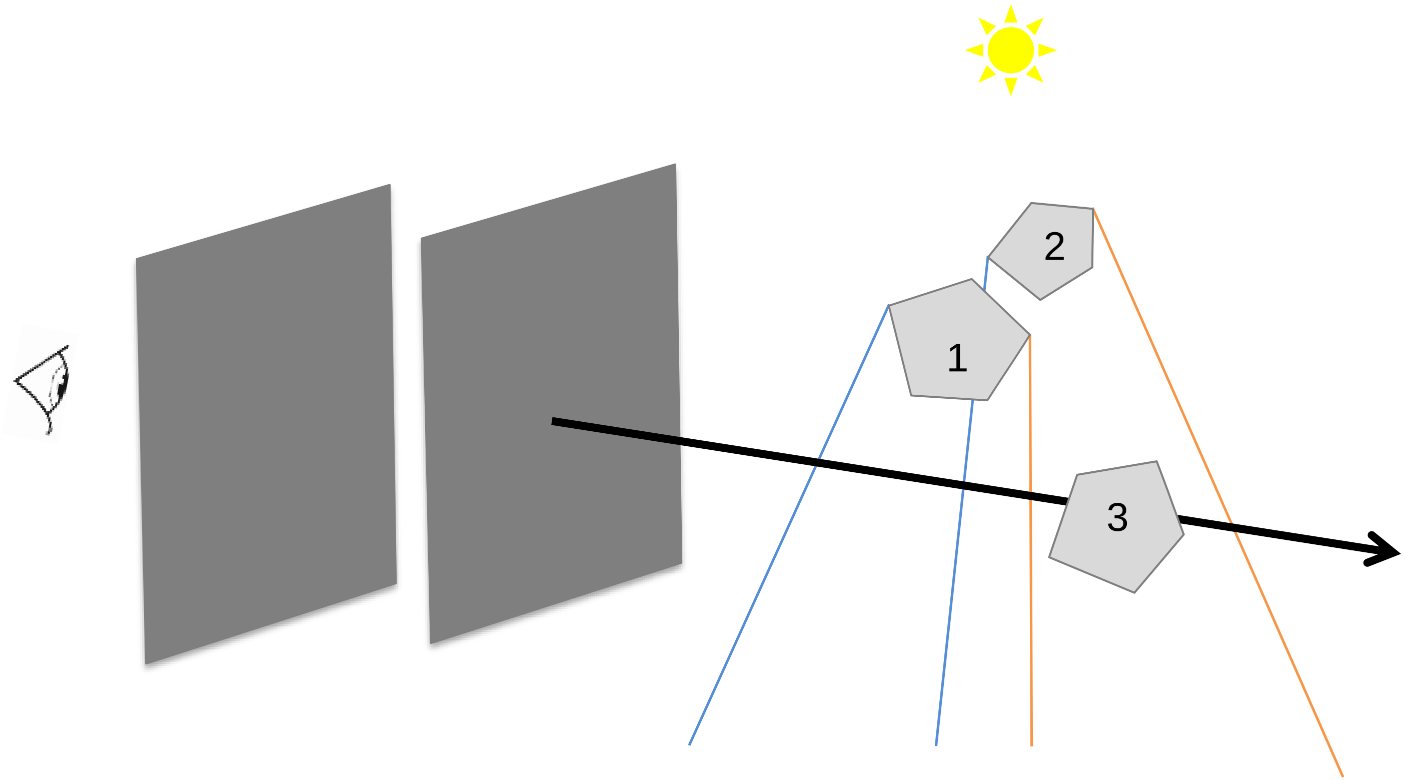

Has-shadow example

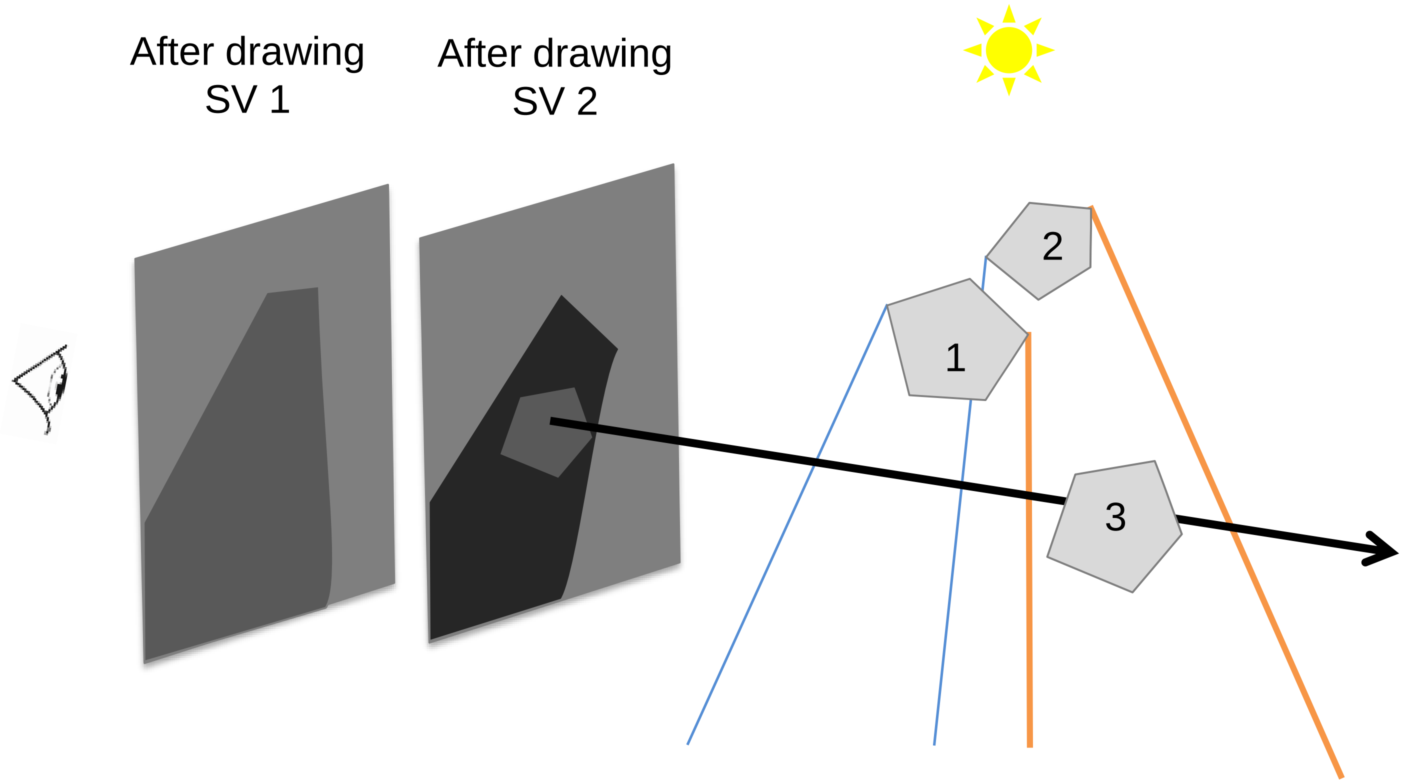

Has-shadow example

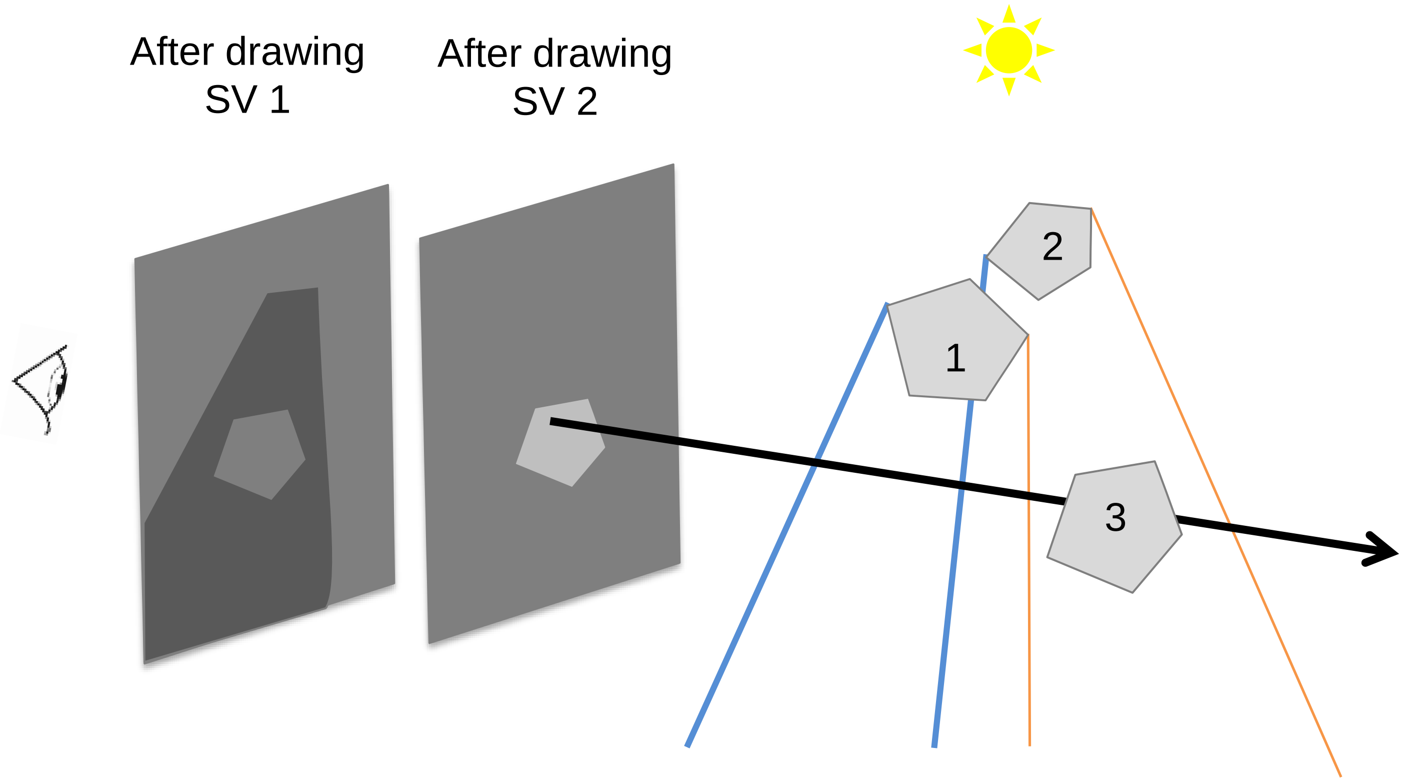

Summary for Sharp Shadows

- Four shadow umbra techniques
- Image space
  - Shadow maps
  - Shadow volumes

Soft Shadows

个

他没讲了

回去看

Soft Shadows

- Source has a finite extend
- Images look a lot more realistic

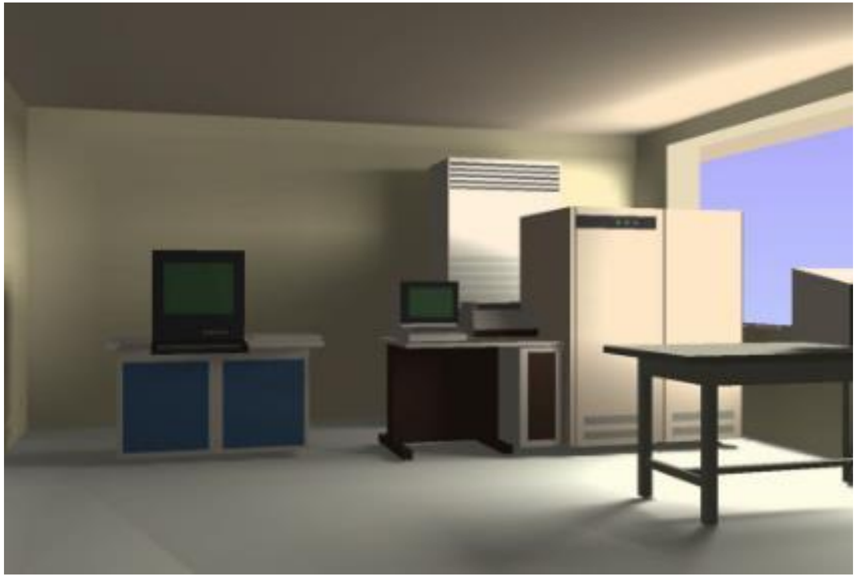

Soft Shadows

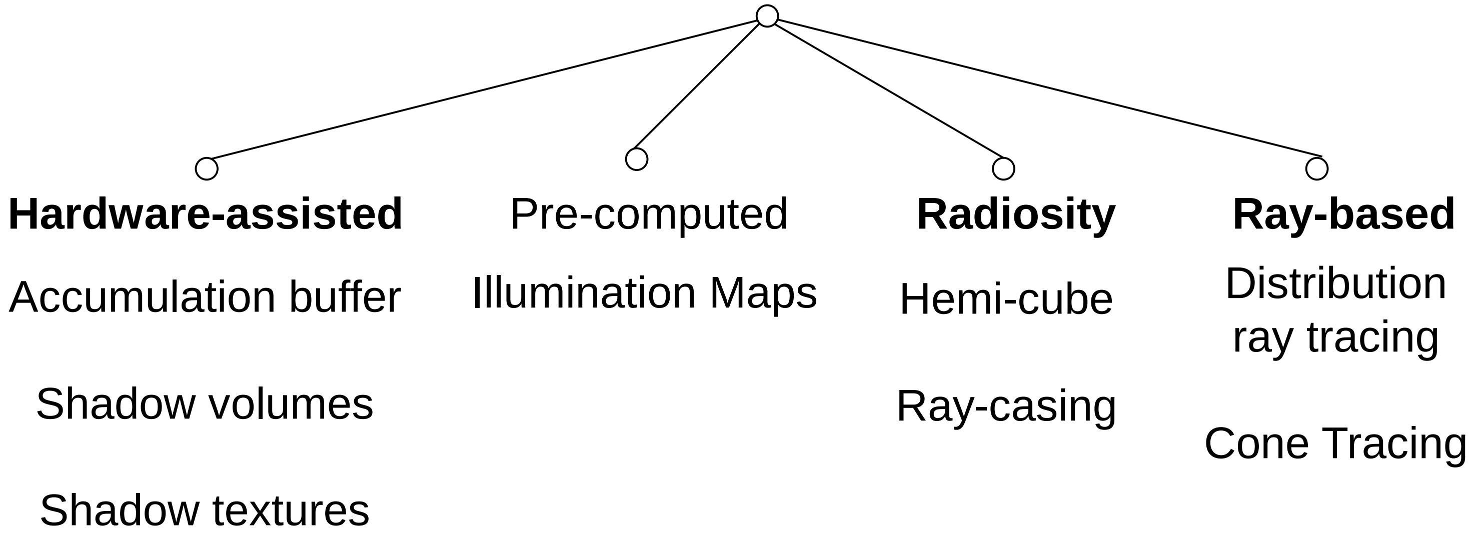

Hardware-assisted

Accumulation buffer

Shadow volumes

Shadow textures

Pre-computed

Illumination Maps

Radiosity

Hemi-cube

Ray-casing

Ray-based

Distribution ray tracing

Cone Tracing

Analytical v. Sampling

- Analytical
  - Find all boundaries within the penumbra. Done almost exclusively for polygonal light sources

上面的 terminology

一半影

- Sampling
  - Approximate solution that treat the light source as a set of points. Any shape source is possible.

在一个 shadowmap 的

: 周别 版块

为一个 人在 shadow 里的

Soft Shadows using Point Sampling

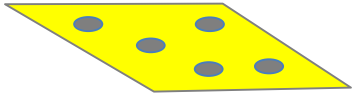

- Place many point lights on an area light
  - – Random positions work just fine
- Render hard shadows from each point light
  - – E.g., using shadow volumes or shadow maps
- Sum up all contributions
  - – Can be done on the GPU (in the frame-buffer)

- Similar to what ray-tracing does to get soft shadows

Example

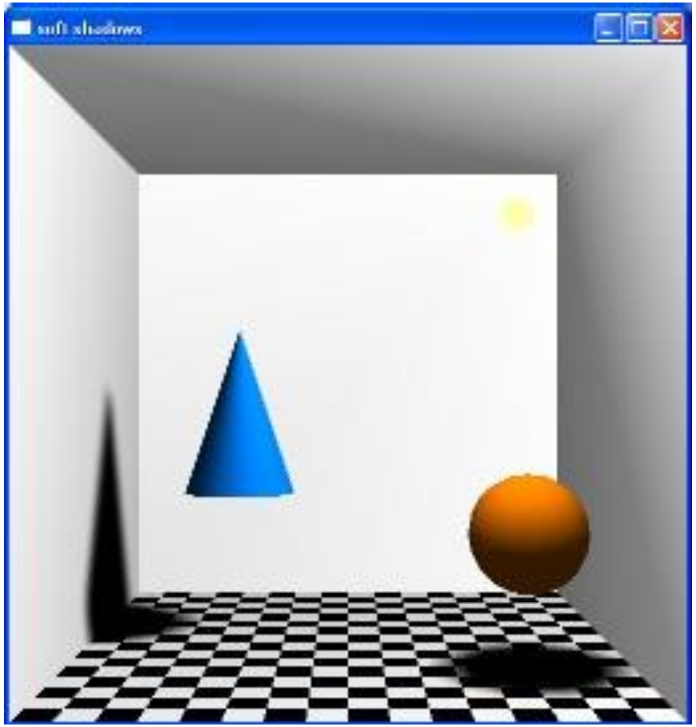

Illumination Maps (Shadow Textures)

- Shadows are pre-computed and stored as textures on the receiving polygons
- Displayed using graphics hardware in real-time
- Often use: Radiosity / Path tracing / Photon mapping
- Sometimes called “baked” lighting, very common in game engines
- Disadvantage: lighting cannot change

Recap

- In order to regain shadows in real-time engines, we have to do a lot of work
- A very large number of shadow algorithms exist
- Many of them are unsuitable for walkthroughs of very complex scenes:
  - with pre-computation methods scene cannot be modified
  - or are to slow (ray-tracing, soft shadows)
- Hard shadows
  - on-the-fly methods (SM and SV) are fast enough